

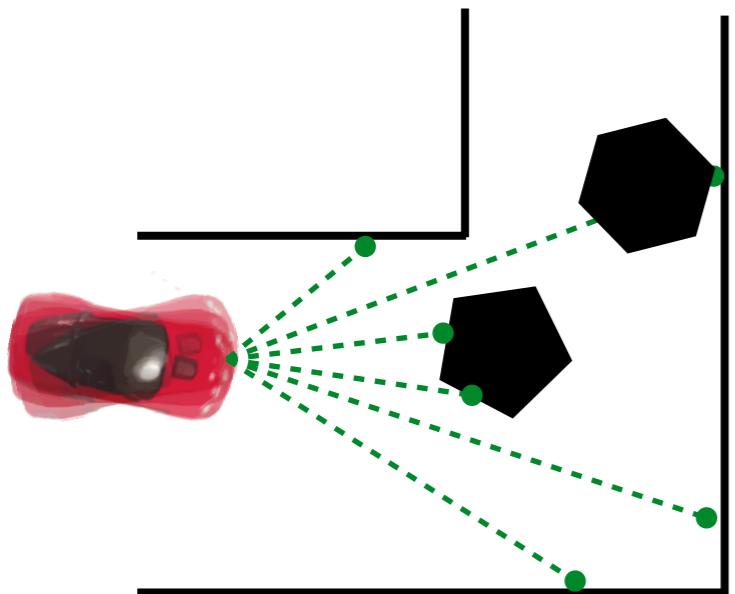
Content heavily adapted from Russ Tedrake

Feedback Control

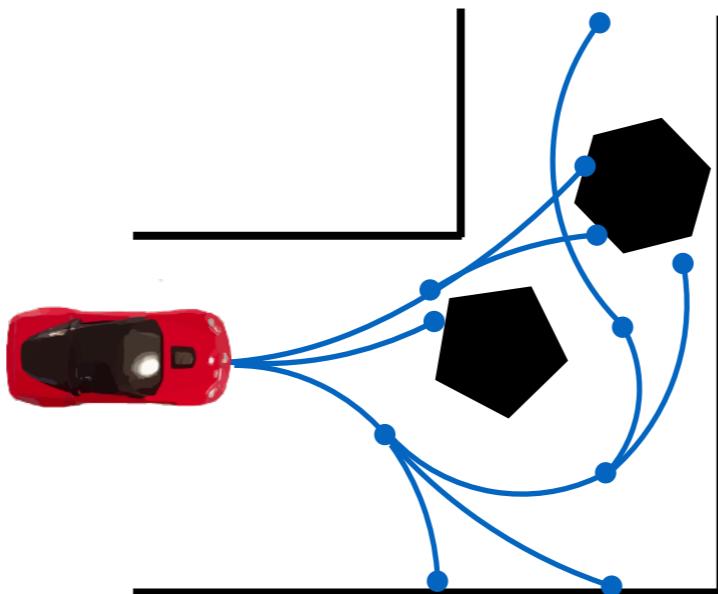
Sanjiban Choudhury

TAs: Matthew Rockett, Gilwoo Lee, Matt Schmittle

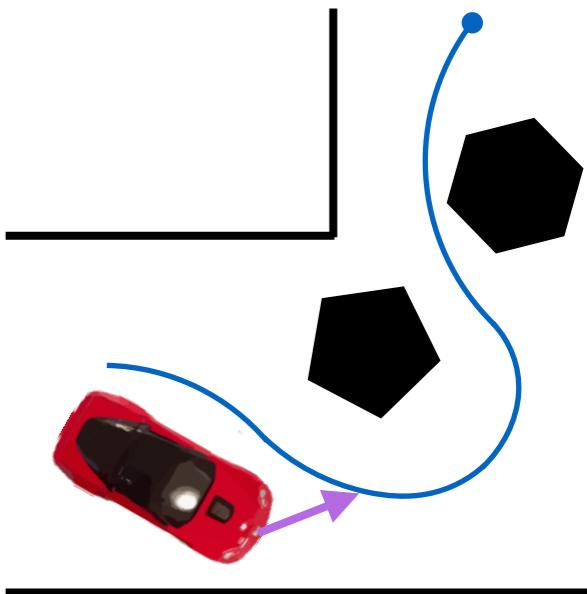
Estimate
state



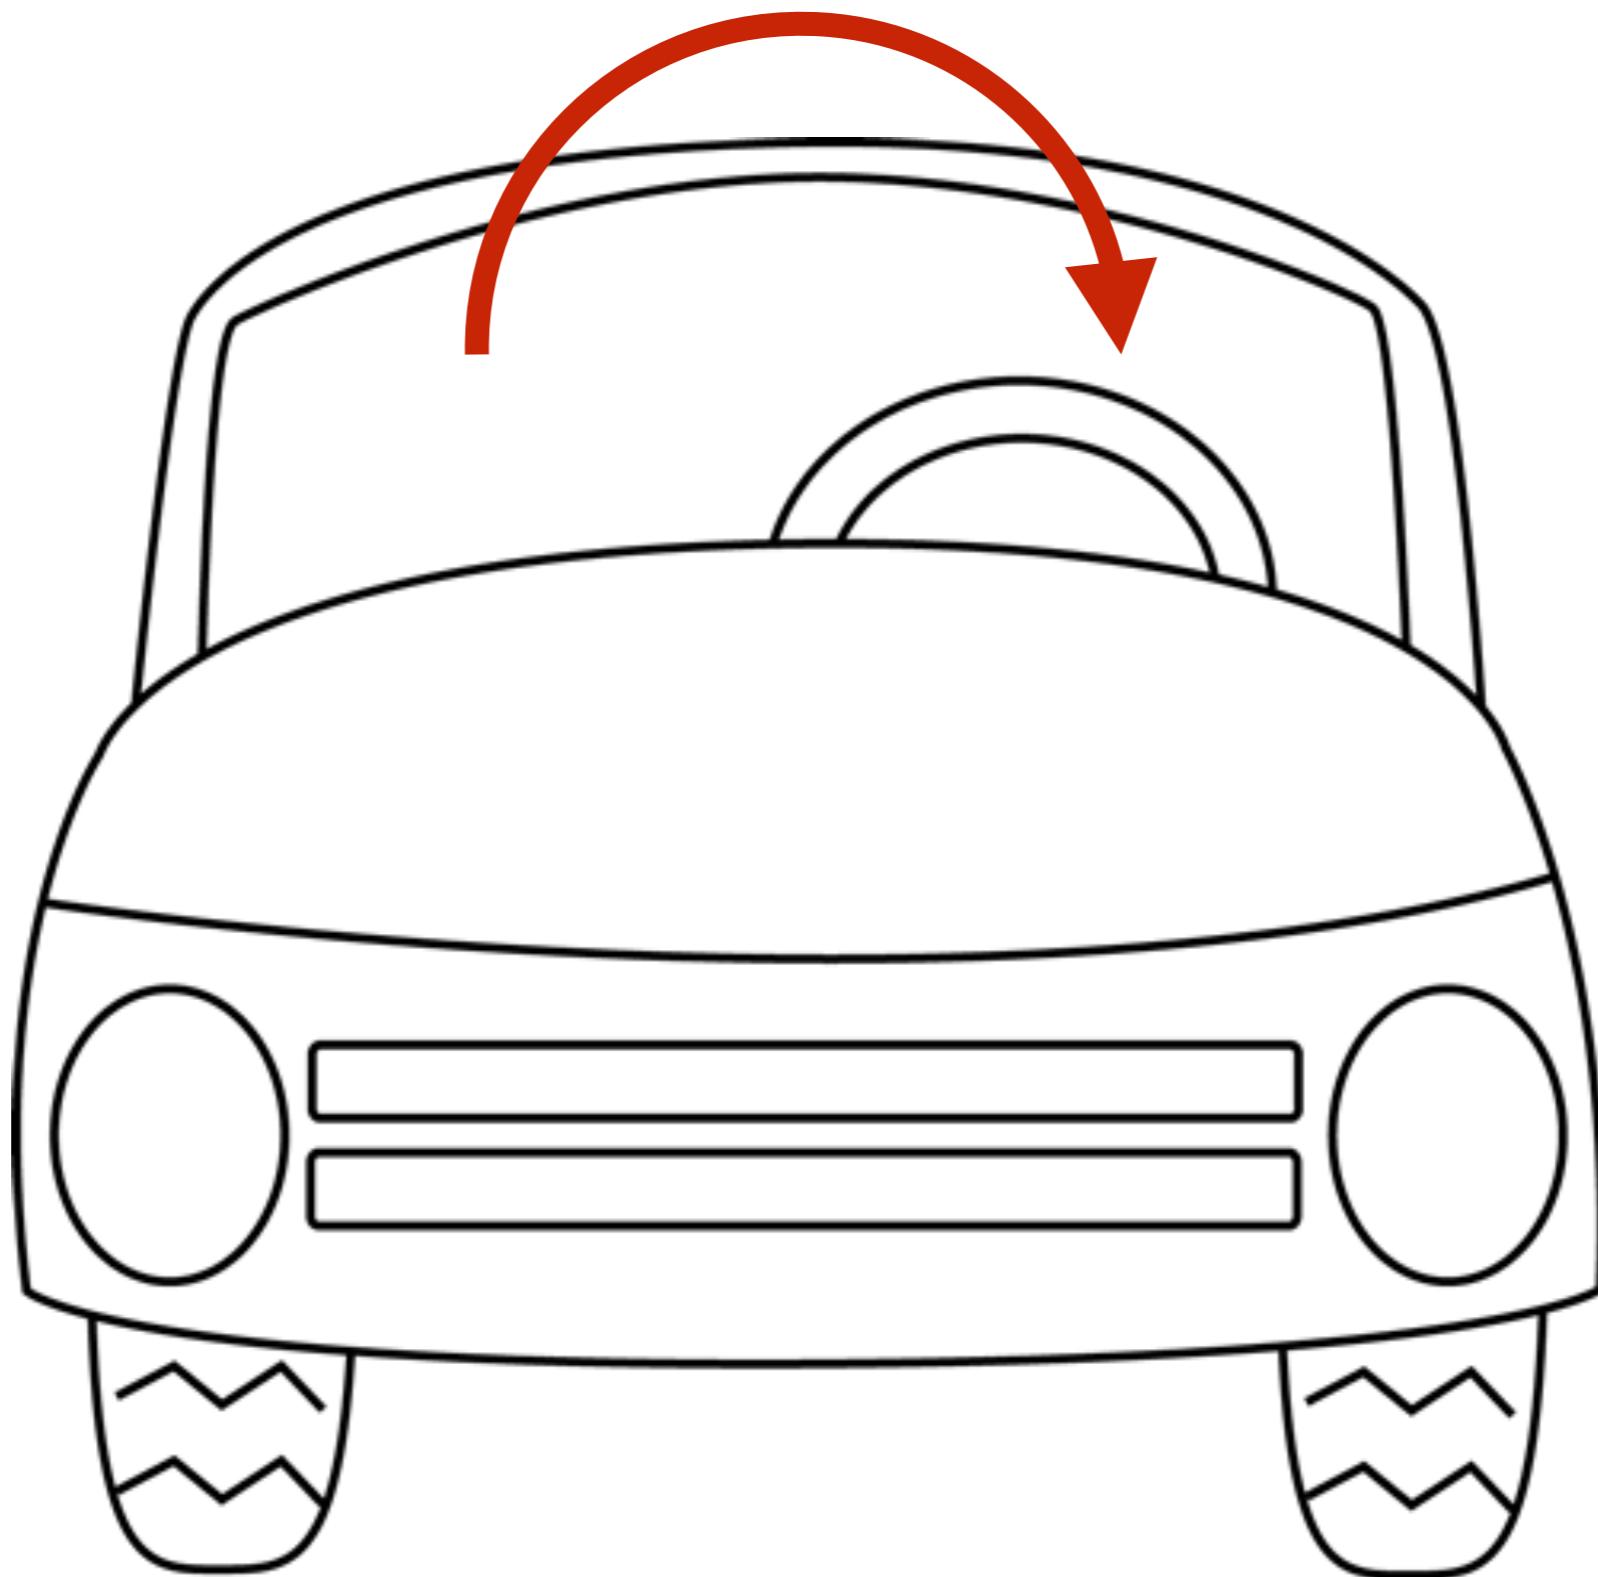
Plan a
sequence of
motions



Control
robot to
follow plan



From perception to control ...



When I think about control ...



<https://www.youtube.com/watch?v=WNR4MqG45pk>

When I think about control ...



<https://www.youtube.com/watch?v=WNR4MqG45pk>

Today's objective

1. Introduce terms and definitions in feedback control
2. Go through challenges in current control research

The control framework

Say we want to get the car to hit a sequence of poses

The control framework

Say we want to get the car to hit a sequence of poses



t_0, x_0, y_0, θ_0

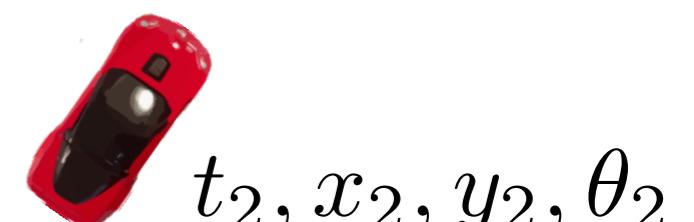
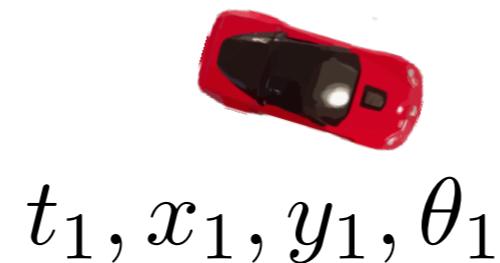
The control framework

Say we want to get the car to hit a sequence of poses

 t_0, x_0, y_0, θ_0  t_1, x_1, y_1, θ_1

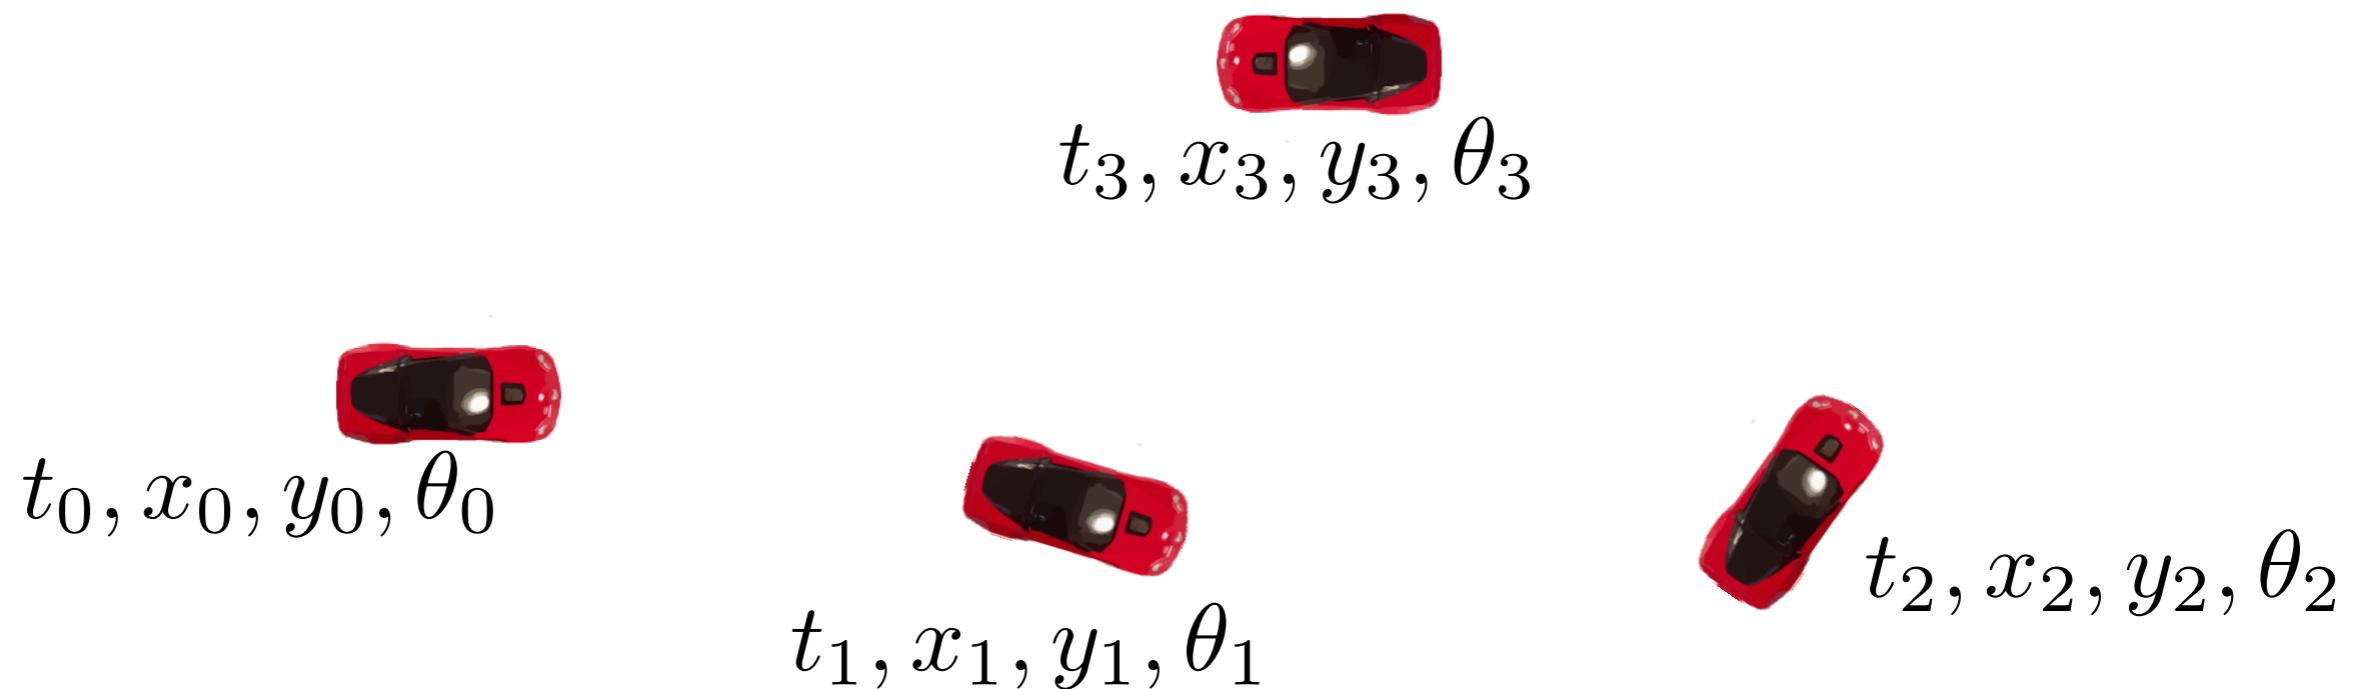
The control framework

Say we want to get the car to hit a sequence of poses



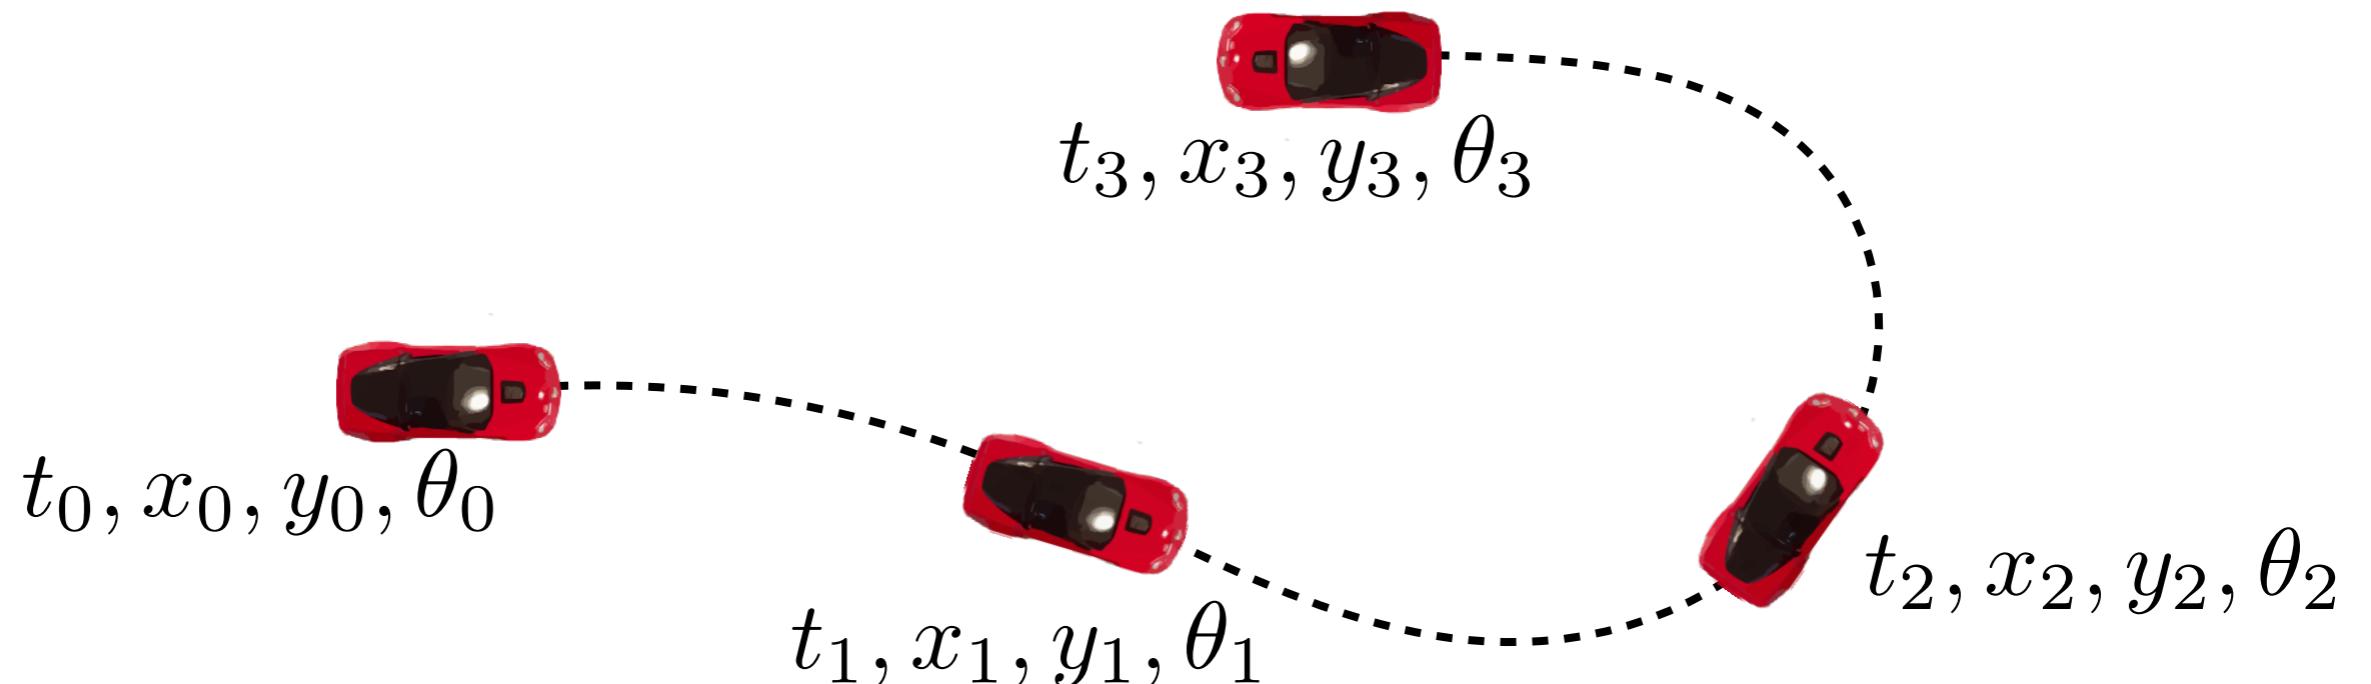
The control framework

Say we want to get the car to hit a sequence of poses



The control framework

Say we want to get the car to hit a sequence of poses



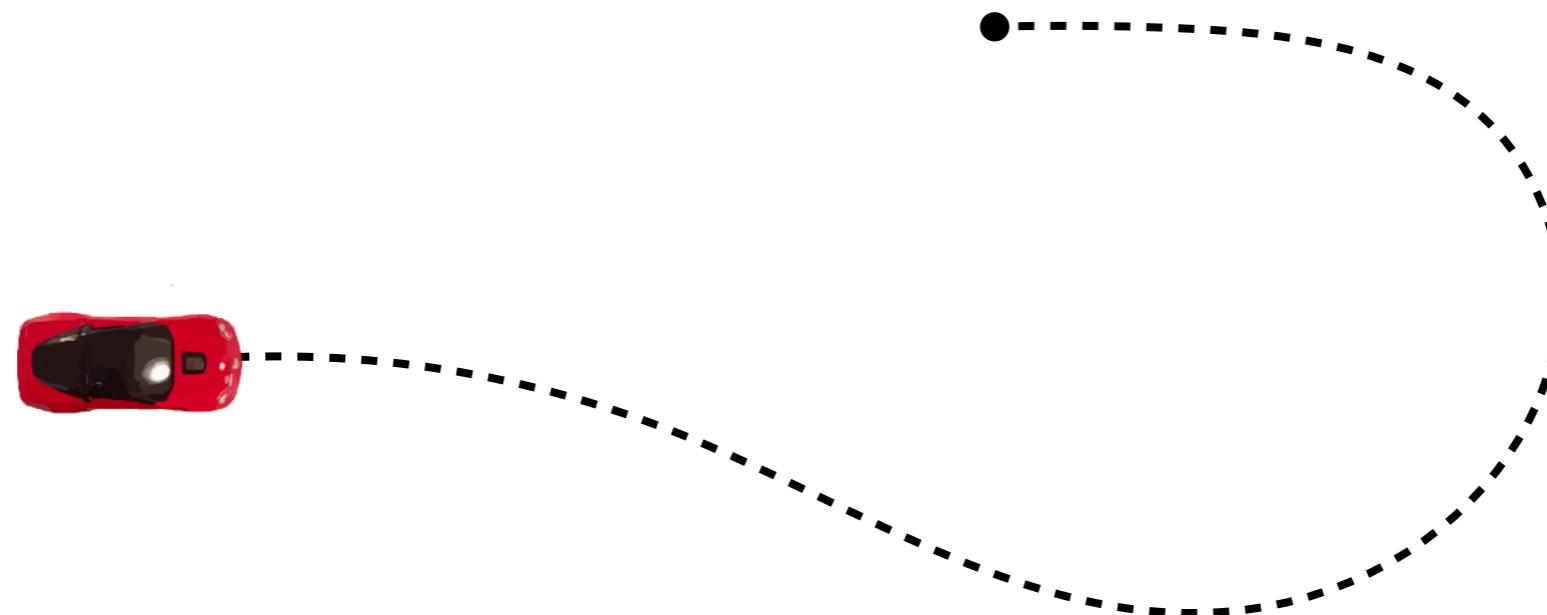
Can express this as wanting to track a **reference** trajectory

$$x(t), y(t), \theta(t)$$

The control framework

Let's say we want to track a reference trajectory

$$x(t), y(t), \theta(t)$$



Objective: Figure out a control trajectory $u(t)$ to achieve this

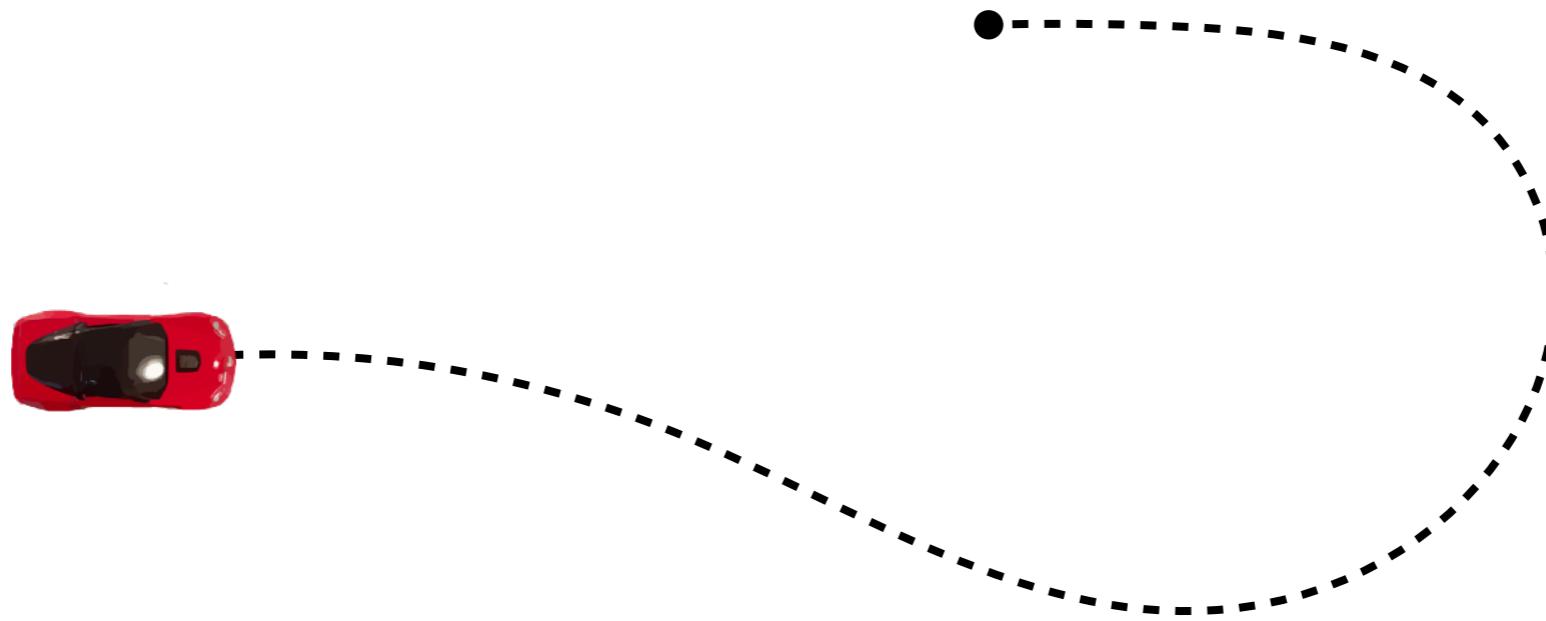
In our case, we will focus on steering angle $\delta(t)$ as control input

Why do we need feedback?

Why do we need feedback?

Let's say we want to track a reference trajectory

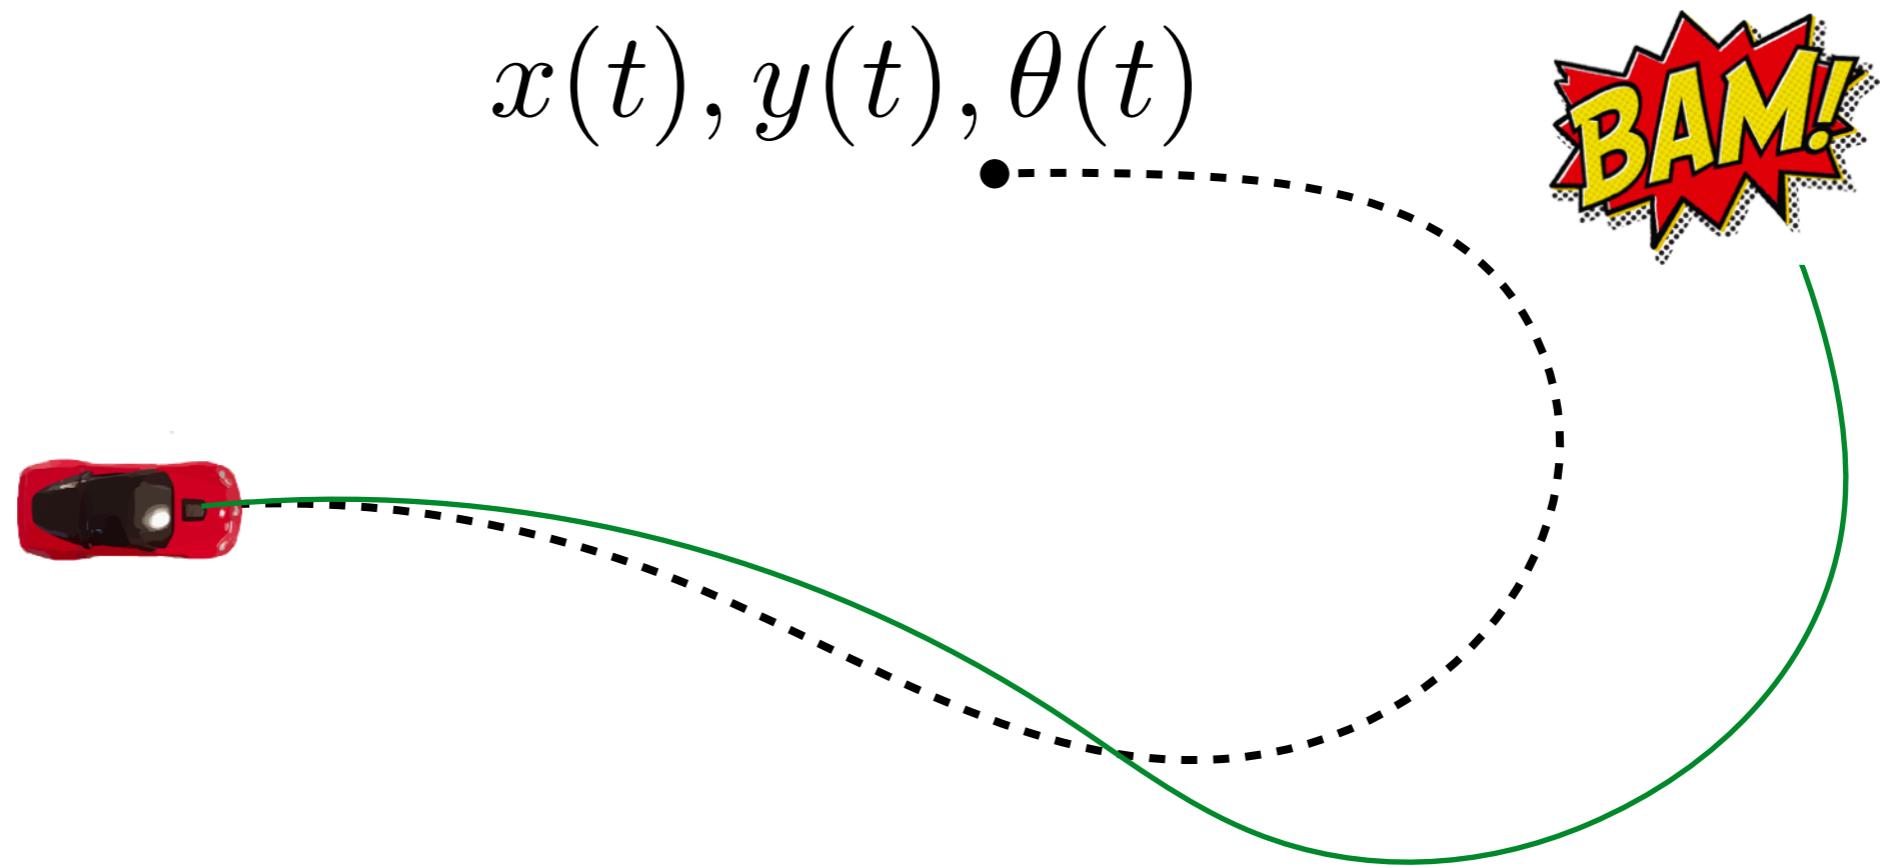
$$x(t), y(t), \theta(t)$$



What if we send out steering angles $\delta(t)$ obtained from kinematic car model?

Why do we need feedback?

Let's say we want to track a reference trajectory



What if we send out steering angles $\delta(t)$ obtained from kinematic car model?

Open loop control leads to accumulating errors!

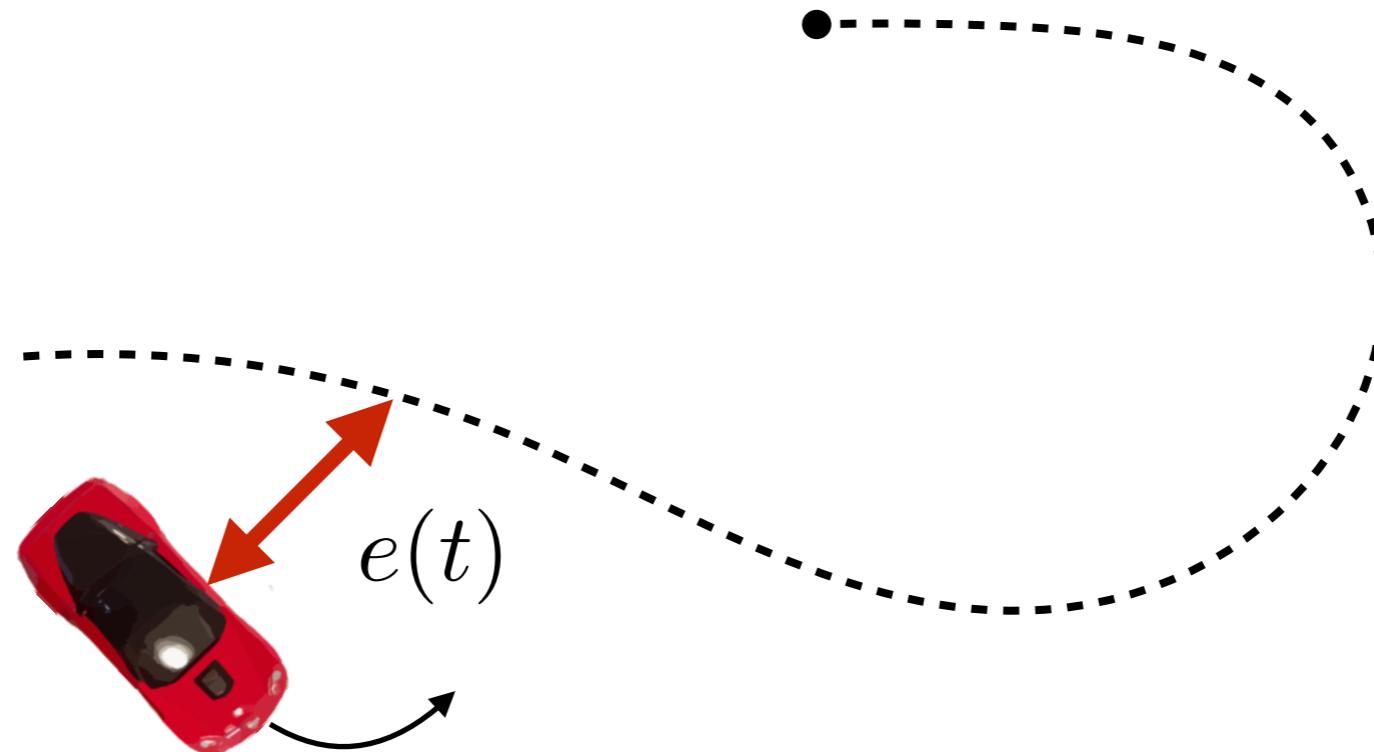
Underlying principle of
feedback control:

Measure error between
reference and state.
Minimize this error.

Measure **error** and minimize it

Error
between
state and
reference

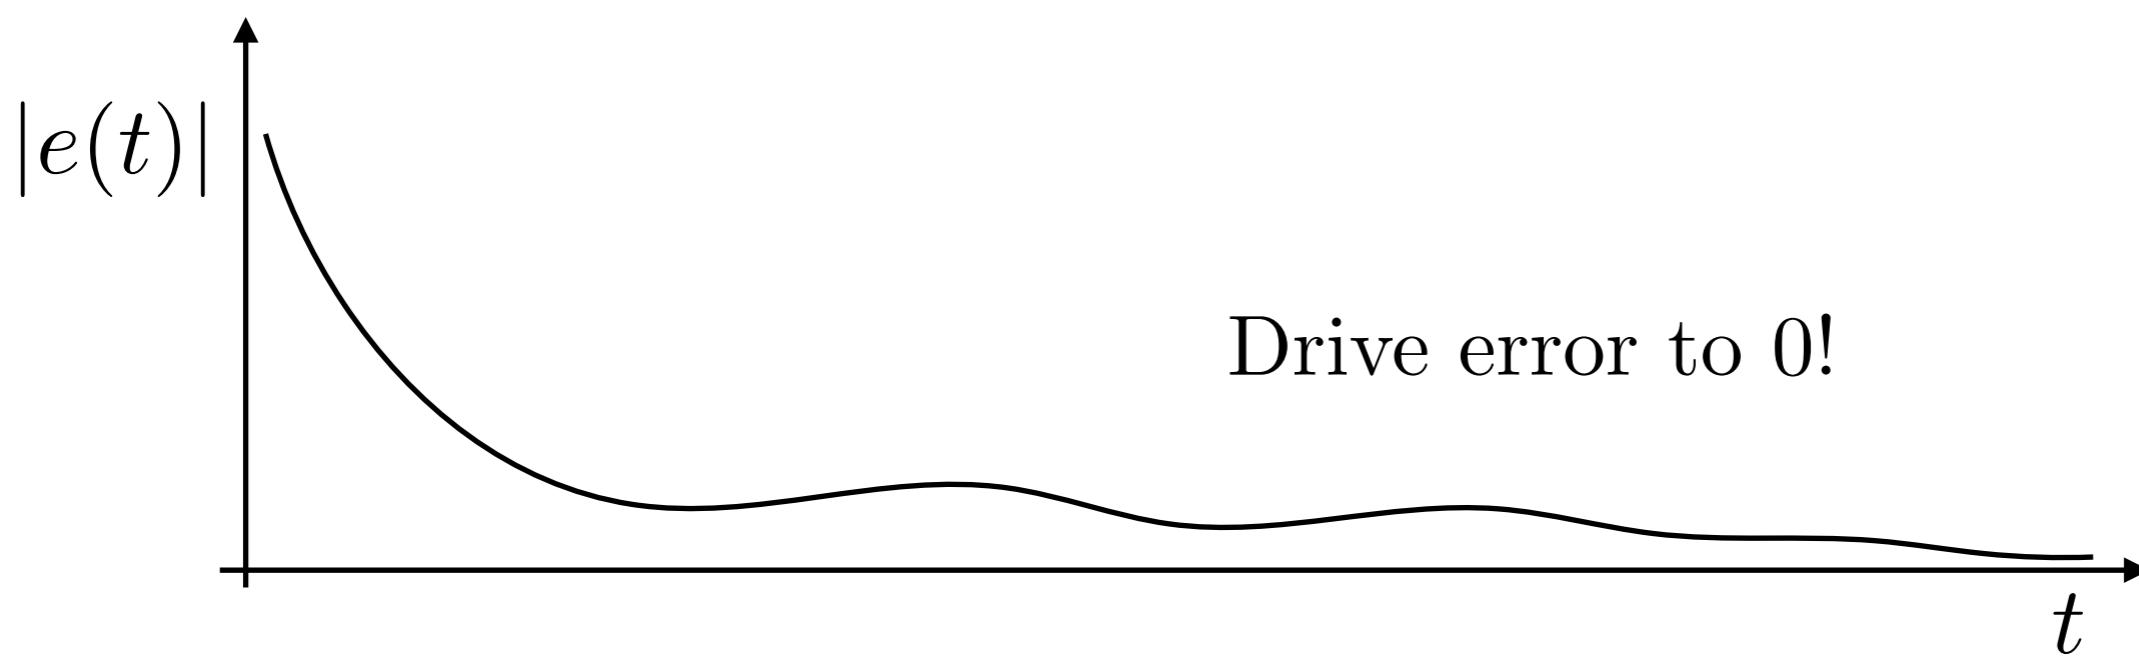
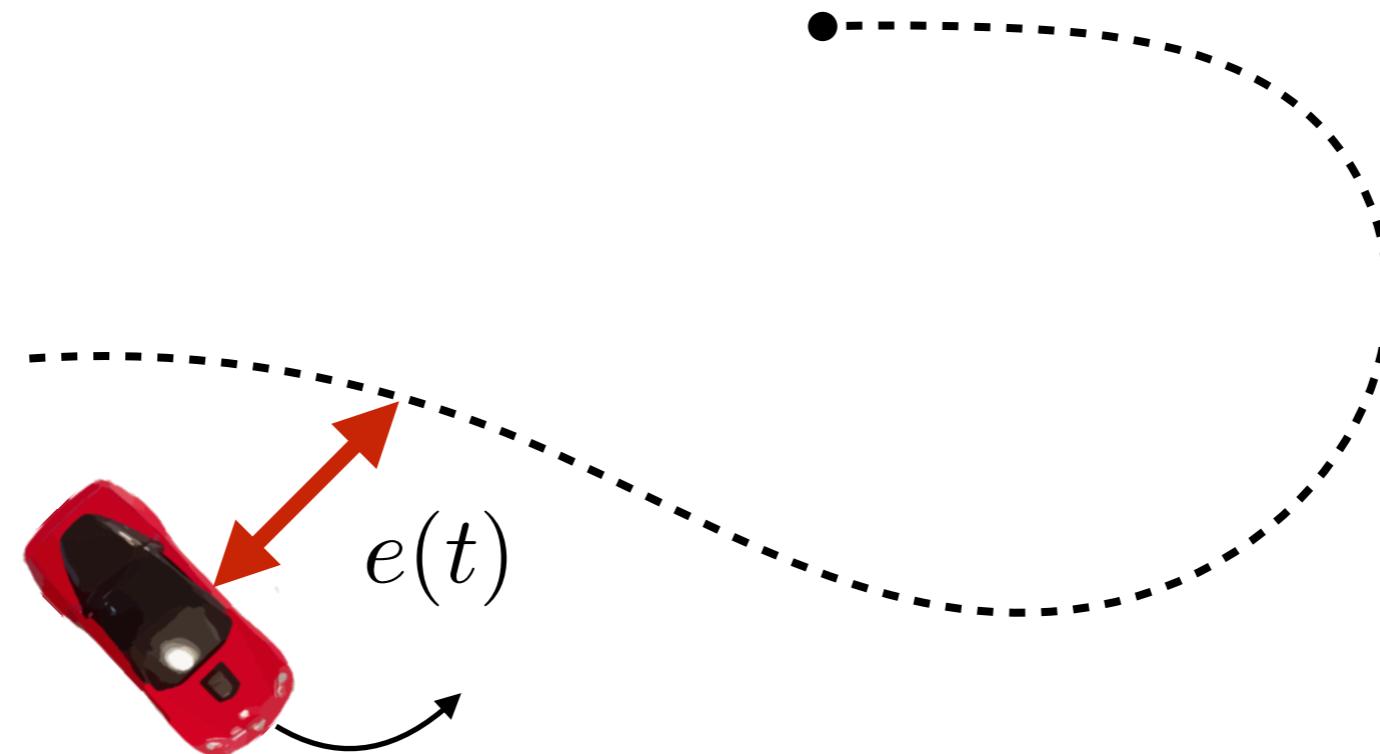
$$x(t), y(t), \theta(t)$$



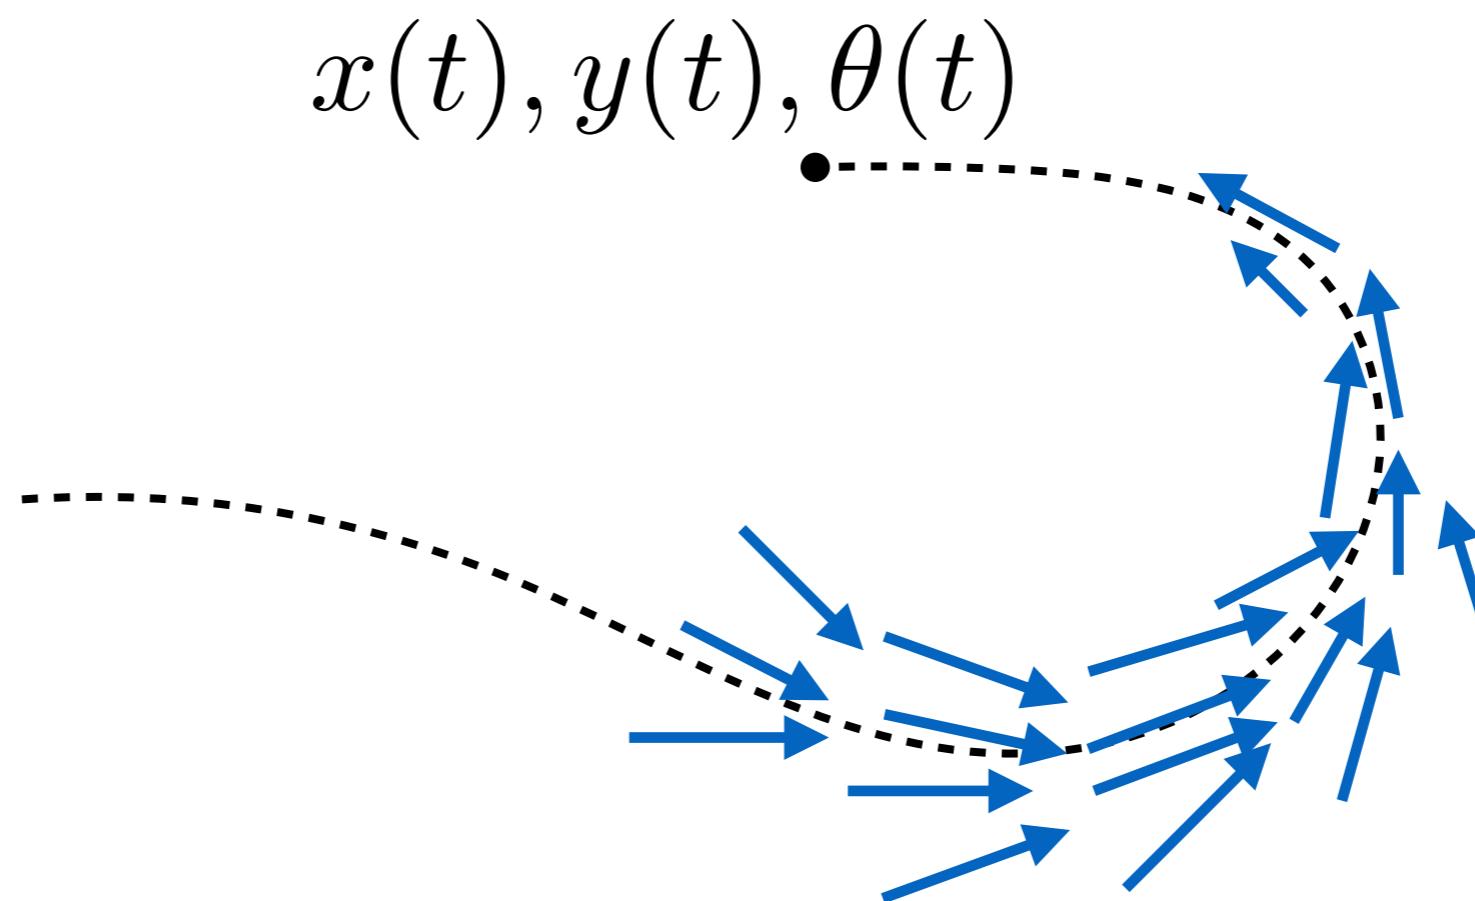
Measure **error** and minimize it

Error
between
state and
reference

$$x(t), y(t), \theta(t)$$

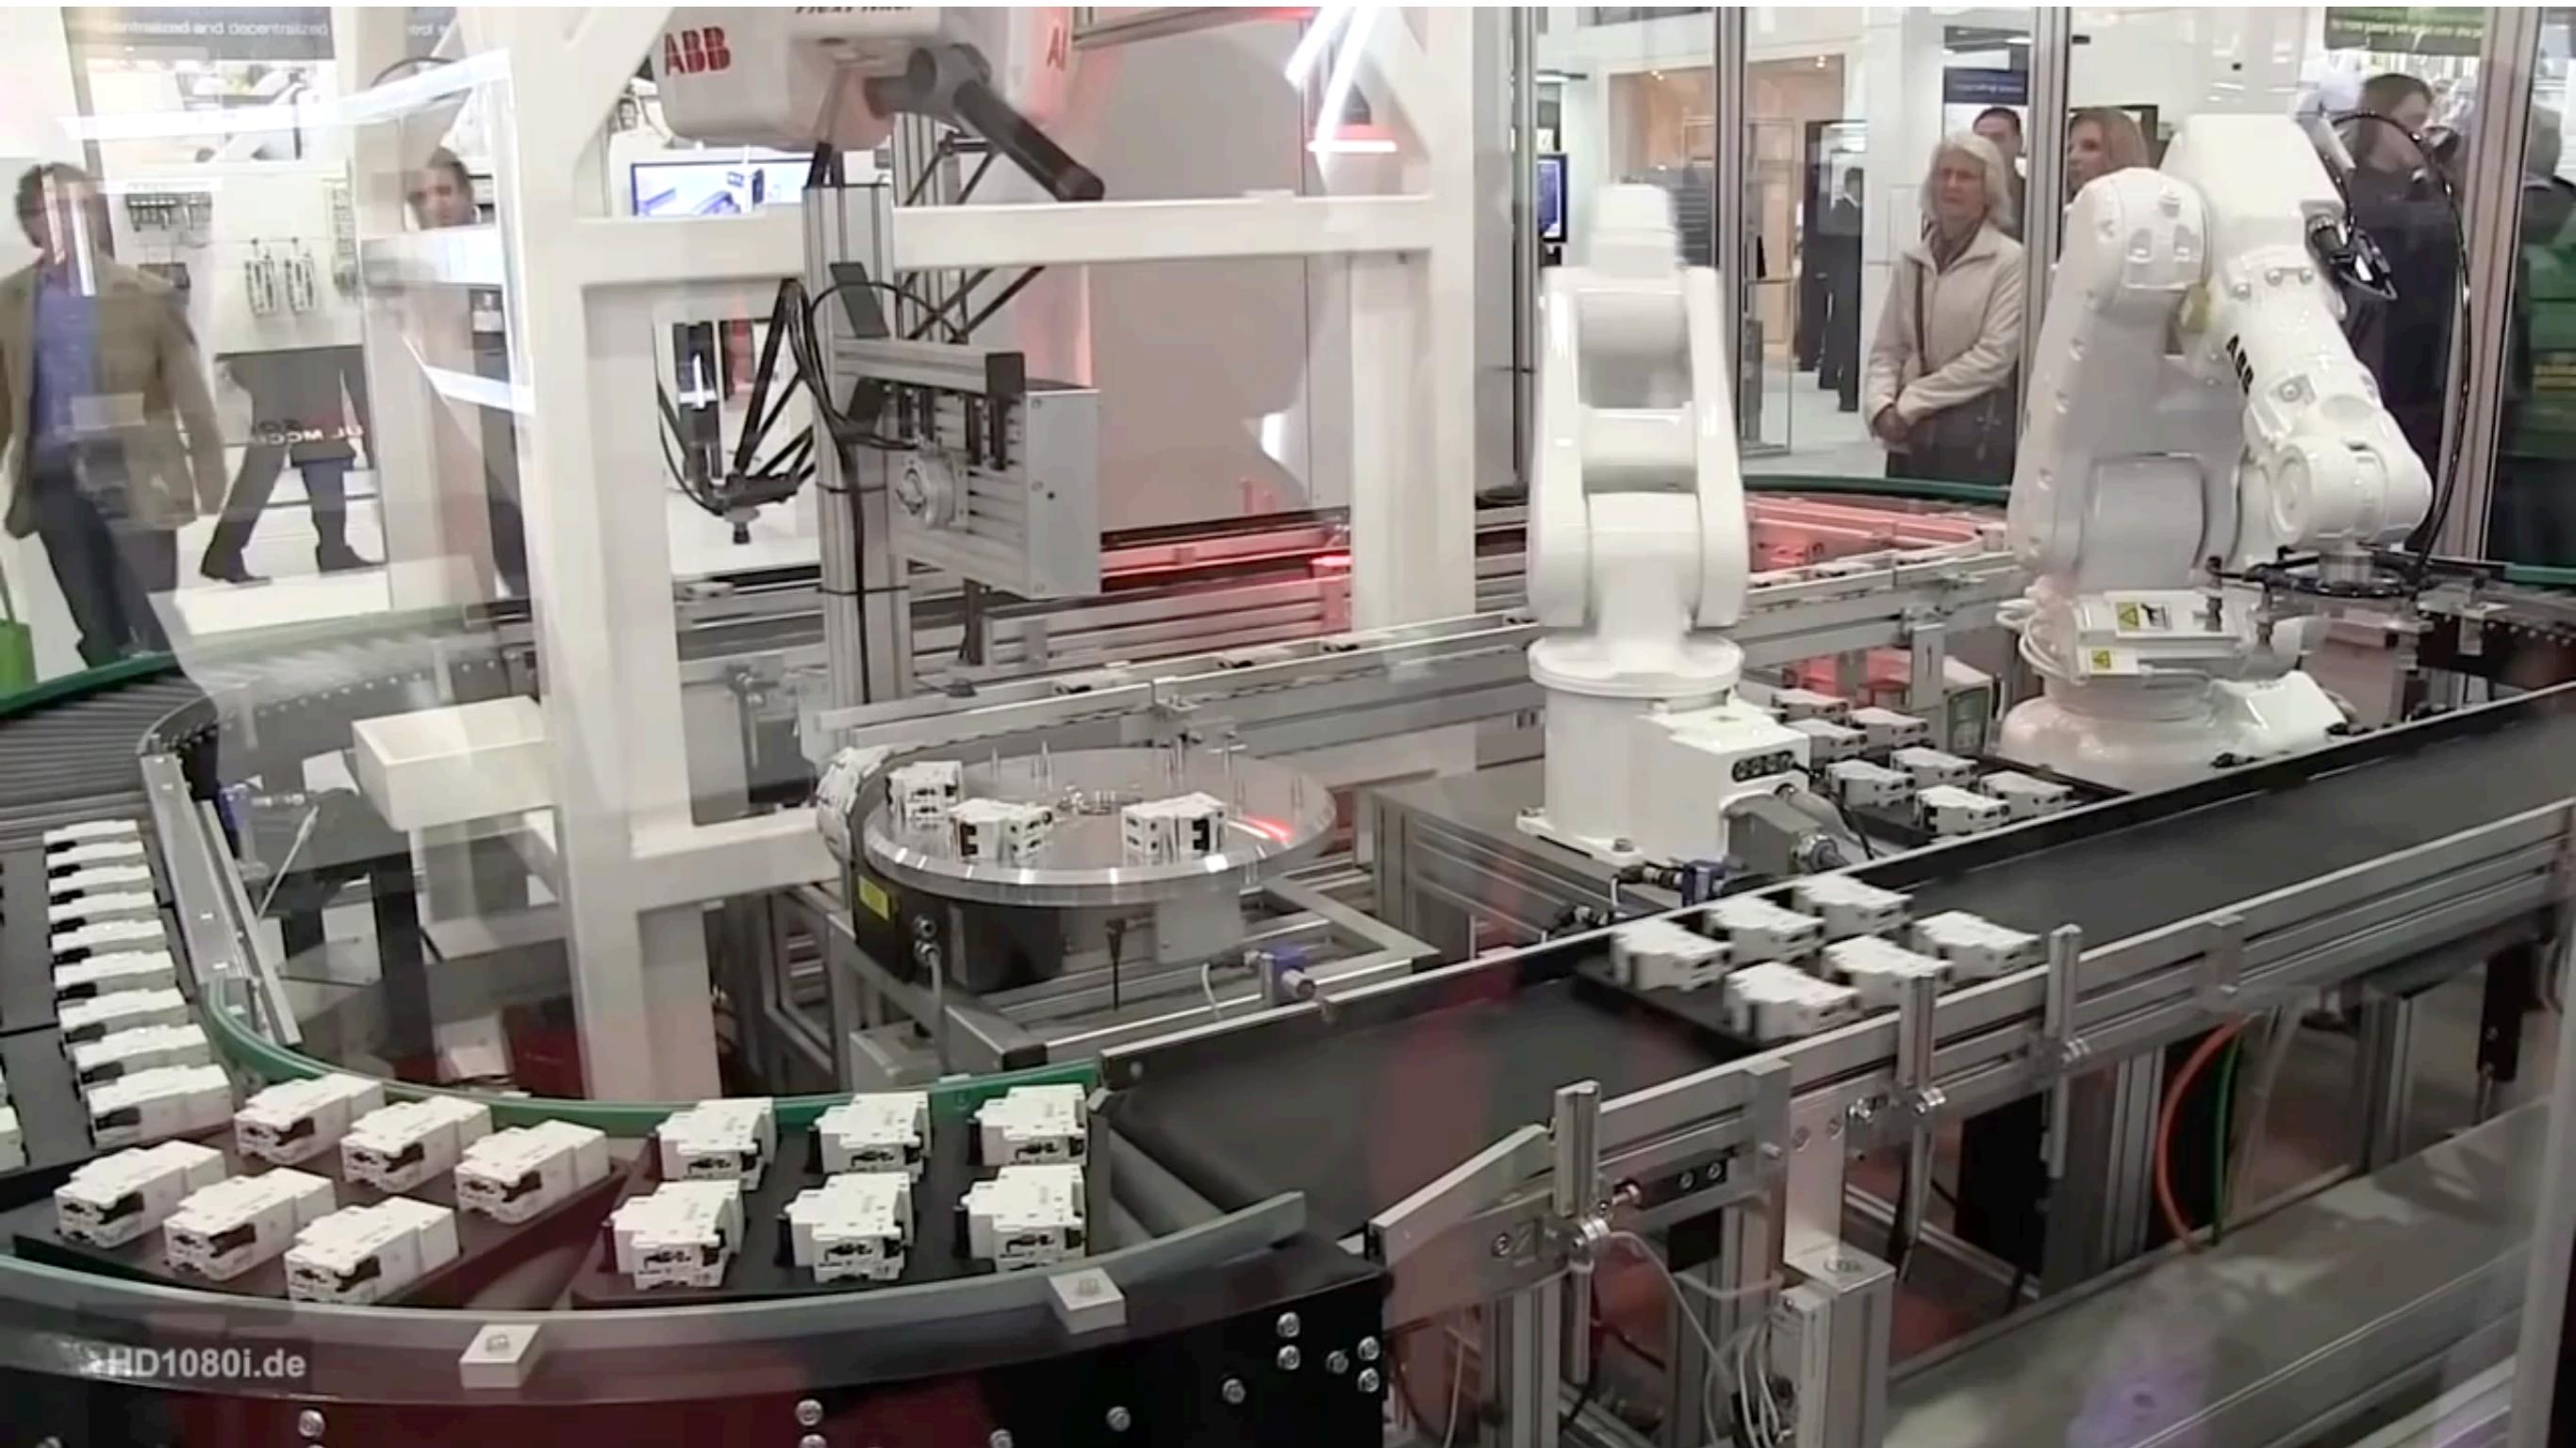


Useful to think of control laws as vector fields



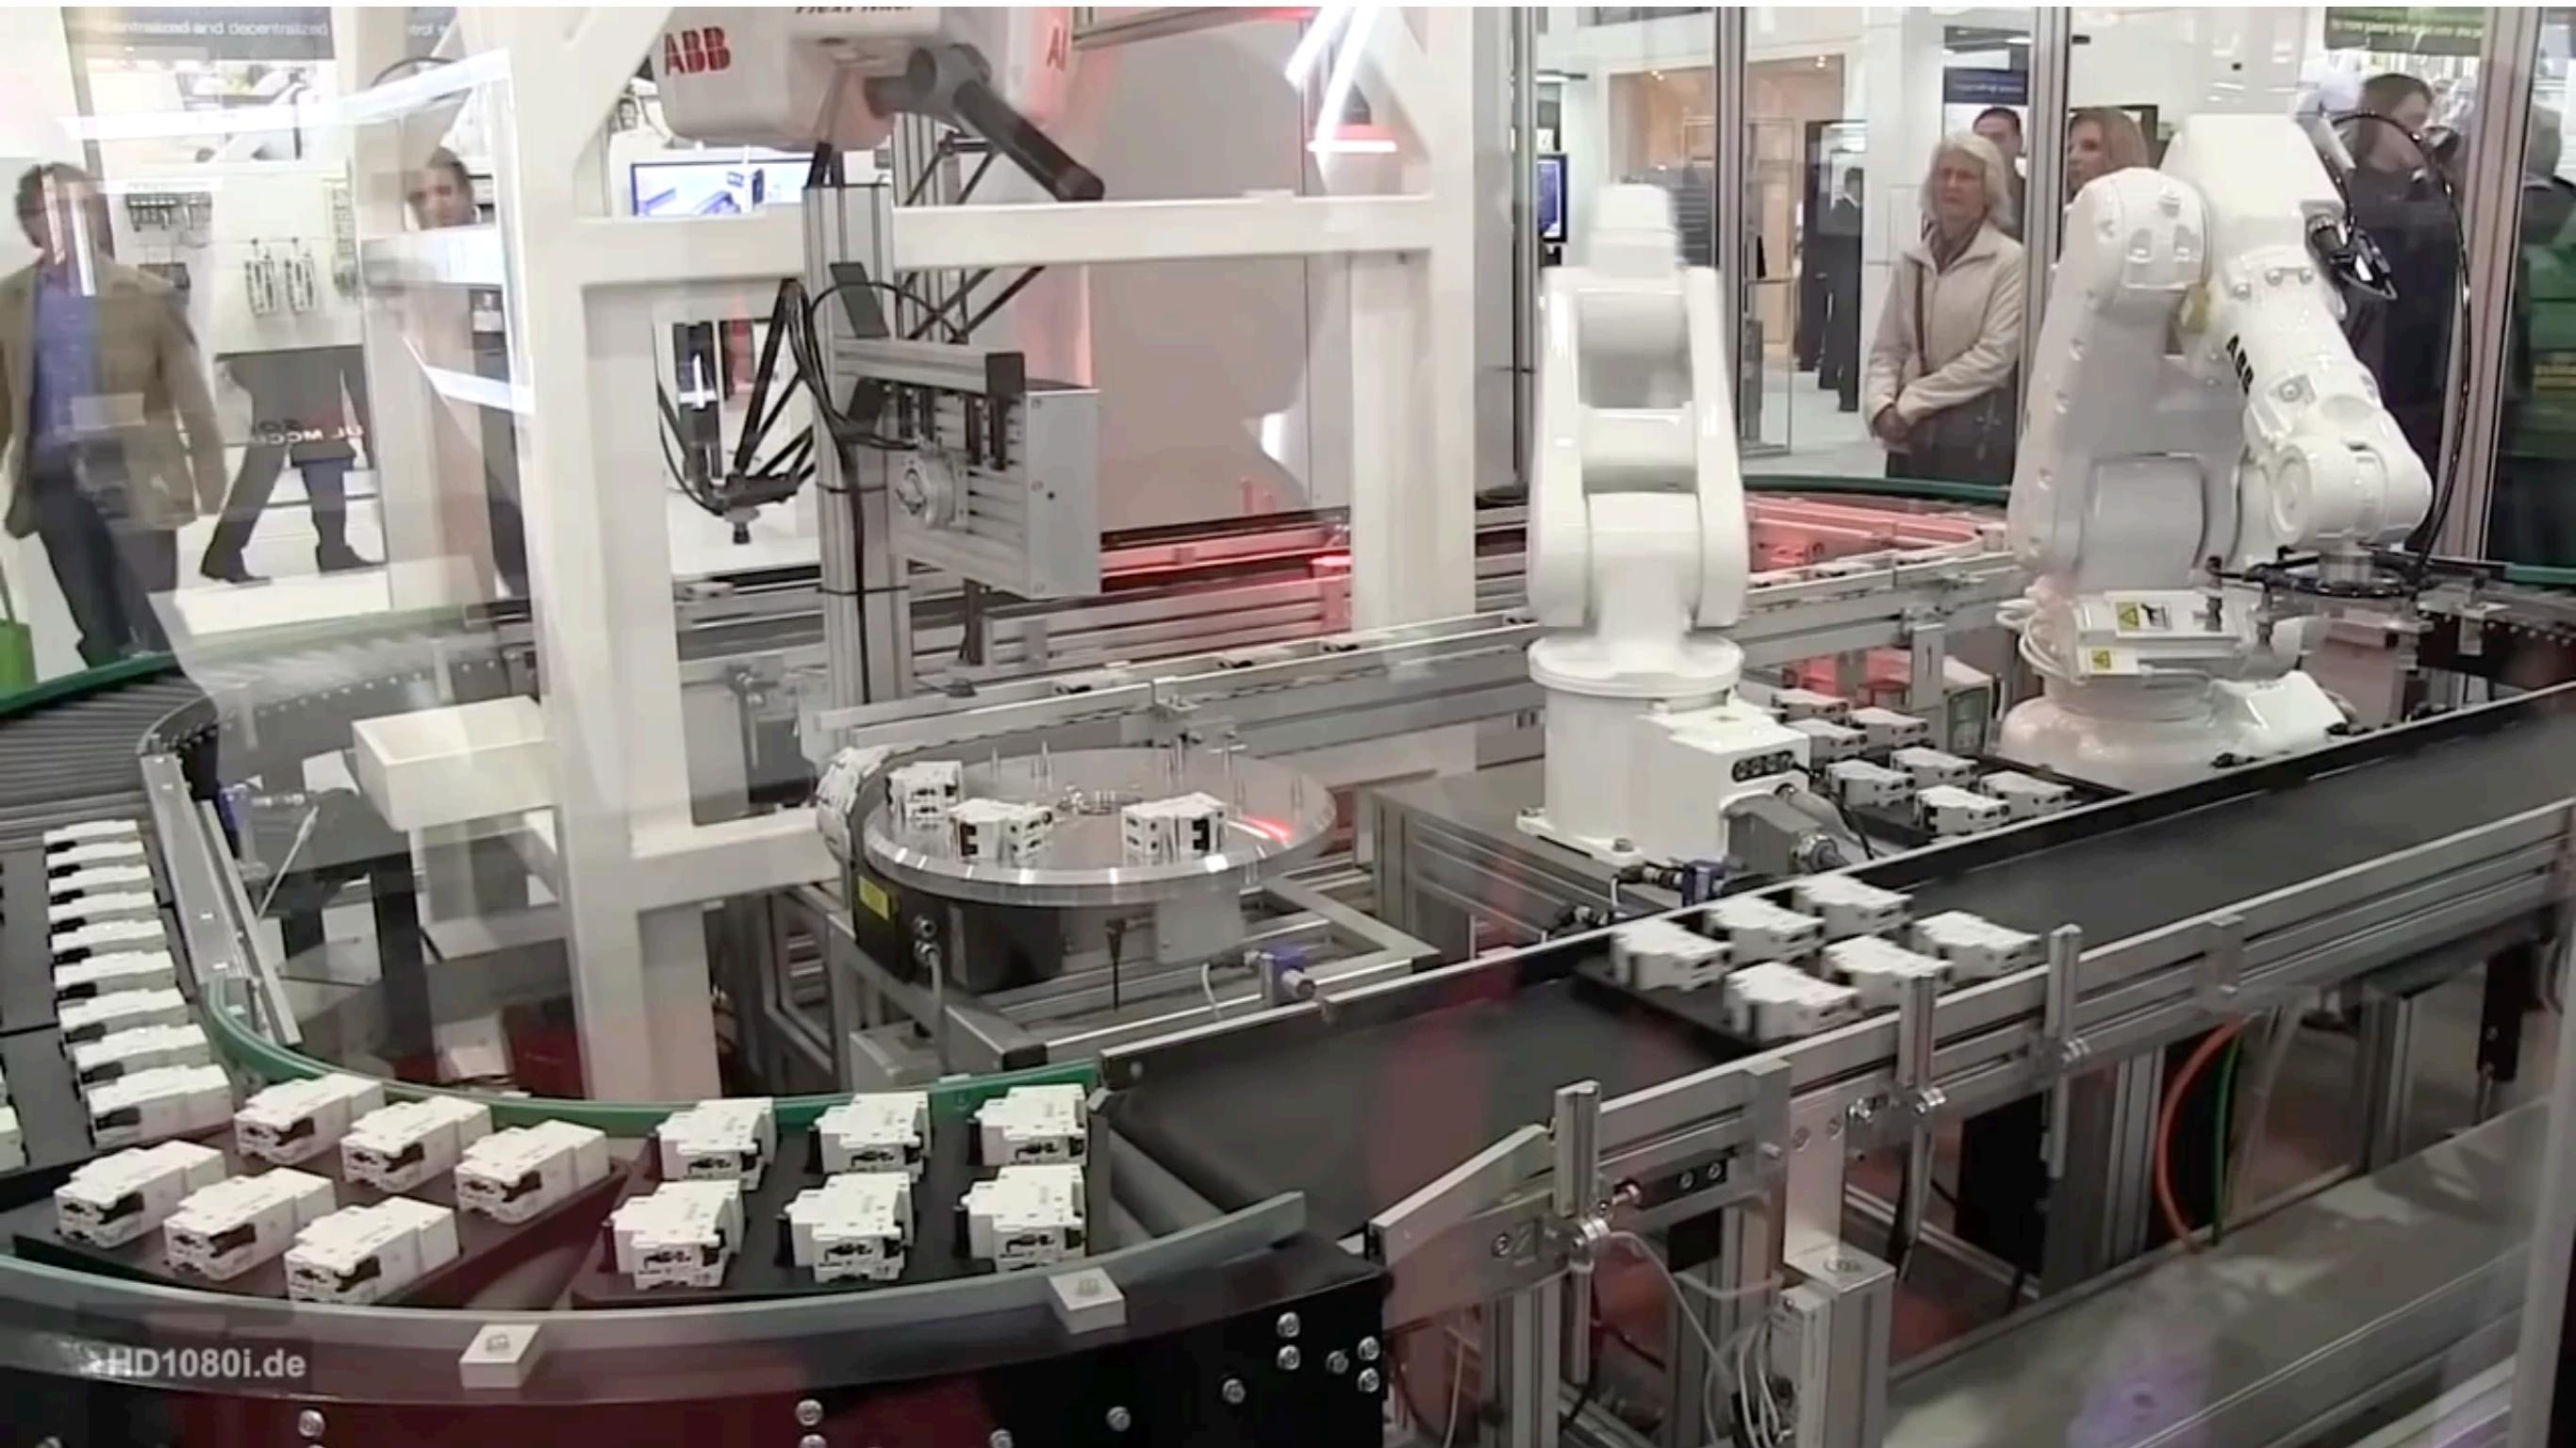
Is this still a research problem?

Industrial robots hard at work



https://www.youtube.com/watch?v=J_8OnDsQVZE&t=315s

Industrial robots hard at work



https://www.youtube.com/watch?v=J_8OnDsQVZE&t=315s

Assumptions made by such controllers

Assumptions made by such controllers

1. Fully actuated: There exists an inverse mapping from reference to control actions

$$\sigma(t) \rightarrow u(t)$$

Assumptions made by such controllers

1. Fully actuated: There exists an inverse mapping from reference to control actions

$$\sigma(t) \rightarrow u(t)$$

2. Almost no execution error or state estimation error

Assumptions made by such controllers

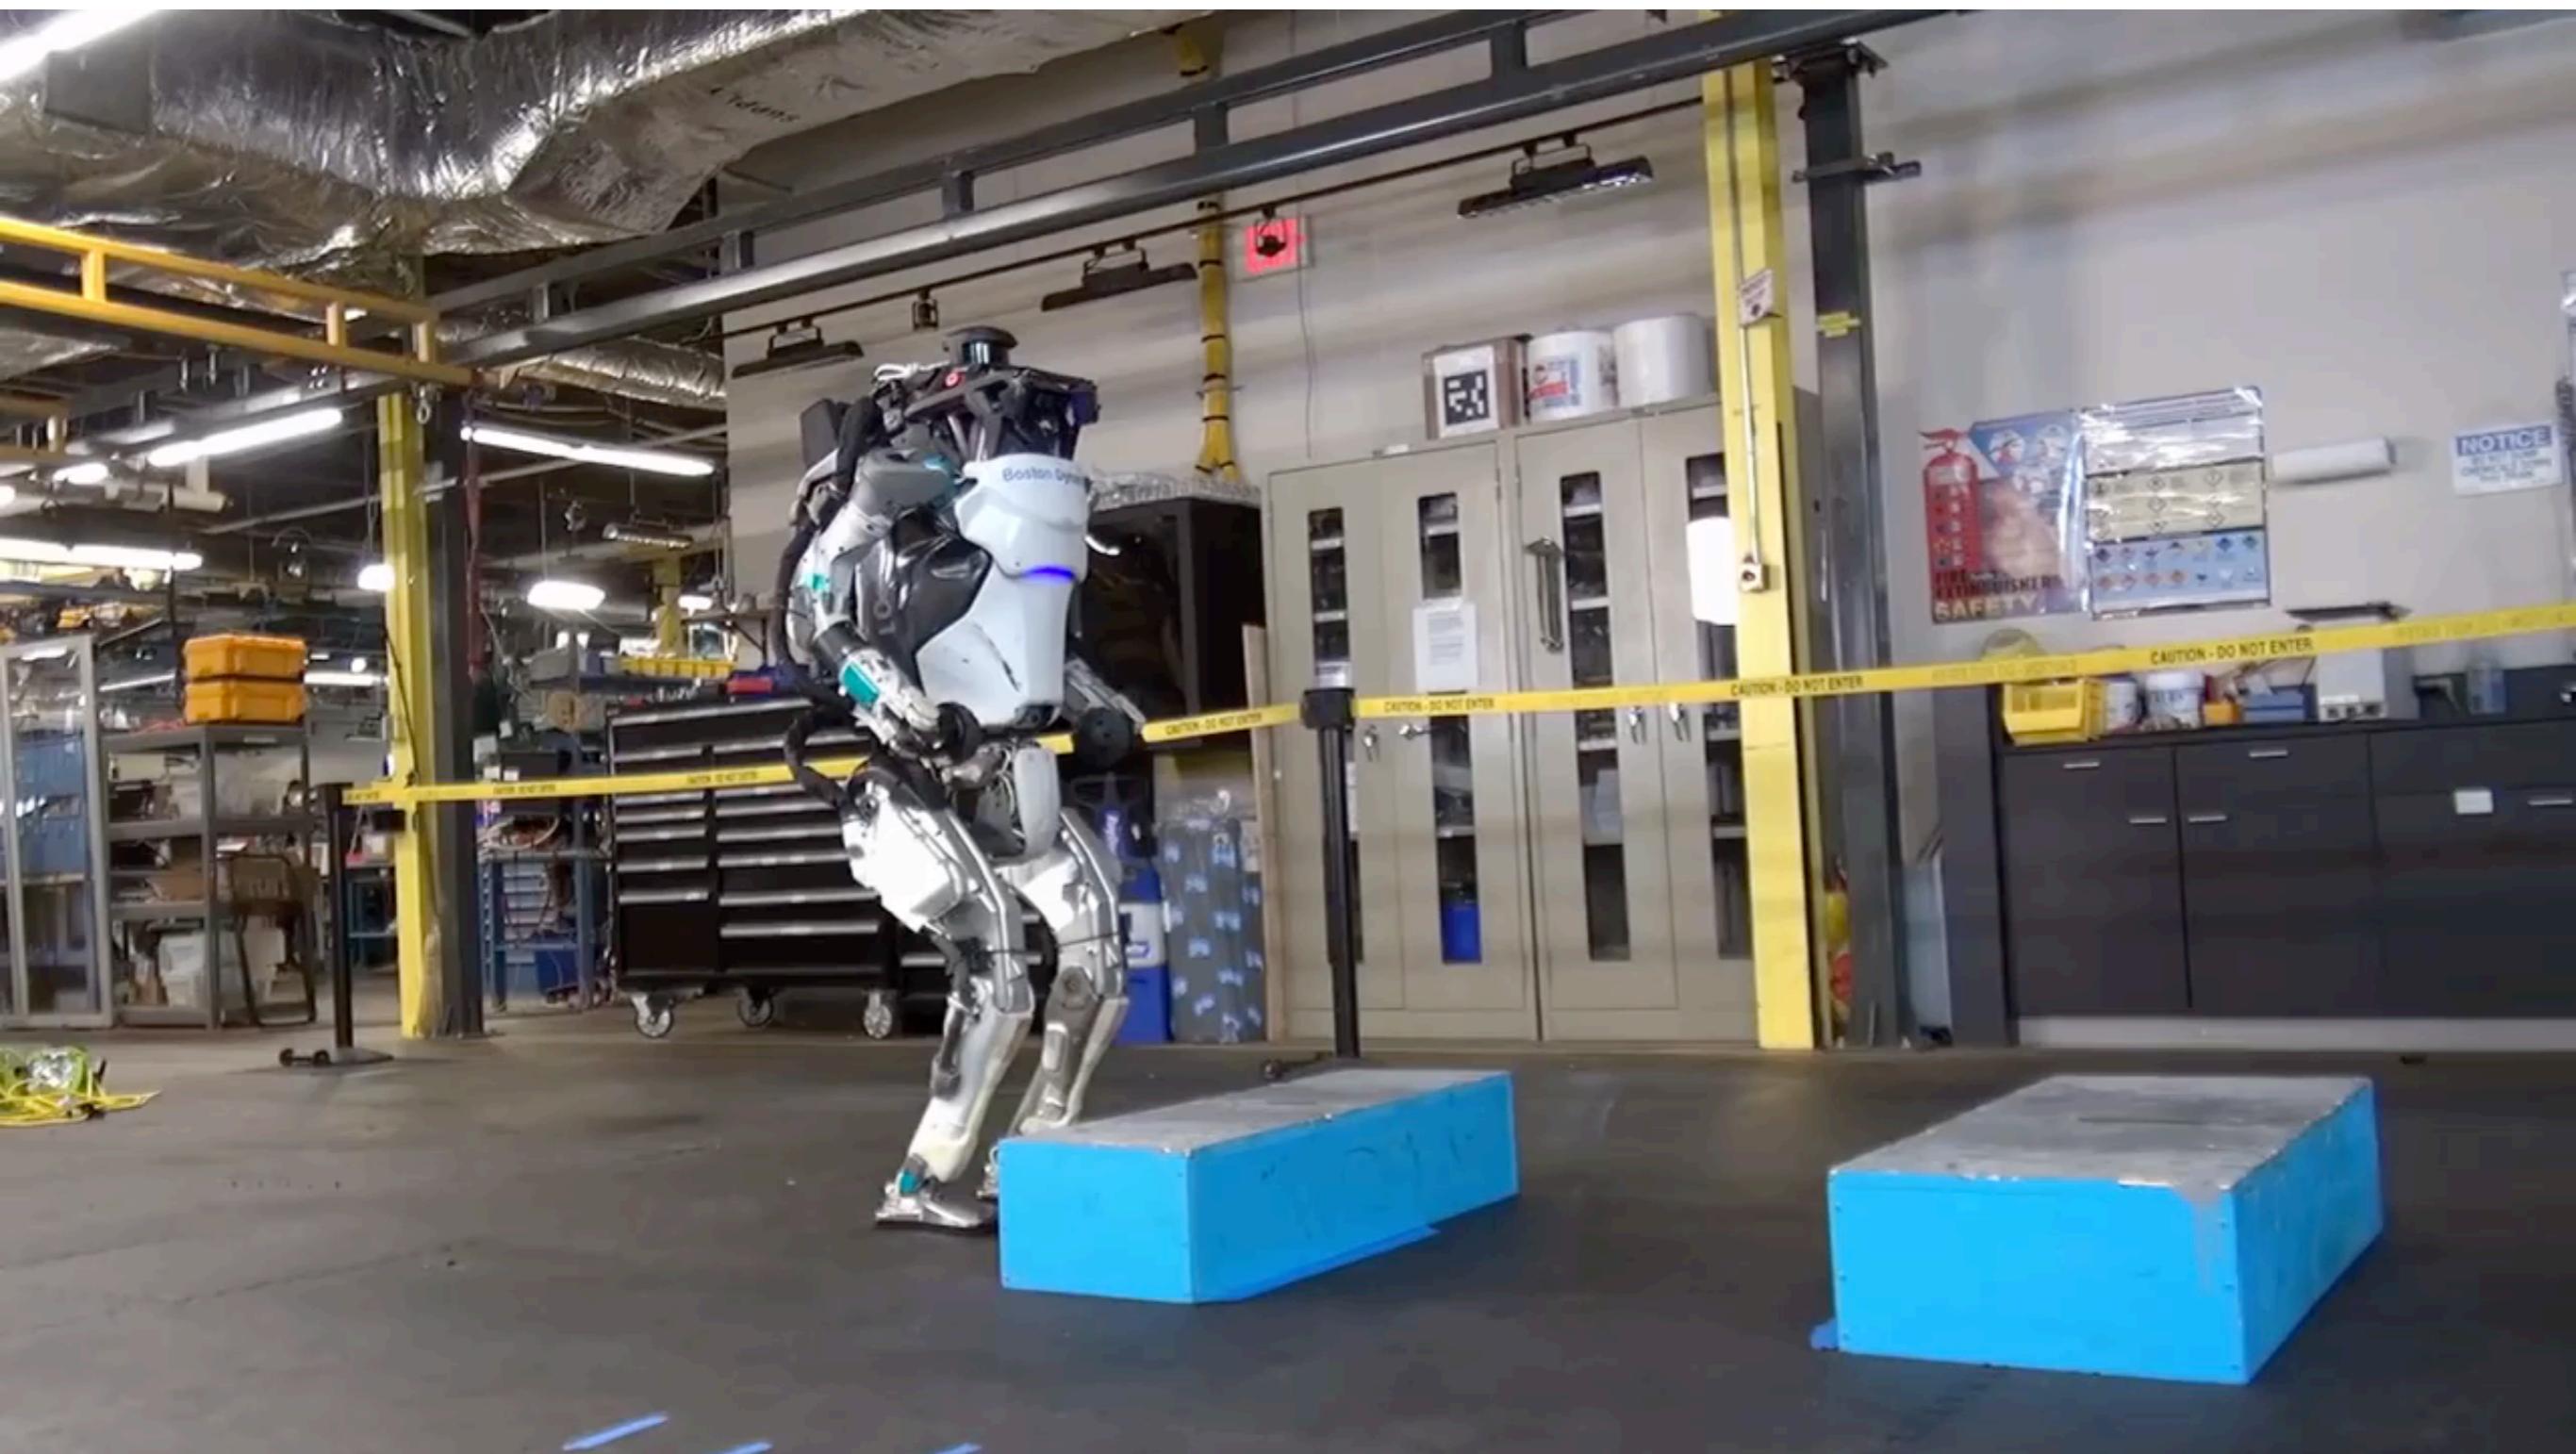
1. Fully actuated: There exists an inverse mapping from reference to control actions

$$\sigma(t) \rightarrow u(t)$$

2. Almost no execution error or state estimation error
3. Enough control authority to clamp down errors / overcome disturbances

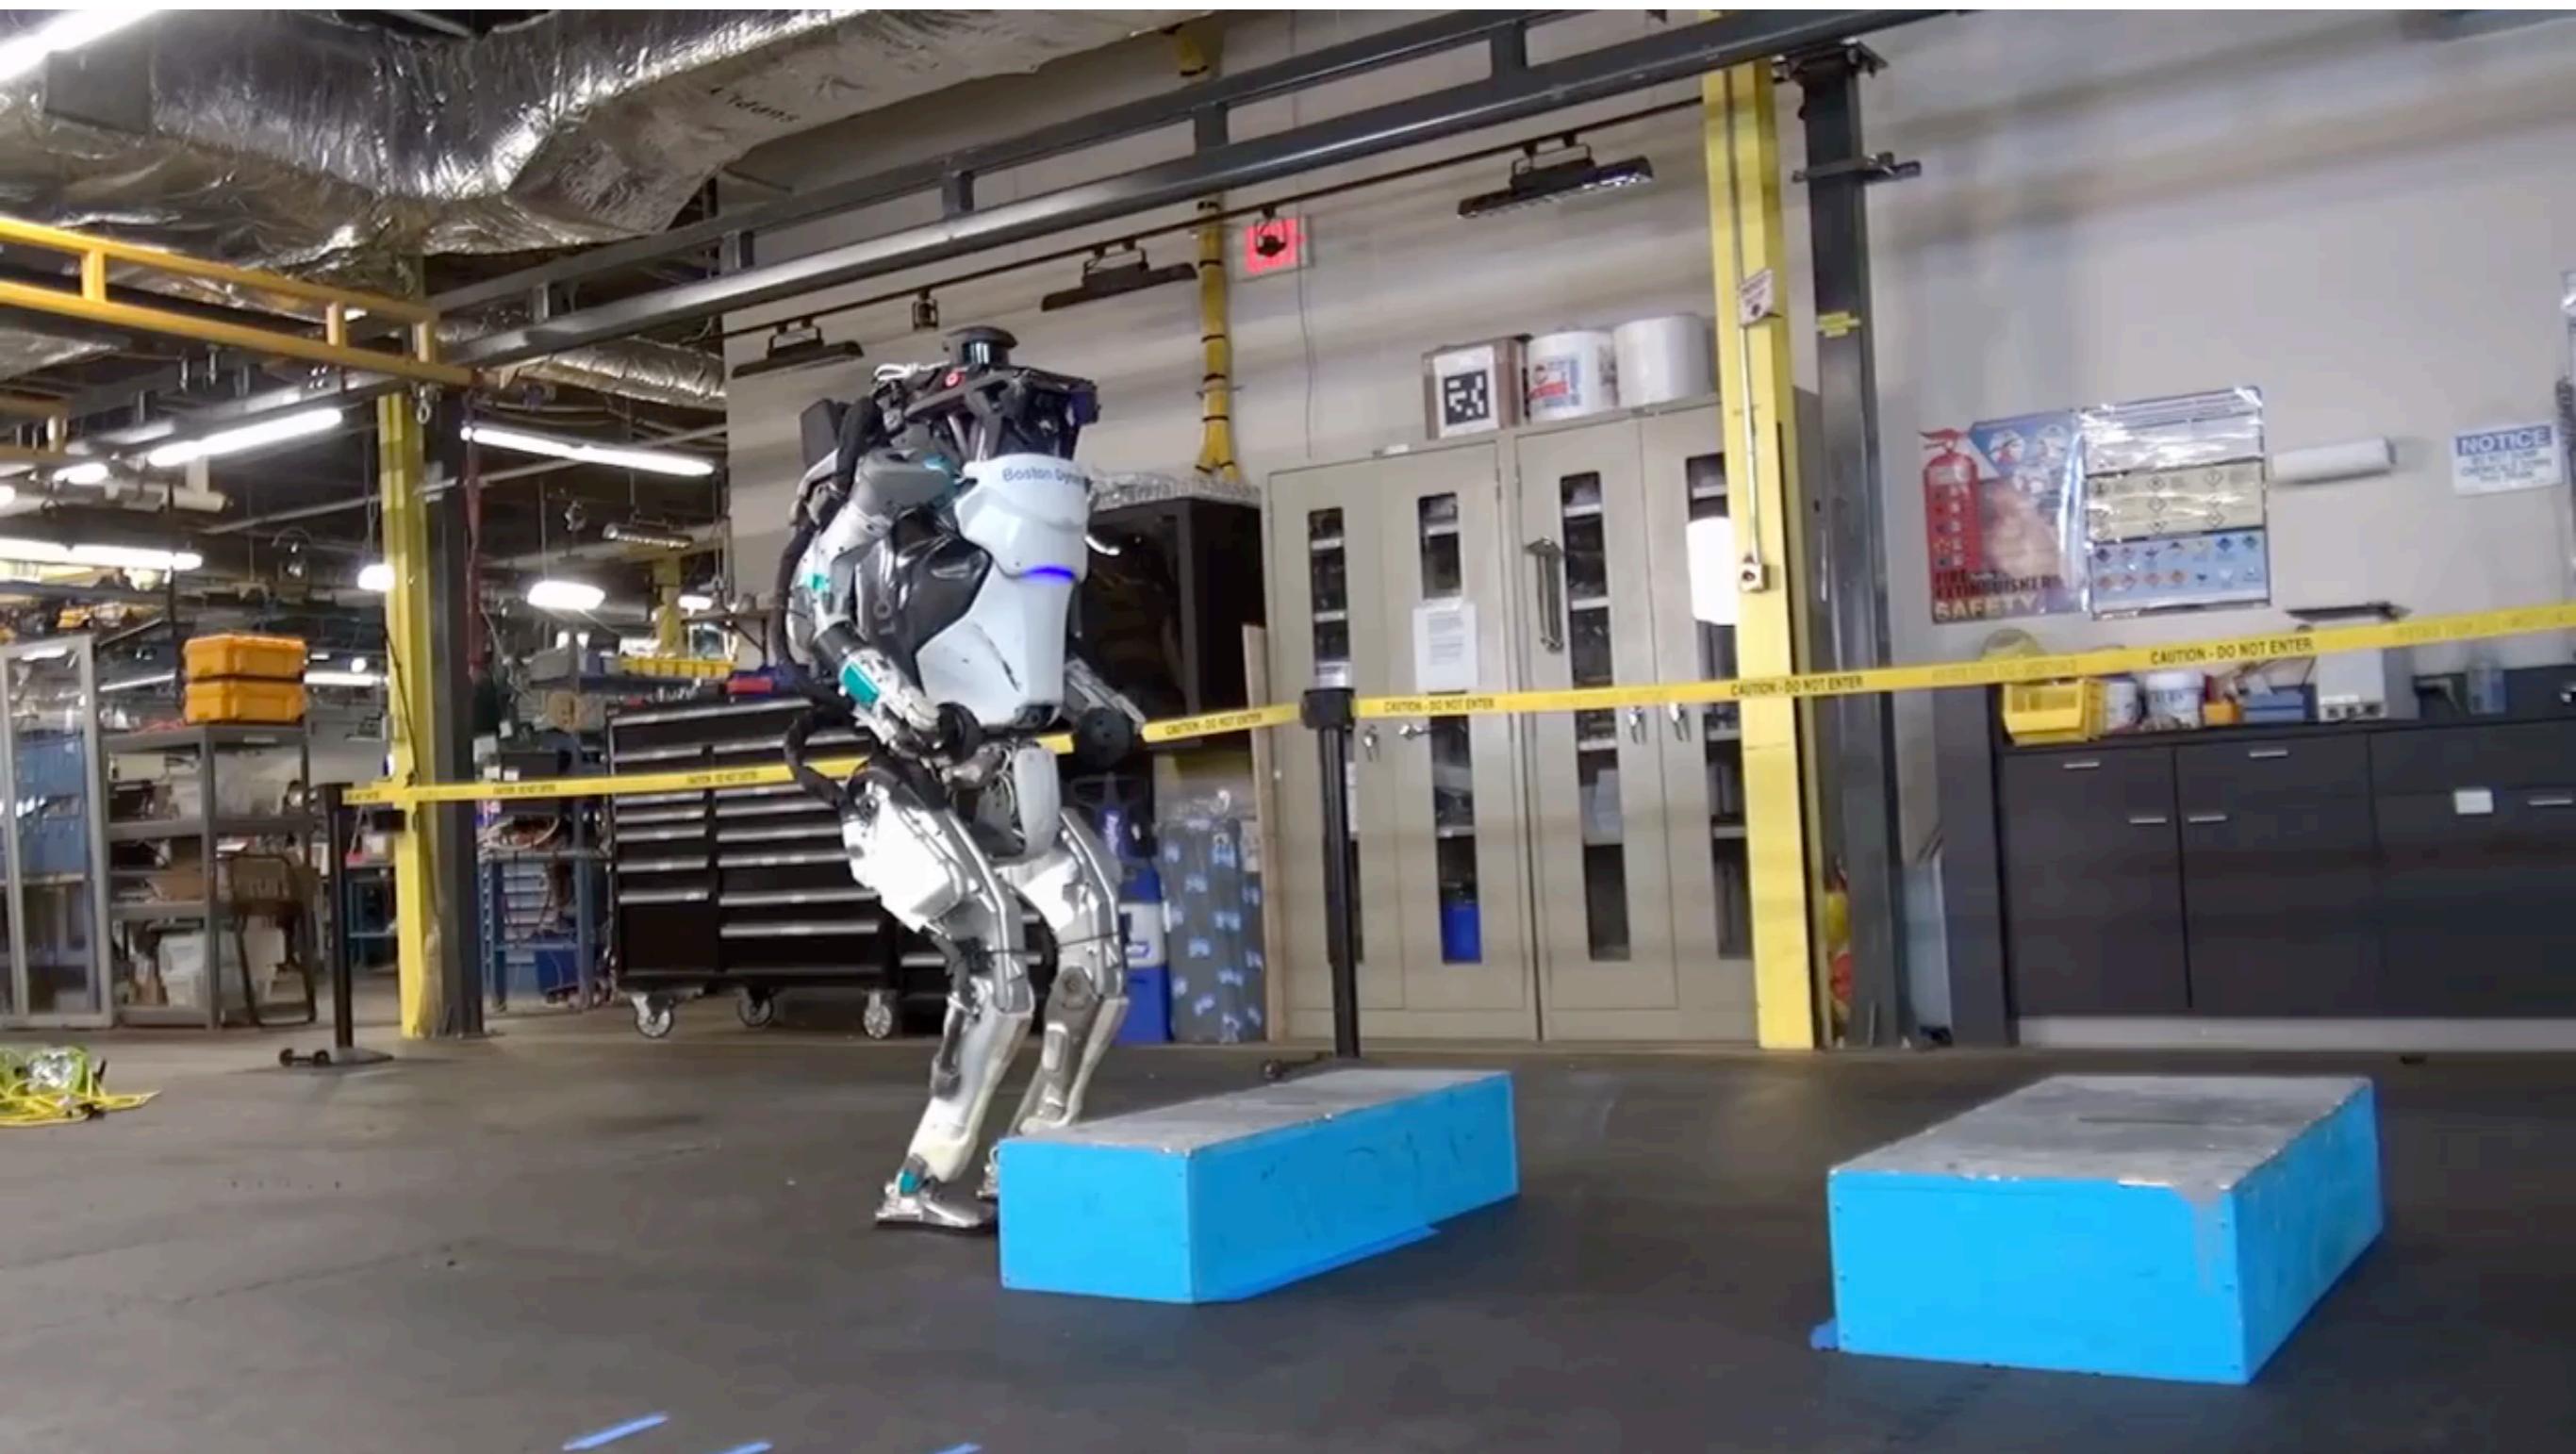
Instead, what are harder
problems we want to think about?

The Atlas robot hard at ... play?



<https://www.youtube.com/watch?v=fRj34o4hN4I>

The Atlas robot hard at ... play?



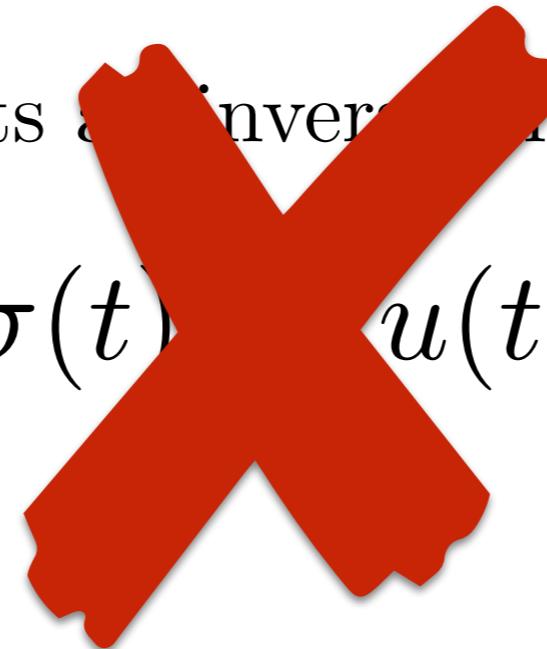
<https://www.youtube.com/watch?v=fRj34o4hN4I>

Why is this a hard problem?

Challenge 1: Underactuated systems

Fully actuated: There exists a universal mapping from reference to control actions

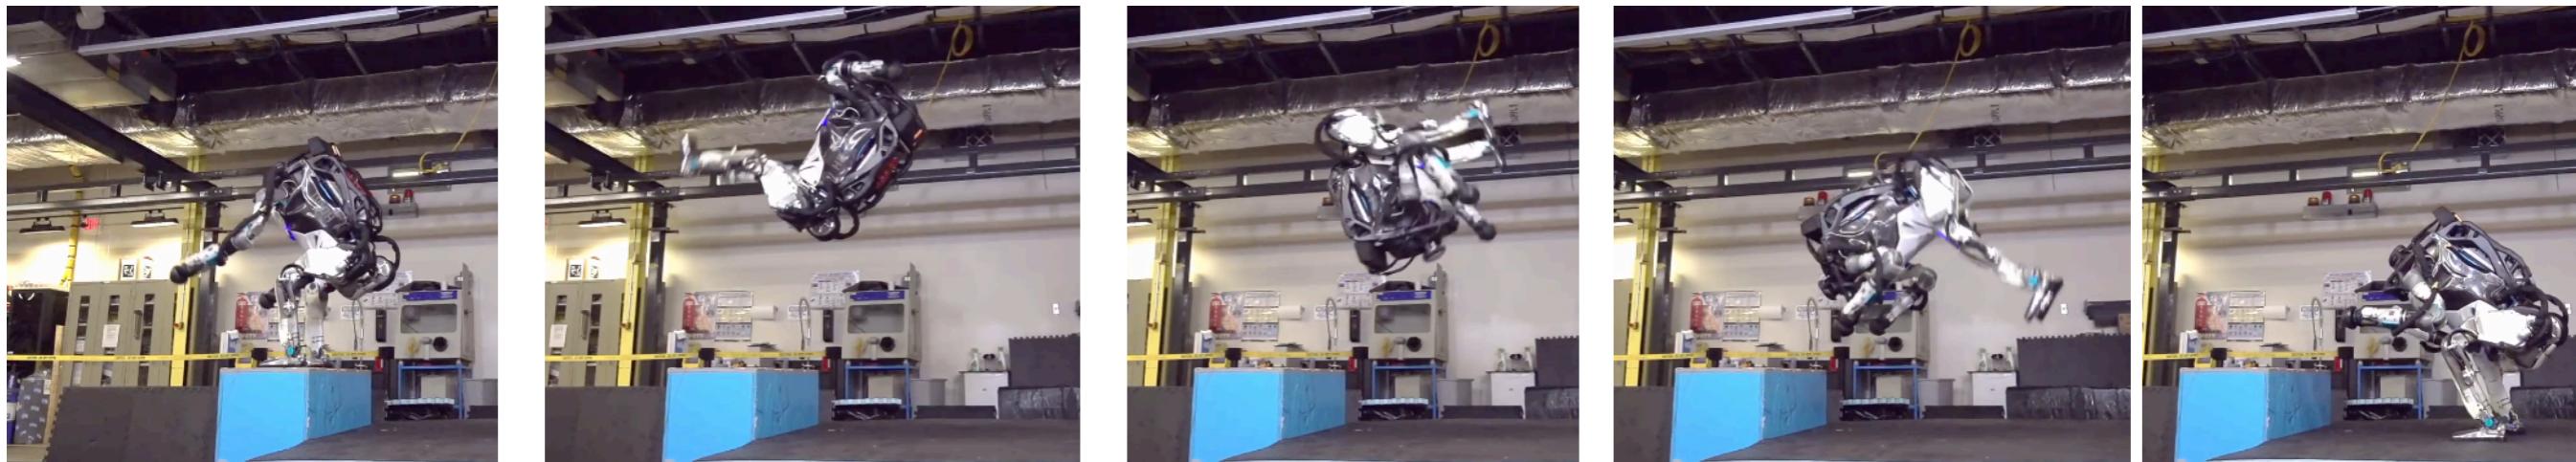
$$\sigma(t) \xrightarrow{\quad} u(t)$$



We don't have full authority to move the system along arbitrary trajectories

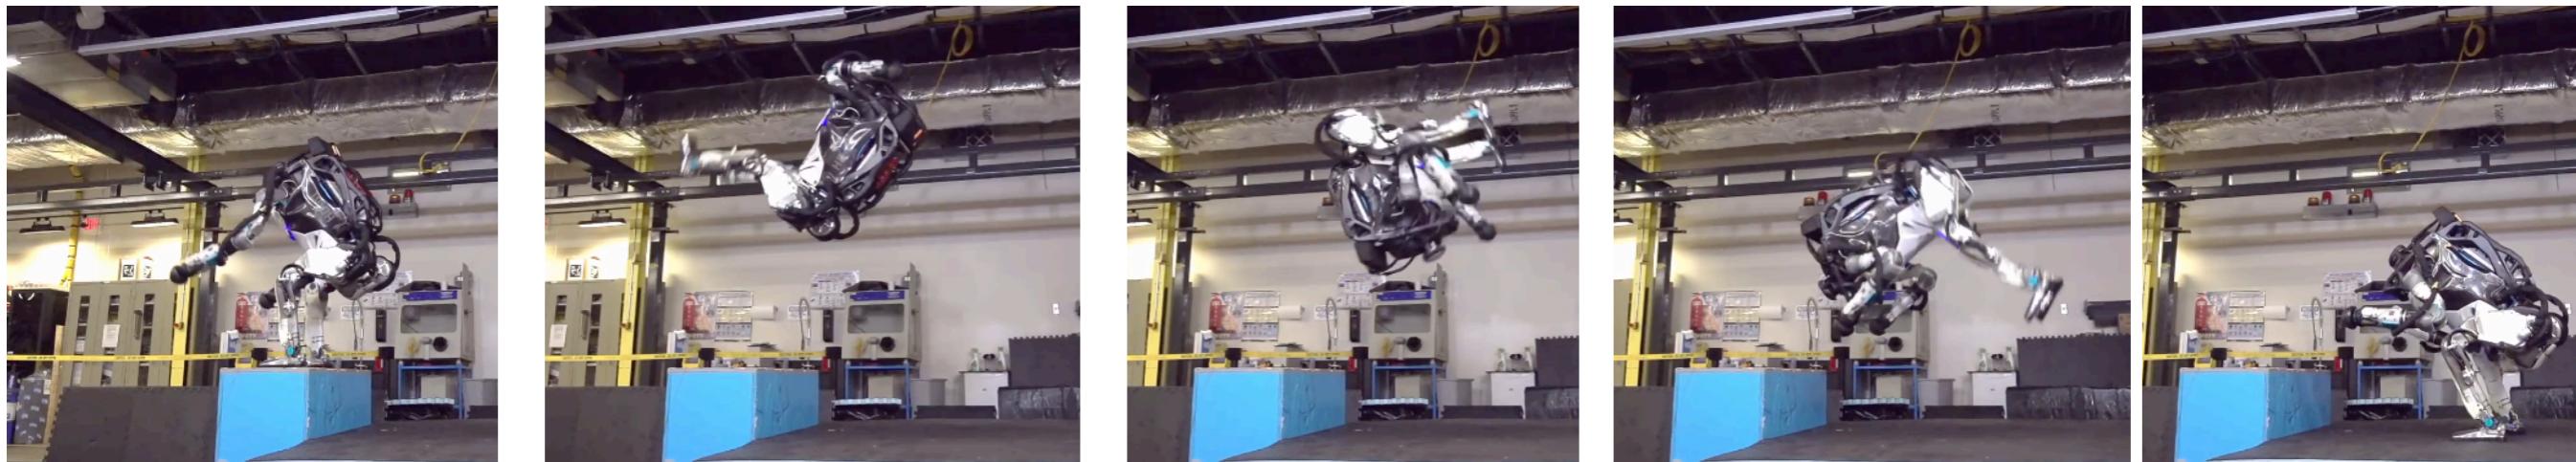
Challenge 1: Underactuated systems

What affects the error between robot state and reference?



Challenge 1: Underactuated systems

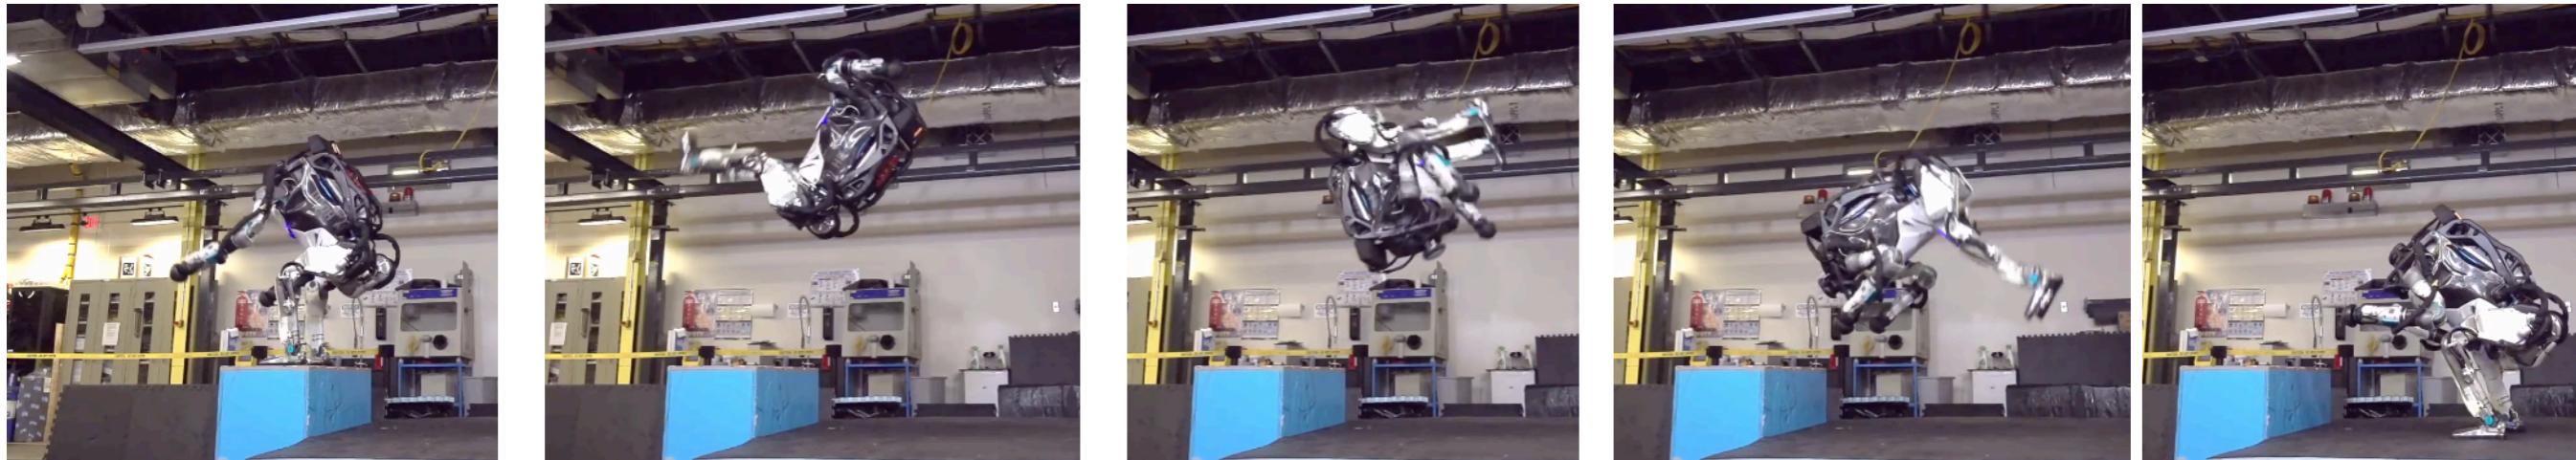
What affects the error between robot state and reference?



Some
initial
motor
thrust ...

Challenge 1: Underactuated systems

What affects the error between robot state and reference?

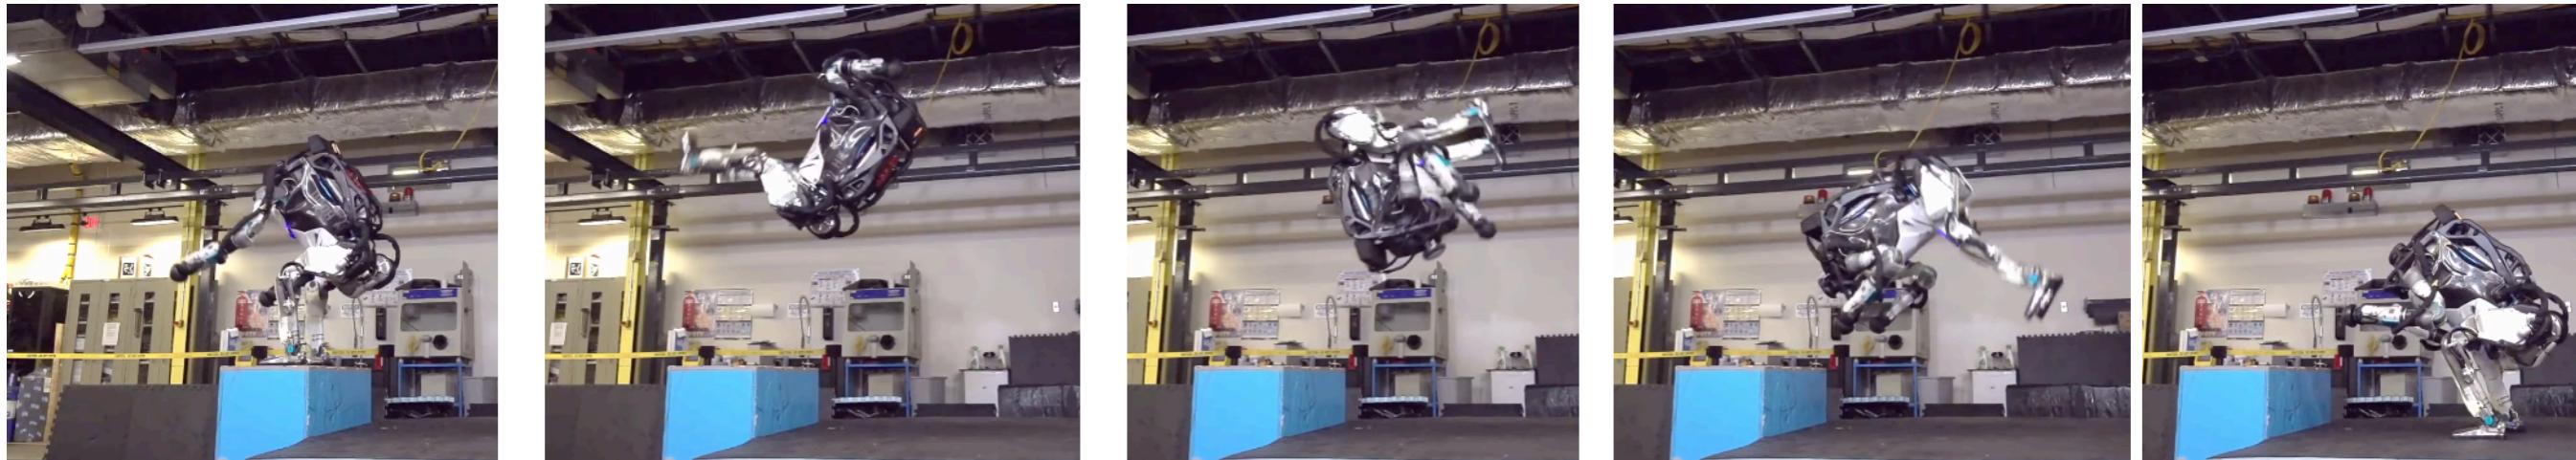


Some
initial
motor
thrust ...

Whole lot of gravity!

Challenge 1: Underactuated systems

What affects the error between robot state and reference?



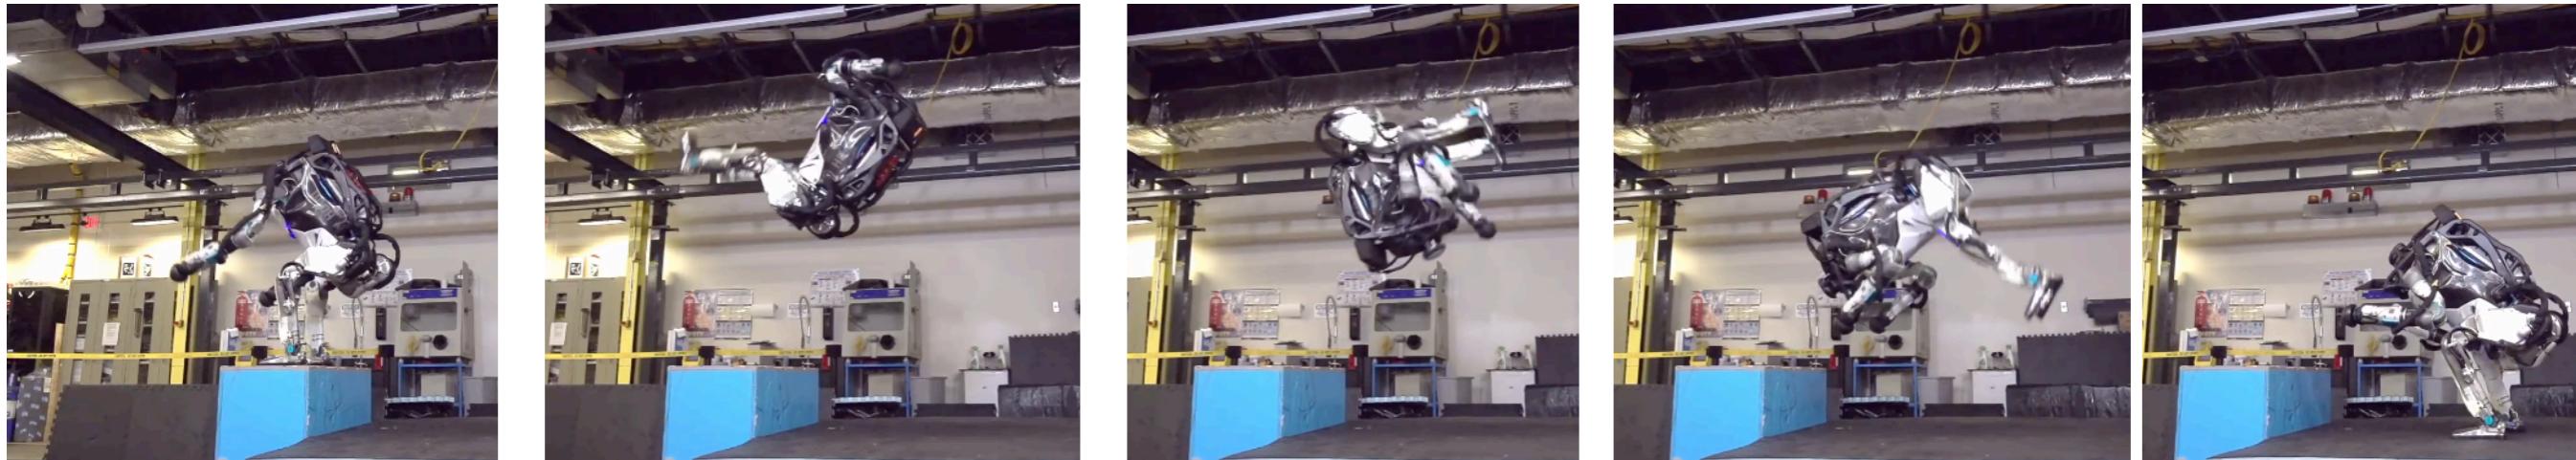
Some
initial
motor
thrust ...

Whole lot of gravity!

Whole lot of momentum!

Challenge 1: Underactuated systems

What affects the error between robot state and reference?



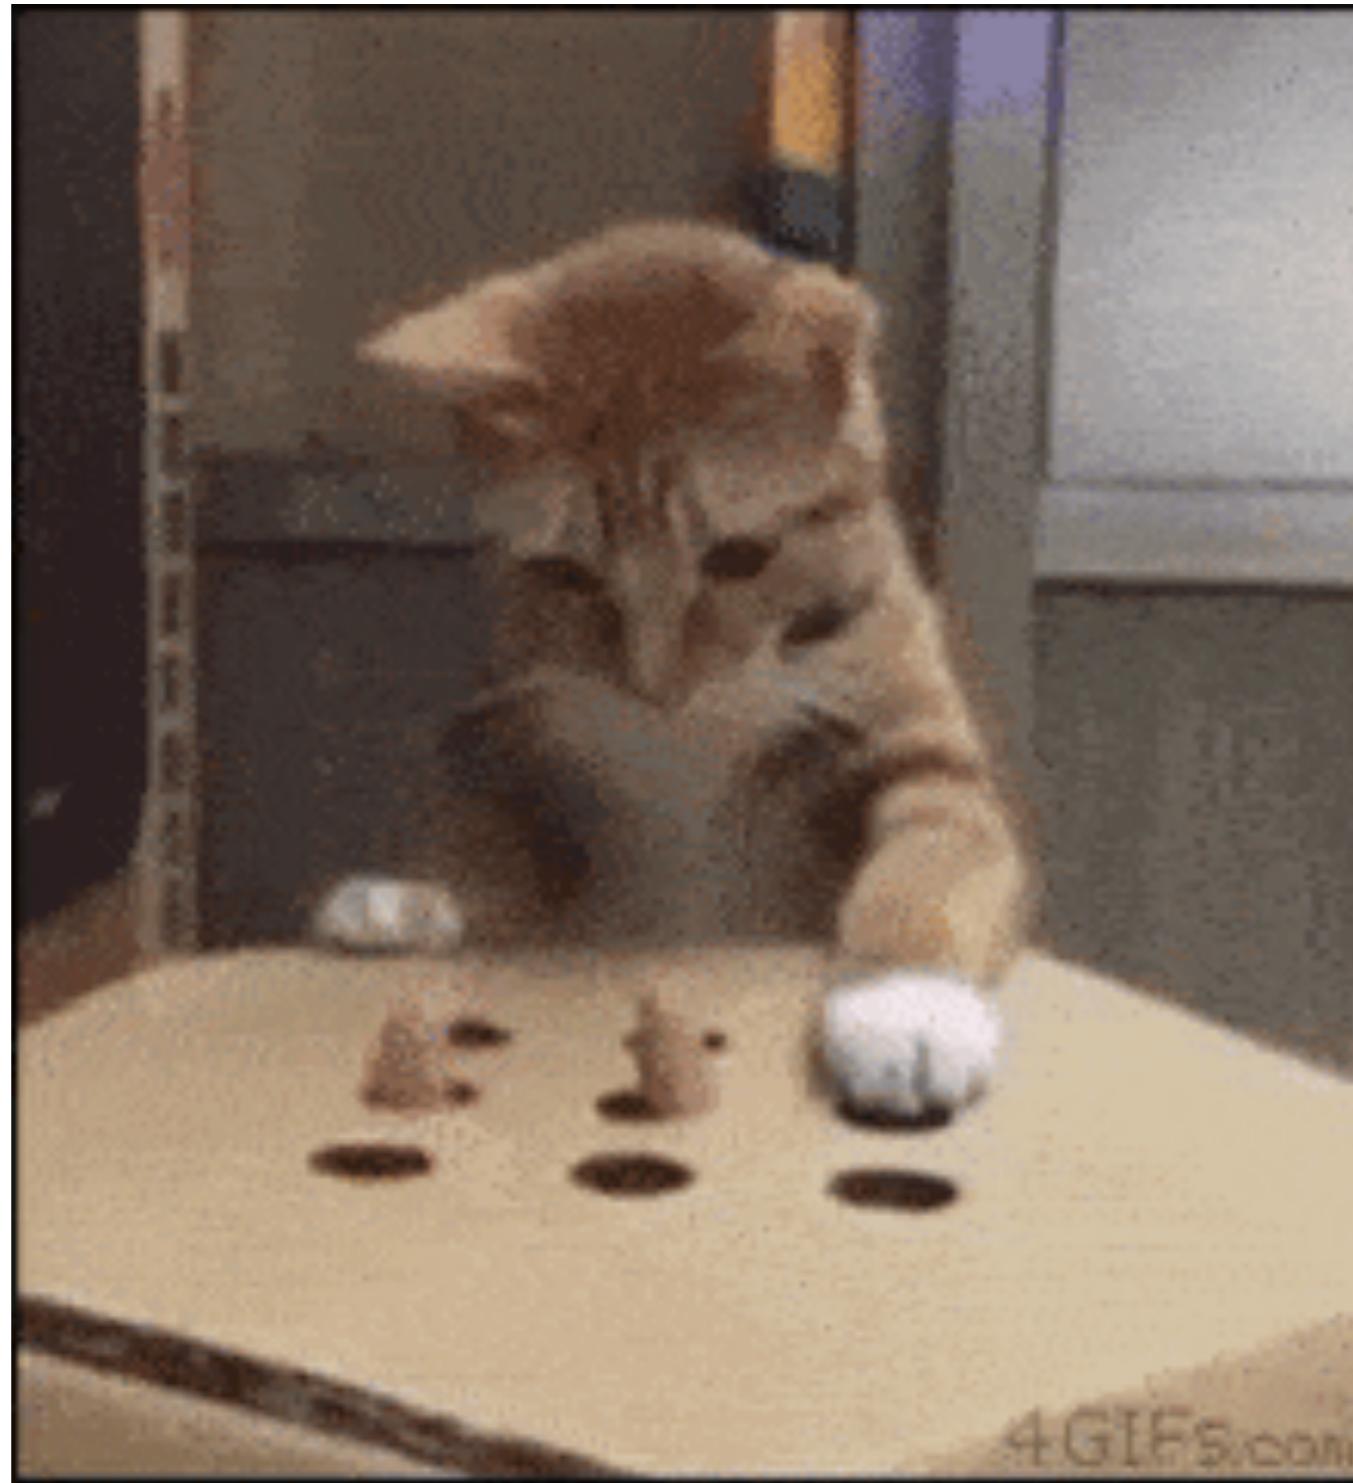
Some
initial
motor
thrust ...

Whole lot of gravity!

Whole lot of momentum!

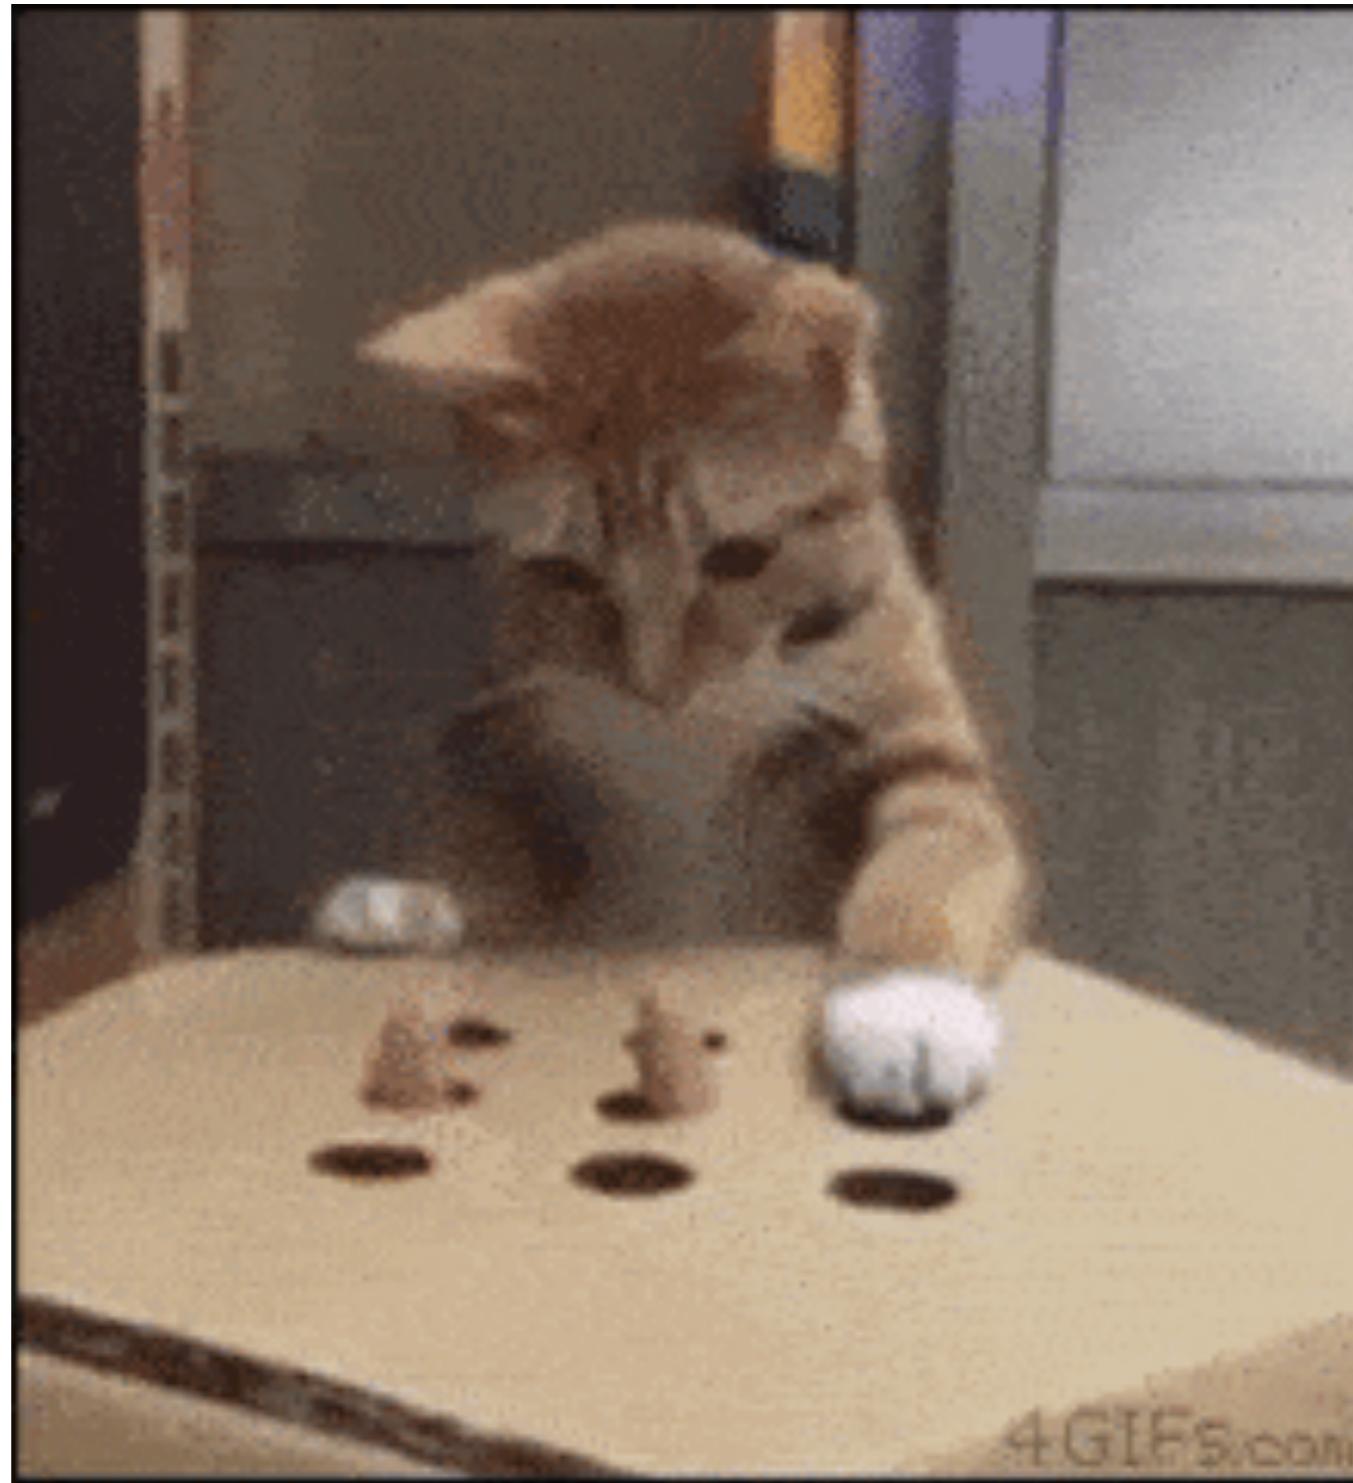
... some precise
control adjustments

Myth: Control is a battle with nature!



Is control like playing whack-a-mole with error terms??

Myth: Control is a battle with nature!



Is control like playing whack-a-mole with error terms??

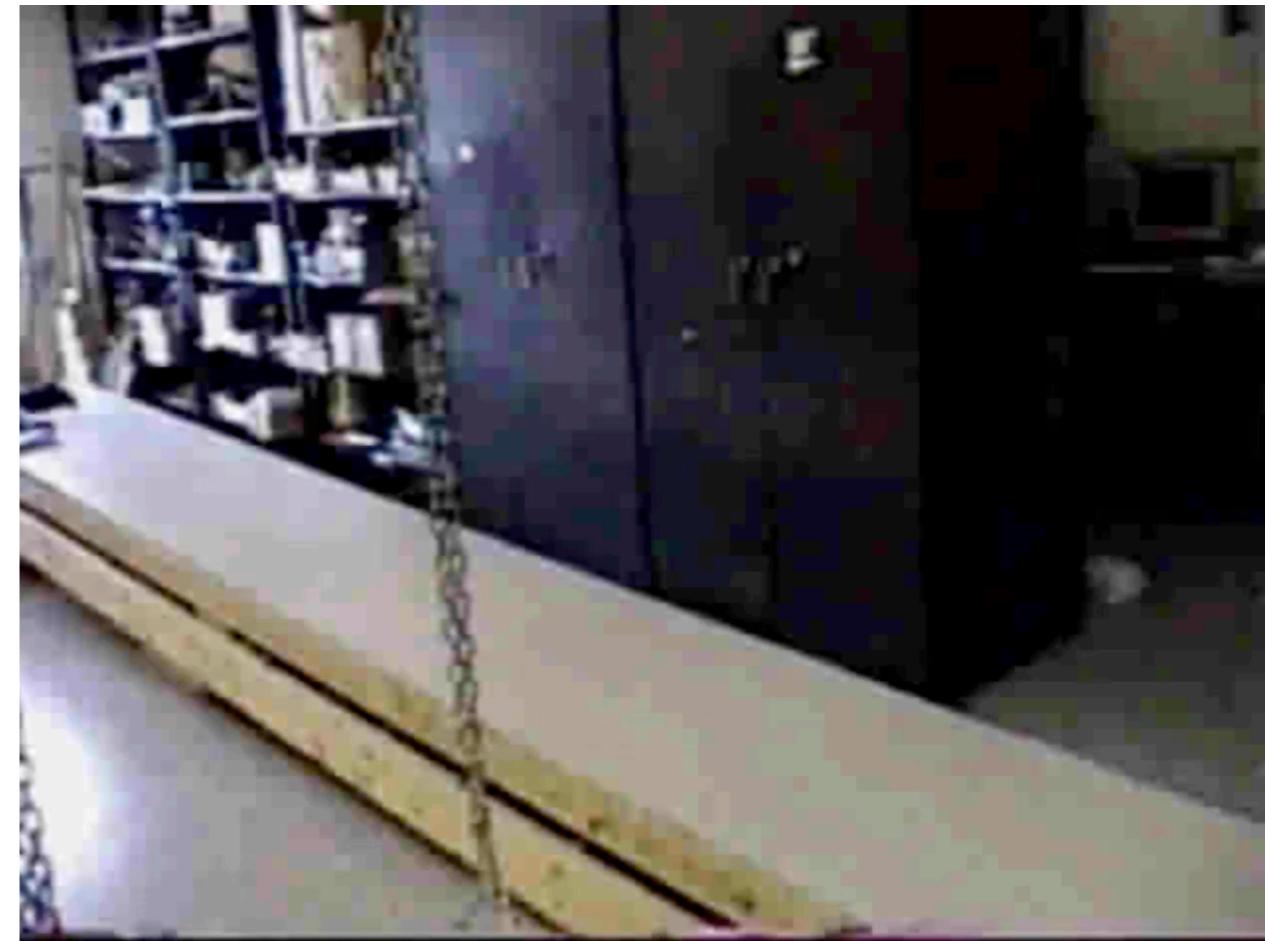
Embrace dynamics instead of fighting it

State-of-the-art humanoid
(at the time)



Honda, 1996

Passive walker
(no motors, powered by gravity!)



Steven H. Collins, Martijn Wisse, and Andy Ruina, 2001

Which one looks more natural? Which one consumes less energy?

Embrace dynamics instead of fighting it

State-of-the-art humanoid
(at the time)



Honda, 1996

Passive walker
(no motors, powered by gravity!)



Steven H. Collins, Martijn Wisse, and Andy Ruina, 2001

Which one looks more natural? Which one consumes less energy?

Question:

If we know the model of our robot,
can't we solve a ginormous
optimization to figure out control ...?

Doing backflips with a helicopter



Redbull Eurocopter BO-105

https://www.youtube.com/watch?v=RGu45s1_QPU

Doing backflips with a helicopter

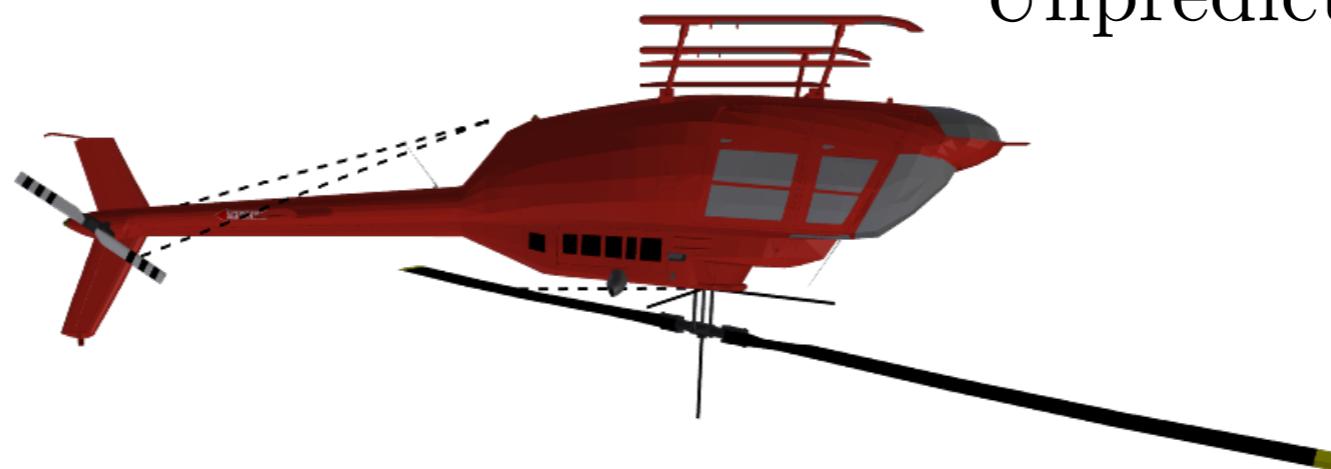


Redbull Eurocopter BO-105

https://www.youtube.com/watch?v=RGu45s1_QPU

And just what is this model ?!?

Nothing
countering
gravity!

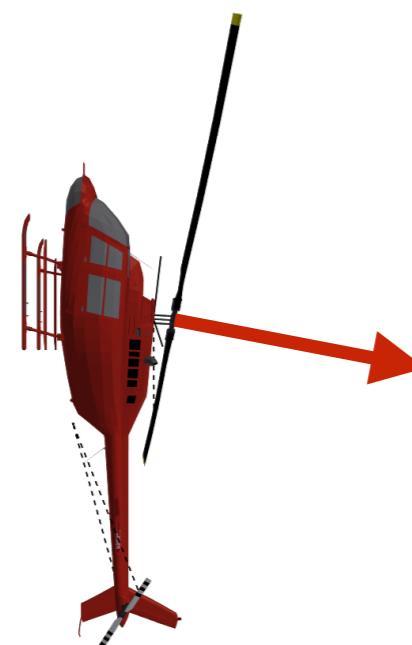


Unpredictable drag forces!

Chaotic vortex around blades!

Hopeless to assume we know exactly how the helicopter
will behave upside down...

Challenge 2: Choose good closed-loop models



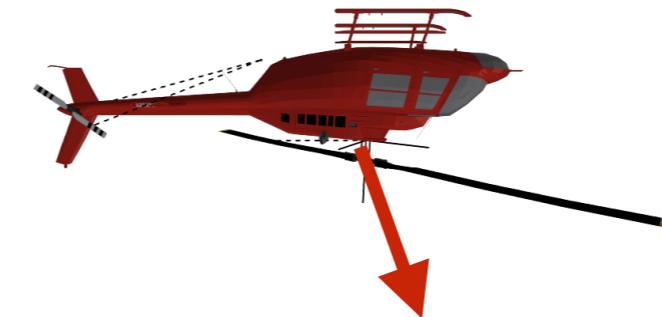
Chaotic
dynamics

+

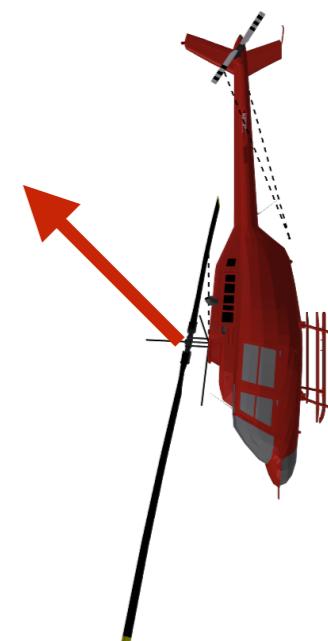


Feedback
control law

=



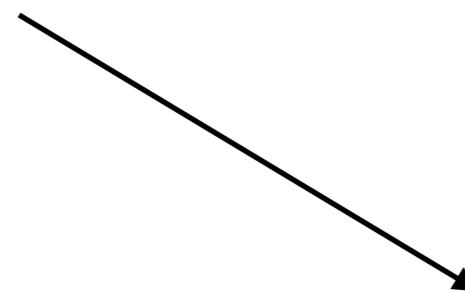
Closed-loop system =
Point mass
with a planar
thrust vector



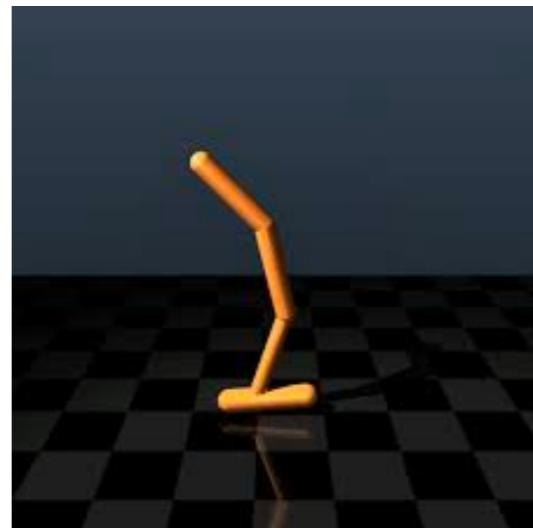
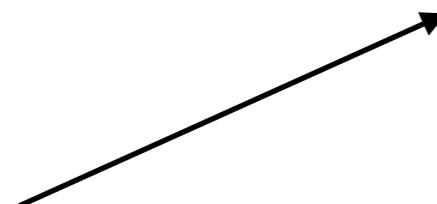
Well-behaved
system

Trotting quadruped = biped = one-leg

Quadruped



Biped



One-leg

Marc Raibert showed how these closed loop system are equivalent

Question:

Is this all offline? Can I pre-compute
a bunch of controllers once and
call it a day?

Figure 8s while drifting



Stanford's MARTY

<https://www.youtube.com/watch?v=nTK56vPb8Zo>

Figure 8s while drifting



Stanford's MARTY

<https://www.youtube.com/watch?v=nTK56vPb8Zo>

Challenge 3: Model changing on the fly!

Run **real-time estimators** for wheel characteristics

Need control laws for all possible model parameters

Other challenges for mobile robot control

1. Unexpected obstacle avoidance

(Lab 2!)

2. Noisy state estimation

(Lab 2 + bimodal estimate from PF)

3. Controllers that guarantee safety

Enough with the motivation ...

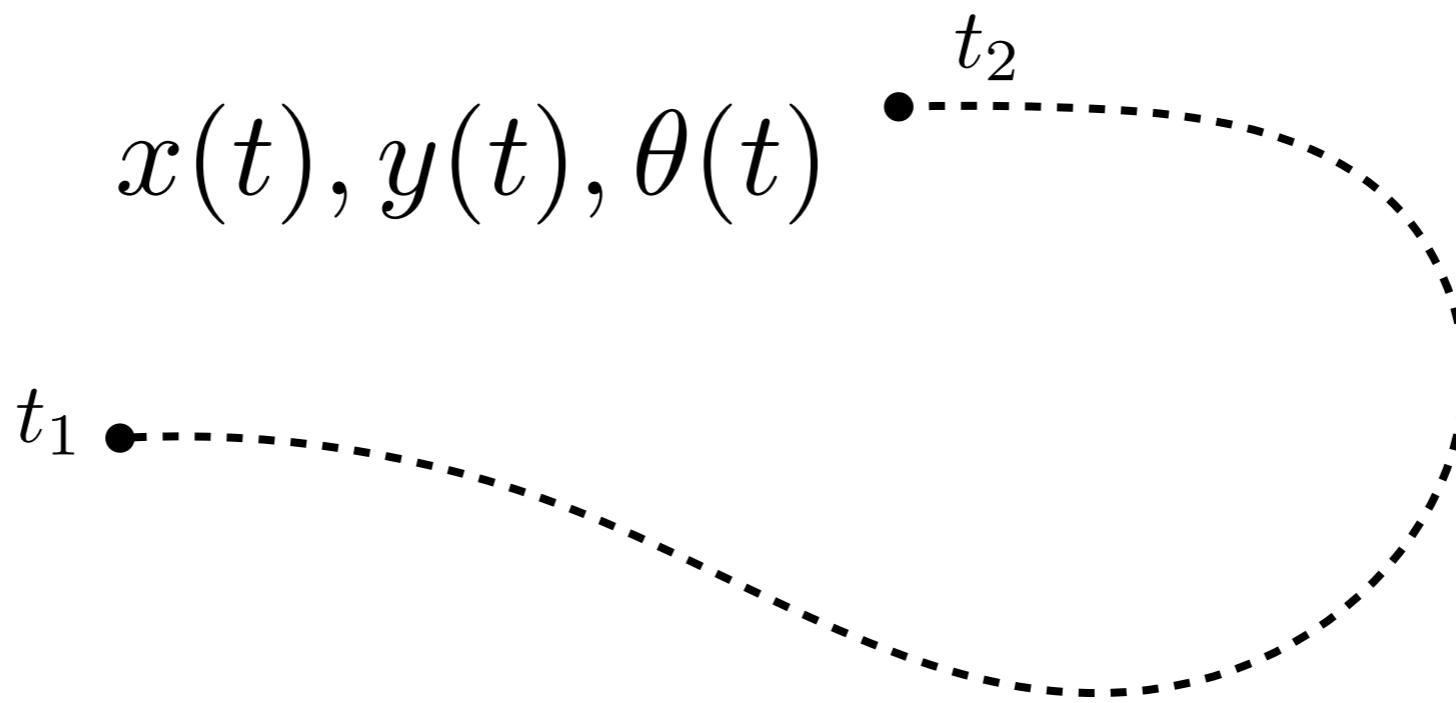
Let's start writing some controllers

A generic template for a controller

1. Get a reference path / trajectory to track
2. Pick a point on the reference
3. Compute error to reference point
4. Compute control law to minimize error

Step 1: Get a reference

Reference can be a **time-parameterized trajectory**



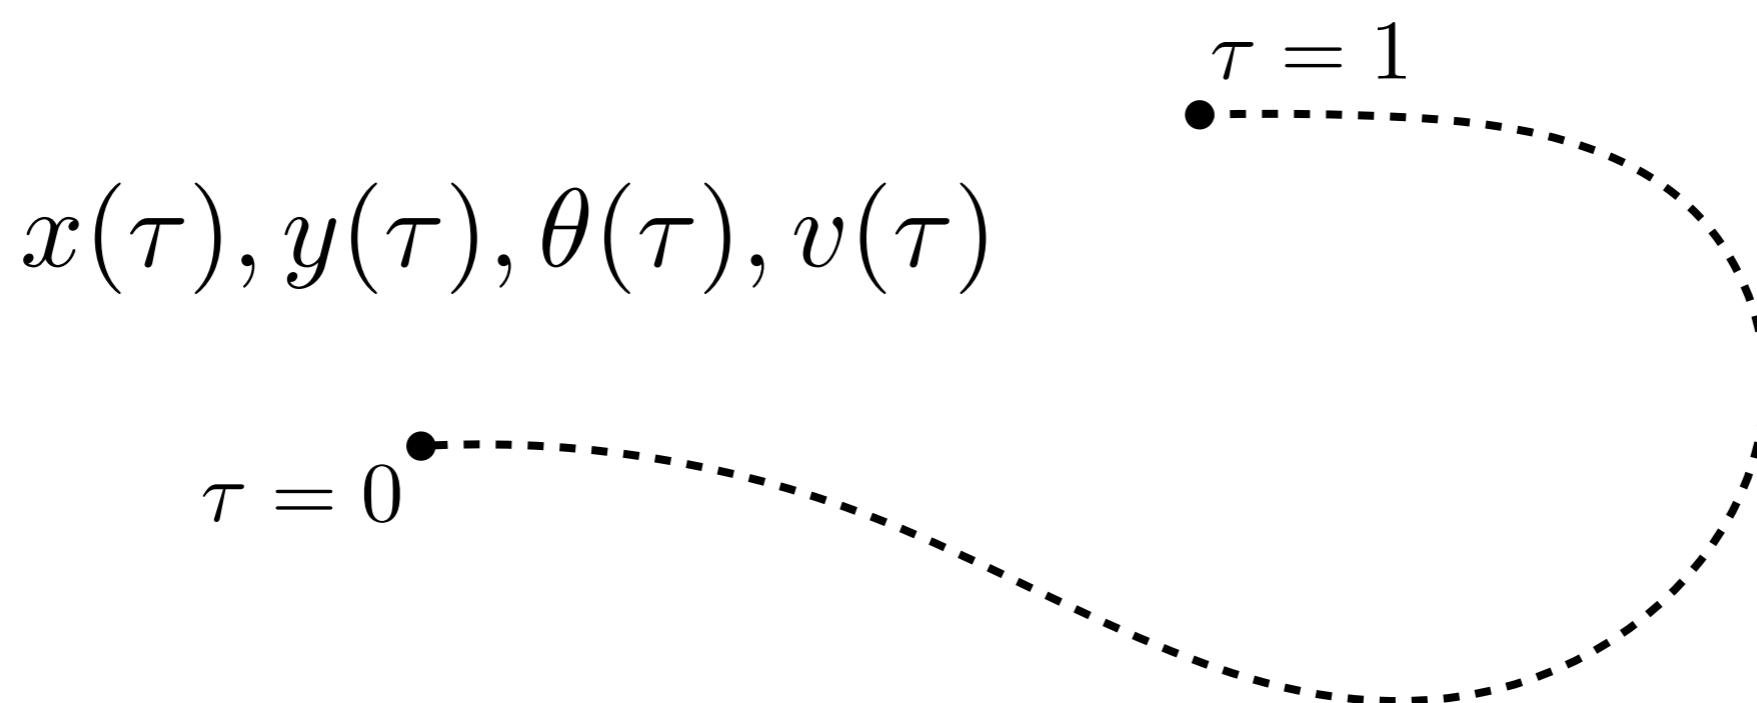
The time index refers to a desired pose we want the system to achieve at a given **time**

Pro: Useful if we want the robot to respect time constraints

Con: Sometimes we care only about deviation from reference

Step 1: Get a reference

Reference can be a index-parameterized path



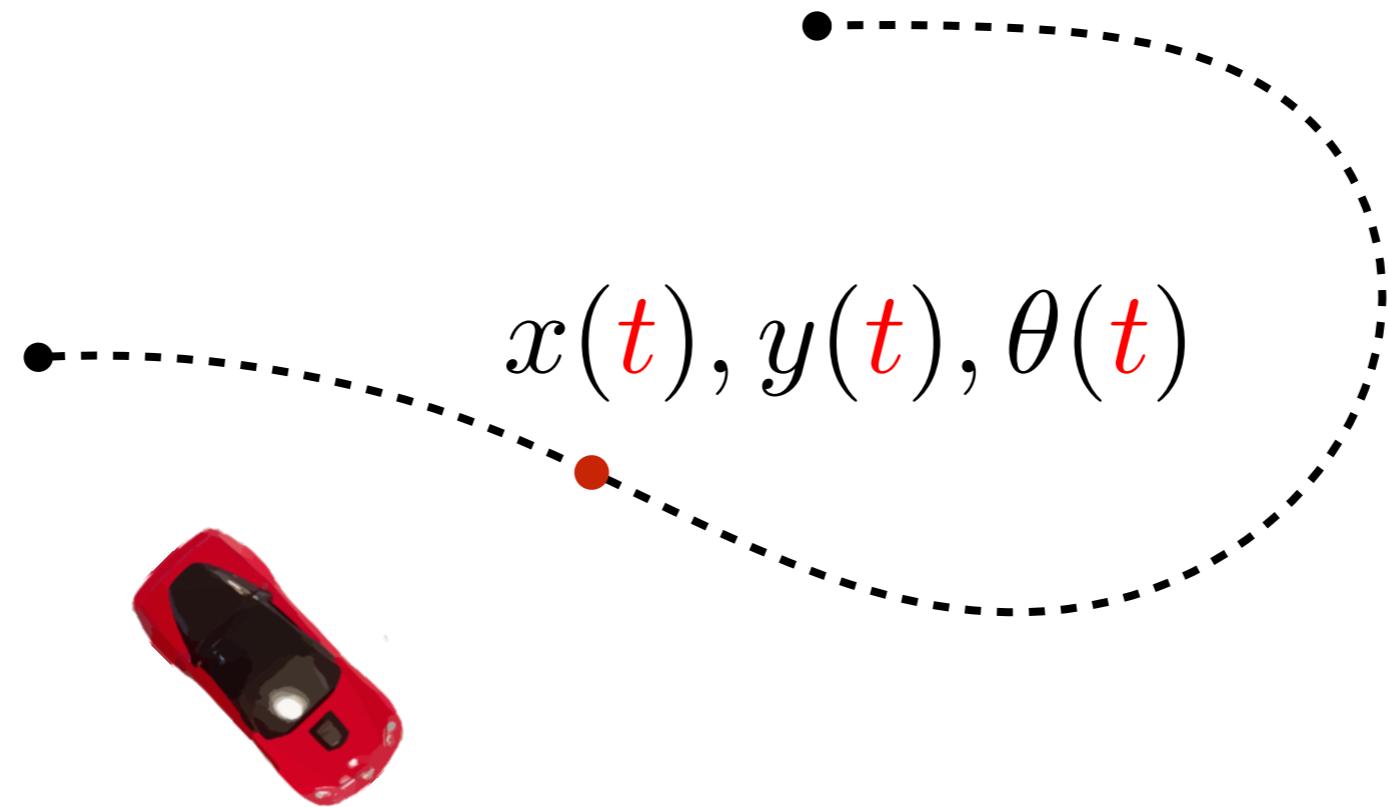
The index is simply a way to access the path; there is no notion of time

Pro: Useful for conveying the shape you want the robot to follow

Con: Can't control when robot will reach a point

Step 2: Pick a point on the reference

Pick a **reference point** on the trajectory / path

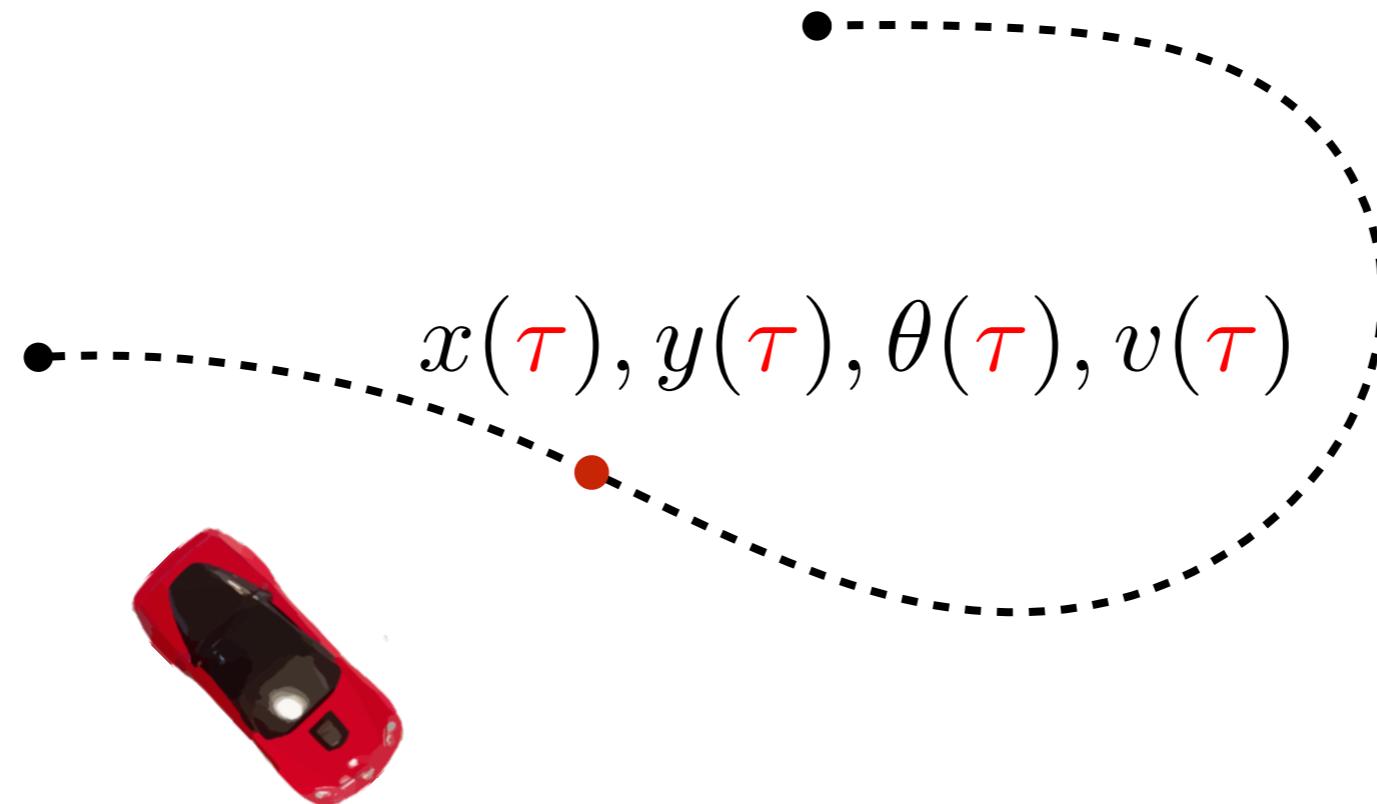


If we are using a time-parameterized trajectory,
the current time is the natural reference

Question: When is this a bad idea?

Step 2: Pick a point on the reference

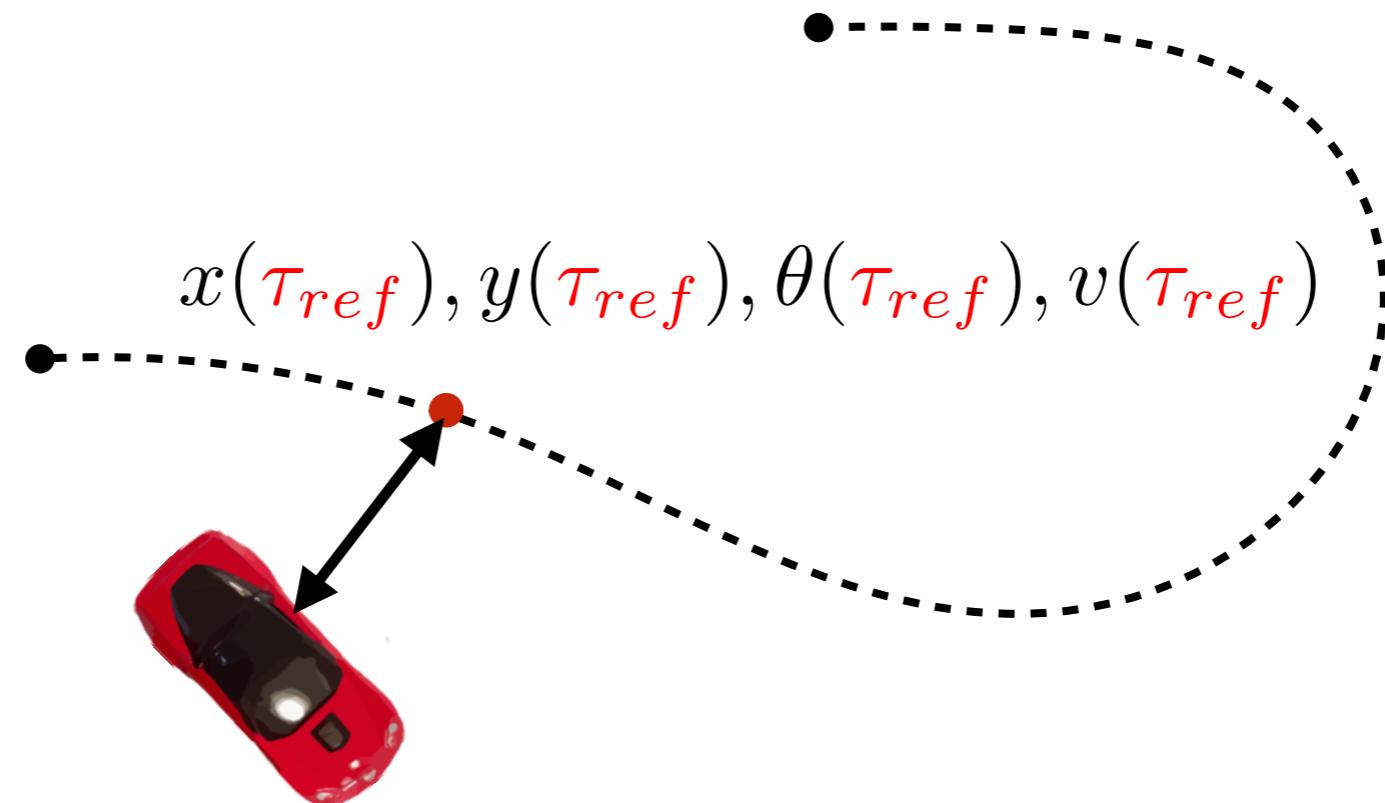
Pick a **reference point** on the trajectory / path



If we are using a index-parameterized trajectory,
there are multiple options

Step 2: Pick a point on the reference

Option 1: Pick the closest point on the path

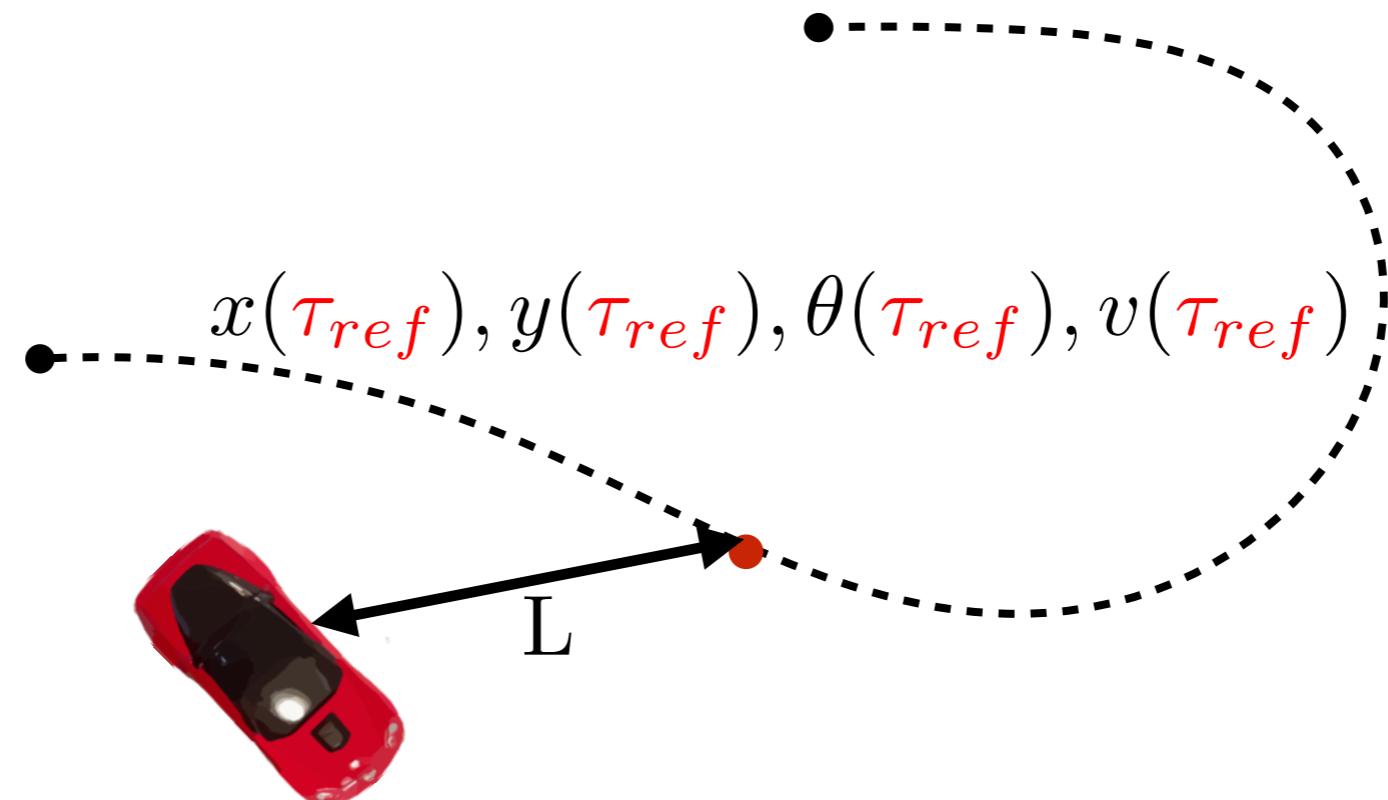


$$\tau_{ref} = \arg \min_{\tau} \| [x \ y]^T - [x(\tau) \ y(\tau)]^T \|$$

Question: When is this a bad idea?

Step 2: Pick a point on the reference

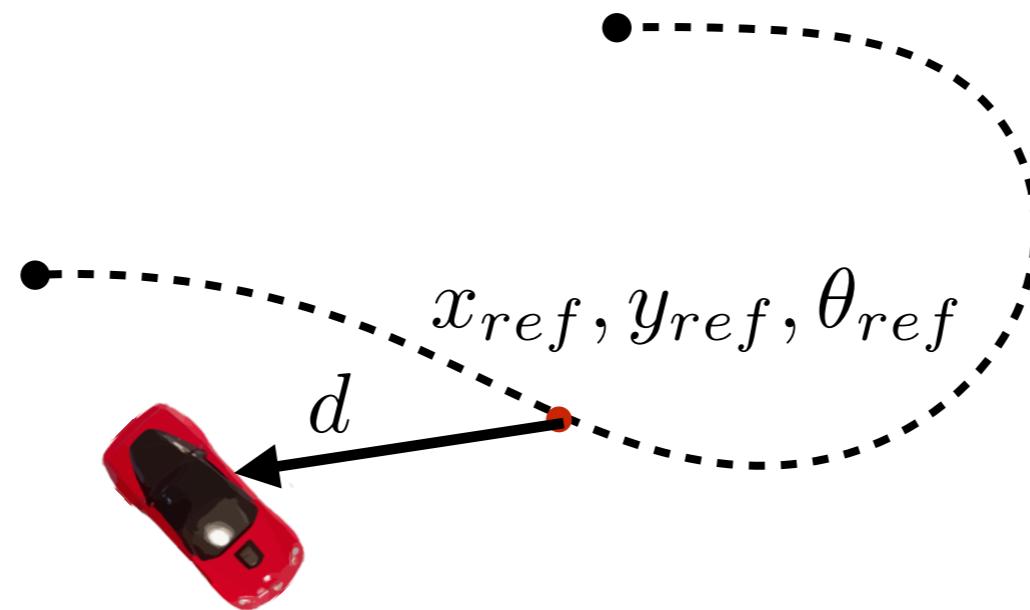
Option 2: Pick a lookahead point on the path



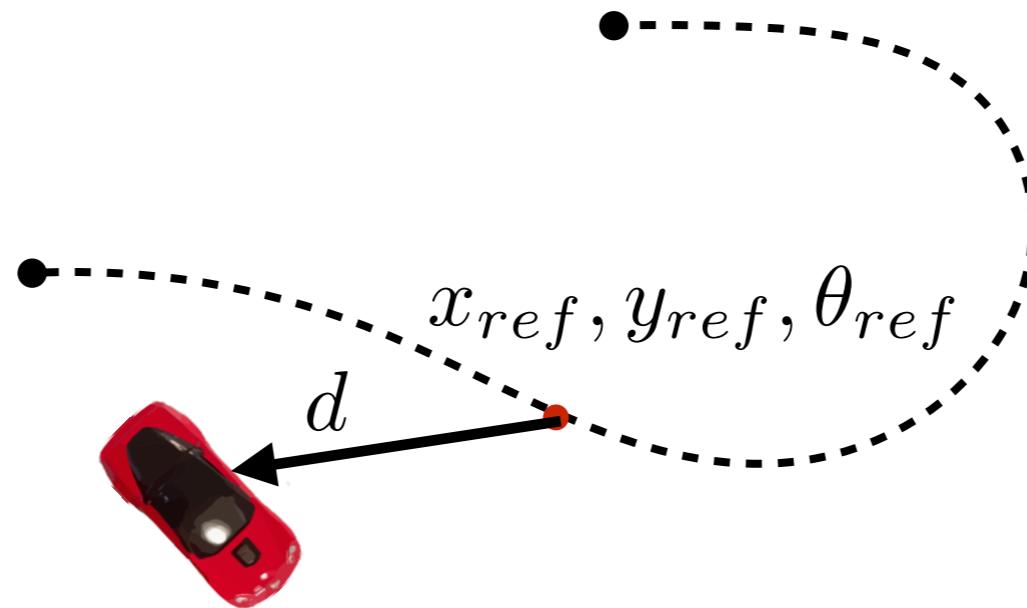
$$\| [x \ y]^T - [x(\tau_{ref}) \ y(\tau_{ref})]^T \| = L$$

Question: What happens when L is too big?

Step 3: Compute error to reference point



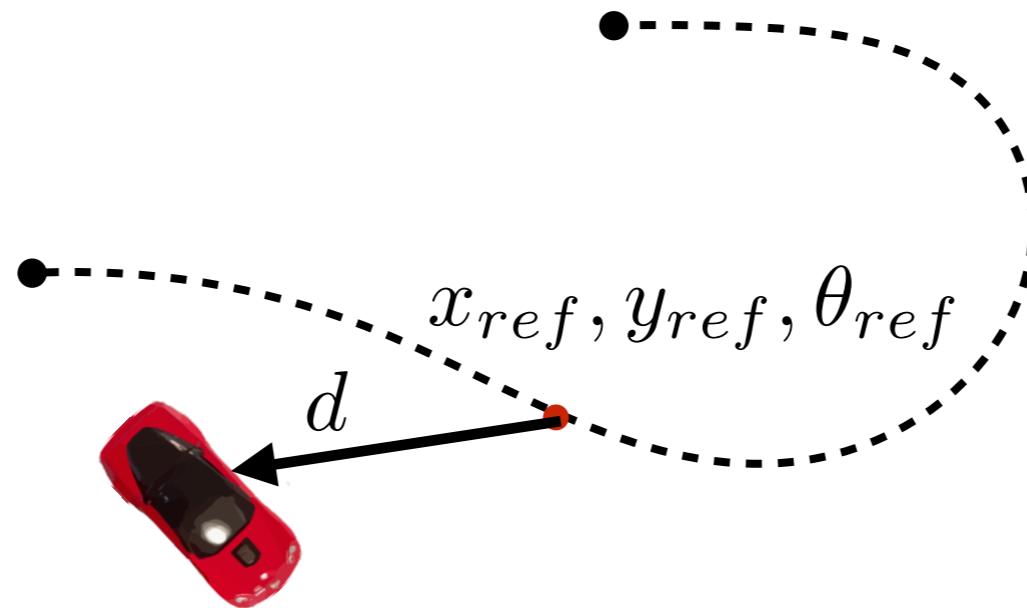
Step 3: Compute error to reference point



1. Define error vector

$$d = \begin{bmatrix} x \\ y \end{bmatrix} - \begin{bmatrix} x_{ref} \\ y_{ref} \end{bmatrix}$$

Step 3: Compute error to reference point



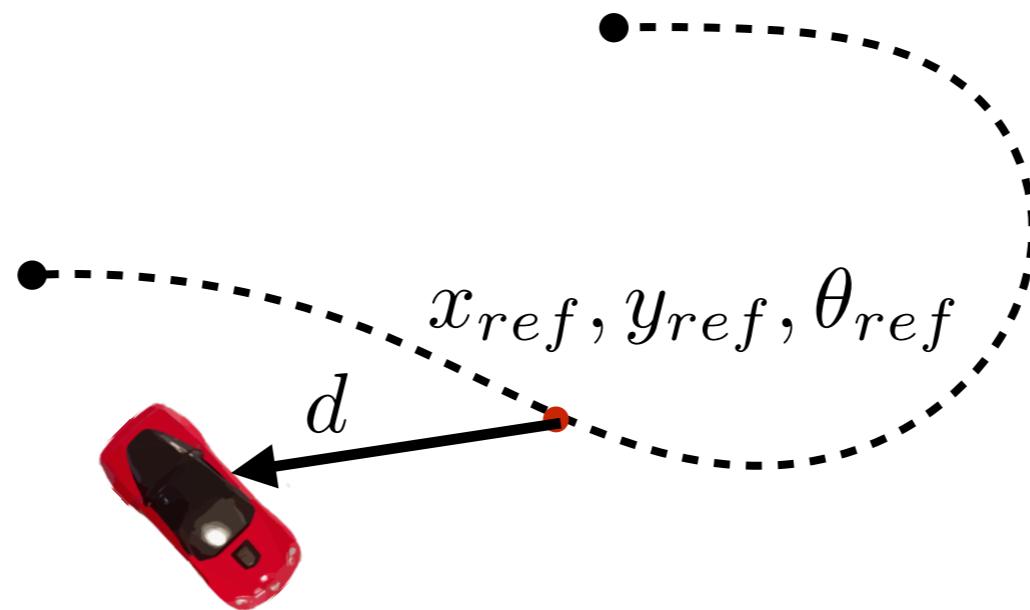
1. Define error vector

$$d = \begin{bmatrix} x \\ y \end{bmatrix} - \begin{bmatrix} x_{ref} \\ y_{ref} \end{bmatrix}$$

2. Rotate error vector to be in the reference frame

$$e = R^T(\theta_{ref}) \left(\begin{bmatrix} x \\ y \end{bmatrix} - \begin{bmatrix} x_{ref} \\ y_{ref} \end{bmatrix} \right)$$

Step 3: Compute error to reference point



1. Define error vector

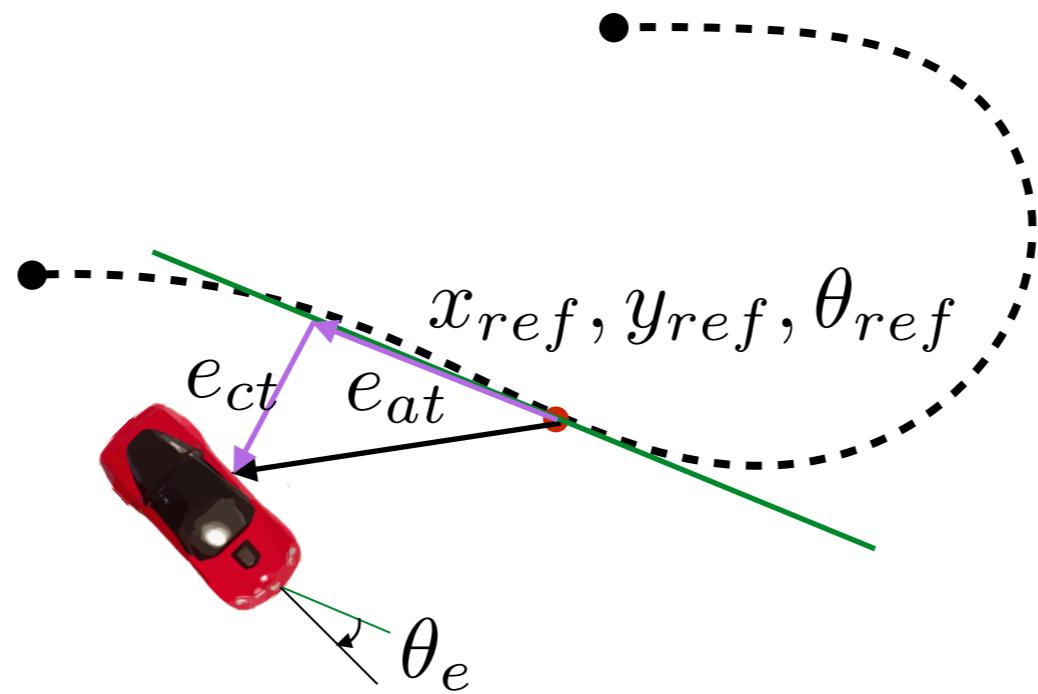
$$d = \begin{bmatrix} x \\ y \end{bmatrix} - \begin{bmatrix} x_{ref} \\ y_{ref} \end{bmatrix}$$

2. Rotate error vector to be in the reference frame

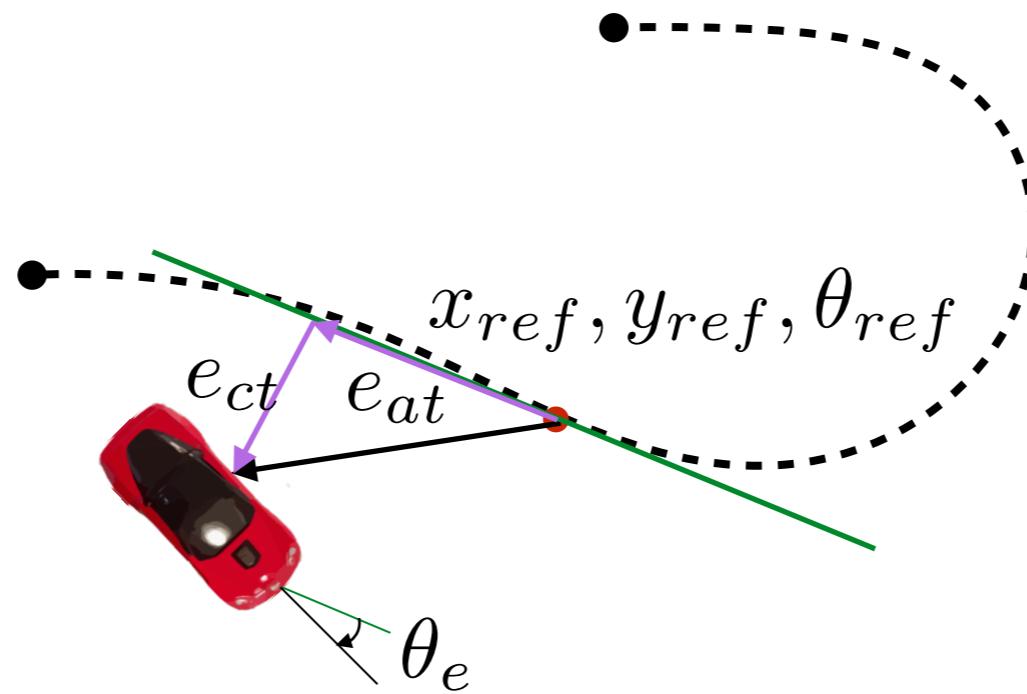
$$e = R^T(\theta_{ref}) \left(\begin{bmatrix} x \\ y \end{bmatrix} - \begin{bmatrix} x_{ref} \\ y_{ref} \end{bmatrix} \right)$$

$$e = \begin{bmatrix} \cos(\theta_{ref}) & \sin(\theta_{ref}) \\ -\sin(\theta_{ref}) & \cos(\theta_{ref}) \end{bmatrix} \left(\begin{bmatrix} x \\ y \end{bmatrix} - \begin{bmatrix} x_{ref} \\ y_{ref} \end{bmatrix} \right)$$

Step 3: Compute error to reference point



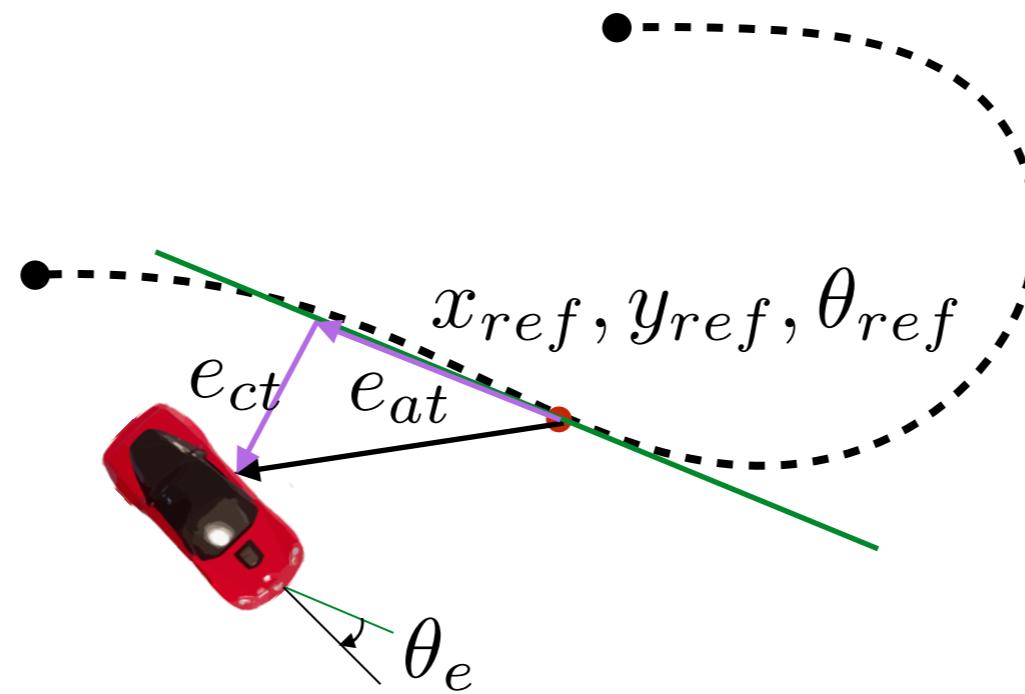
Step 3: Compute error to reference point



Error has two components: along-track and cross-track

$$e = \begin{bmatrix} e_{at} \\ e_{ct} \end{bmatrix} \quad \begin{array}{l} \text{(along track)} \\ \text{(cross track)} \end{array}$$

Step 3: Compute error to reference point



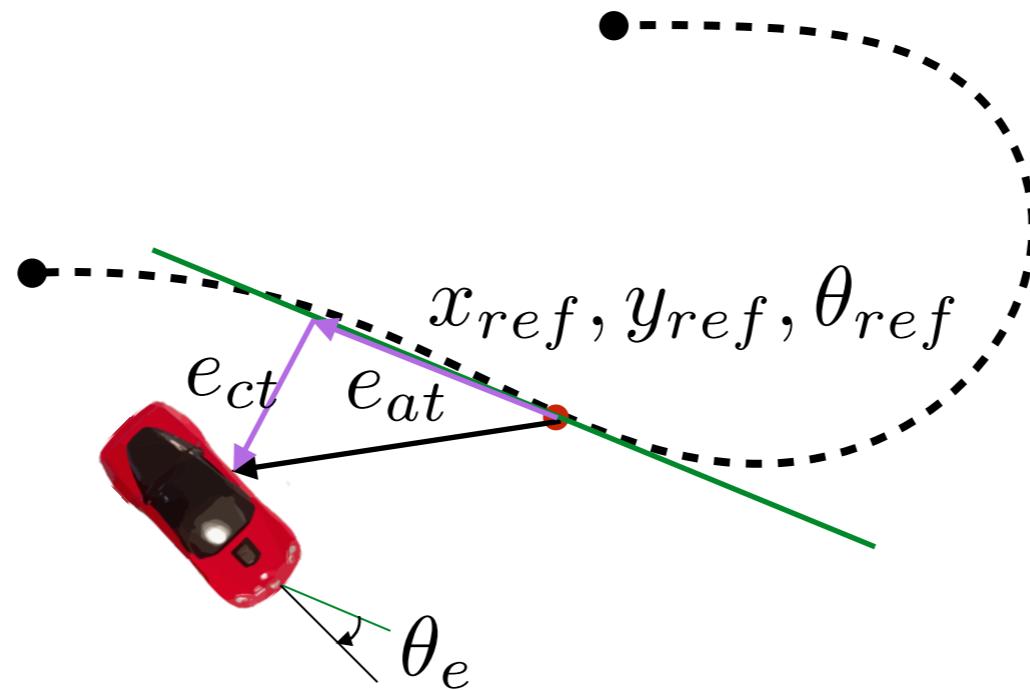
Error has two components: along-track and cross-track

$$e = \begin{bmatrix} e_{at} \\ e_{ct} \end{bmatrix} \quad \begin{array}{l} \text{(along track)} \\ \text{(cross track)} \end{array}$$

$$\theta_e = \theta - \theta_{ref} \quad \text{(heading error)}$$

Let's consider only the cross track error for now

Step 3: Compute error to reference point

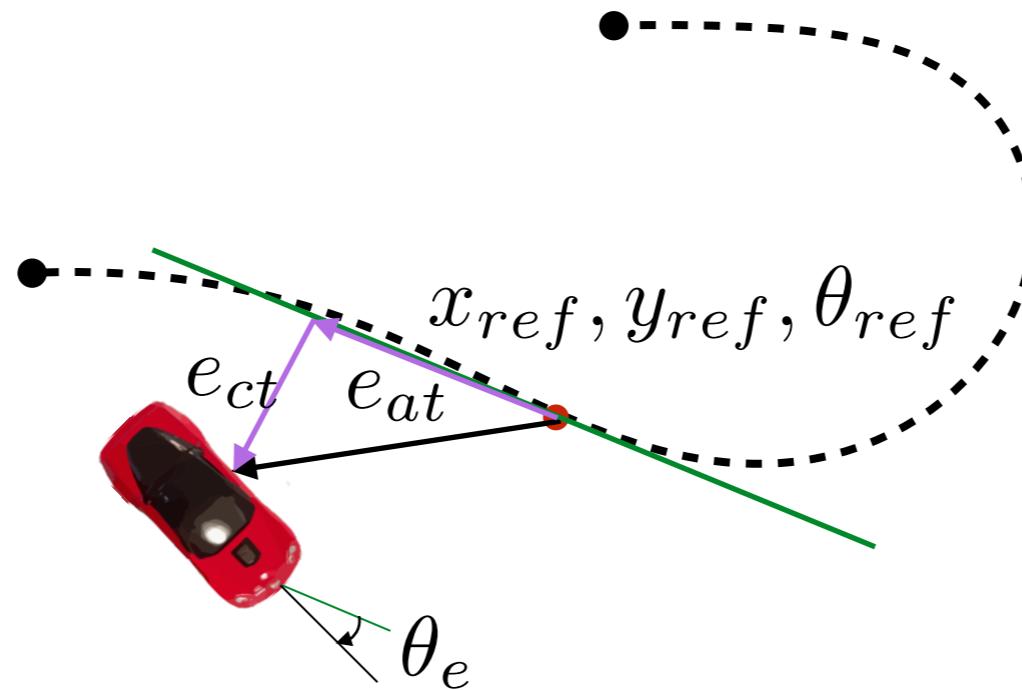


Cross-track error

$$e_{ct} = -\sin(\theta_{ref})(x - x_{ref}) + \cos(\theta_{ref})(y - y_{ref})$$

$$= V \sin(\theta_e)$$

Step 3: Compute error to reference point



Cross-track error

$$e_{ct} = -\sin(\theta_{ref})(x - x_{ref}) + \cos(\theta_{ref})(y - y_{ref})$$

Derivative of cross-track error

$$\begin{aligned}\dot{e}_{ct} &= -\sin(\theta_{ref})\dot{x} + \cos(\theta_{ref})\dot{y} \\ &= -\sin(\theta_{ref})V \cos(\theta) + \cos(\theta_{ref})V \sin(\theta) \\ &= V \sin(\theta - \theta_{ref}) = V \sin(\theta_e)\end{aligned}$$

Step 4: Compute control law

Compute control action based on instantaneous error

$$u = K(\mathbf{x}, e)$$

control

state error

(steering angle, speed)

Step 4: Compute control law

Compute control action based on instantaneous error

$$u = K(\mathbf{x}, e)$$

control

state error

(steering angle, speed)

Apply control action, robot moves a bit, compute new error, repeat

Step 4: Compute control law

Compute control action based on instantaneous error

$$u = K(\mathbf{x}, e)$$

control

state error

(steering angle, speed)

Apply control action, robot moves a bit, compute new error, repeat

Different laws have different trade-offs

We assume that control speed is set to be equal to reference speed

Hence control = steering angle

Different control laws

1. Bang-bang control
2. PID control
3. Pure-pursuit control
4. Lyapunov control
5. Linear Quadratic Regulator (LQR)
6. Model Predictive Control (MPC)