

Linear Quadratic Regulator

Sanjiban Choudhury

Different control laws

1. Bang-bang control

2. PID control

3. Pure-pursuit control

4. Lyapunov control

5. LQR

6. MPC

Recap of controllers

PID / Pure pursuit: Worked well, no provable guarantees

Lyapunov: Provable stability, convergence rate depends on gains

Table of controllers

	Control Law	Uses model	Stability Guarantee	Minimize Cost
PID	$u = K_p e + \dots$	No	No	No
Pure Pursuit	$u = \tan^{-1} \left(\frac{2B \sin \alpha}{L} \right)$	Circular arcs	Yes - with assumptions	No
Lyapunov	$u = \tan^{-1} \left(-\frac{k_1 e_{ct} B}{\theta_e} \sin \theta_e - \frac{B}{V} k_2 \theta_e \right)$	Non-linear	Yes	No


Is stability enough?

$$\lim_{t \rightarrow \infty} e(t) = 0$$


Is *stability* enough of a guarantee?




Is *stability* enough of a guarantee?



Is **stability** enough of a guarantee?




Control action changes abruptly - why is this bad?



Is *stability* enough of a guarantee?


What if we just choose really small gains?

Is **stability** enough of a guarantee?



What if we just choose really small gains?

Is **stability** enough of a guarantee?



What if we just choose really small gains?

Stability guarantees that the error will go to zero ...
but can take arbitrary long time

Question:

How do we trade-off both
driving error to zero
AND
keeping control action small?

Key Idea:

Turn the problem into an
optimization

$$\min_{u(t)} \int_0^{\infty} (w_1 e(t)^2 + w_2 u(t)^2) dt$$

Optimal Control

A fundamental framework:


Linear Quadratic Regulator

Trivia! :) (from <http://www.uta.edu/utari/acs/history.htm>)

In 1960 three major papers were published by R. Kalman and coworkers...


1. One of these [Kalman and Bertram 1960], publicized the vital work of Lyapunov in the time-domain control of nonlinear systems.
2. The next [Kalman 1960a] discussed the optimal control of systems, providing the design equations for the linear quadratic regulator (LQR).
3. The third paper [Kalman 1960b] discussed optimal filtering and estimation theory, providing the design equations for the discrete Kalman filter.

LQR flying RC helicopters



(Excellent work by Pieter Abeel et al. https://people.eecs.berkeley.edu/~pabbeel/autonomous_helicopter.html)

LQR flying RC helicopters



(Excellent work by Pieter Abeel et al. https://people.eecs.berkeley.edu/~pabbeel/autonomous_helicopter.html)

A fundamental framework:

Linear Quadratic Regulator

The Linear Quadratic Regulator (LQR) Problem

Given:

The Linear Quadratic Regulator (LQR) Problem

Given:

1. **Linear** dynamical system

$$x_{t+1} = Ax_t + Bu_t$$

The Linear Quadratic Regulator (LQR) Problem

Given:

1. **Linear** dynamical system

$$x_{t+1} = Ax_t + Bu_t$$

2. A reference state which we are **regulating** around

$$x_{ref} = 0$$

The Linear Quadratic Regulator (LQR) Problem

Given:

1. **Linear** dynamical system

$$x_{t+1} = Ax_t + Bu_t$$

2. A reference state which we are **regulating** around

$$x_{ref} = 0$$

3. A **quadratic** cost function to minimize

$$\begin{aligned} c(x_t, u_t) &= (x_t - x_{ref})^T Q (x_t - x_{ref}) + u_t^T R u_t \\ &= x_t^T Q x_t + u_t^T R u_t \end{aligned}$$

The Linear Quadratic Regulator (LQR) Problem

Given:

1. **Linear** dynamical system

$$x_{t+1} = Ax_t + Bu_t$$

2. A reference state which we are **regulating** around

$$x_{ref} = 0$$


3. A **quadratic** cost function to minimize

$$\begin{aligned} c(x_t, u_t) &= (x_t - x_{ref})^T Q (x_t - x_{ref}) + u_t^T R u_t \\ &= x_t^T Q x_t + u_t^T R u_t \end{aligned}$$

Goal: Compute control actions to minimize cumulative cost (value)

$$J = \sum_{t=0}^{T-1} x_t^T Q x_t + u_t^T R u_t$$


Example: Inverted Pendulum



Example: Inverted Pendulum

Equations of motion

$$ml^2\ddot{\theta} - mgl \sin \theta = \tau$$




Example: Inverted Pendulum


Equations of motion

$$ml^2\ddot{\theta} - mgl \sin \theta = \tau$$

$$\ddot{\theta} = \frac{g}{l} \sin \theta + \frac{1}{ml^2} \tau$$



Example: Inverted Pendulum



Equations of motion


$$ml^2\ddot{\theta} - mgl \sin \theta = \tau$$

$$\ddot{\theta} = \frac{g}{l} \sin \theta + \frac{1}{ml^2} \tau$$

$$\approx \frac{g}{l} \theta + \frac{1}{ml^2} \tau$$

Example: Inverted Pendulum

Equations of motion



$$ml^2\ddot{\theta} - mgl \sin \theta = \tau$$

$$\ddot{\theta} = \frac{g}{l} \sin \theta + \frac{1}{ml^2} \tau$$


$$\approx \frac{g}{l} \theta + \frac{1}{ml^2} \tau$$

$$\begin{bmatrix} \theta \\ \dot{\theta} \end{bmatrix}_{t+1} = \begin{bmatrix} 1 + \frac{1}{2} \frac{g}{l} \Delta t^2 & \Delta t \\ \frac{g}{l} \Delta t & 1 \end{bmatrix} \begin{bmatrix} \theta \\ \dot{\theta} \end{bmatrix}_t + \begin{bmatrix} \frac{1}{2} \Delta t^2 \\ \Delta t \end{bmatrix} \frac{\tau}{ml^2}$$


A

B


Get to (0,0) while minimizing cost




Get to (0,0) while minimizing cost




Get to (0,0) while minimizing cost




Get to (0,0) while minimizing cost




Get to (0,0) while minimizing cost




Get to (0,0) while minimizing cost




Get to (0,0) while minimizing cost



Get to (0,0) while minimizing cost



Observation: Cost-to-go is not uniform



How do we solve for controls?

How do we solve for controls?

Dynamic programming to the rescue!

How do we solve for controls?

Dynamic programming to the rescue!

Recall the Bellman function that relates value at consecutive time steps

$$\begin{aligned} J(x_t, t) &= \min_{u_t} c(x_t, u_t) + J(x_{t+1}, t+1) \\ &= \min_{u_t} x_t^T Q x_t + u_t^T R u_t + J(x_{t+1}, t+1) \end{aligned}$$

How do we solve for controls?

Dynamic programming to the rescue!

Start from timestep $T-1$ and solve backwards

How do we solve for controls?

Dynamic programming to the rescue!

Start from timestep $T-1$ and solve backwards

|
|
|
|
|
|
|
|
|
|
|

|
|
|
|
|
|
|
|
|
|
|

$T-3$

|
|
|
|
|
|
|
|
|
|
|

$T-2$


|
|
|
|
|
|
|
|
|
|
|

$T-1$

How do we solve for controls?

Dynamic programming to the rescue!


Start from timestep $T-1$ and solve backwards



How do we solve for controls?

Dynamic programming to the rescue!


Start from timestep $T-1$ and solve backwards




How do we solve for controls?

Dynamic programming to the rescue!

Start from timestep $T-1$ and solve backwards




Last time step T-1



Last time step T-1

We have only 1 term in the cost function

$$J(x_{T-1}, u_{T-1}) = \min_{u_T} x_{T-1}^T Q x_{T-1} + u_{T-1}^T R u_{T-1}$$




Last time step T-1

We have only 1 term in the cost function

$$J(x_{T-1}, u_{T-1}) = \min_{u_T} x_{T-1}^T Q x_{T-1} + u_{T-1}^T R u_{T-1}$$

To minimize cost, set control to 0

$$u_{T-1} = 0$$



Last time step T-1

We have only 1 term in the cost function

$$J(x_{T-1}, u_{T-1}) = \min_{u_T} x_{T-1}^T Q x_{T-1} + u_{T-1}^T R u_{T-1}$$

To minimize cost, set control to 0


$$u_{T-1} = 0$$

The cost function is a **quadratic**

$$\begin{aligned} J(x_{T-1}, u_{T-1}) &= x_{T-1}^T Q x_{T-1} \\ &= x_{T-1}^T V_{T-1} x_{T-1} \end{aligned}$$



(this is a value matrix)



Last time step T-1

We have only 1 term in the cost function

$$J(x_{T-1}, u_{T-1}) = \min_{u_T} x_{T-1}^T Q x_{T-1} + u_{T-1}^T R u_{T-1}$$


To minimize cost, set control to 0

$$u_{T-1} = 0$$


The cost function is a **quadratic**

$$\begin{aligned} J(x_{T-1}, u_{T-1}) &= x_{T-1}^T Q x_{T-1} \\ &= x_{T-1}^T V_{T-1} x_{T-1} \end{aligned}$$

↑
(this is a value matrix)




Last time step T-2



Last time step T-2


$$J(x_{T-2}, T-2) = \min_{u_{T-2}} c(x_{T-2}, u_{T-2}) + J(x_{T-1}, T-1)$$



Last time step T-2

$$J(x_{T-2}, T-2) = \min_{u_{T-2}} c(x_{T-2}, u_{T-2}) + J(x_{T-1}, T-1)$$

$$= \min_{u_{T-2}} x_{T-2}^T Q x_{T-2} + u_{T-2}^T R u_{T-2} + x_{T-1}^T V_{T-1} x_{T-1}$$




Last time step T-2

$$J(x_{T-2}, T-2) = \min_{u_{T-2}} c(x_{T-2}, u_{T-2}) + J(x_{T-1}, T-1)$$

$$= \min_{u_{T-2}} x_{T-2}^T Q x_{T-2} + u_{T-2}^T R u_{T-2} + x_{T-1}^T V_{T-1} x_{T-1}$$

Solve for control at timestep T-1



Last time step T-2

$$J(x_{T-2}, T-2) = \min_{u_{T-2}} c(x_{T-2}, u_{T-2}) + J(x_{T-1}, T-1)$$


$$= \min_{u_{T-2}} x_{T-2}^T Q x_{T-2} + u_{T-2}^T R u_{T-2} + x_{T-1}^T V_{T-1} x_{T-1}$$

Solve for control at timestep T-1

$$u_{T-2} = - \underbrace{(R + B^T V_{T-1} B)^{-1} B^T V_{T-1} A}_{K_{T-2}} x_{T-2}$$


K_{T-2}

Observation: Control law is linear!



Key insight: Value function is always quadratic


Plug back control in the value function (cumulative cost)



Key insight: Value function is always quadratic

Plug back control in the value function (cumulative cost)


$$J(x_{T-2}, T-2) = x_{T-2}^T \underbrace{(Q + K_{T-2}^T R K_{T-2} + (A + B K_{T-2})^T V_{T-1} (A + B K_{T-2}))}_{V_{T-2}} x_{T-2}$$



Key insight: Value function is always quadratic

Plug back control in the value function (cumulative cost)

$$J(x_{T-2}, T-2) = x_{T-2}^T \underbrace{(Q + K_{T-2}^T R K_{T-2} + (A + B K_{T-2})^T V_{T-1} (A + B K_{T-2}))}_{V_{T-2}} x_{T-2}$$



We can derive this relation at **ALL time steps**

$$K_t = -(R + B^T V_{t+1} B)^{-1} B^T V_{t+1} A$$

$$V_t = Q + K_t^T R K_t + (A + B K_t)^T V_{t+1} (A + B K_t)$$

Current
cost

Action
cost

Closed
loop
dynamics

Future
value func

Closed
loop
dynamics

The LQR algorithm

Algorithm OptimalValue(A, B, Q, R, t, T)

if $t = T - 1$ **then**

return Q

end

else

$V_{t+1} = \text{OptimalValue}(A, B, Q, R, t + 1, T)$


$K_t = -(B^T V_{t+1} B + R)^{-1} B^T V_{t+1} A$

return $V_t = Q + K_t^T R K_t + (A + B K_t)^T V_{t+1} (A + B K_t)$


end

(Courtesy Drew Bagnell)


Contours of value function (T-1)




Contours of value function (T-2)



Contours of value function (many steps)



How does the value function evolve?



What if my time horizon is very very
very large?

Convergence of value iteration

Convergence of value iteration

Theorem: If the system is stabilizable, then the value V will converge

Convergence of value iteration

Theorem: If the system is stabilizable, then the value V will converge

$$V = Q + K^T R K + (A + B K)^T V (A + B K)$$

$$K = -(R + B^T V B)^{-1} B^T V A$$

Discrete Algebraic Ricatti Equation (DARE)

Convergence of value iteration

Theorem: If the system is stabilizable, then the value V will converge

$$V = Q + K^T R K + (A + B K)^T V (A + B K)$$

$$K = -(R + B^T V B)^{-1} B^T V A$$

Discrete Algebraic Ricatti Equation (DARE)

How do I solve? Can iterate over V / use eigen value decomposition [1]

Type into MATLAB: `dare(A,B,Q,R)`

[1] Arnold, W.F., III and A.J. Laub, "Generalized Eigenproblem Algorithms and Software for Algebraic Riccati Equations," *Proc. IEEE*®, 72 (1984), pp. 1746-1754.


So, can this controller
stabilize inverted pendulum for all
angles?

So, can this controller
stabilize inverted pendulum for all
angles?

No!

Linearization error is too large when angle is large

Instead, can we use LQR to track reference trajectory?



Yes

But, first we need to talk
about time-varying systems

LQR for Time-Varying Dynamical Systems

$$x_{t+1} = A_t x_t + B_t u_t$$

LQR for Time-Varying Dynamical Systems

$$x_{t+1} = A_t x_t + B_t u_t$$

$$c(x_t, u_t) = x_t^T Q_t x_t + u_t^T R_t u_t$$

LQR for Time-Varying Dynamical Systems

$$x_{t+1} = A_t x_t + B_t u_t$$

$$c(x_t, u_t) = x_t^T Q_t x_t + u_t^T R_t u_t$$

Straight forward to get LQR equations

$$K_t = -(R_t + B_t^T V_{t+1} B_t)^{-1} B_t^T V_{t+1} A_t$$


$$V_t = Q_t + K_t^T R_t K_t + (A_t + B_t K_t)^T V_{t+1} (A_t + B_t K_t)$$

Why do we care about time-varying?


Ans: Linearization about a trajectory

Linearize about a time-varying trajectory


Linearize about a time-varying trajectory




Linearize about a time-varying trajectory



Linearize about a time-varying trajectory




Linearize about a time-varying trajectory



$$\dot{x} = f(x, u)$$

$$A_t = \left. \frac{\partial f}{\partial x} \right|_{x_{ref}(t)}$$

Linearize about a time-varying trajectory




$$\dot{x} = f(x, u)$$

$$A_t = \left. \frac{\partial f}{\partial x} \right|_{x_{ref}(t)}$$


$$B_t = \left. \frac{\partial f}{\partial u} \right|_{u_{ref}(t)}$$

Linearize about a time-varying trajectory



$$x_{t+1} = A_t x_t + B_t u_t + x_t^{off}$$

Linearize about a time-varying trajectory



$$\dot{x} = f(x, u)$$

$$A_t = \left. \frac{\partial f}{\partial x} \right|_{x_{ref}(t)}$$

$$B_t = \left. \frac{\partial f}{\partial u} \right|_{u_{ref}(t)}$$

$$x_{t+1} = A_t x_t + B_t u_t + x_t^{off}$$

Making an affine system linear

Making an affine system linear

$$x_{t+1} = A_t x_t + B_t u_t + x_t^{off}$$

Making an affine system linear

$$x_{t+1} = A_t x_t + B_t u_t + x_t^{off}$$

Homogenous coordinates

$$\tilde{x} = \begin{pmatrix} x \\ 1 \end{pmatrix}$$

Making an affine system linear

$$x_{t+1} = A_t x_t + B_t u_t + x_t^{off}$$

Homogenous coordinates

$$\tilde{x} = \begin{pmatrix} x \\ 1 \end{pmatrix}$$

$$\tilde{x}_{t+1} = \begin{pmatrix} A_t & x_t^{off} \\ 0 & 1 \end{pmatrix} + \begin{pmatrix} B_t \\ 0 \end{pmatrix} u_t$$

Making an affine system linear

$$x_{t+1} = A_t x_t + B_t u_t + x_t^{off}$$

Homogenous coordinates $\tilde{x} = \begin{pmatrix} x \\ 1 \end{pmatrix}$

$$\tilde{x}_{t+1} = \begin{pmatrix} A_t & x_t^{off} \\ 0 & 1 \end{pmatrix} + \begin{pmatrix} B_t \\ 0 \end{pmatrix} u_t$$

Similarly you can transform cost function


$$c(\tilde{x}_t, u_t) = \tilde{x}_t^T \tilde{Q}_t \tilde{x}_t + u_t^T R_t u_t$$

Shape of the value function changes along trajectory


$$x_{ref}(t) \quad u_{ref}(t)$$




Shape of the value function changes along trajectory



Shape of the value function changes along trajectory



Shape of the value function changes along trajectory



Questions

Questions

1. Can we solve LQR for **continuous** time dynamics?

Questions

1. Can we solve LQR for **continuous** time dynamics?

Yes! Refer to Continuous Algebraic Ricatti Equations (CARE)

Questions

1. Can we solve LQR for **continuous** time dynamics?

Yes! Refer to Continuous Algebraic Ricatti Equations (CARE)

2. Can LQR handle **arbitrary costs** (not just tracking)?

Questions

1. Can we solve LQR for **continuous** time dynamics?

Yes! Refer to Continuous Algebraic Ricatti Equations (CARE)

2. Can LQR handle **arbitrary costs** (not just tracking)?

Yes! We will talk about iterative LQR next class

Questions

1. Can we solve LQR for **continuous** time dynamics?

Yes! Refer to Continuous Algebraic Ricatti Equations (CARE)

2. Can LQR handle **arbitrary costs** (not just tracking)?

Yes! We will talk about iterative LQR next class

3. What if I want to penalize control **derivatives**?

Questions

1. Can we solve LQR for **continuous** time dynamics?

Yes! Refer to Continuous Algebraic Ricatti Equations (CARE)

2. Can LQR handle **arbitrary costs** (not just tracking)?

Yes! We will talk about iterative LQR next class

3. What if I want to penalize control **derivatives**?

No problem! Add control as part of state space

Questions

1. Can we solve LQR for **continuous** time dynamics?

Yes! Refer to Continuous Algebraic Ricatti Equations (CARE)

2. Can LQR handle **arbitrary costs** (not just tracking)?

Yes! We will talk about iterative LQR next class

3. What if I want to penalize control **derivatives**?

No problem! Add control as part of state space

4. Can we handle **noisy** dynamics?

Questions

1. Can we solve LQR for **continuous** time dynamics?

Yes! Refer to Continuous Algebraic Ricatti Equations (CARE)

2. Can LQR handle **arbitrary costs** (not just tracking)?

Yes! We will talk about iterative LQR next class

3. What if I want to penalize control **derivatives**?

No problem! Add control as part of state space

4. Can we handle **noisy** dynamics?

Yes! Gaussian noise does not change the answer

Trivia: Duality between control and estimation

R. Kalman “A new approach to linear filtering and prediction problems.” (1960)

**linear-quadratic
regulator**

**Kalman-Bucy
filter**

V

Σ

(state variance)

A

A^\top

(dynamics)

B

H^\top

(measurement)

R

DD^\top

(dynamics noise)

Q

CC^\top

(motion noise)

t

$t_f - t$

(Table from E.Todorov “General duality between optimal control and estimation”, CDC, 2008)