

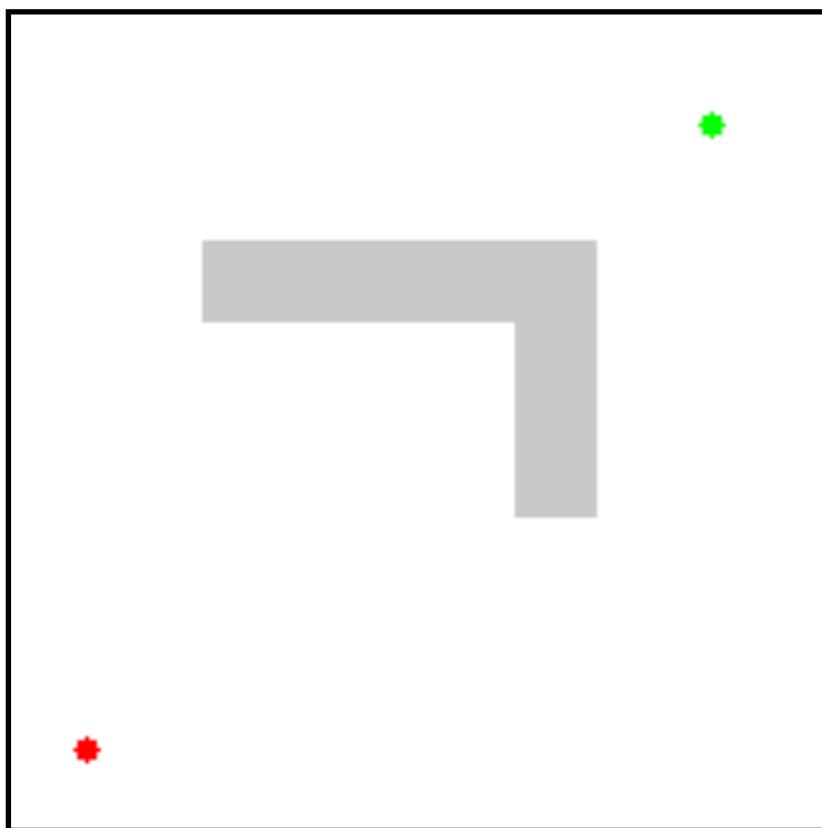
# Lazy Search

Sanjiban Choudhury

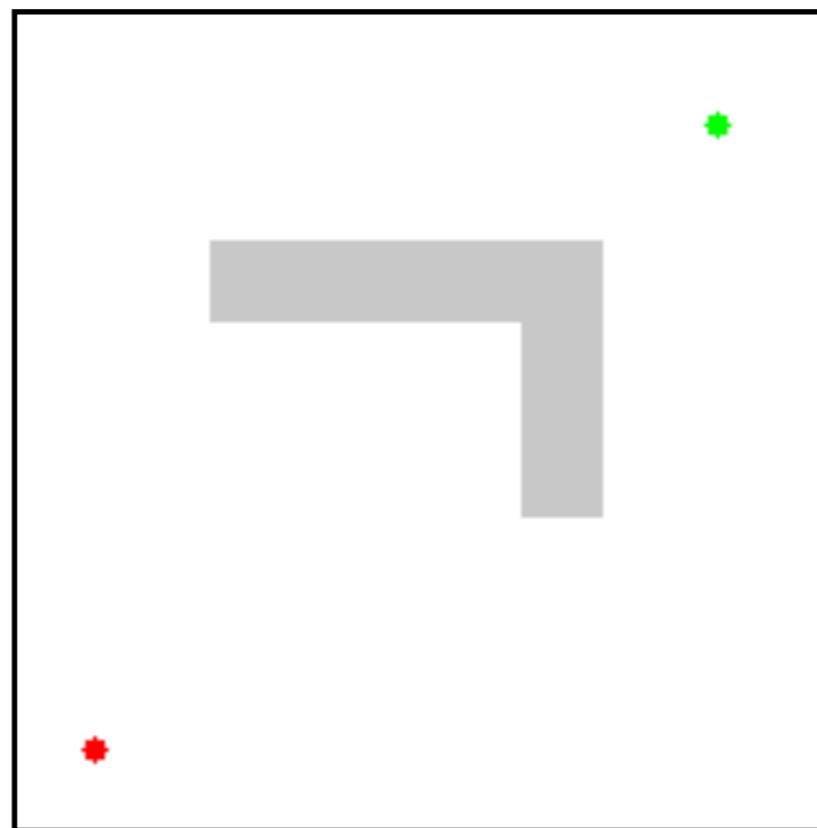
TAs: Matthew Rockett, Gilwoo Lee, Matt Schmittle

# High-order bit

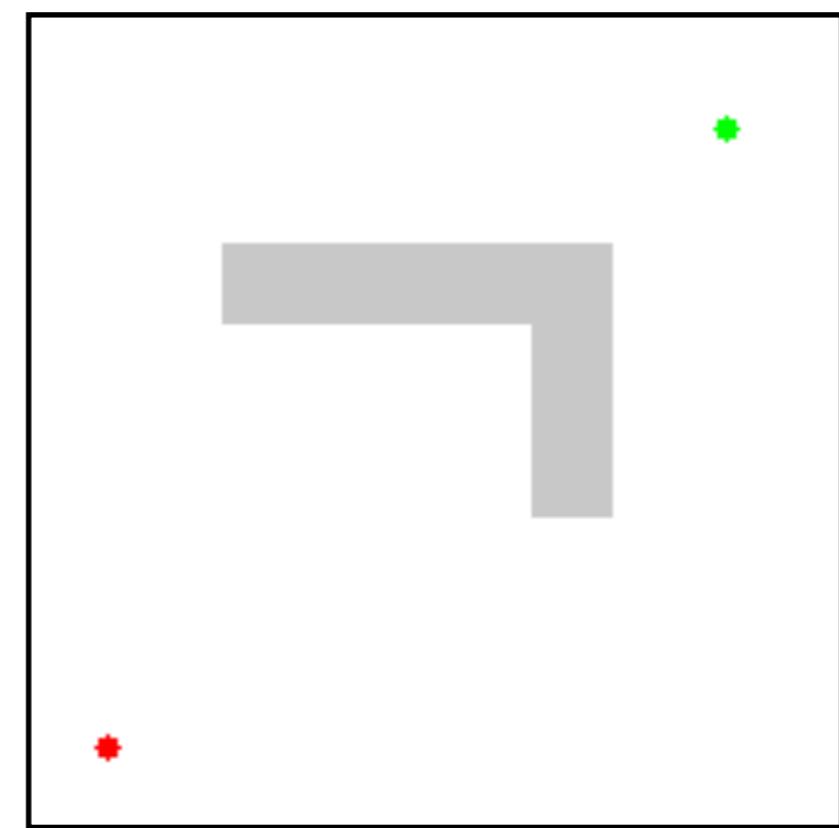
Expansion of a search wavefront from start to goal



Dijkstra



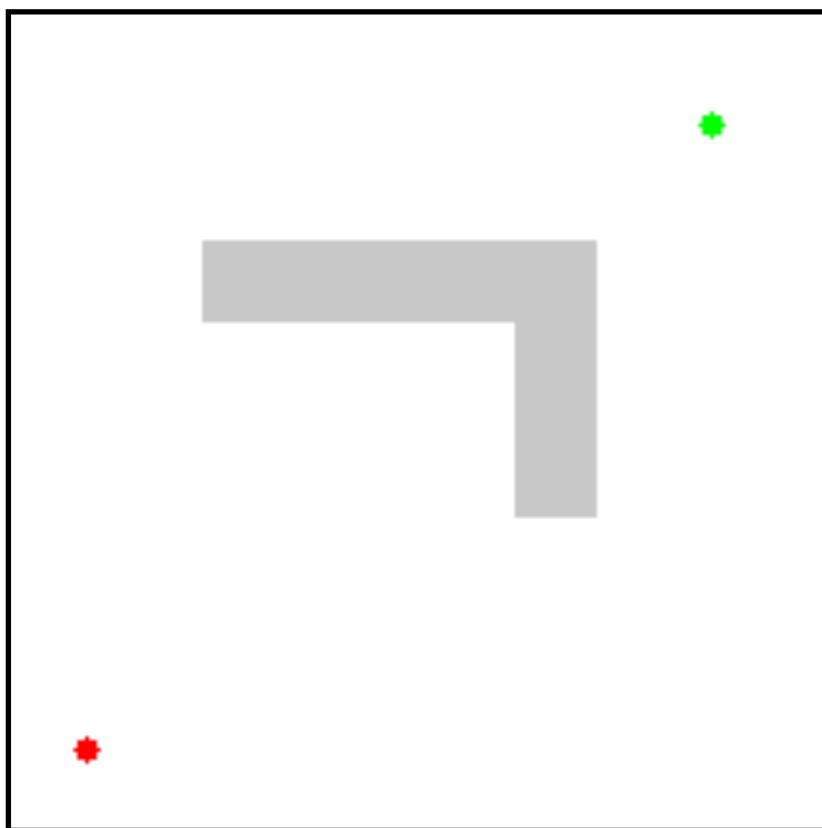
$A^*$



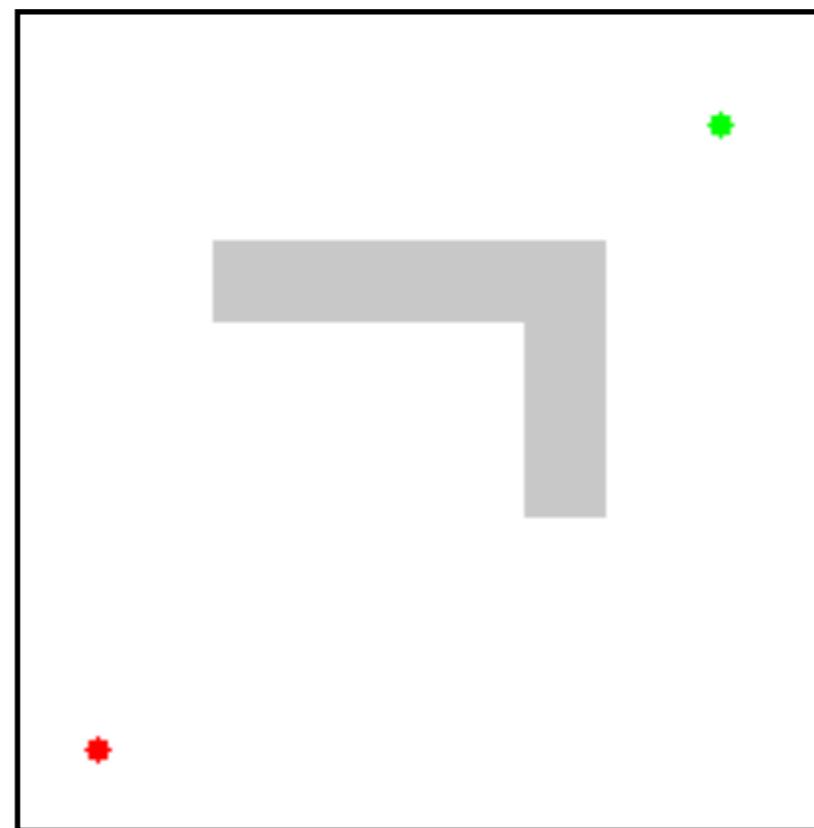
Weighted  $A^*$

# High-order bit

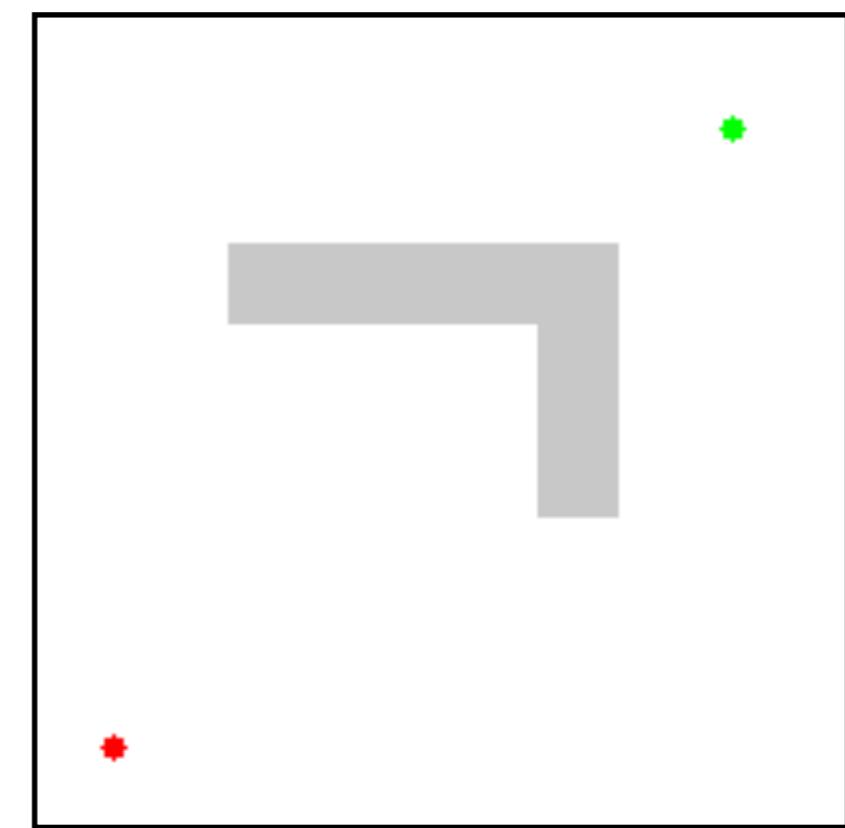
Expansion of a search wavefront from start to goal



Dijkstra



$A^*$



Weighted  $A^*$

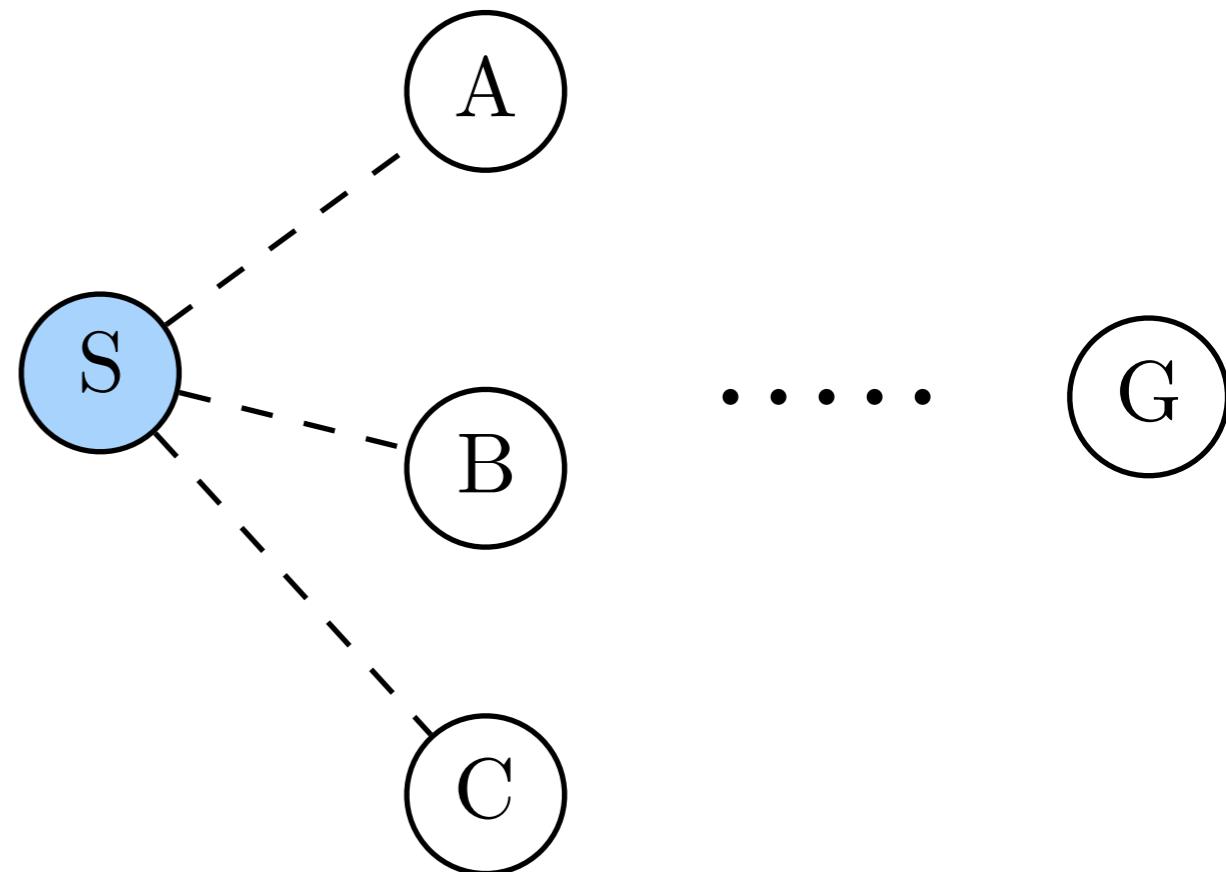
# What do we want?

1. Search to systematically reason over the space of paths
2. Find a (near)-optimal path quickly  
(minimize planning effort)

# Best First Search

(Explore the graph by expanding/processing promising nodes)

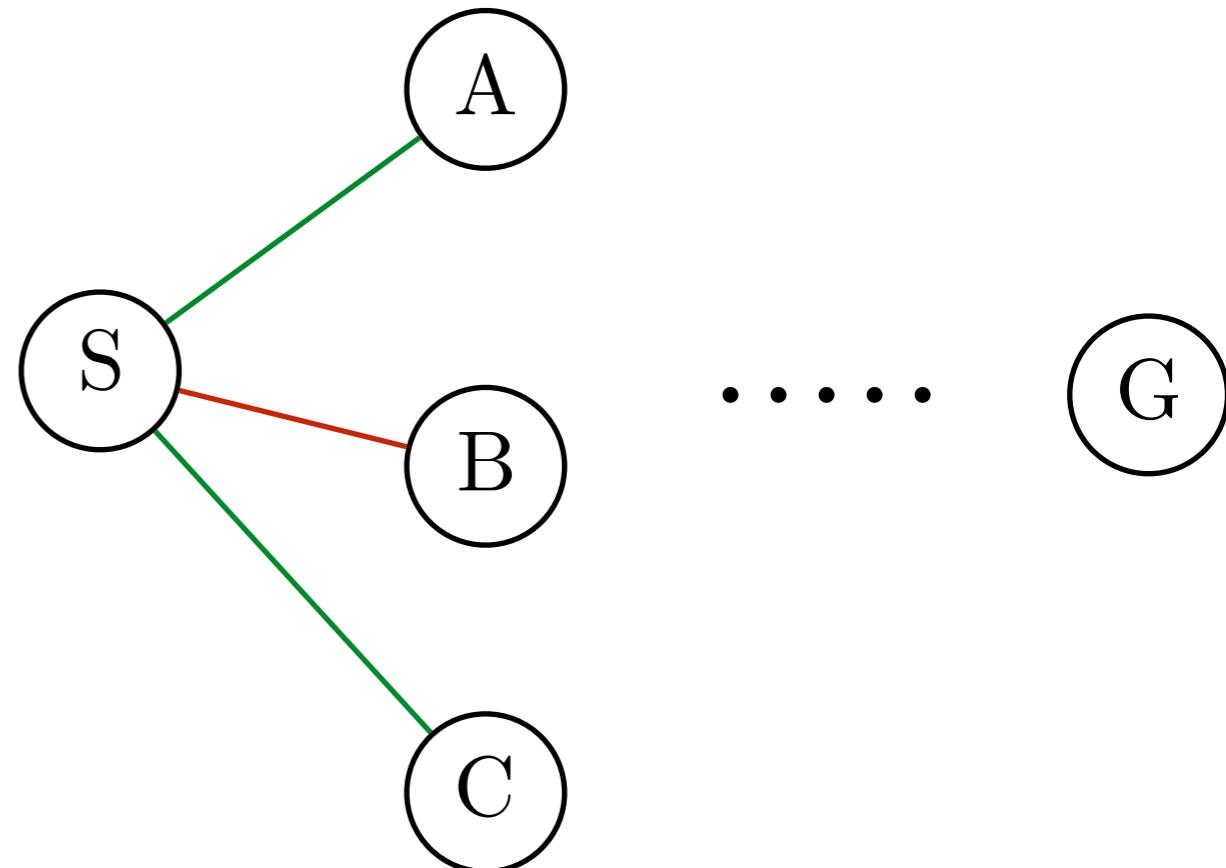
Element (Node)	Priority Value (f-value)
Node S	$f(S)$



# Best First Search

(Explore the graph by expanding/processing promising nodes)

Element (Node)	Priority Value (f-value)
Node S	$f(S)$
Node A	$f(A)$
Node C	$f(C)$



# Different $f(s)$ leads to different algorithms

Algorithm	$f(s)$	Optimality	Efficiency
Dijkstra	$f(s) = g(s)$	Yes	Poor
A*	$f(s) = g(s) + h(s)$	Yes* (if $h(s)$ is admissible)	Good* (better if $h(s)$ is consistent)

# What is a heuristic $h(s)$ ?

So far, we have been thinking of heuristics as an  
estimate of cost-to-go to goal

# What is a heuristic $h(s)$ ?

So far, we have been thinking of heuristics as an  
estimate of cost-to-go to goal

Goal is to come up with cheap-to-compute estimates

# What is a heuristic $h(s)$ ?

So far, we have been thinking of heuristics as an  
**estimate of cost-to-go** to goal

Goal is to come up with **cheap-to-compute** estimates

- External knowledge about metrics
  - Euclidean distance admissible because triangle inequality

# What is a heuristic $h(s)$ ?

So far, we have been thinking of heuristics as an  
**estimate of cost-to-go** to goal

Goal is to come up with **cheap-to-compute** estimates

- External knowledge about metrics
  - Euclidean distance admissible because triangle inequality
- Solving a relaxation of the problem (which is easier)
  - Ignore obstacles, ignore dynamics

# What is a heuristic $h(s)$ ?

So far, we have been thinking of heuristics as an  
**estimate of cost-to-go** to goal

Goal is to come up with **cheap-to-compute** estimates

- External knowledge about metrics
  - Euclidean distance admissible because triangle inequality
- Solving a relaxation of the problem (which is easier)
  - Ignore obstacles, ignore dynamics
- Statistical knowledge of the problem
  - Keep arms tucked in, don't drive headfirst into a wall

# What is a heuristic $h(s)$ ?

So far, we have been thinking of heuristics as an  
estimate of cost-to-go to goal

# What is a heuristic $h(s)$ ?

So far, we have been thinking of heuristics as an  
estimate of cost-to-go to goal

But another way to think of heuristics is as a ranking function

# What is a heuristic $h(s)$ ?

So far, we have been thinking of heuristics as an  
estimate of cost-to-go to goal

But another way to think of heuristics is as a ranking function

Even if estimates are off, as long as a heuristic is ranking  
good states better than bad states,  
it's useful

# The case for inadmissible search

# The case for inadmissible search

Let's say you have a  $h(s)$  that is very very admissible

# The case for inadmissible search

Let's say you have a  $h(s)$  that is very very admissible

But it is actually good at ranking good vs bad states

# The case for inadmissible search

Let's say you have a  $h(s)$  that is very very admissible

But it is actually good at ranking good vs bad states

$f(s) = g(s) + h(s)$  might not be best at ranking states. Why?

# The case for inadmissible search

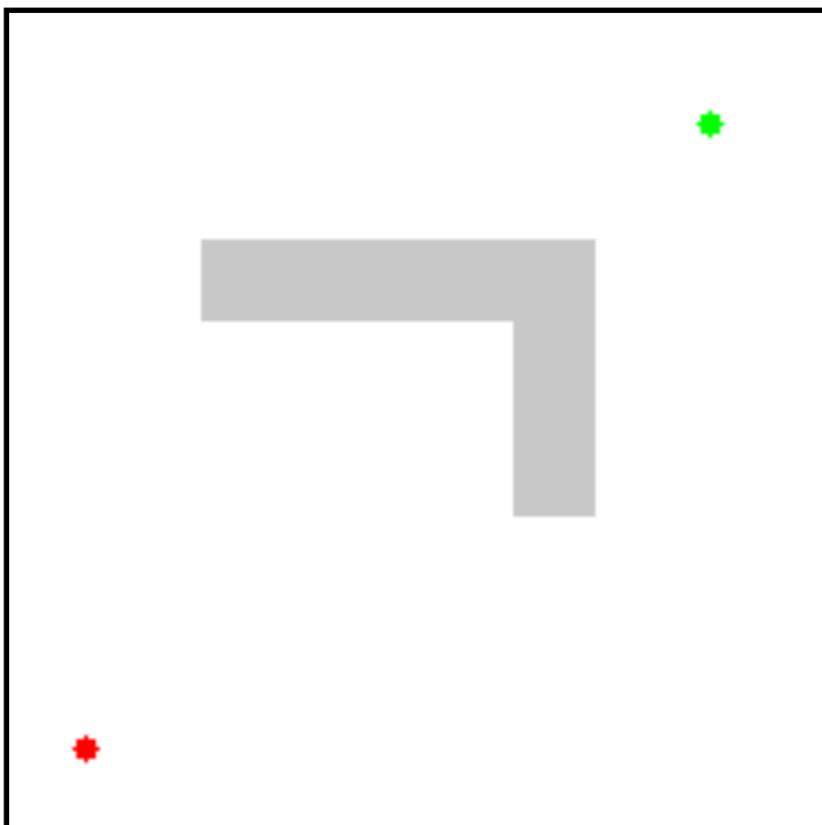
Let's say you have a  $h(s)$  that is very very admissible

But it is actually good at ranking good vs bad states

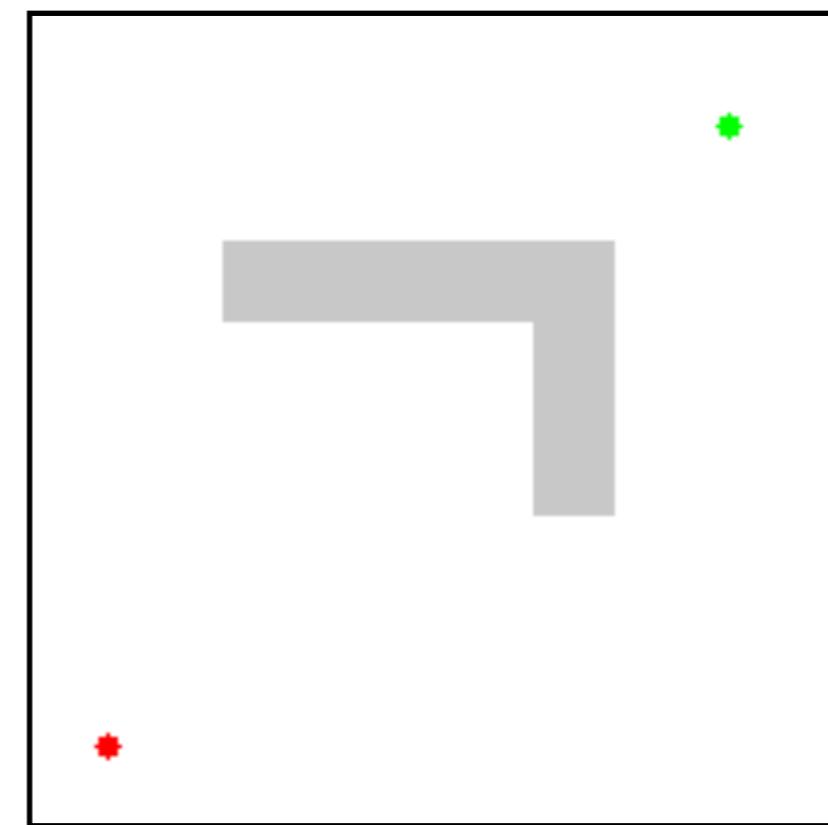
$f(s) = g(s) + h(s)$  might not be best at ranking states. Why?

$f(s) = g(s) + 10h(s)$  might be much better. Why?

# Side-effects of inflating heuristics

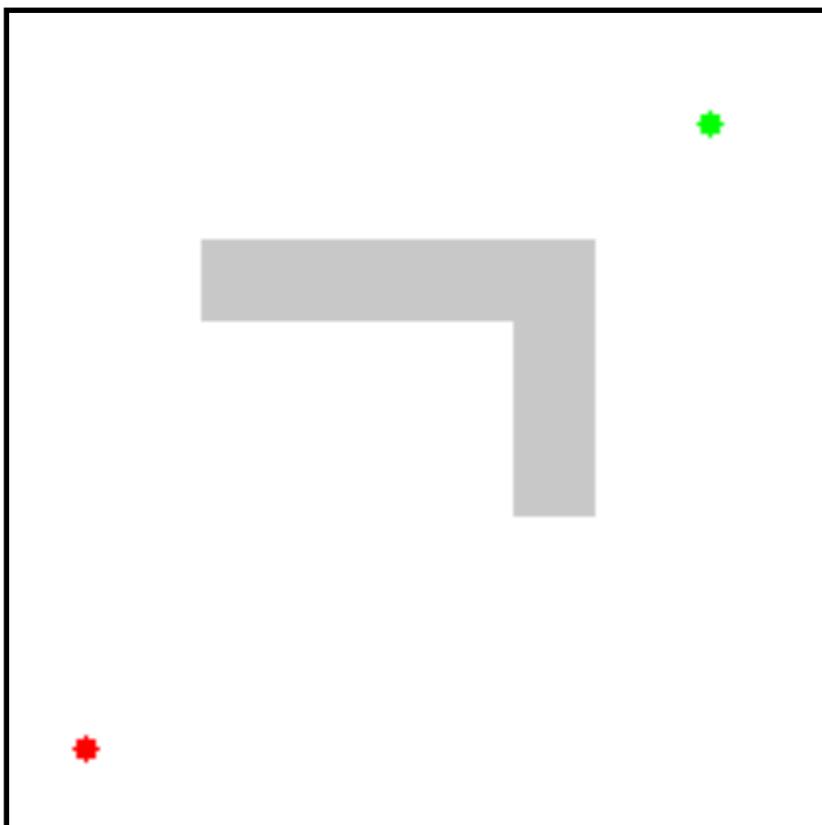


$$f(s) = g(s) + h(s)$$

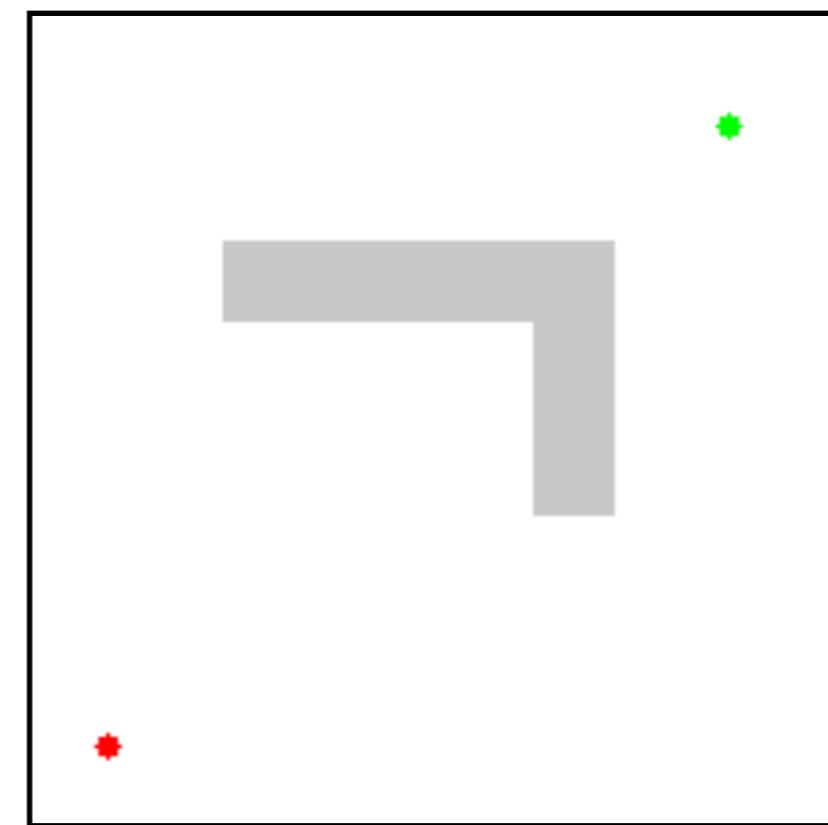


$$f(s) = g(s) + 5h(s)$$

# Side-effects of inflating heuristics



$$f(s) = g(s) + h(s)$$



$$f(s) = g(s) + 5h(s)$$

Can we bound the solution quality?

$$\epsilon \geq 1$$

Can we bound the solution quality?

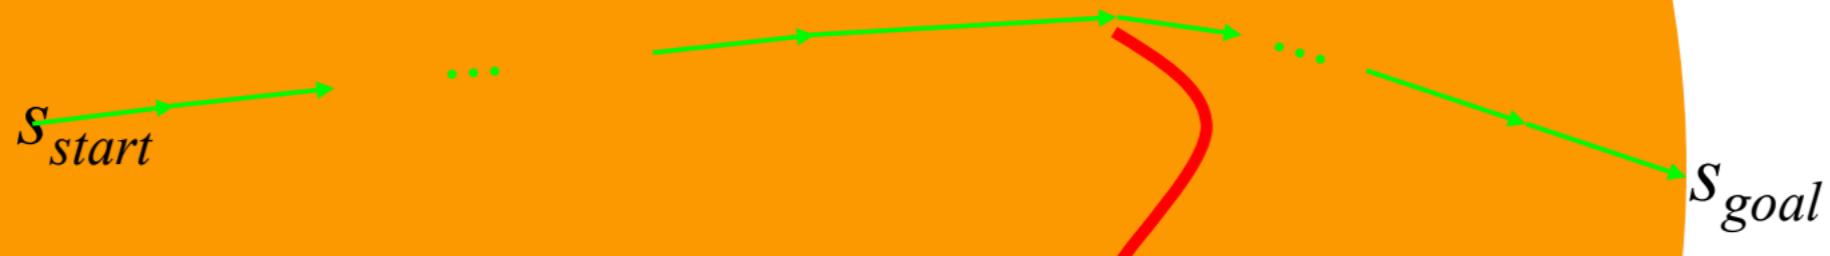
$$g(s) + \epsilon h(s) \quad \epsilon \geq 1$$

results in

$$\text{cost(solution)} \leq \epsilon \text{ cost(solution)}$$

- Dijkstra's: expands states in the order of  $f = g$  values

*What are the states expanded?*



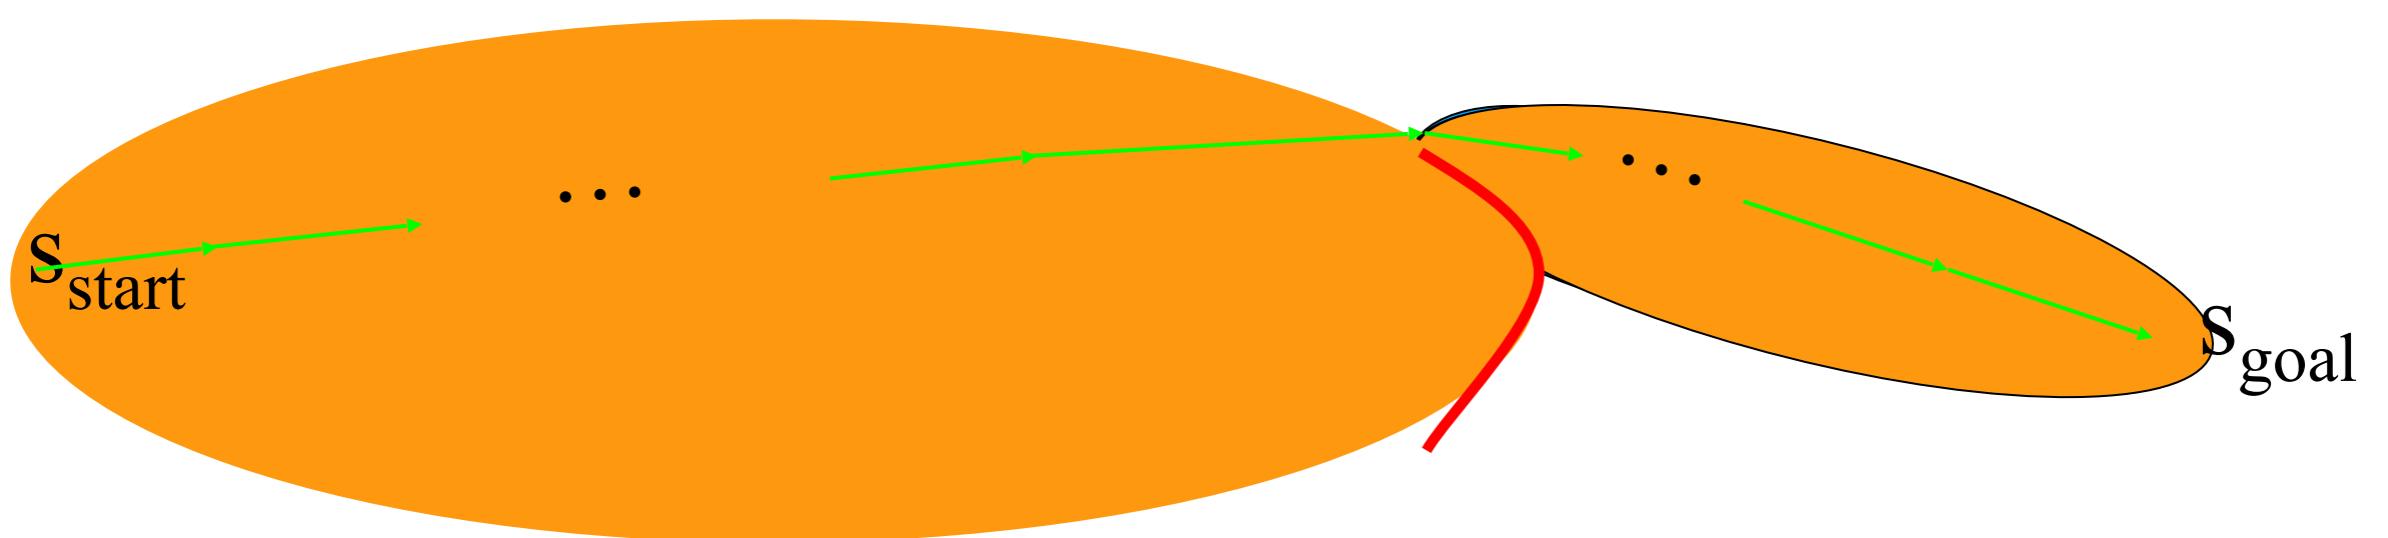
# Effect of the Heuristic Function

A\* Search: expands states in the order of  $f = g+h$  values



# Effect of the Heuristic Function

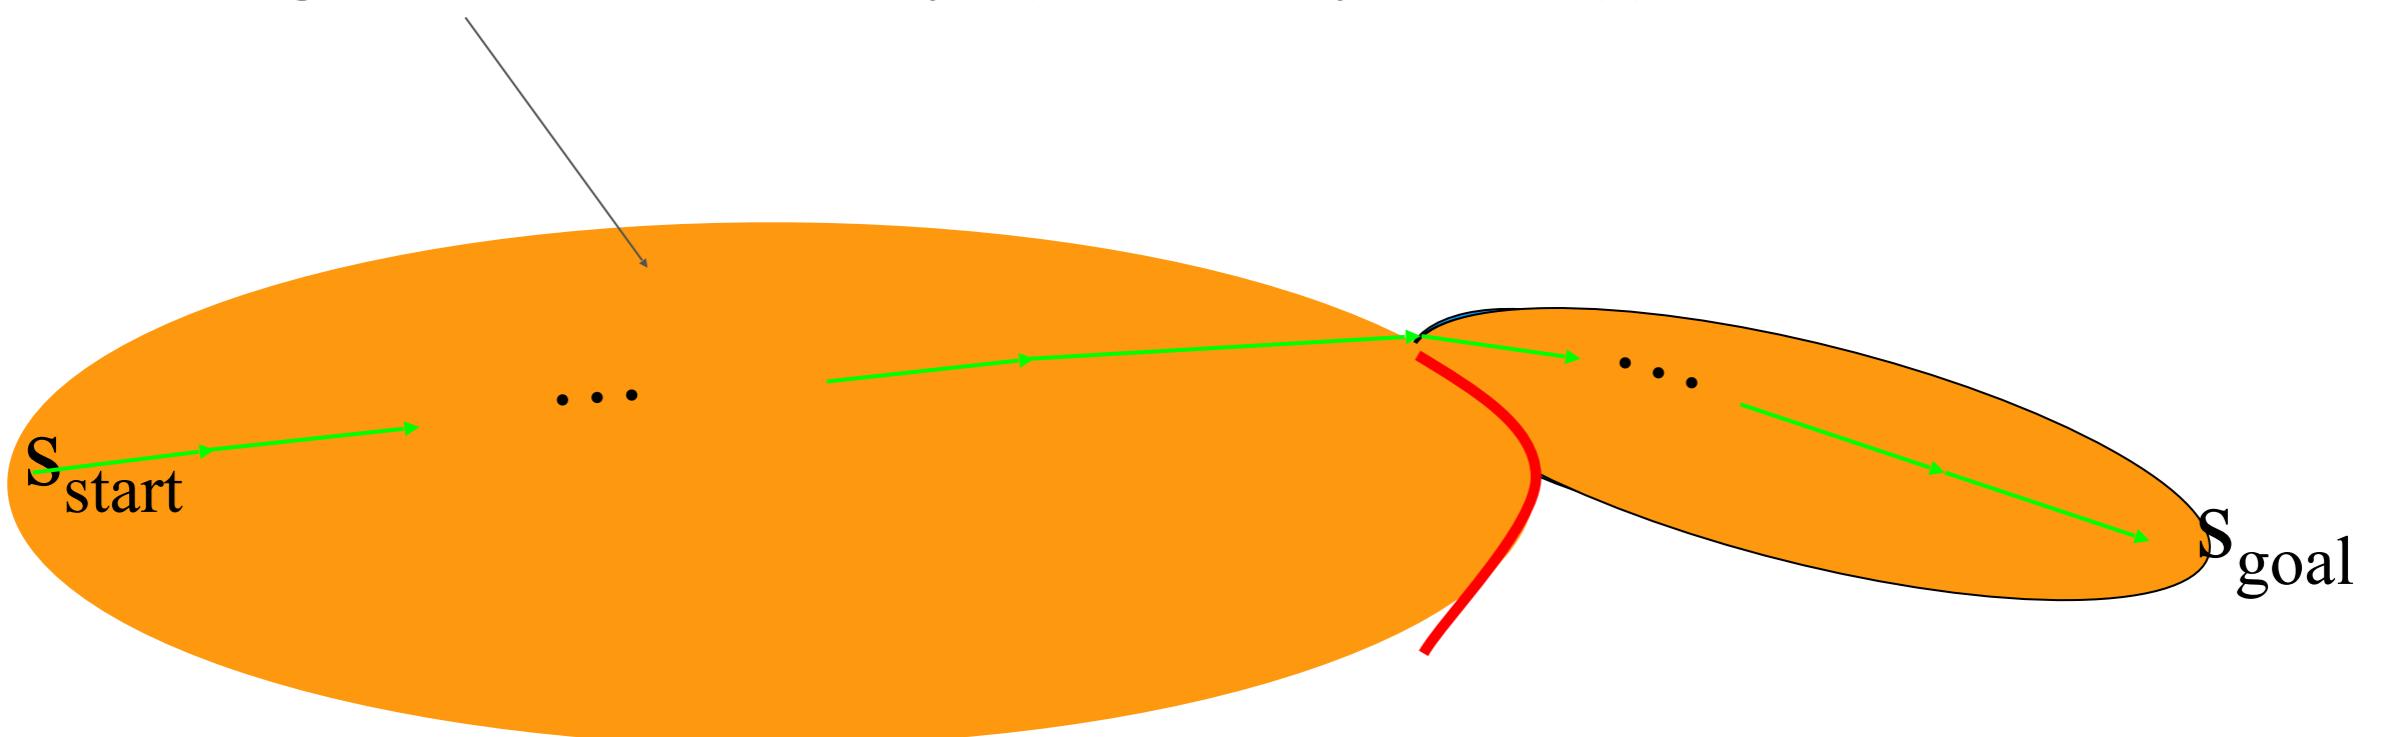
A\* Search: expands states in the order of  $f = g+h$  values



# Effect of the Heuristic Function

A\* Search: expands states in the order of  $f = g+h$  values

for large problems this results in A\* quickly  
running out of memory (memory:  $O(n)$ )



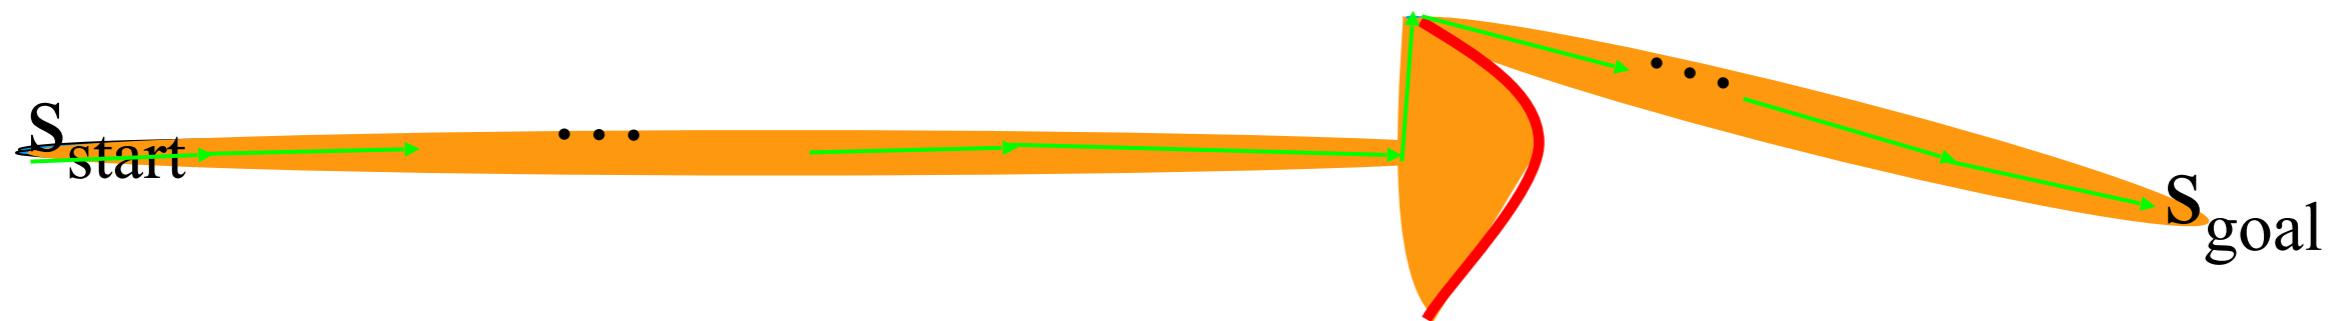
# Effect of the Heuristic Function

Weighted A\* Search: expands states in the order of  $f = g + \varepsilon h$  values,  $\varepsilon > 1$  = bias towards states that are closer to goal



# Effect of the Heuristic Function

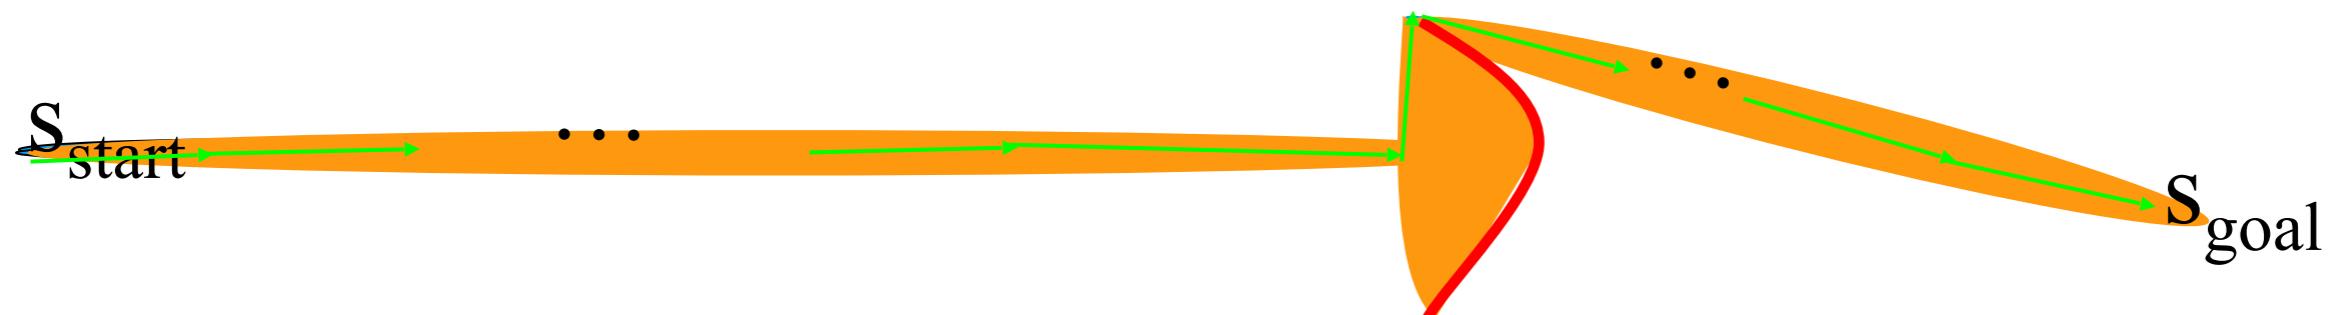
Weighted A\* Search: expands states in the order of  $f = g + \varepsilon h$  values,  $\varepsilon > 1$  = bias towards states that are closer to goal



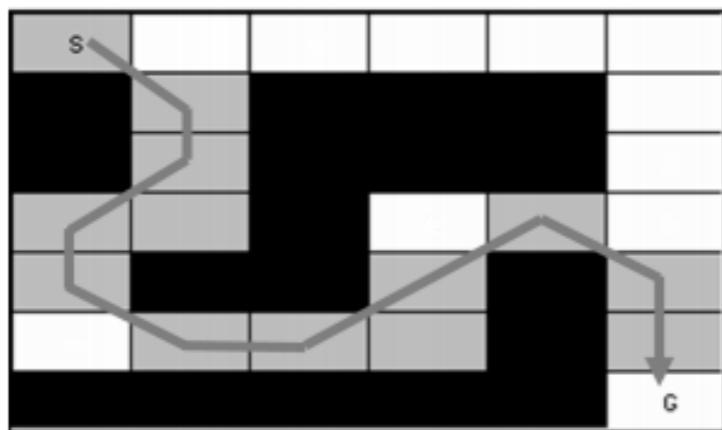
# Effect of the Heuristic Function

Weighted A\* Search: expands states in the order of  $f = g + \varepsilon h$  values,  $\varepsilon > 1$  = bias towards states that are closer to goal

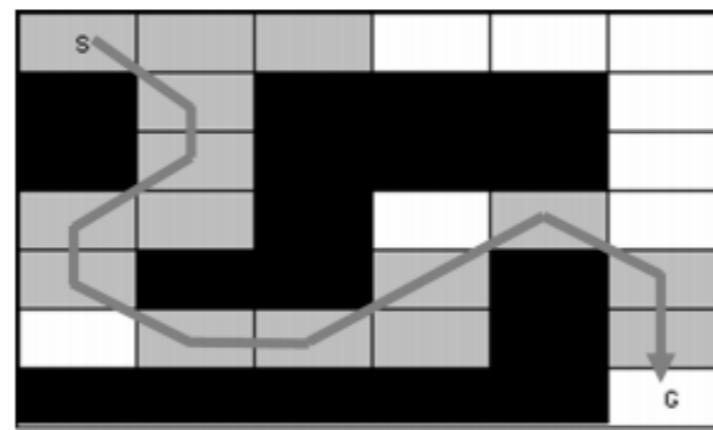
solution is always  $\varepsilon$ -suboptimal:  
 $\text{cost}(\text{solution}) \leq \varepsilon \cdot \text{cost}(\text{optimal solution})$



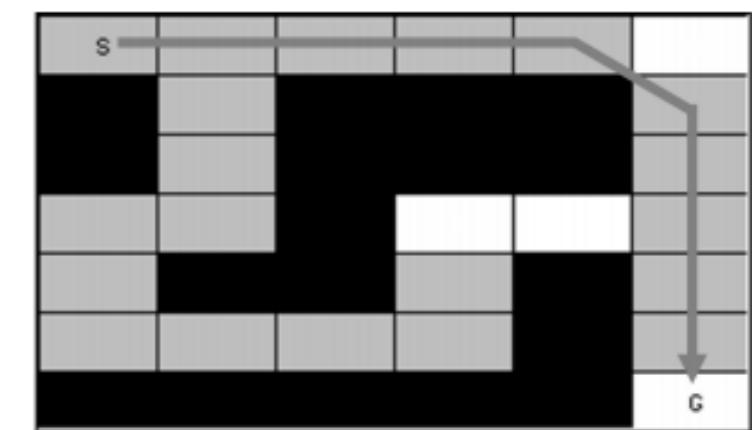
# Effect of the Heuristic Function



$\epsilon = 2.5$



$\epsilon = 1.5$



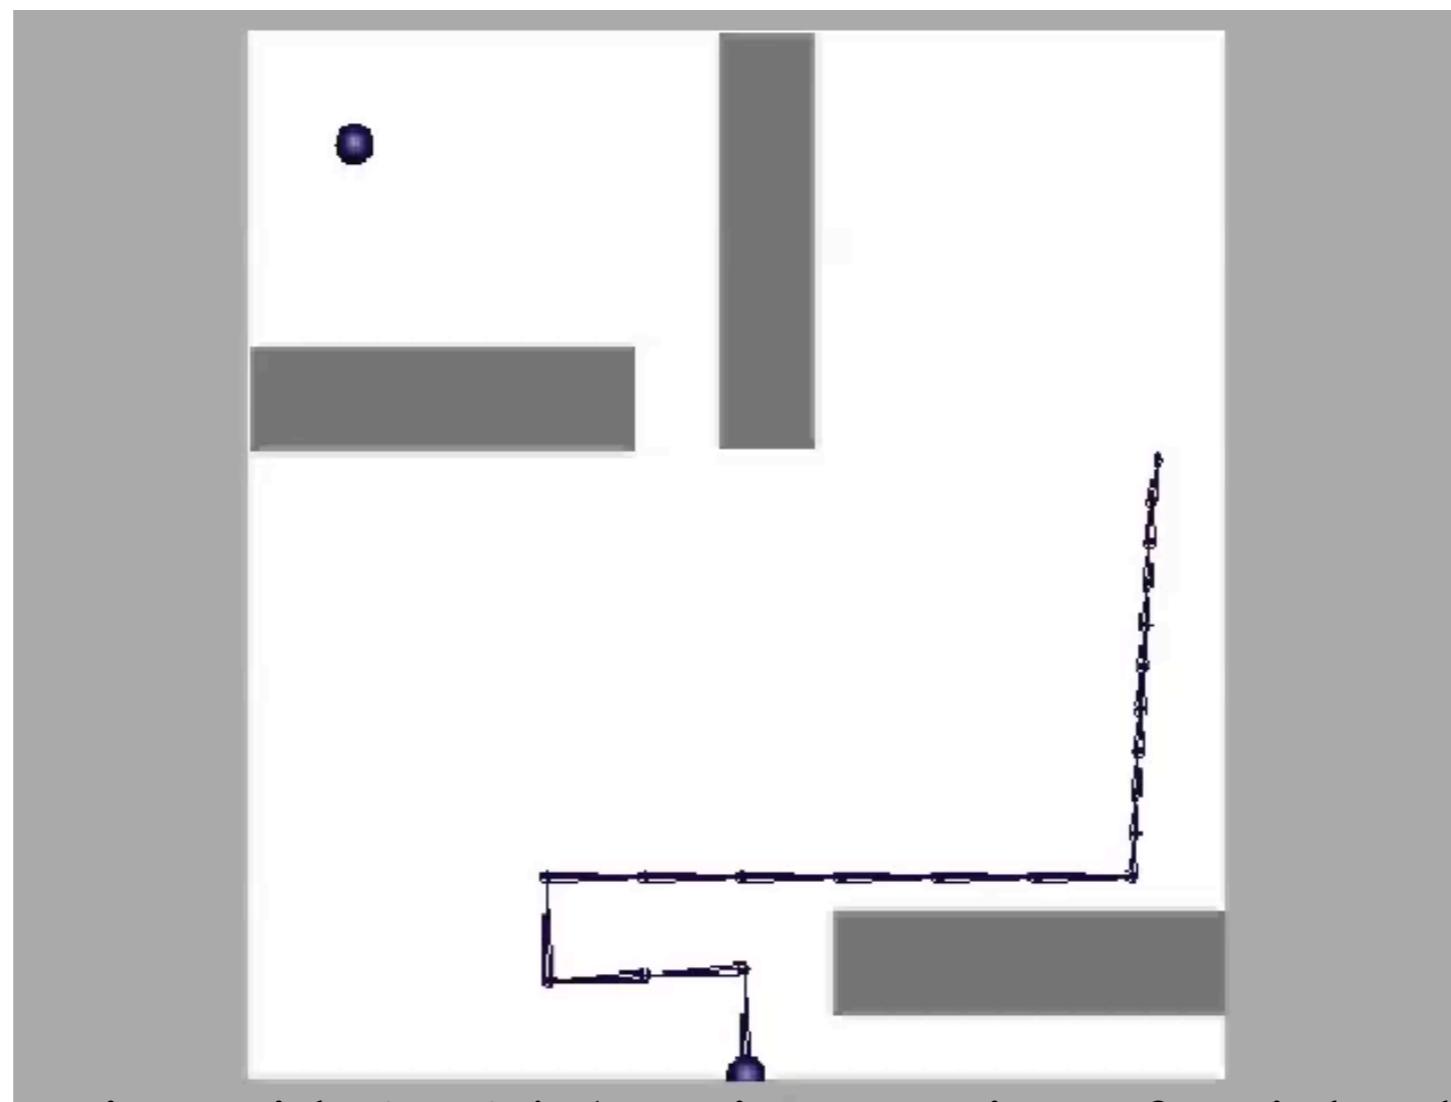
$\epsilon = 1.0$  (optimal search)

Courtesy Max Likhachev

# Effect of the Heuristic Function

Weighted A\* Search: expands states in the order of  $f = g + \varepsilon h$  values,  $\varepsilon > 1$  = bias towards states that are closer to goal

20DOF simulated robotic arm  
state-space size: over  $10^{26}$  states



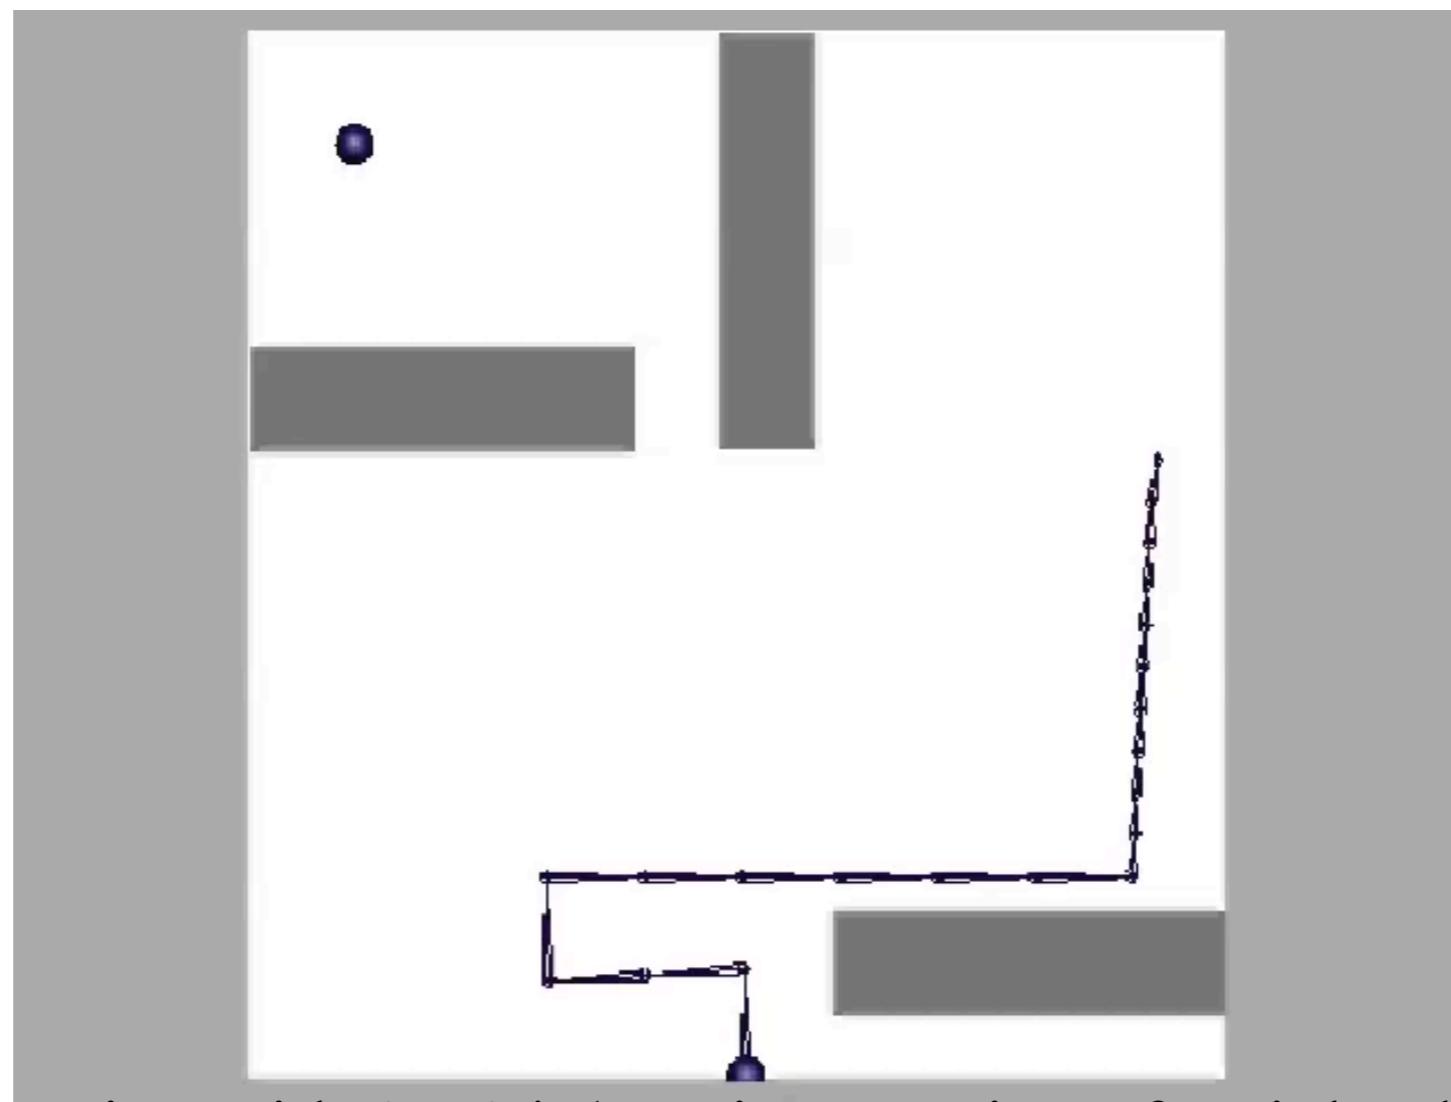
planning with ARA\* (anytime version of weighted A\*)

Courtesy Max Likhachev

# Effect of the Heuristic Function

Weighted A\* Search: expands states in the order of  $f = g + \varepsilon h$   
values,  $\varepsilon > 1$  = bias towards states that are closer to goal

20DOF simulated robotic arm  
state-space size: over  $10^{26}$  states



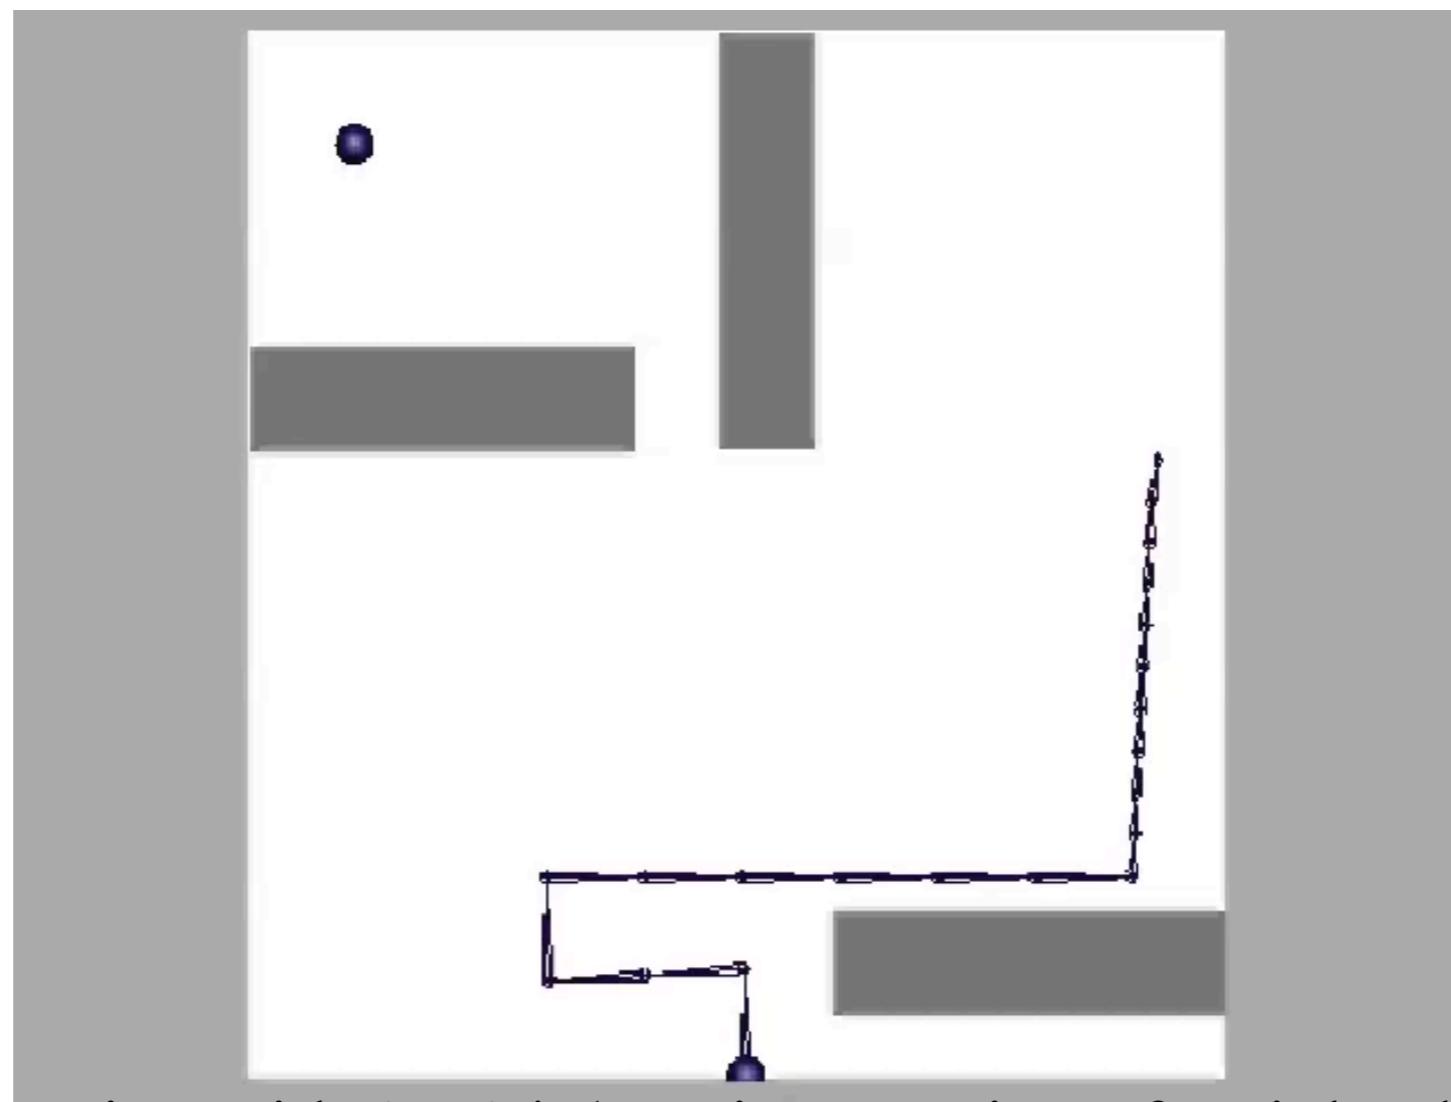
planning with ARA\* (anytime version of weighted A\*)

Courtesy Max Likhachev

# Effect of the Heuristic Function

Weighted A\* Search: expands states in the order of  $f = g + \varepsilon h$   
values,  $\varepsilon > 1$  = bias towards states that are closer to goal

20DOF simulated robotic arm  
state-space size: over  $10^{26}$  states



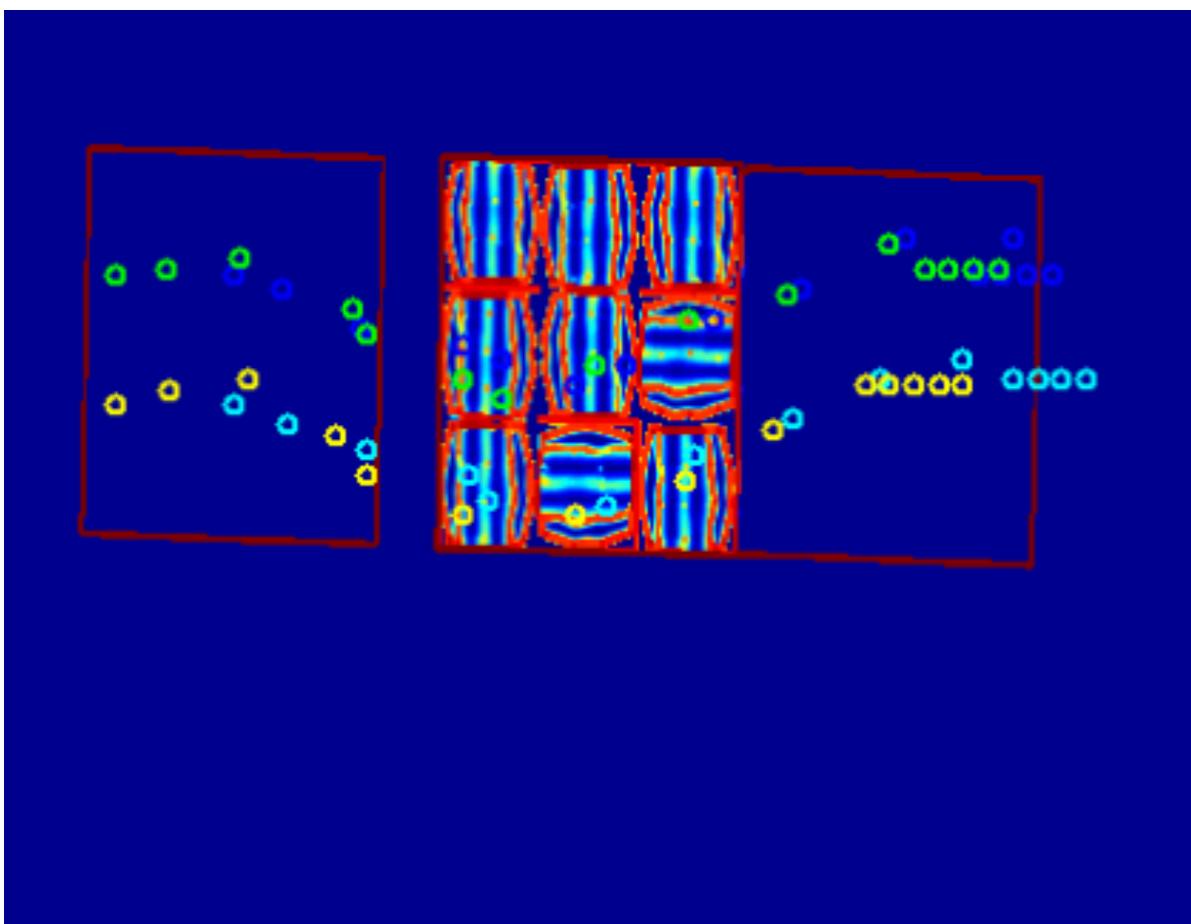
planning with ARA\* (anytime version of weighted A\*)

Courtesy Max Likhachev

# Effect of the Heuristic Function

Courtesy Max Likhachev

- planning in 8D ( $x, y$  for each foothold)
- heuristic is Euclidean distance from the center of the body to the goal location
- cost of edges based on kinematic stability of the robot and quality of footholds



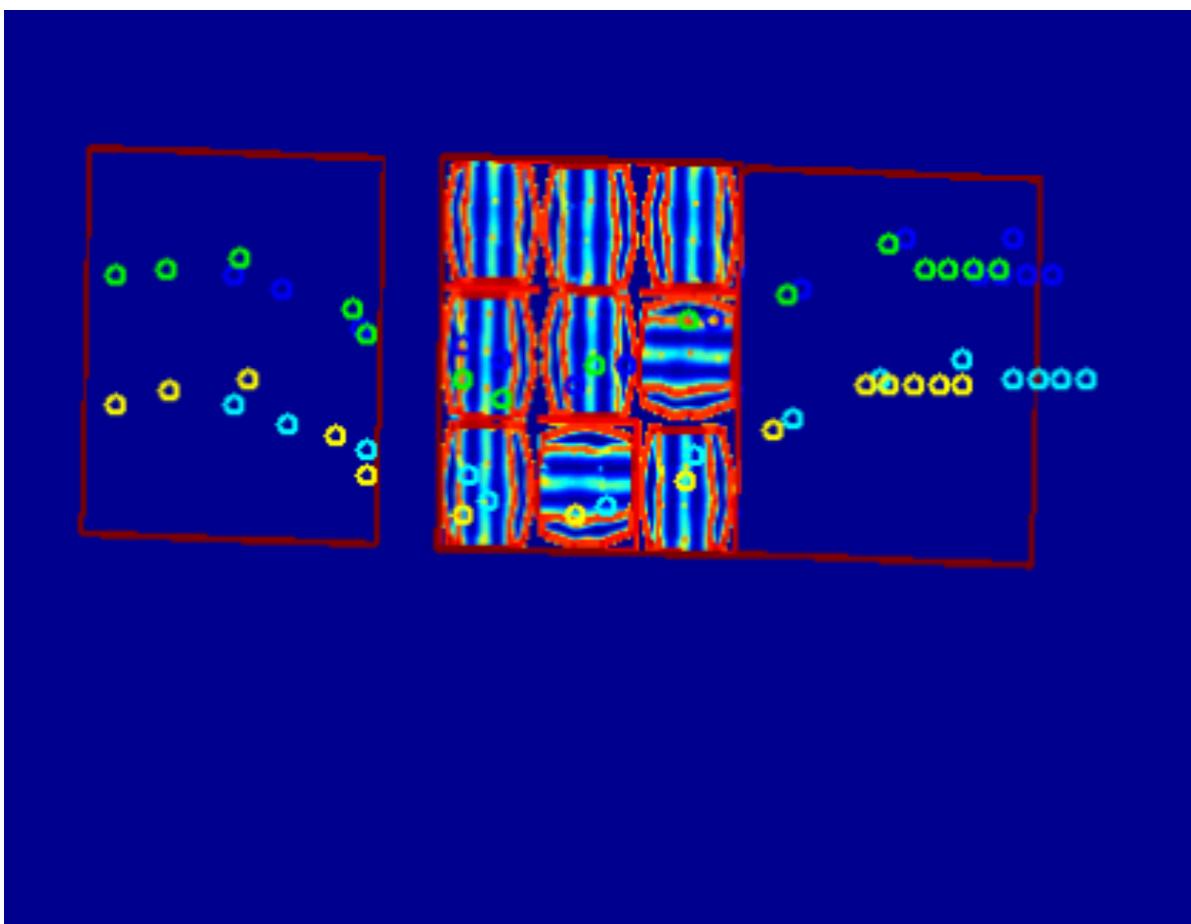
Uses R\* - A randomized version of weighted A\*

Joint work between Max Likhachev, Subhrajit Bhattacharya, Joh Bohren, Sachin Chitta, Daniel D. Lee, Aleksandr Kushleyev, and Paul Vernaza

# Effect of the Heuristic Function

Courtesy Max Likhachev

- planning in 8D ( $x, y$  for each foothold)
- heuristic is Euclidean distance from the center of the body to the goal location
- cost of edges based on kinematic stability of the robot and quality of footholds



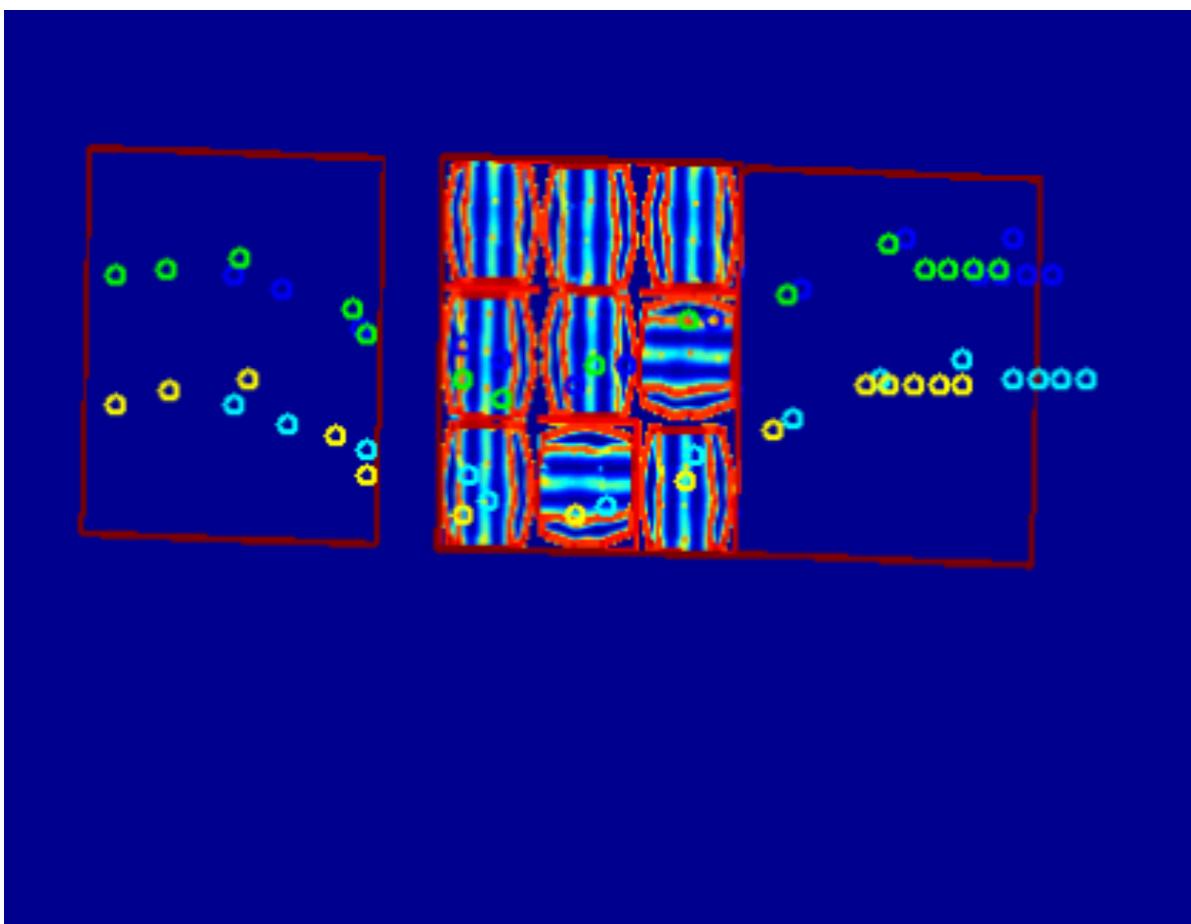
Uses R\* - A randomized version of weighted A\*

Joint work between Max Likhachev, Subhrajit Bhattacharya, Joh Bohren, Sachin Chitta, Daniel D. Lee, Aleksandr Kushleyev, and Paul Vernaza

# Effect of the Heuristic Function

Courtesy Max Likhachev

- planning in 8D ( $x, y$  for each foothold)
- heuristic is Euclidean distance from the center of the body to the goal location
- cost of edges based on kinematic stability of the robot and quality of footholds



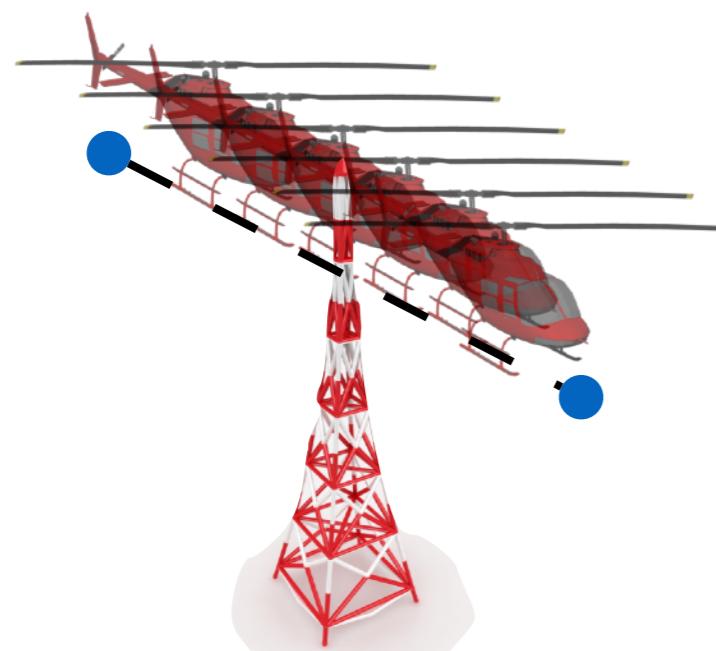
Uses R\* - A randomized version of weighted A\*

Joint work between Max Likhachev, Subhrajit Bhattacharya, Joh Bohren, Sachin Chitta, Daniel D. Lee, Aleksandr Kushleyev, and Paul Vernaza

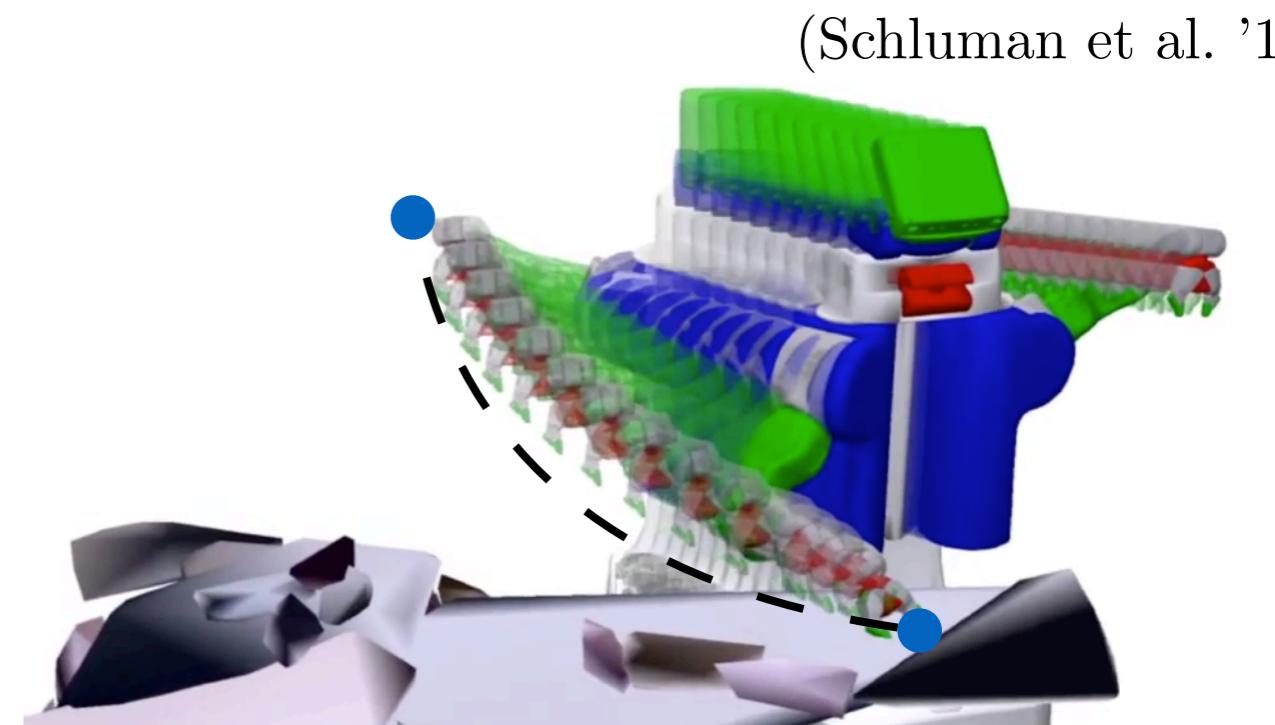
But is the number of expansions  
really what we want to minimize in  
motion planning?

What is the most expensive step?

# Edge evaluation is expensive

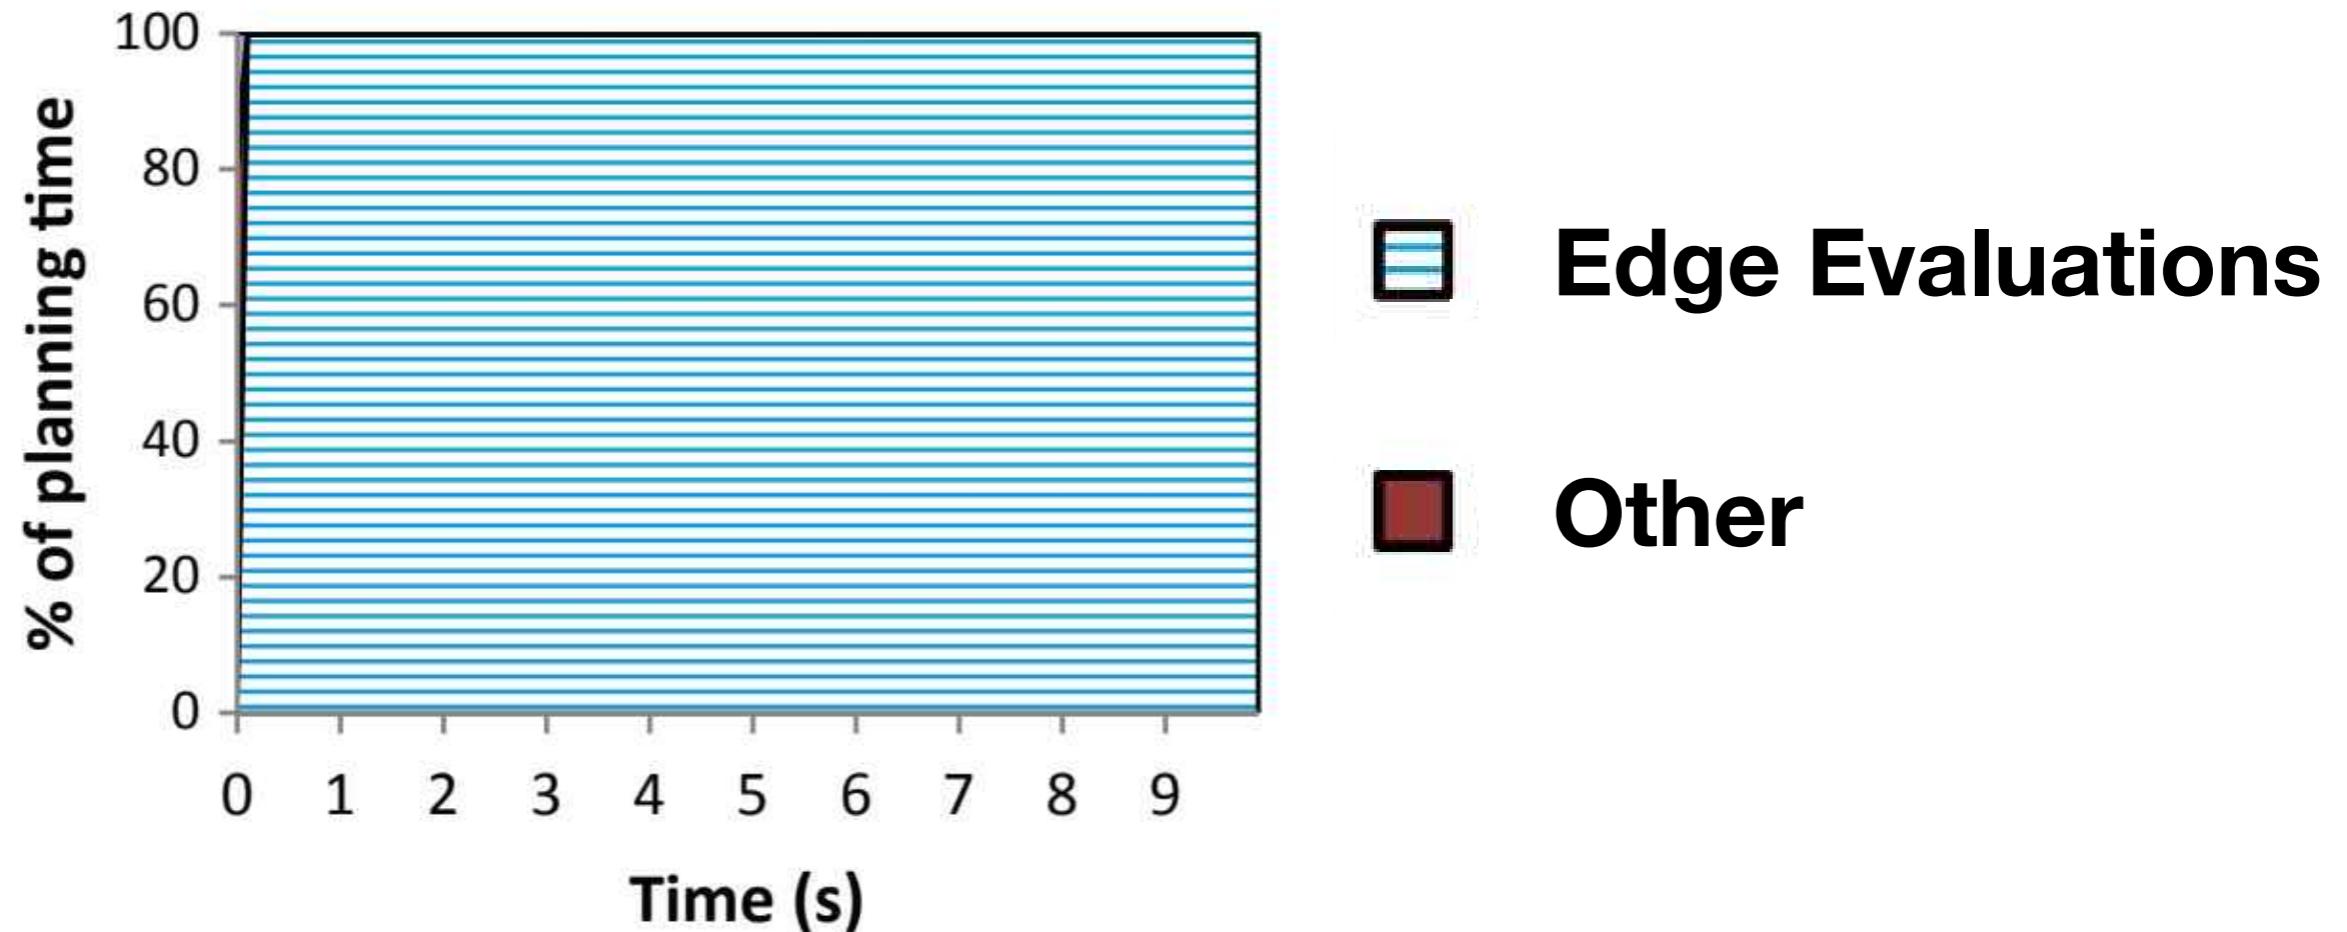


Check if helicopter  
intersects with tower



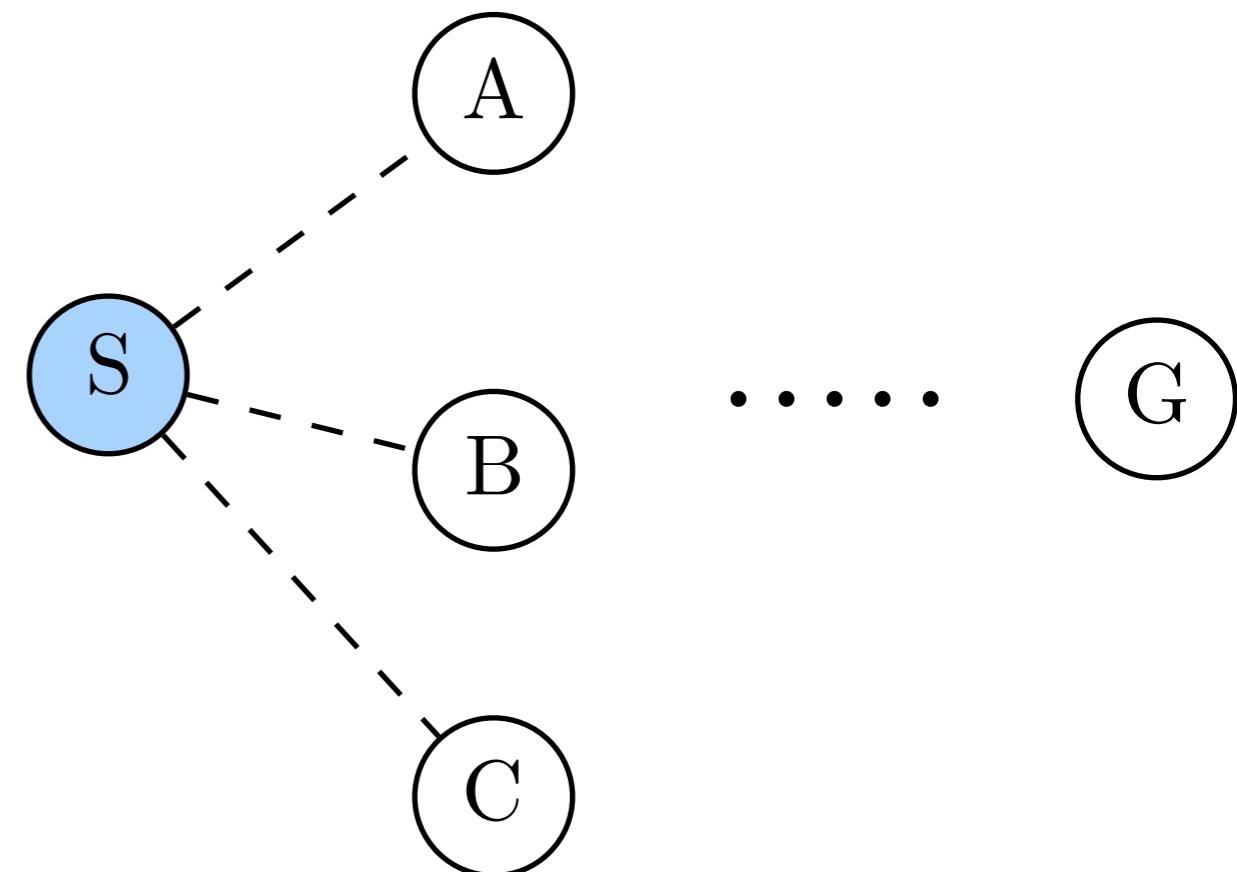
Check if manipulator  
intersects with table

# Edge evaluation dominates planning time



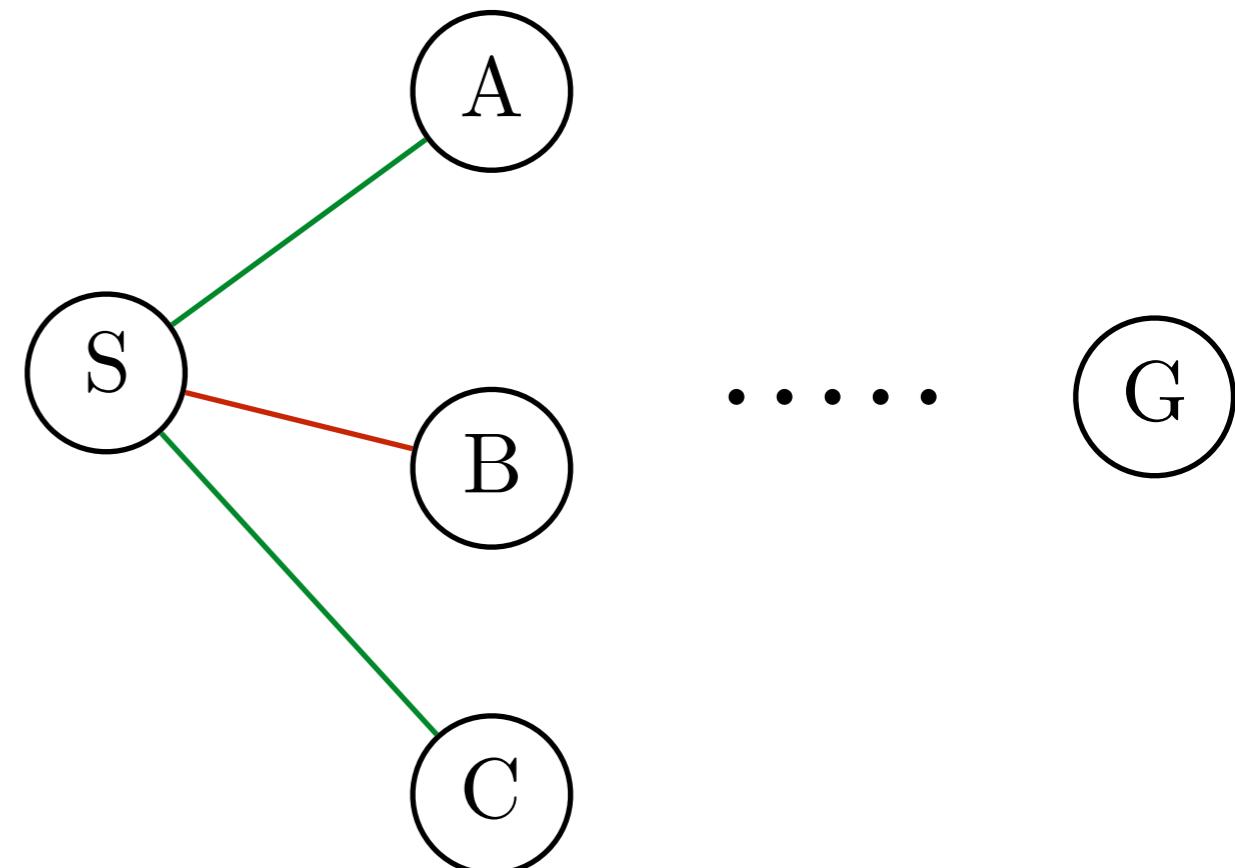
# Let's revisit Best First Search

Element (Node)	Priority Value (f-value)
Node S	$f(S)$



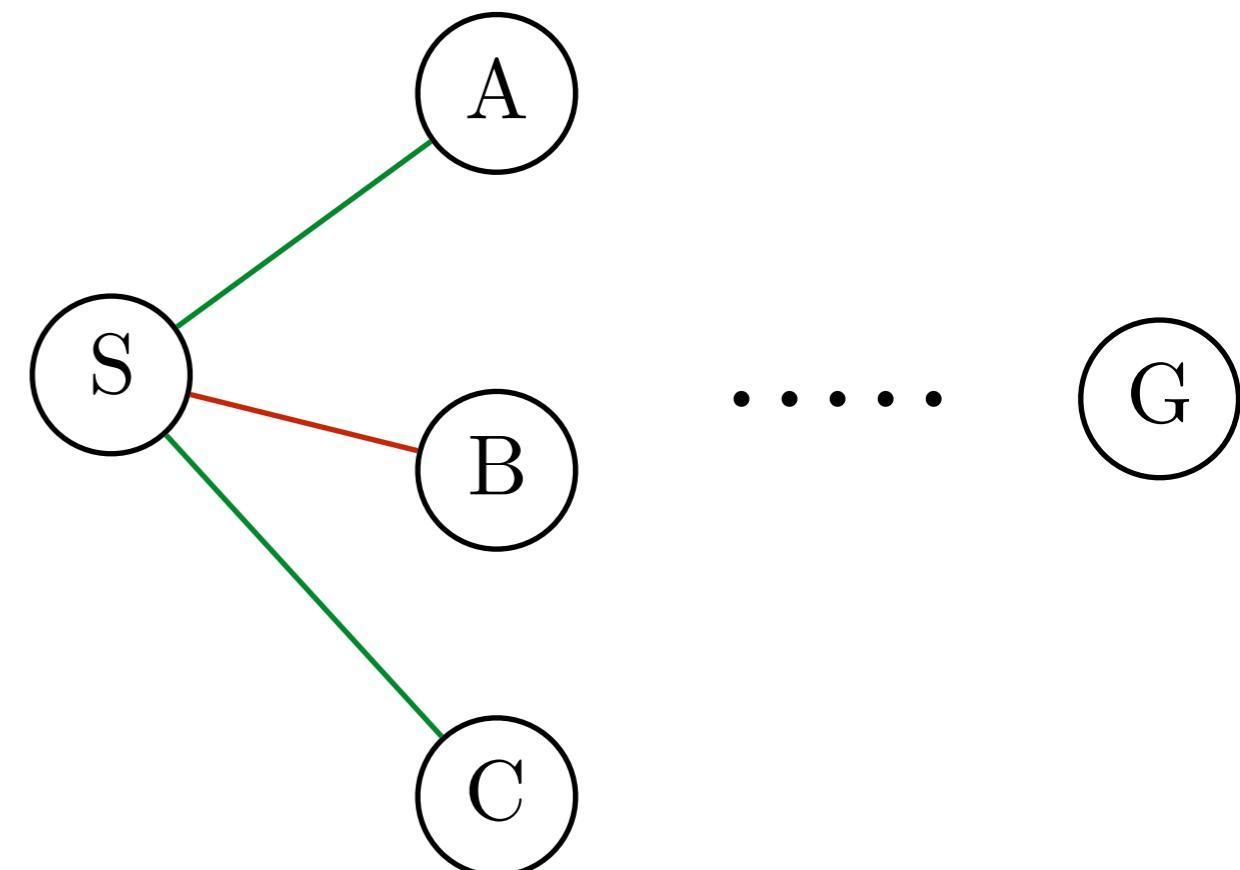
# Let's revisit Best First Search

Element (Node)	Priority Value (f-value)
Node S	$f(S)$
Node A	$f(A)$
Node C	$f(C)$



# What if we never use C? Wasted collision check!

Element (Node)	Priority Value (f-value)
Node S	$f(S)$
Node A	$f(A)$
Node C	$f(C)$



The provable virtue of laziness:

Take the thing that's **expensive**  
(collision checking)

and

**procrastinate** as long as possible  
till you have to evaluate it!

# Lazy (weighted) A\*

Cohen, Phillips, and Likhachev 2014

# Lazy (weighted) A\*

Cohen, Phillips, and Likhachev 2014

Key Idea:

# Lazy (weighted) A\*

Cohen, Phillips, and Likhachev 2014

Key Idea:

1. When expanding a node, **don't collision check** edge to successors  
(be optimistic and assume the edge will be valid)

# Lazy (weighted) A\*

Cohen, Phillips, and Likhachev 2014

## Key Idea:

1. When expanding a node, **don't collision check** edge to successors  
(be optimistic and assume the edge will be valid)
2. When expanding a node, collision check the **edge to parent**  
(expansion means this node is good and worth the effort)

# Lazy (weighted) A\*

Cohen, Phillips, and Likhachev 2014

## Key Idea:

1. When expanding a node, **don't collision check** edge to successors  
(be optimistic and assume the edge will be valid)
2. When expanding a node, collision check the **edge to parent**  
(expansion means this node is good and worth the effort)
3. Important: OPEN list will have **multiple copies** of a node  
(multiple candidate parents since we haven't collision check)

# Lazy A\*

Cohen, Phillips, and Likhachev 2014

## Non lazy A\*

while( $s_{goal}$  is not expanded)

remove  $s$  with the smallest  
 $[f(s) = g(s) + h(s)]$  from  $OPEN$ ;

insert  $s$  into  $CLOSED$ ;

for every successor  $s'$  of  $s$  such  
that  $s'$  not in  $CLOSED$

if *edge* ( $s, s'$ ) in collision

$c(s, s') = \infty$

if  $g(s') > g(s) + c(s, s')$

$g(s') = g(s) + c(s, s');$

insert  $s'$  into  $OPEN$ ;

# Lazy A\*

Cohen, Phillips, and Likhachev 2014

## Non lazy A\*

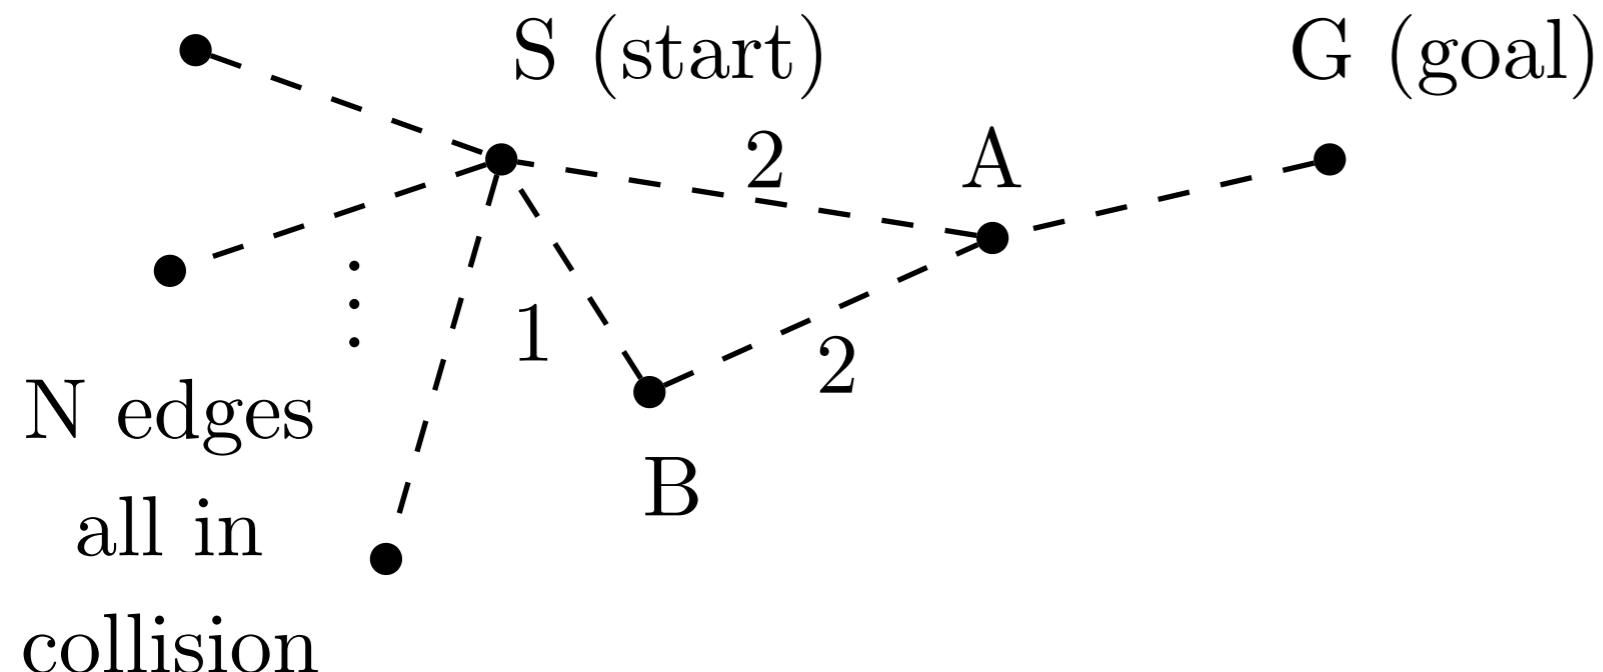
```
while( $s_{goal}$  is not expanded)
    remove  $s$  with the smallest
        [ $f(s) = g(s) + h(s)$ ] from  $OPEN$ ;
    insert  $s$  into  $CLOSED$ ;
    for every successor  $s'$  of  $s$  such
        that  $s'$  not in  $CLOSED$ 
        if  $edge(s, s')$  in collision
             $c(s, s') = \infty$ 
        if  $g(s') > g(s) + c(s, s')$ 
             $g(s') = g(s) + c(s, s');$ 
        insert  $s'$  into  $OPEN$ ;
```

## Lazy A\*

```
while( $s_{goal}$  is not expanded)
    remove  $s$  with the smallest
        [ $f(s) = g(s) + h(s)$ ] from  $OPEN$ ;
    if  $s$  is in  $CLOSED$ 
        continue;
    if  $edge(parent(s), s)$  in collision
        continue;
    insert  $s$  into  $CLOSED$ ;
    for every successor  $s'$  of  $s$  such
        that  $s'$  not in  $CLOSED$ 
        no collision checking of edge
        if  $g(s') > g(s) + c(s, s')$ 
             $g(s') = g(s) + c(s, s');$ 
        insert  $s'$  into  $OPEN$ ; // multiple
                                copies
```

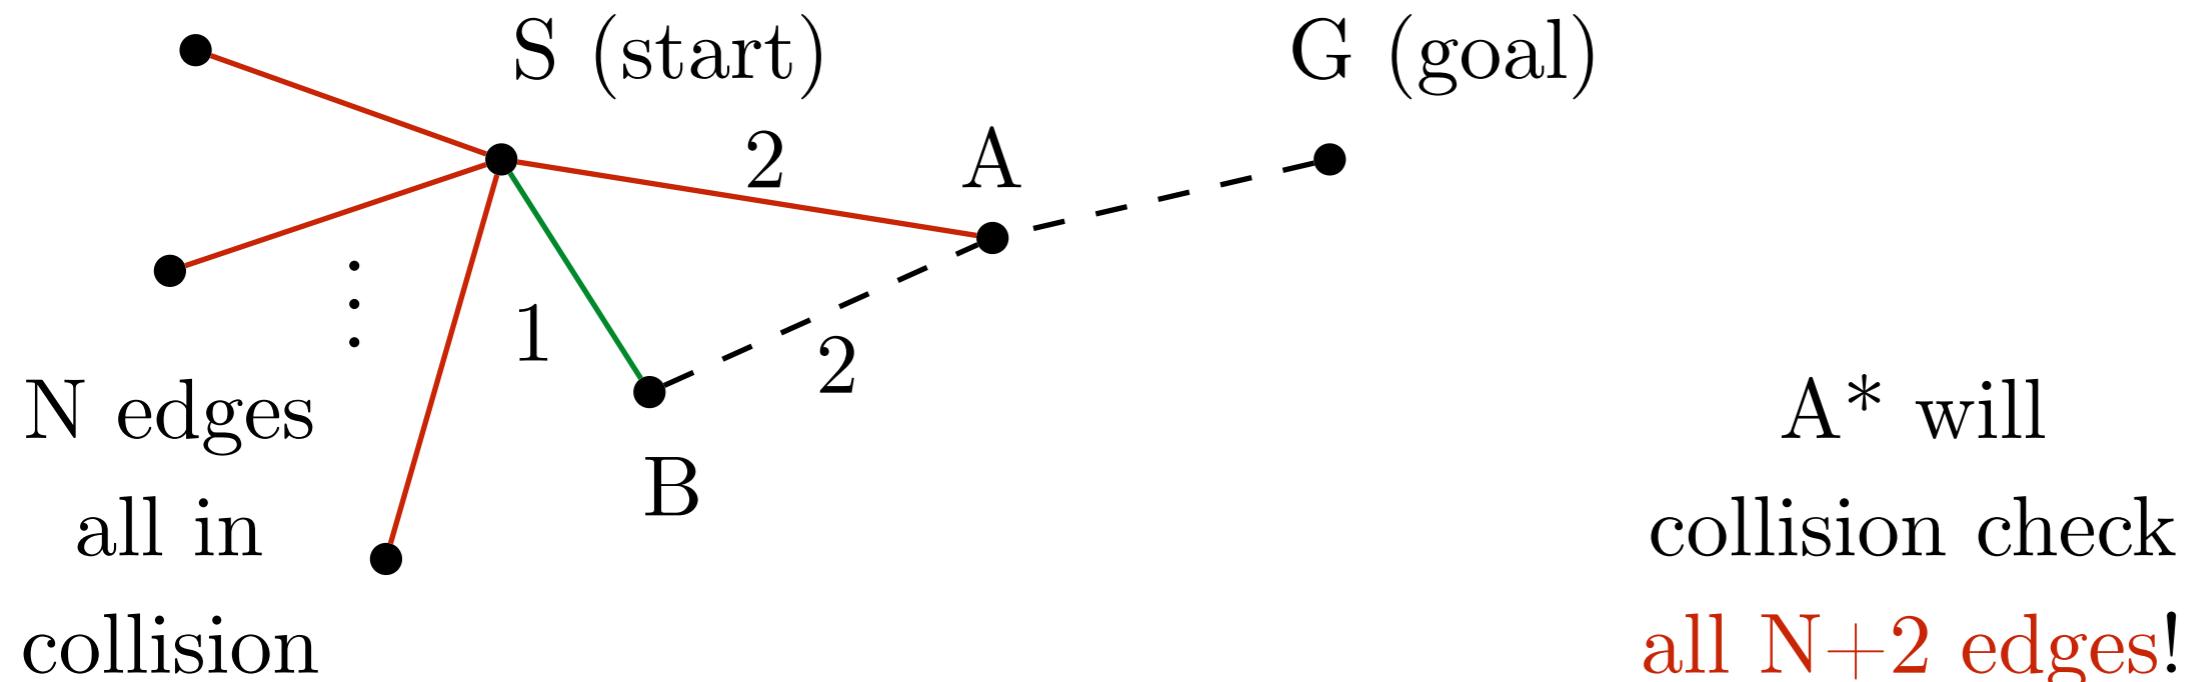
# A\*

Let's say S-A is in collision and true shortest path is S-B-A-G



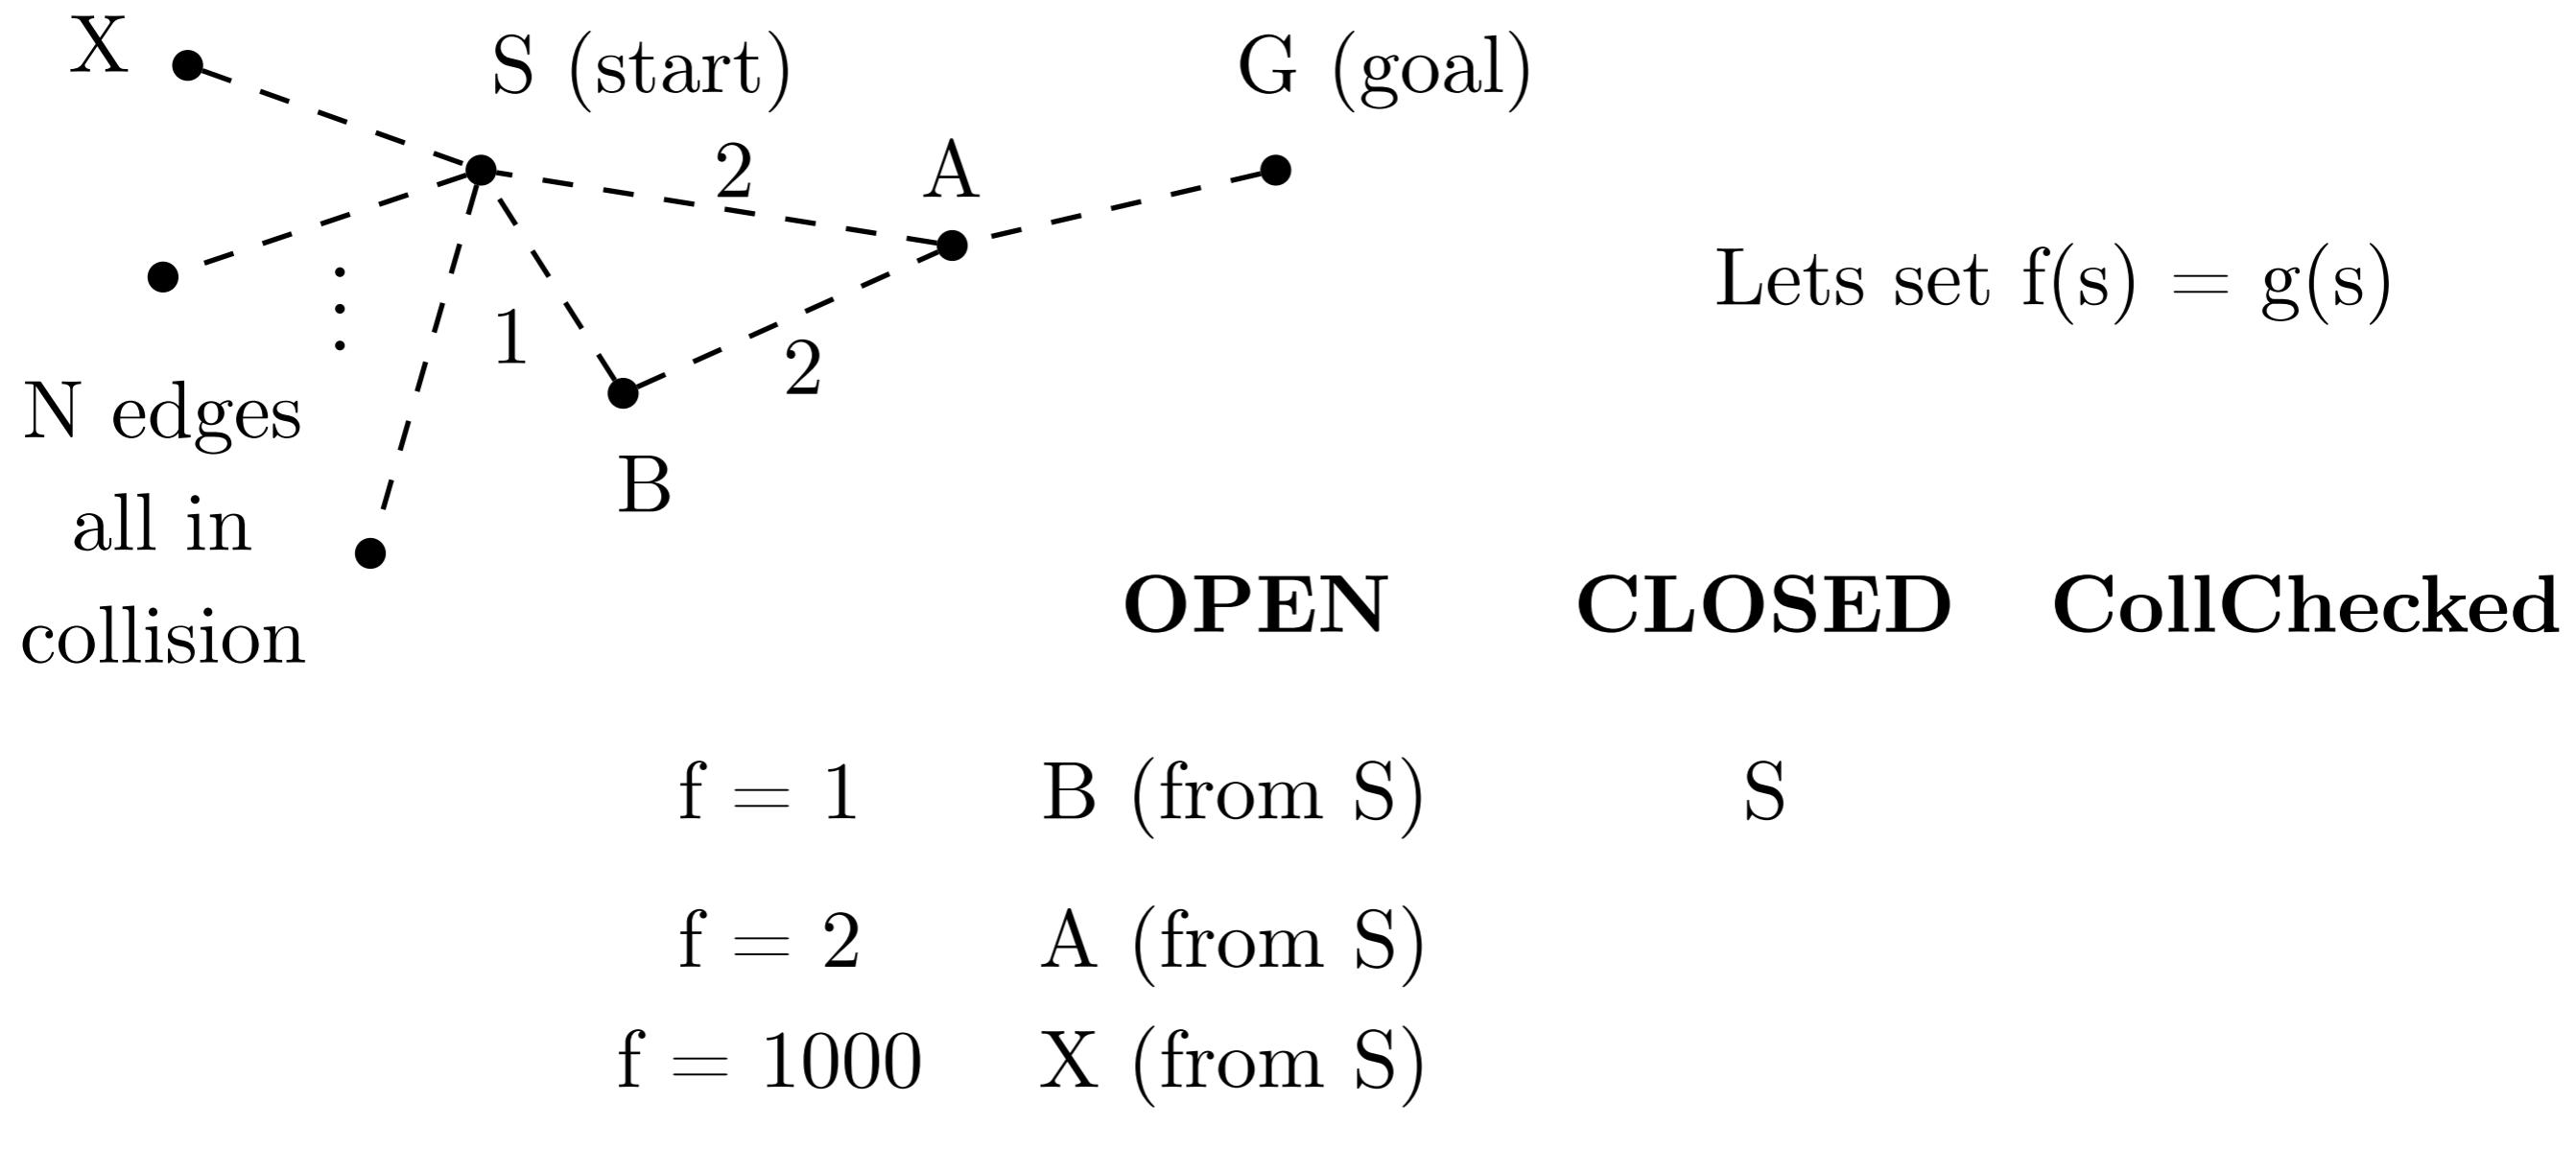
# A\*

Let's say S-A is in collision and true shortest path is S-B-A-G



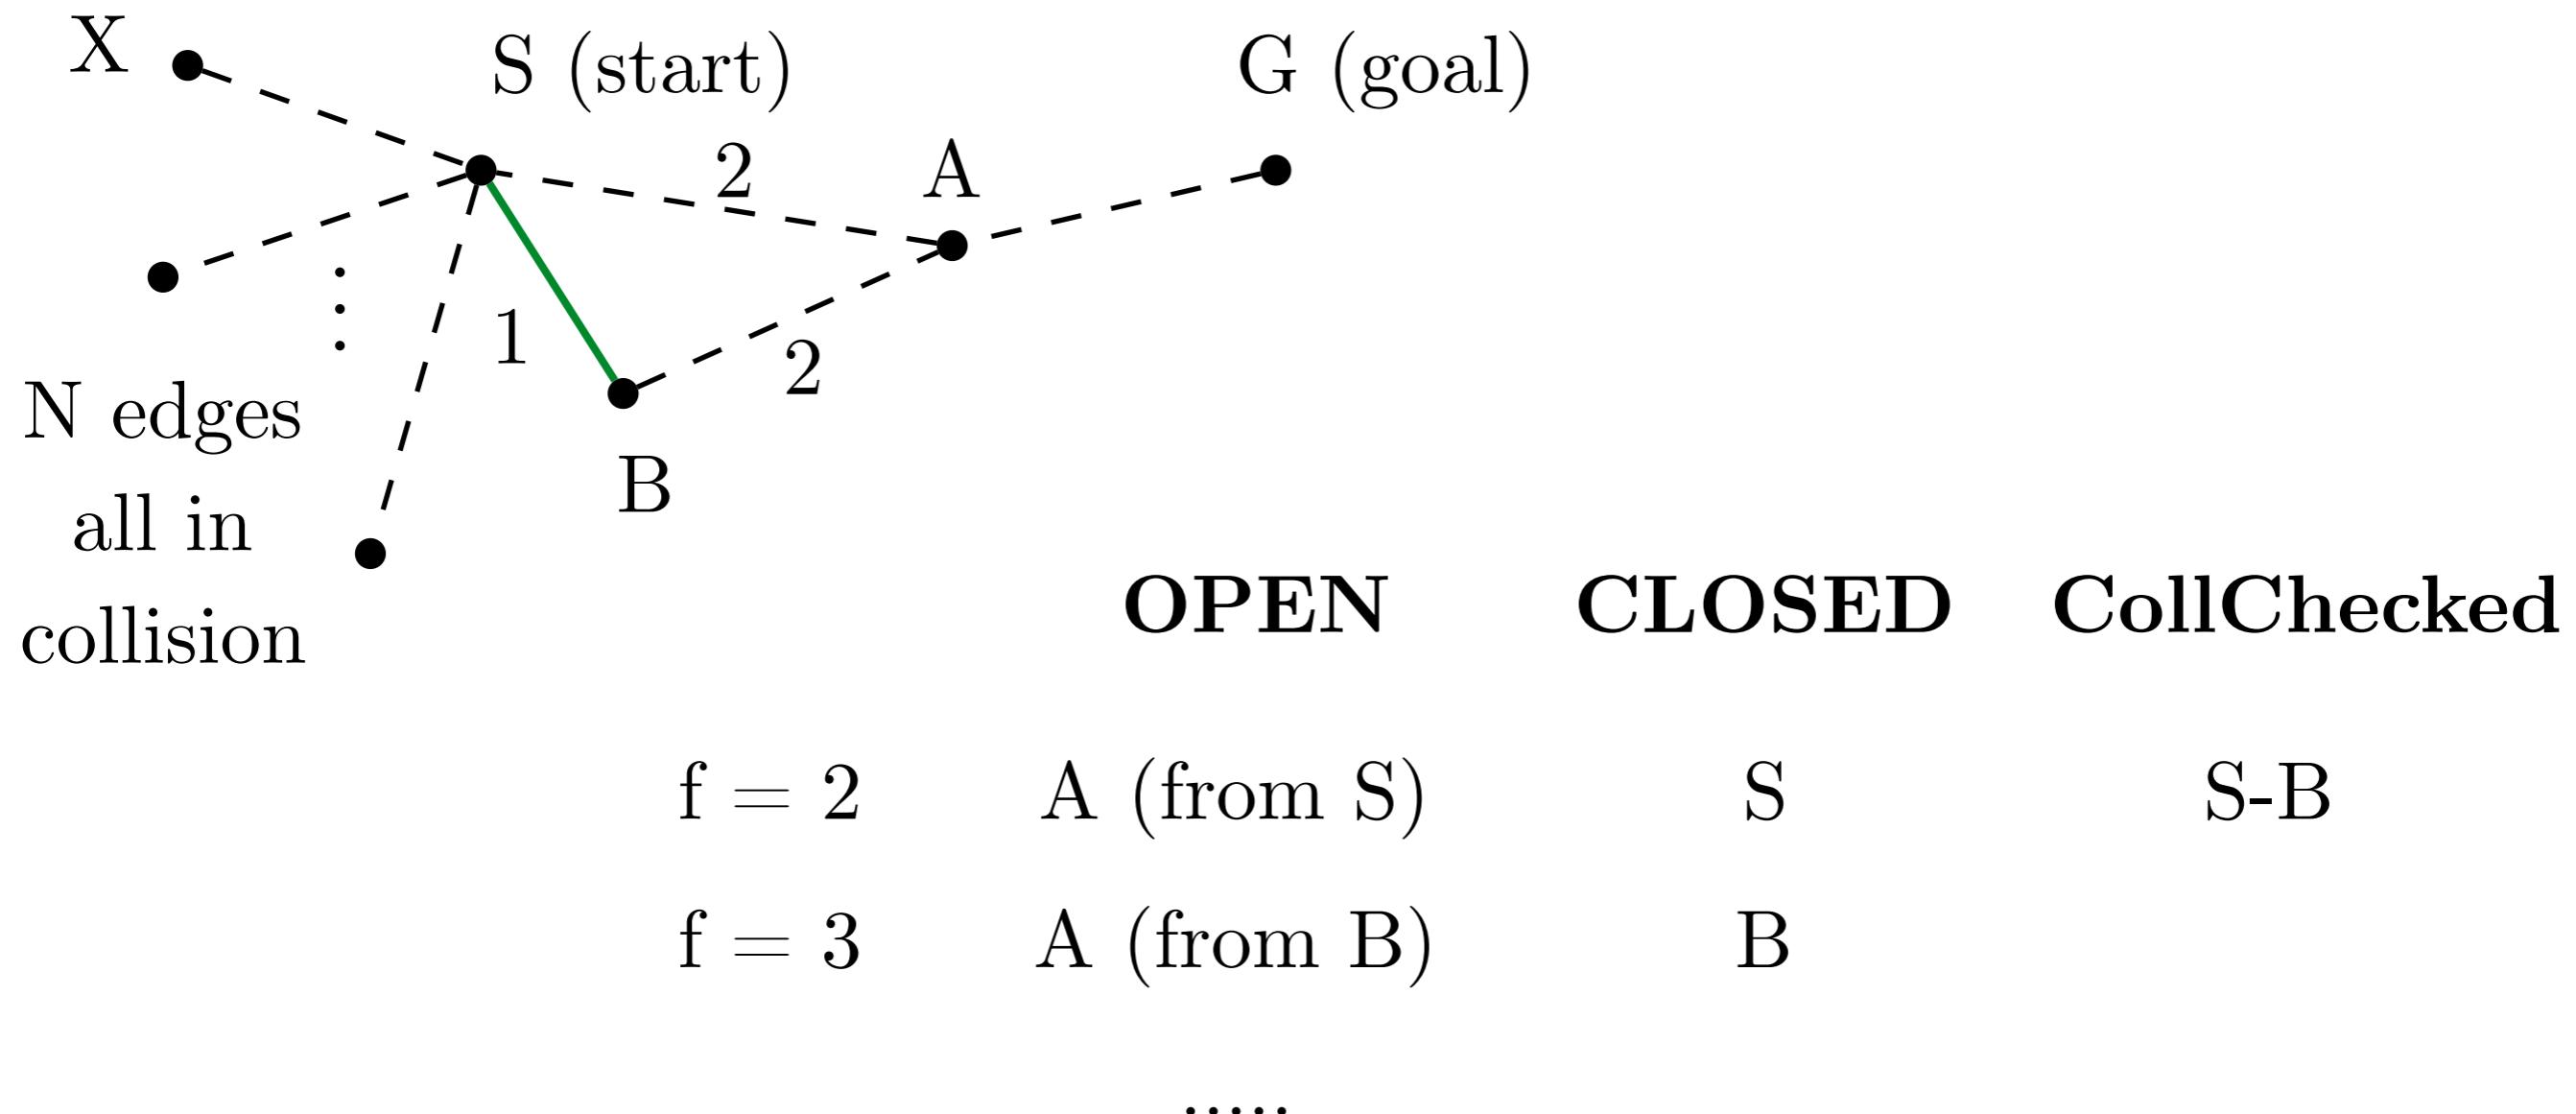
# Lazy A\*

Let's say S-A is in collision and true shortest path is S-B-A-G



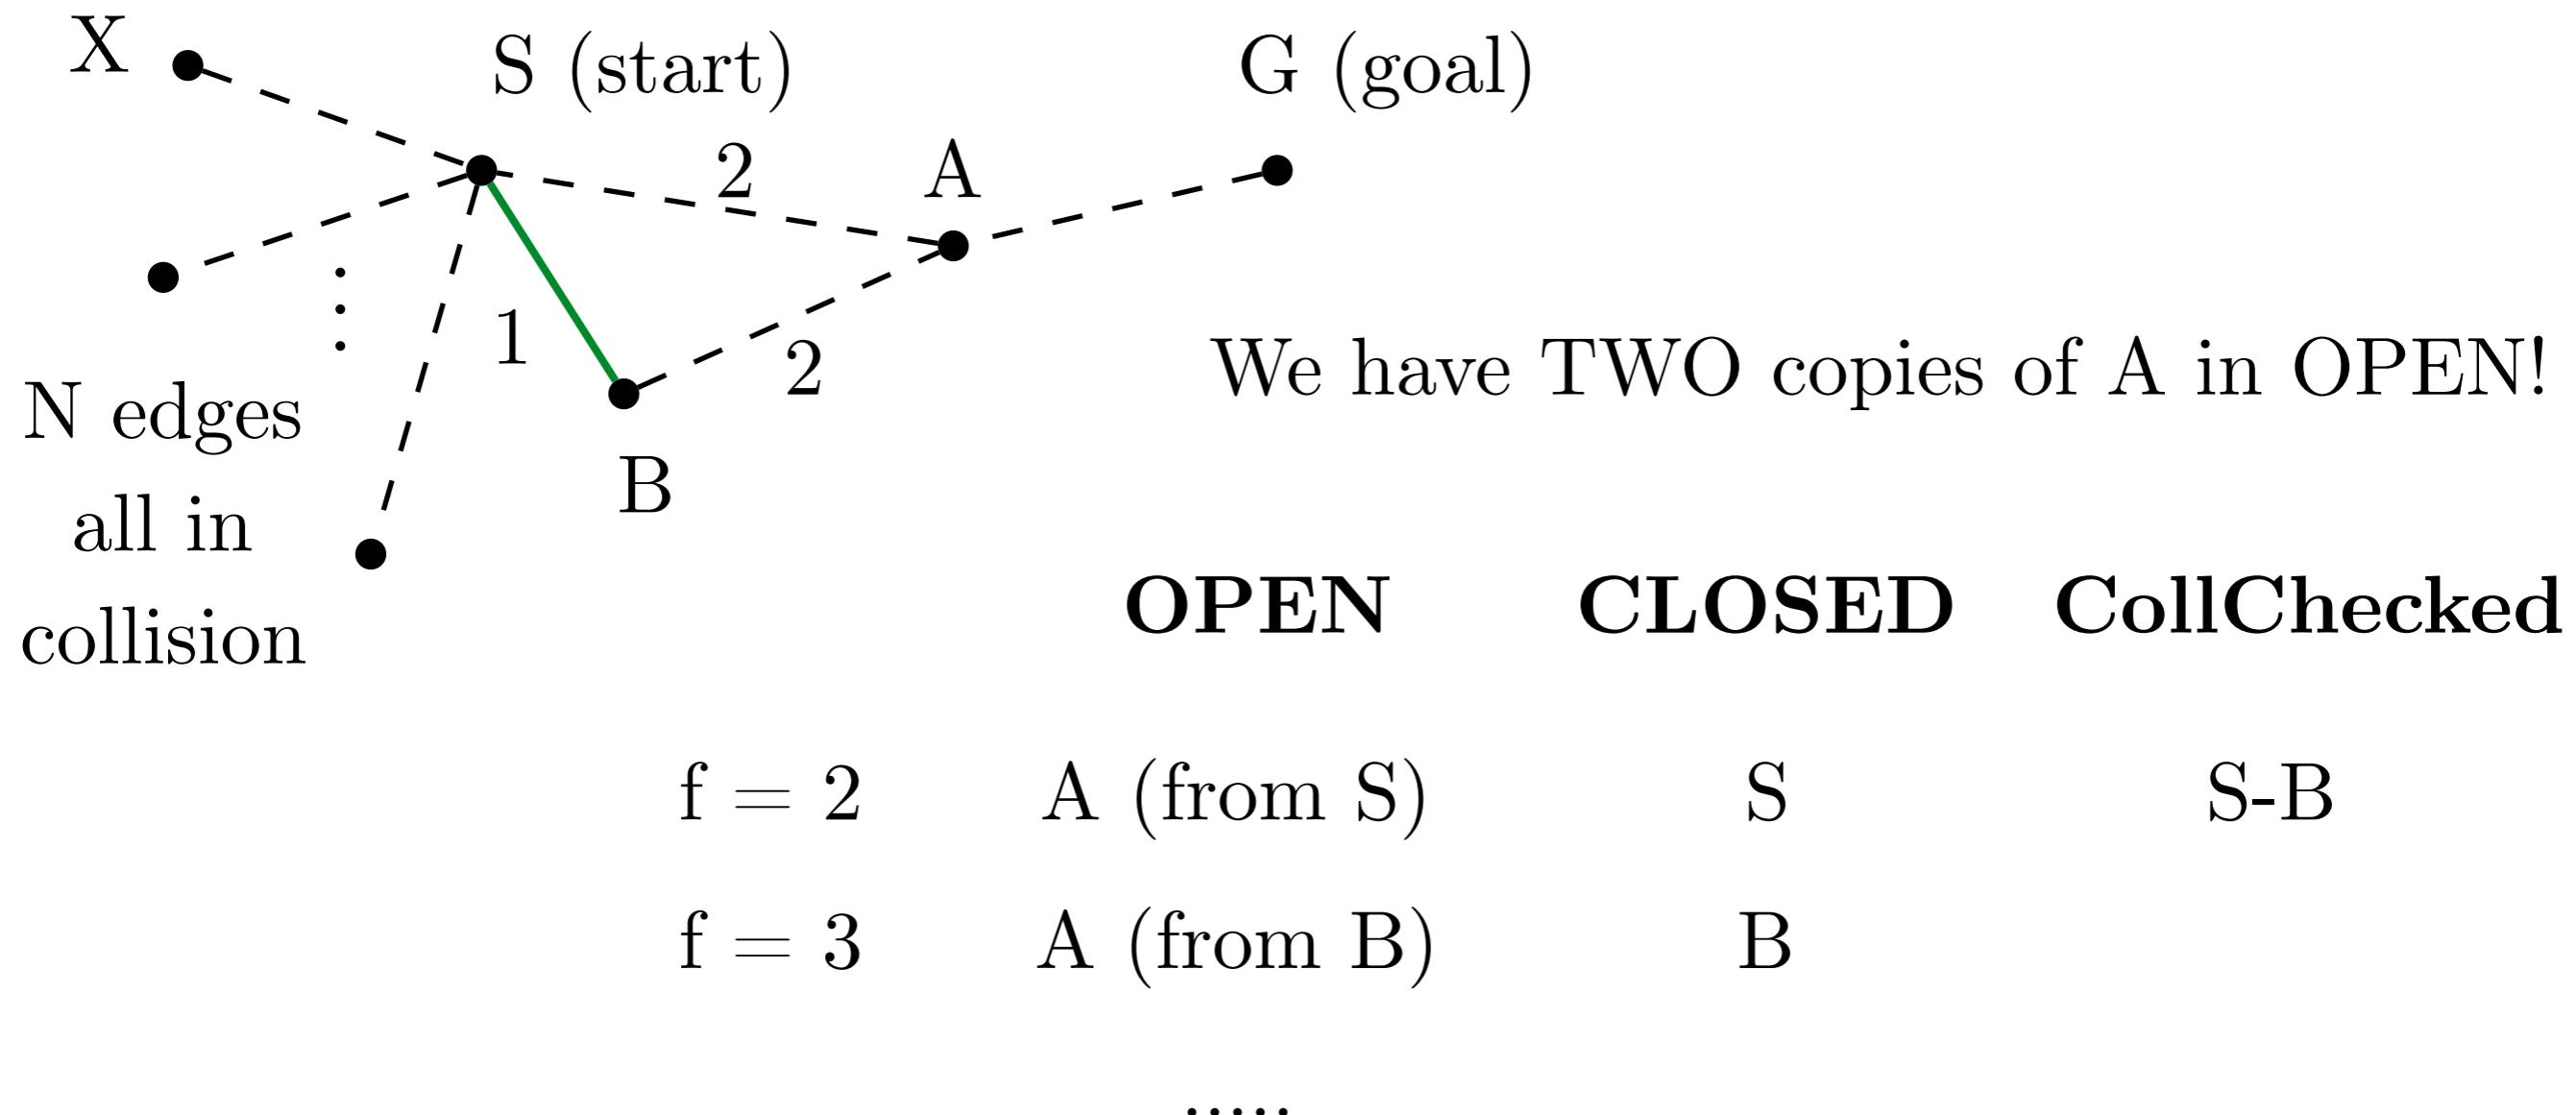
# Lazy A\*

Let's say S-A is in collision and true shortest path is S-B-A-G



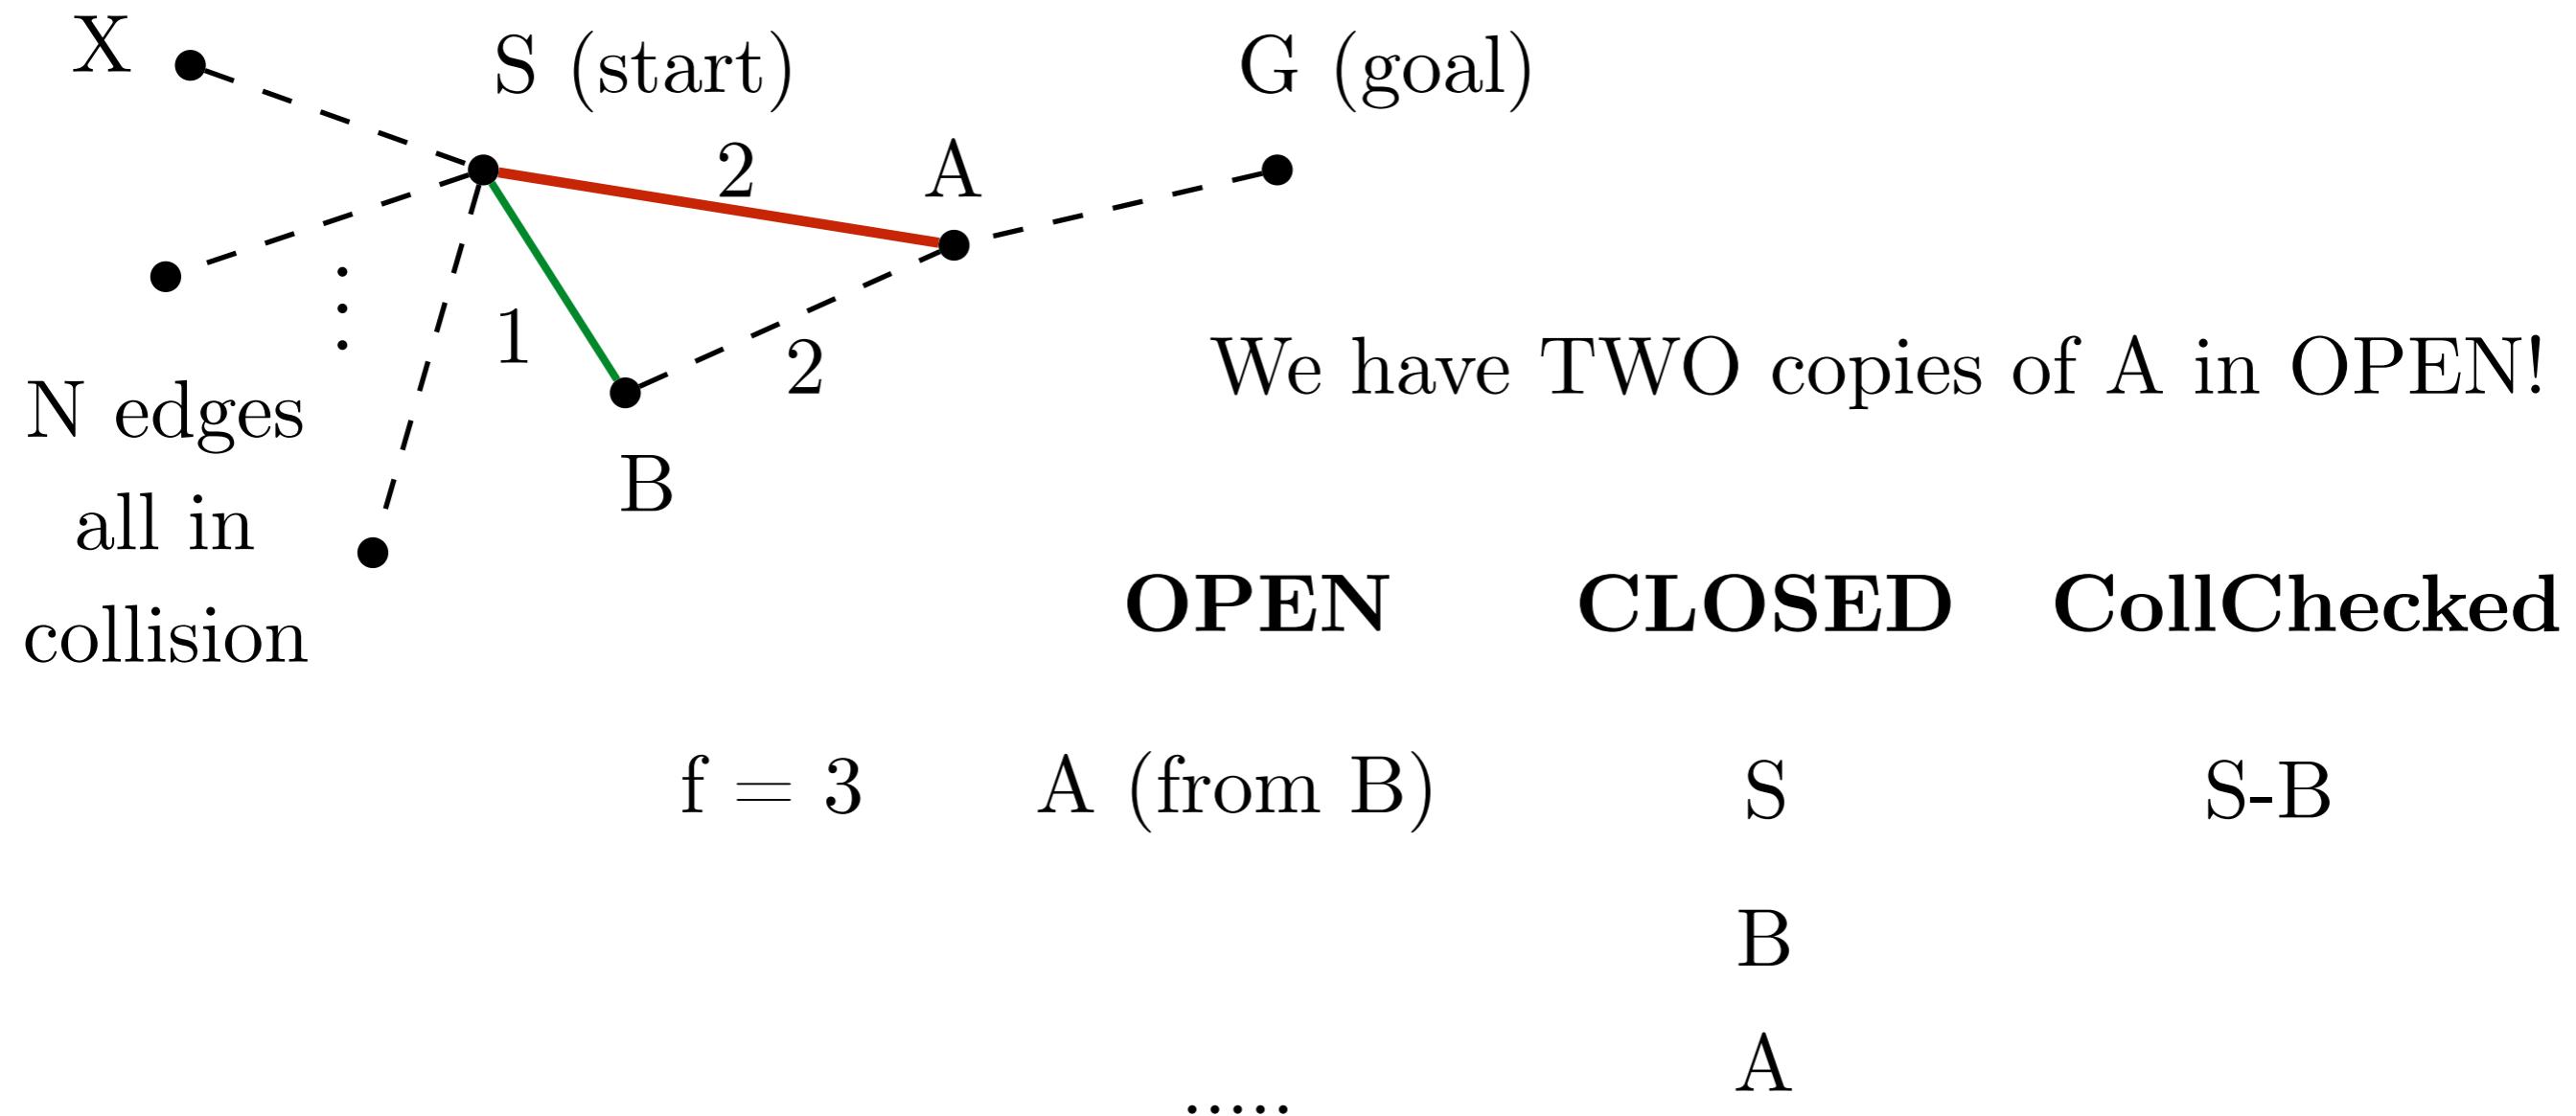
# Lazy A\*

Let's say S-A is in collision and true shortest path is S-B-A-G



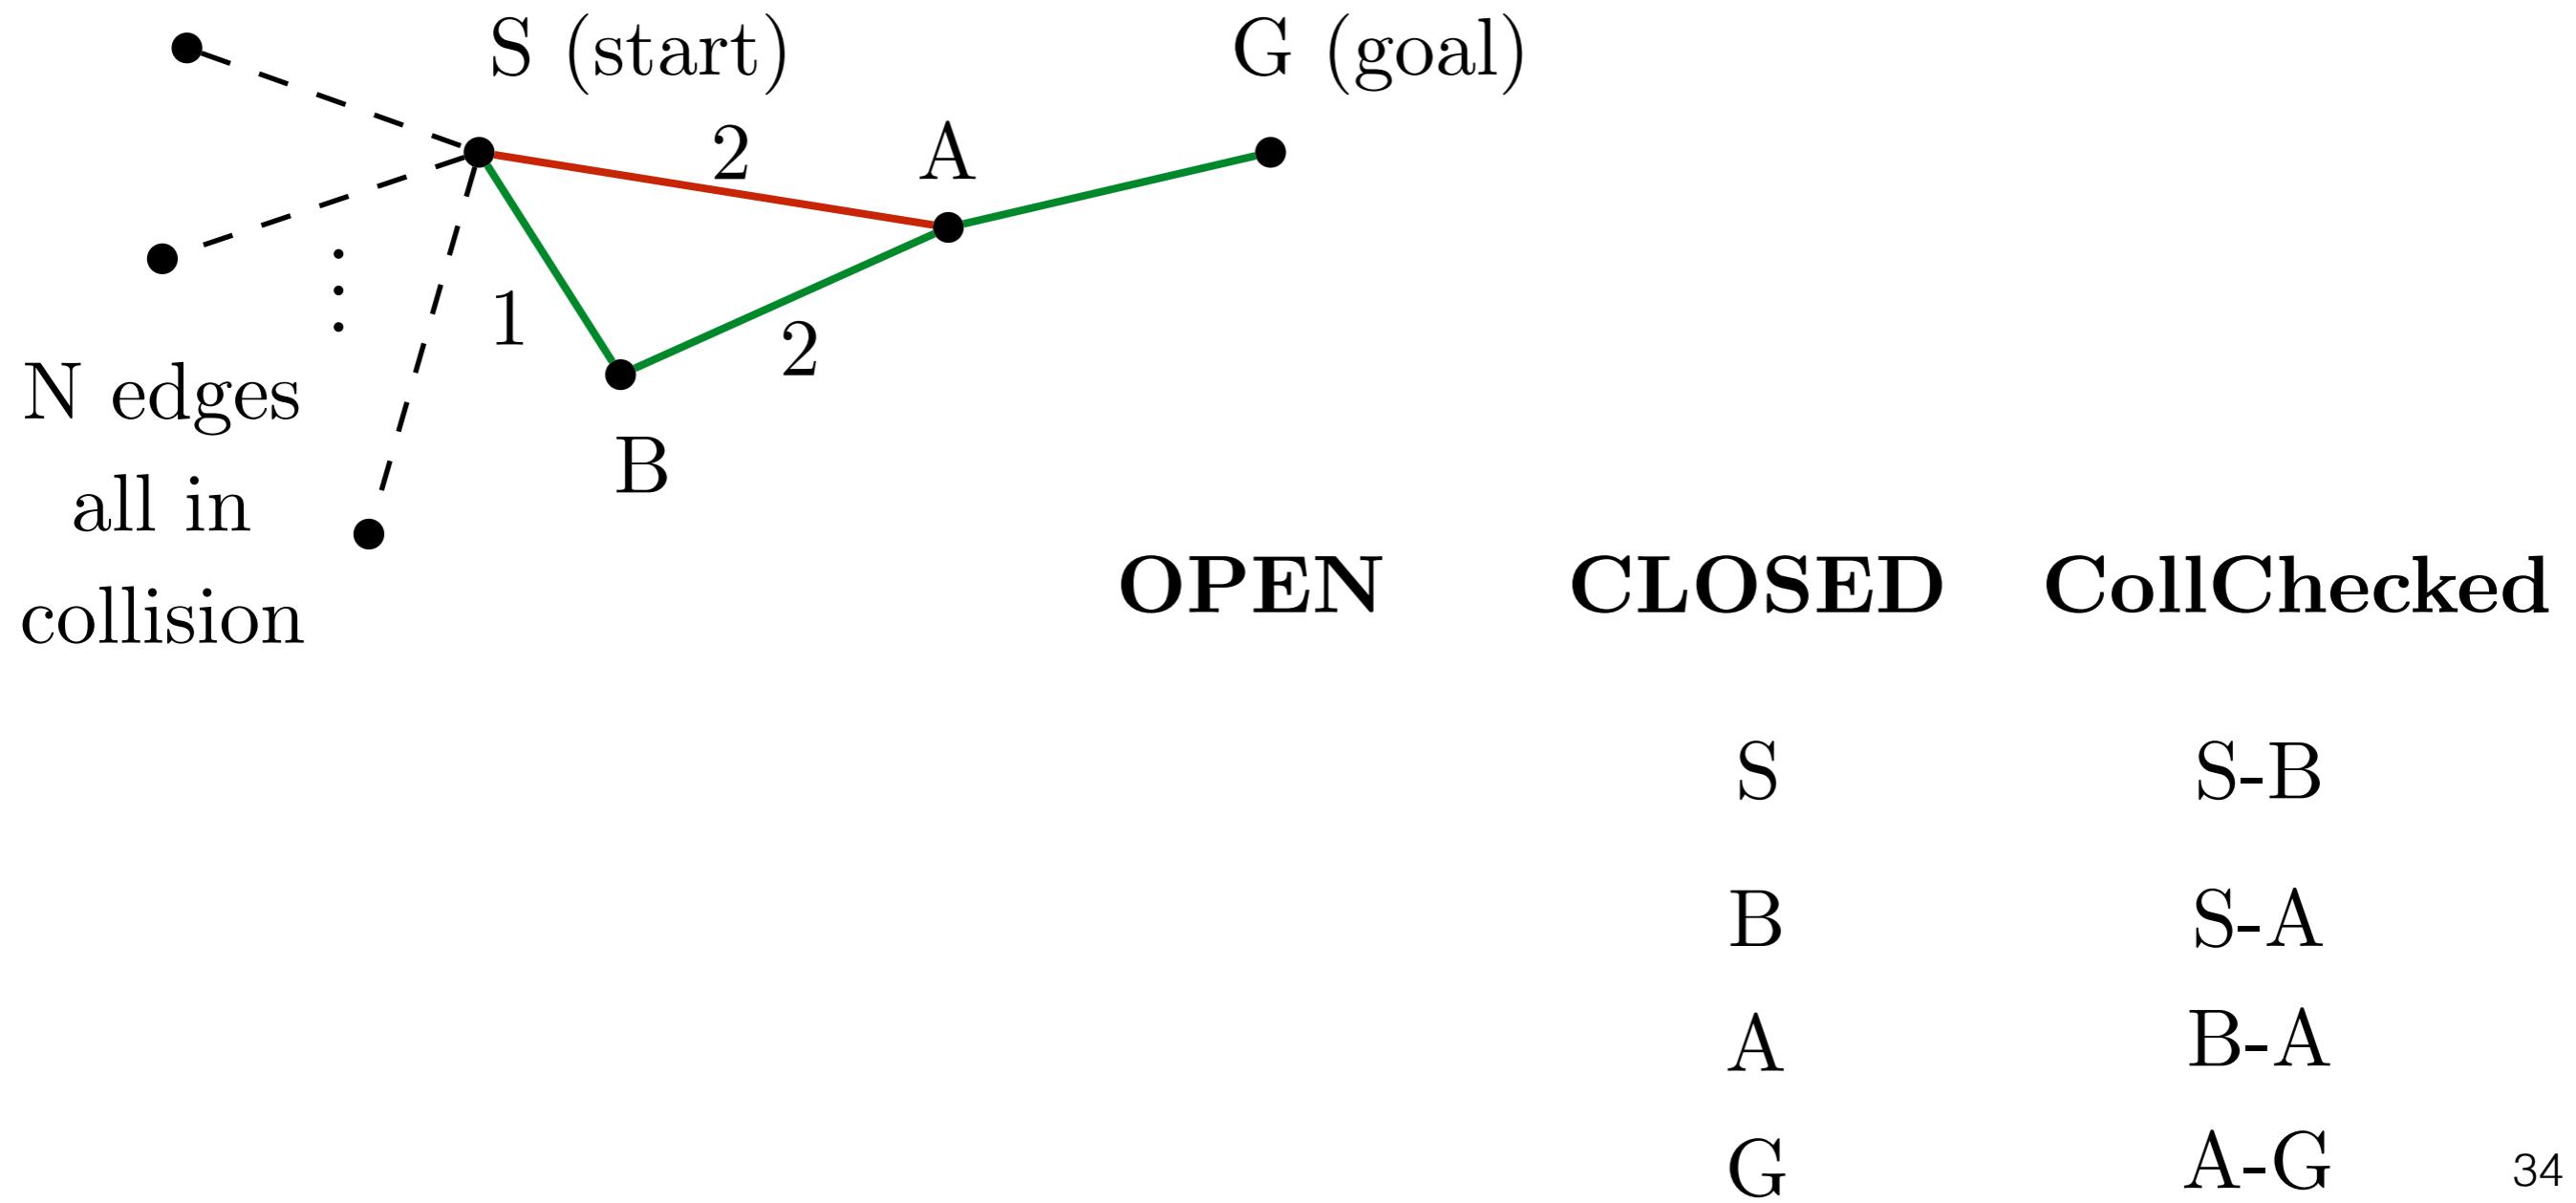
# Lazy A\*

Let's say S-A is in collision and true shortest path is S-B-A-G

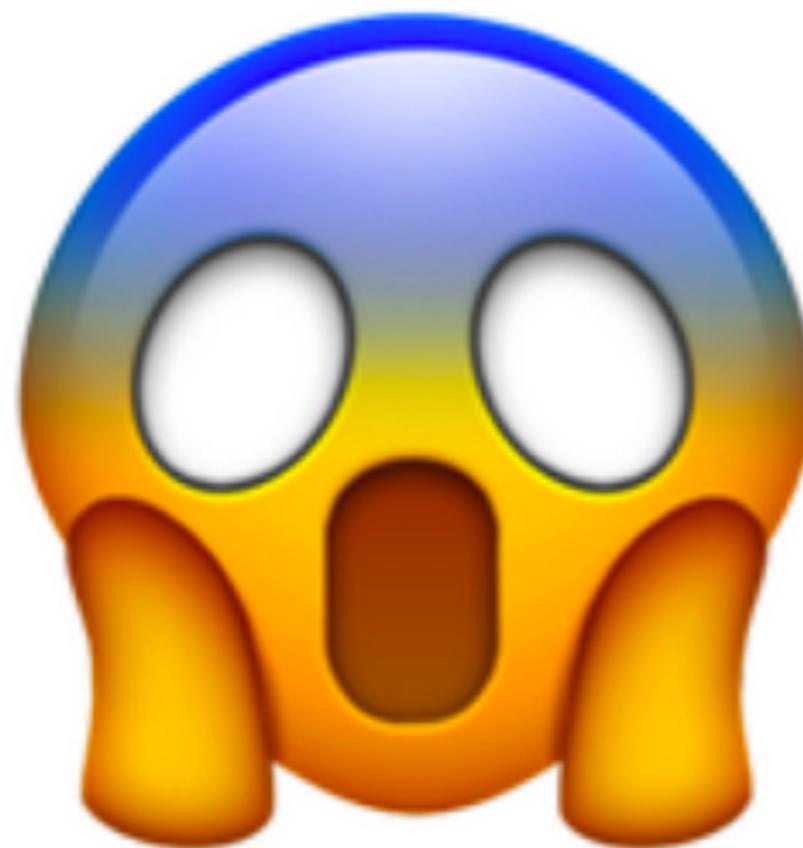


# Lazy A\*

Let's say S-A is in collision and true shortest path is S-B-A-G



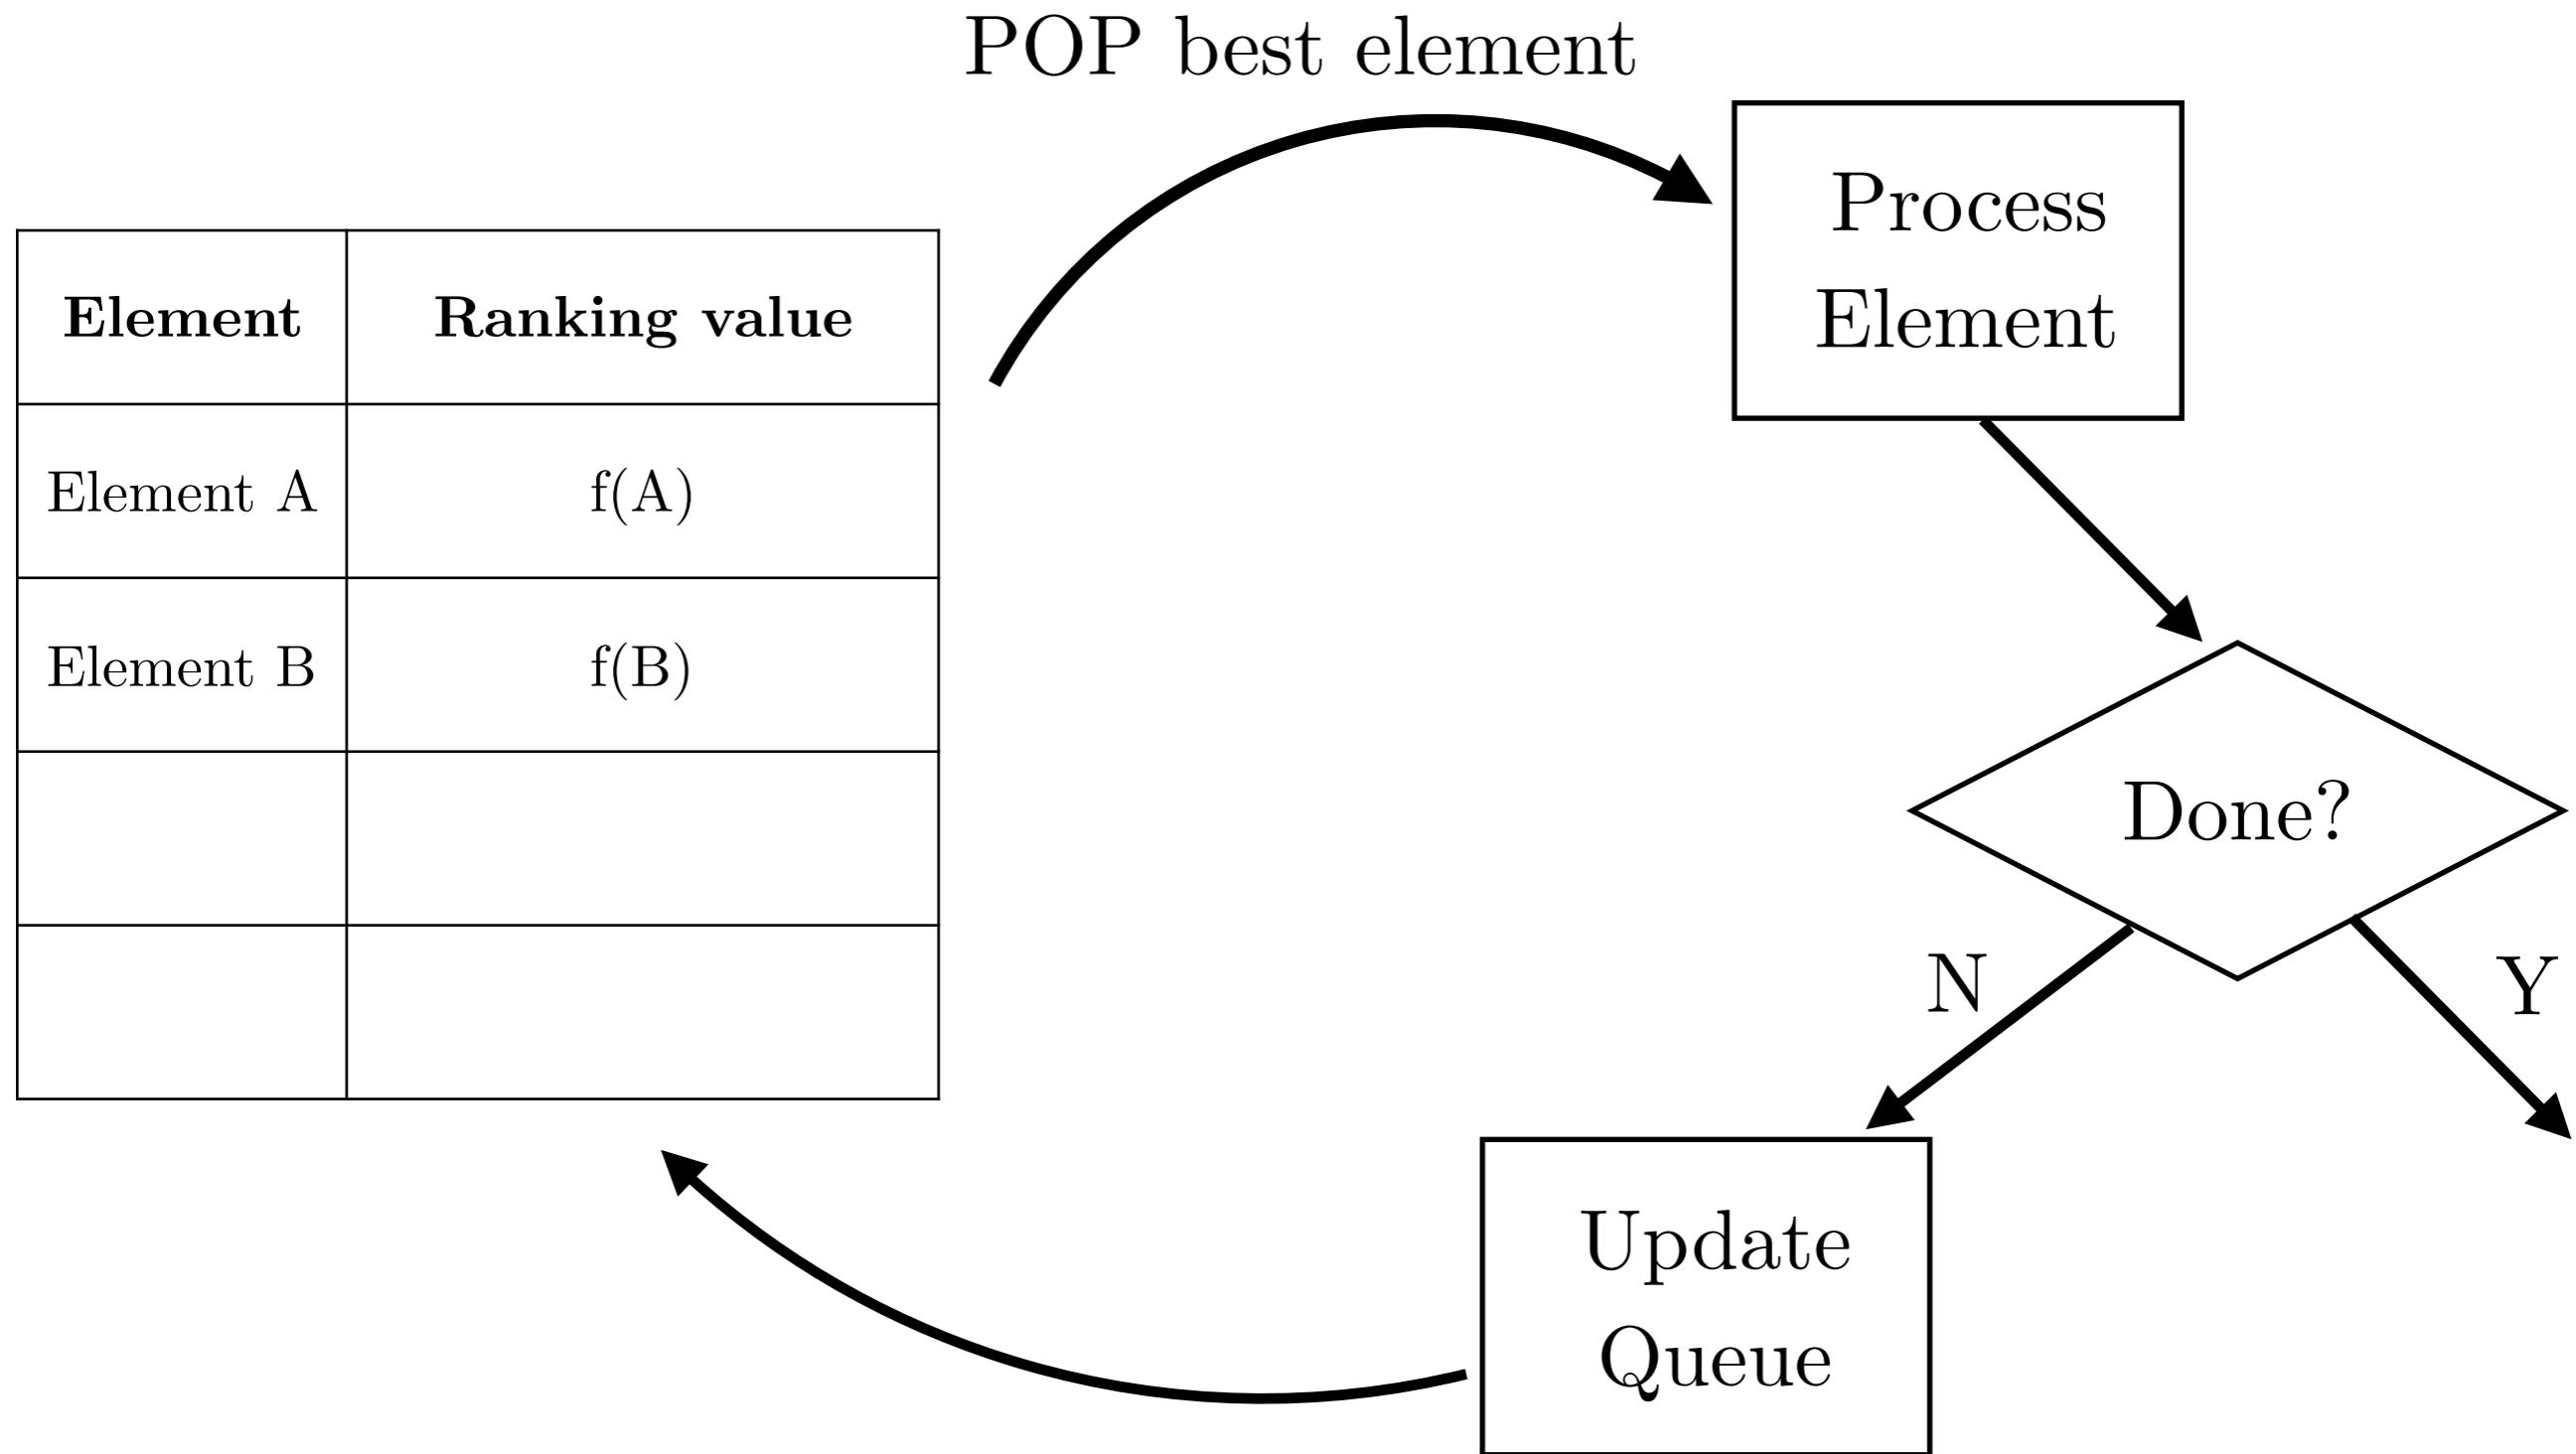
Yet another algorithm to  
remember....!?!?!



# Best First Search explains it all



# Best First Search as a Meta++ algorithm



If we want to minimize the  
number of vertices expanded,  
then  
elements of the queue should be  
candidate vertices to expand!

If we want to minimize the  
number of edges evaluated,  
then

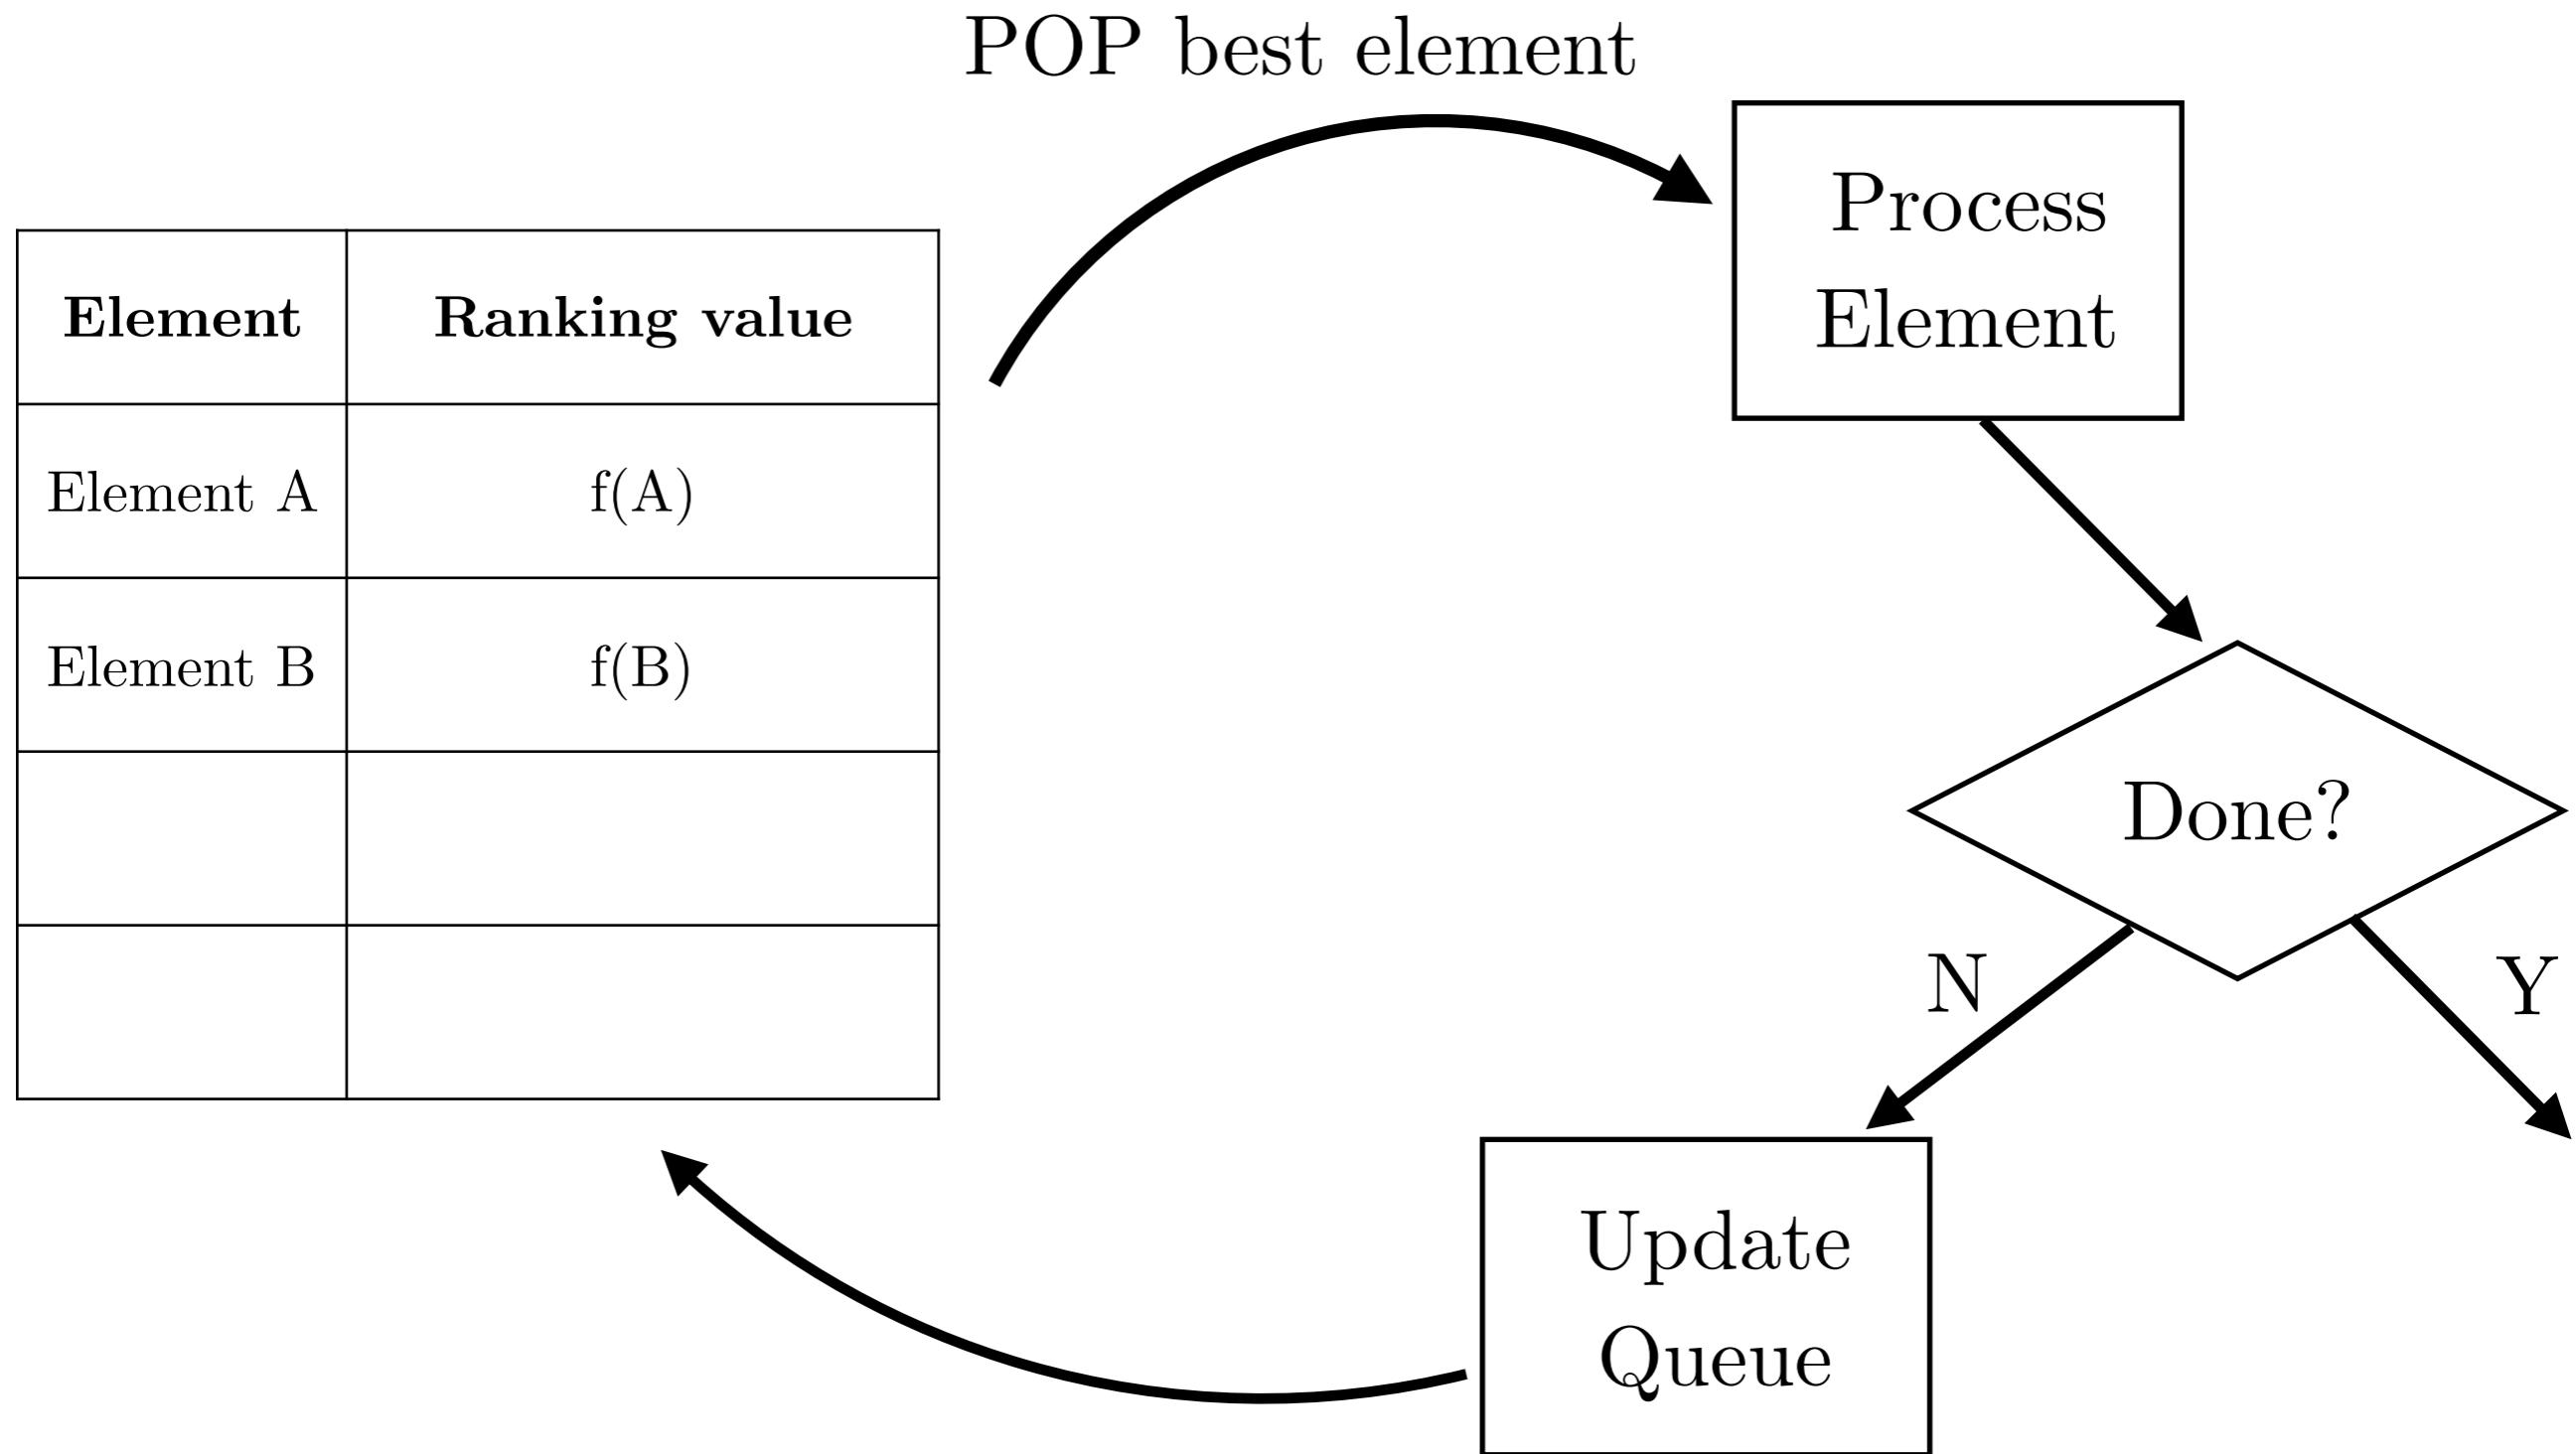
.....

If we want to minimize the  
number of edges evaluated,  
then

.....

elements of the queue should be  
candidate edges to evaluate!

# Best First Search as a Meta++ algorithm



# How do we define a queue over edges???

OPEN list of A\*

Element (Vertex)	Value (f-value of path through vertex)
Vertex A	$f(A) = g(A) + h(A)$
Vertex B	$f(B) = g(B) + h(B)$

# How do we define a queue over edges???

OPEN list of A\*

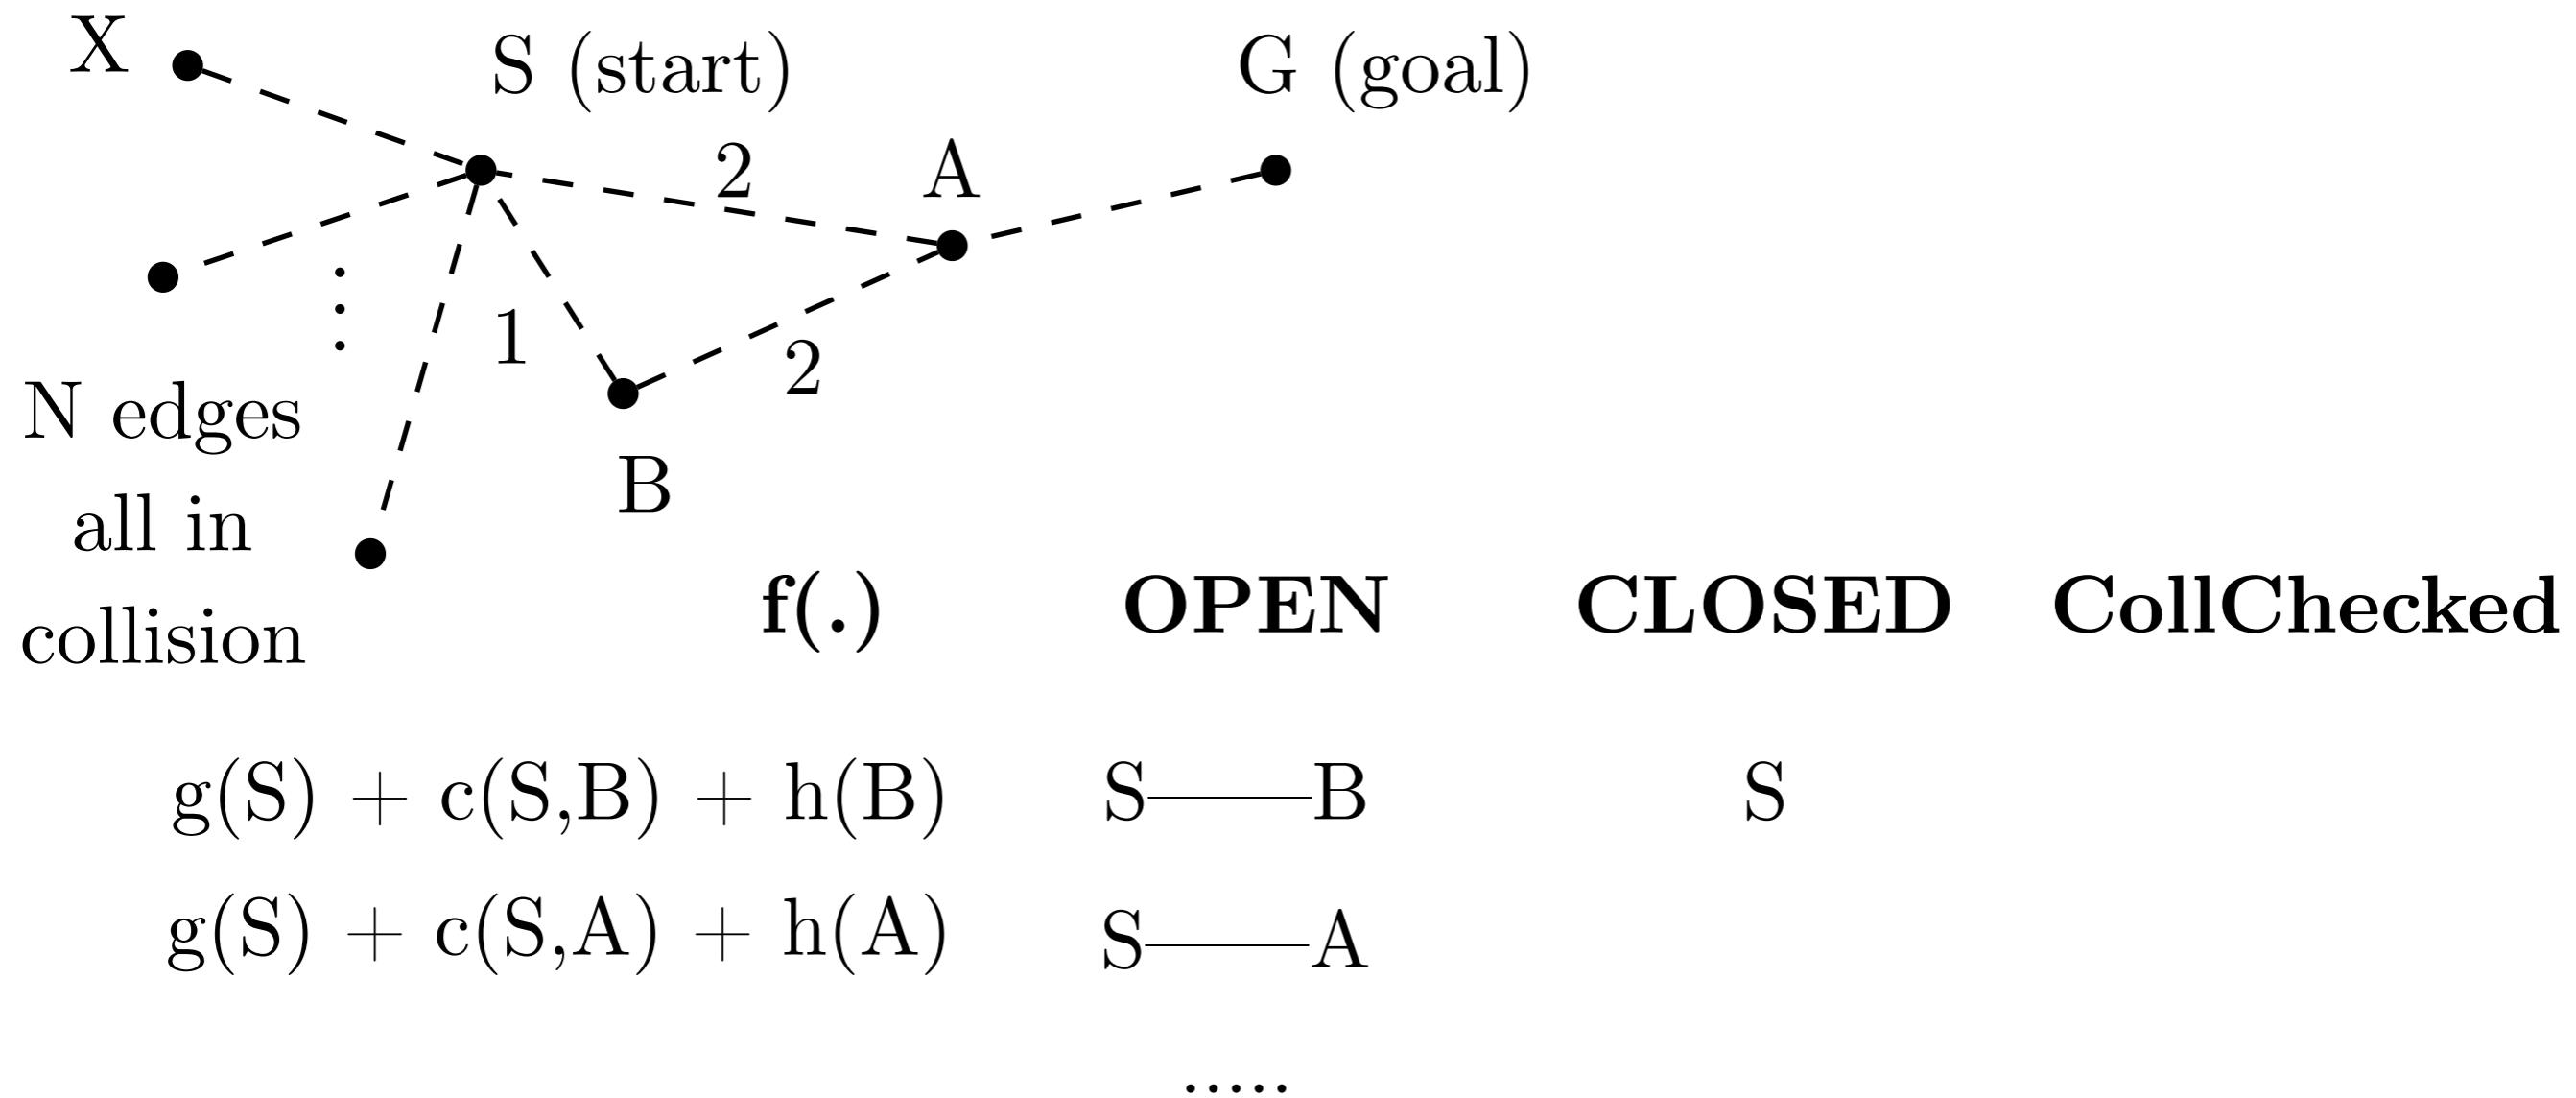
Element (Vertex)	Value (f-value of path through vertex)
Vertex A	$f(A) = g(A) + h(A)$
Vertex B	$f(B) = g(B) + h(B)$

OPEN list of Lazy A\*

Element (Edge)	Value (f-value of path through edge)
Edge (X,Y)	$f(X,Y) = g(X) + c(X,Y) + h(Y)$
Edge (P,Q)	$f(P,Q) = g(P) + c(P,Q) + h(Q)$
Edge (W,Y)	$f(W,Y) = g(W) + c(W,Y) + h(Y)$

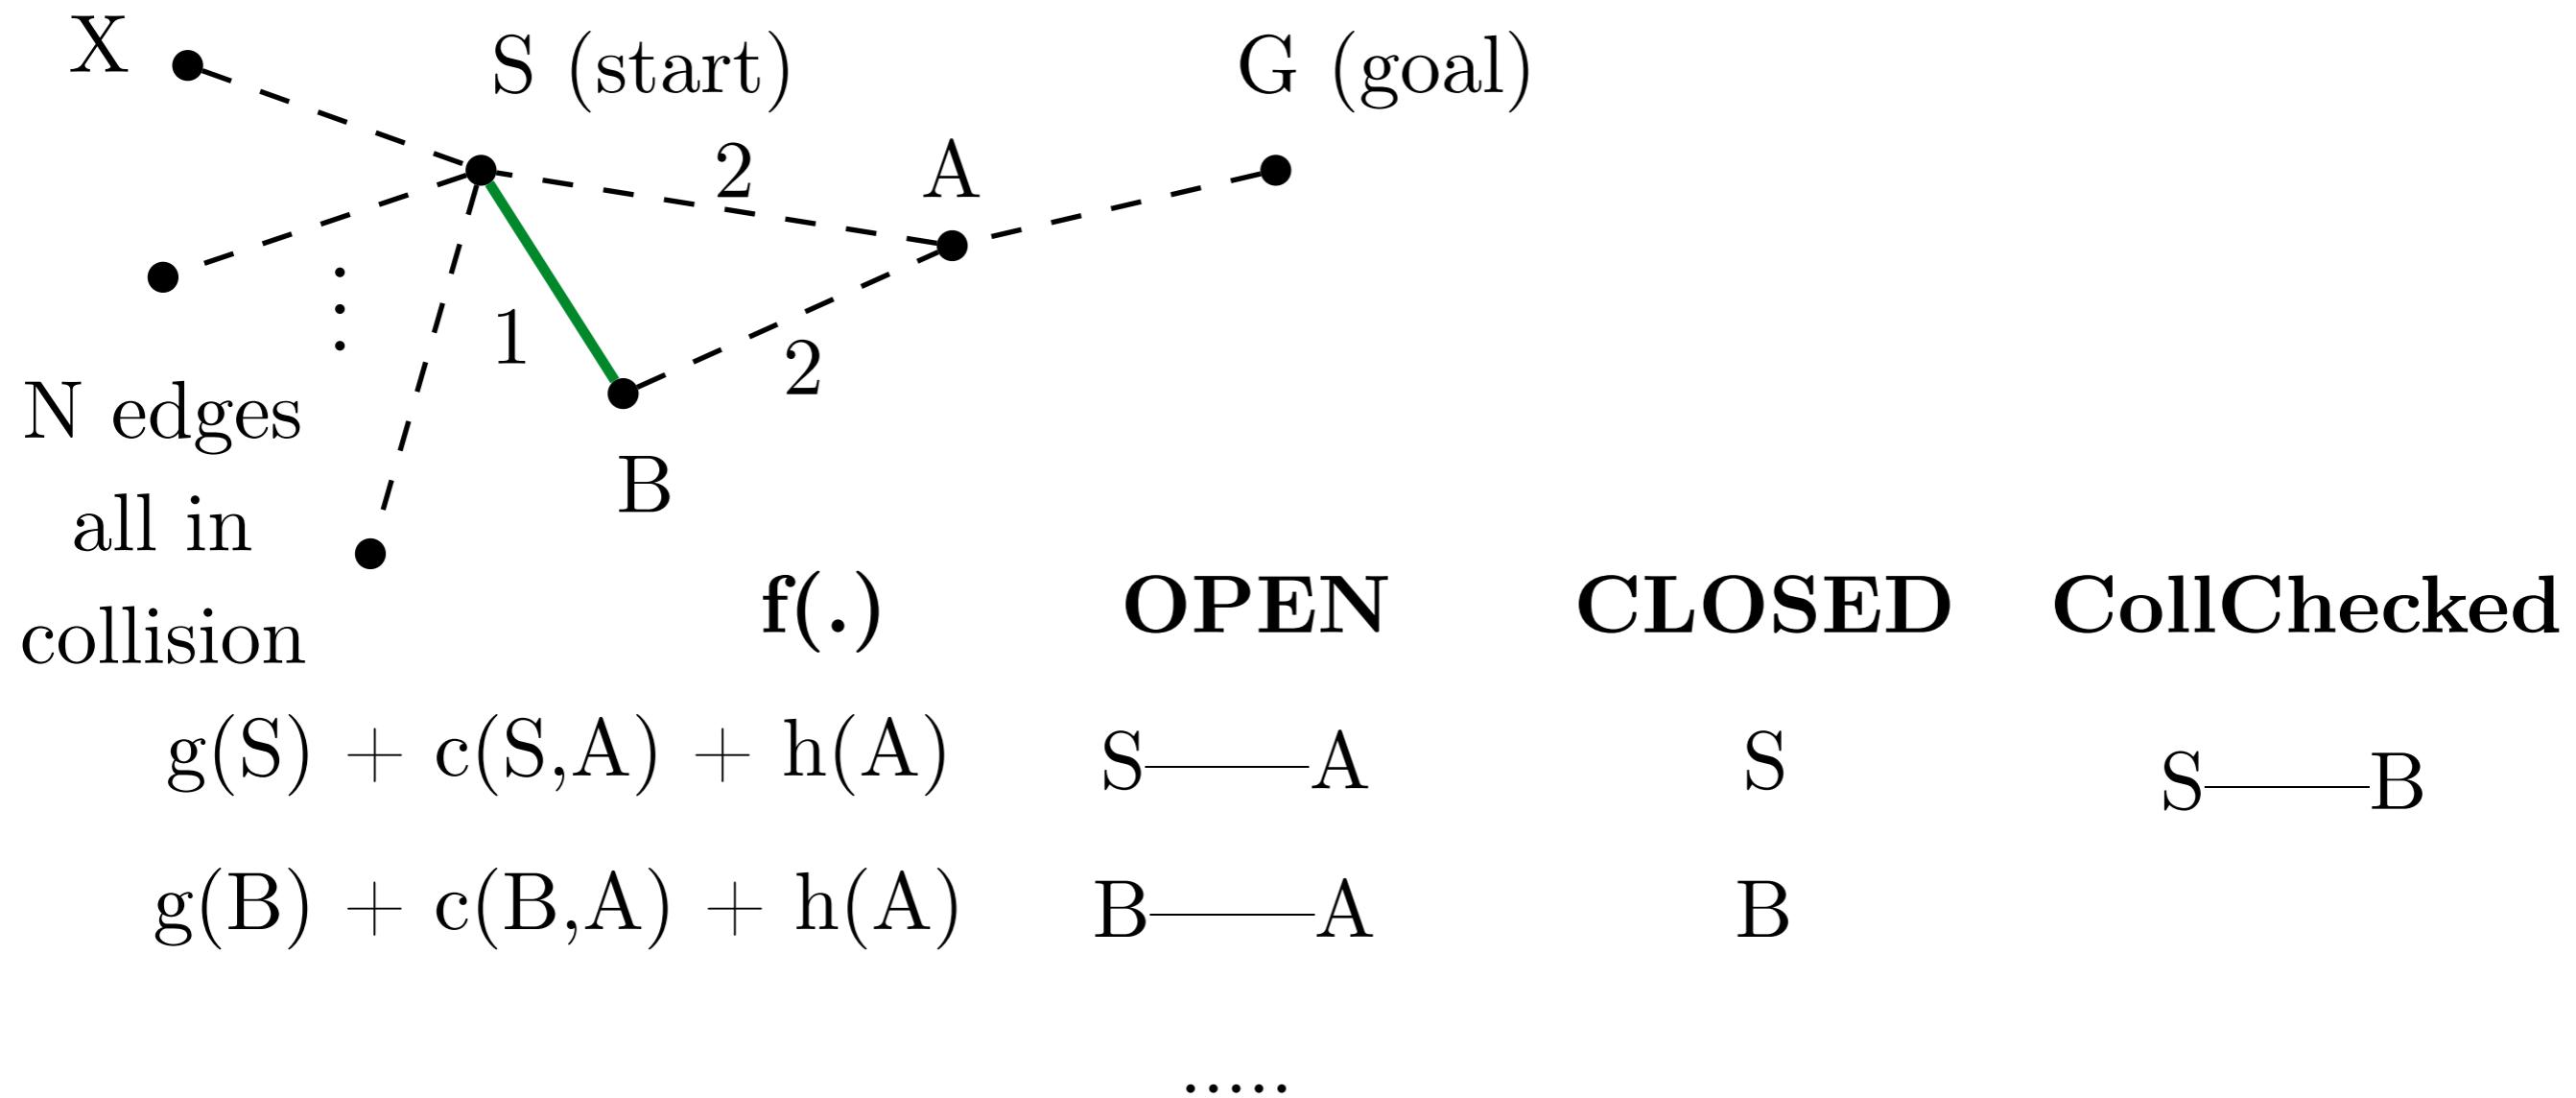
# Lazy A\*

Let's say S-A is in collision and true shortest path is S-B-A-G



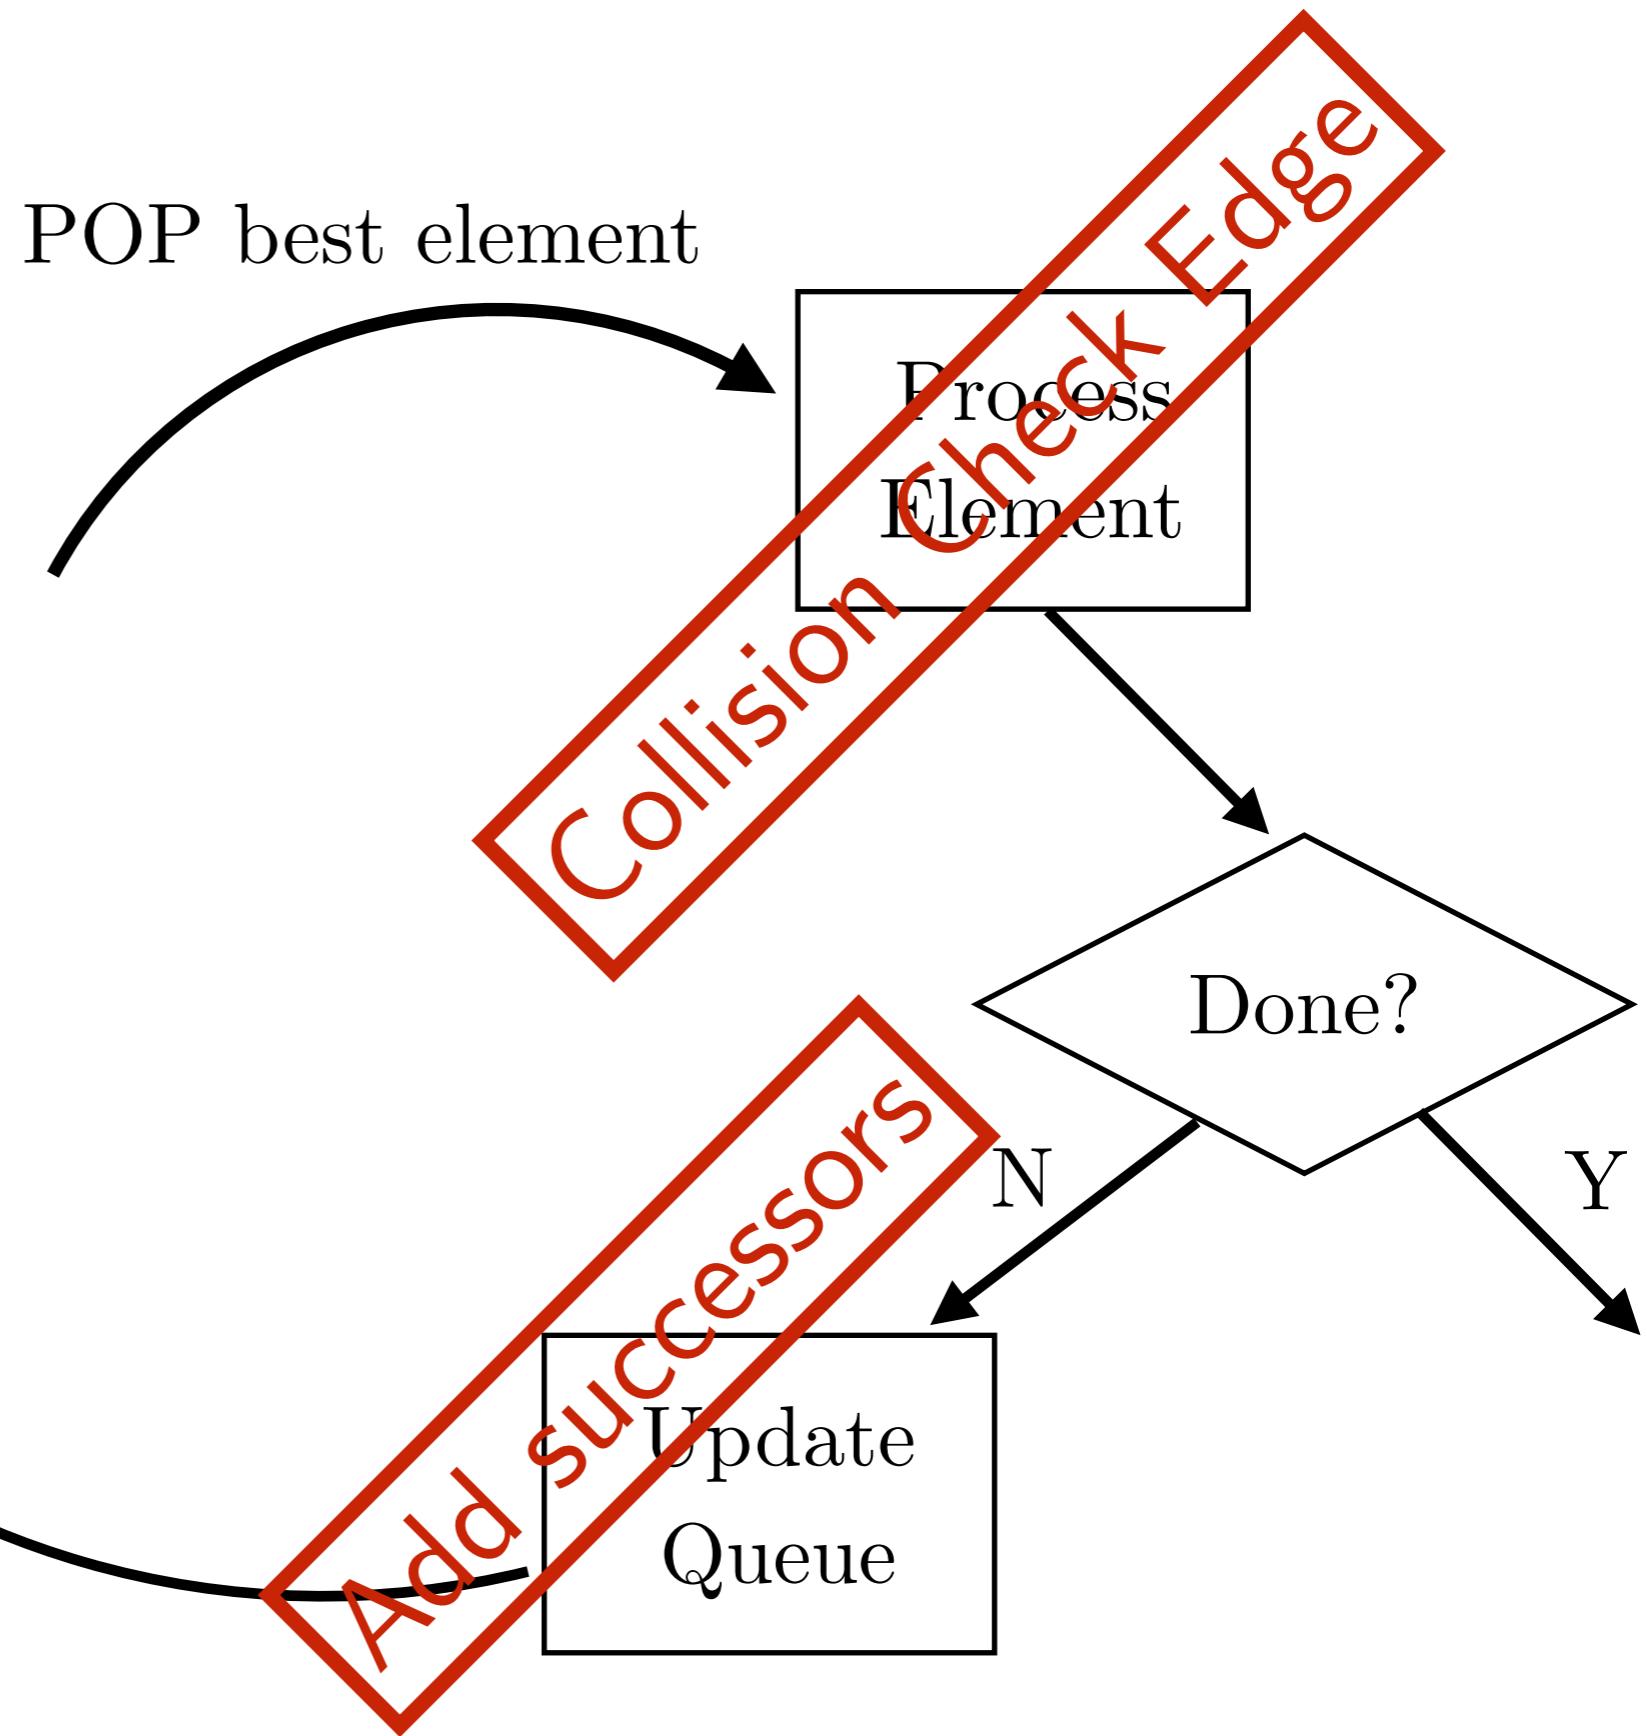
# Lazy A\*

Let's say S-A is in collision and true shortest path is S-B-A-G



# Lazy A\* as BFS

Element	Ranking value
Element A	$f(A)$
Element B	$f(B)$
	Queue over edges



What is the laziest that we can  
be?

What is the laziest that we can  
be?

# LazySP

(Lazy Shortest Path)

Dellin and Srinivasa, 2016

First Provably Edge-Optimal A\*-like Search Algorithm

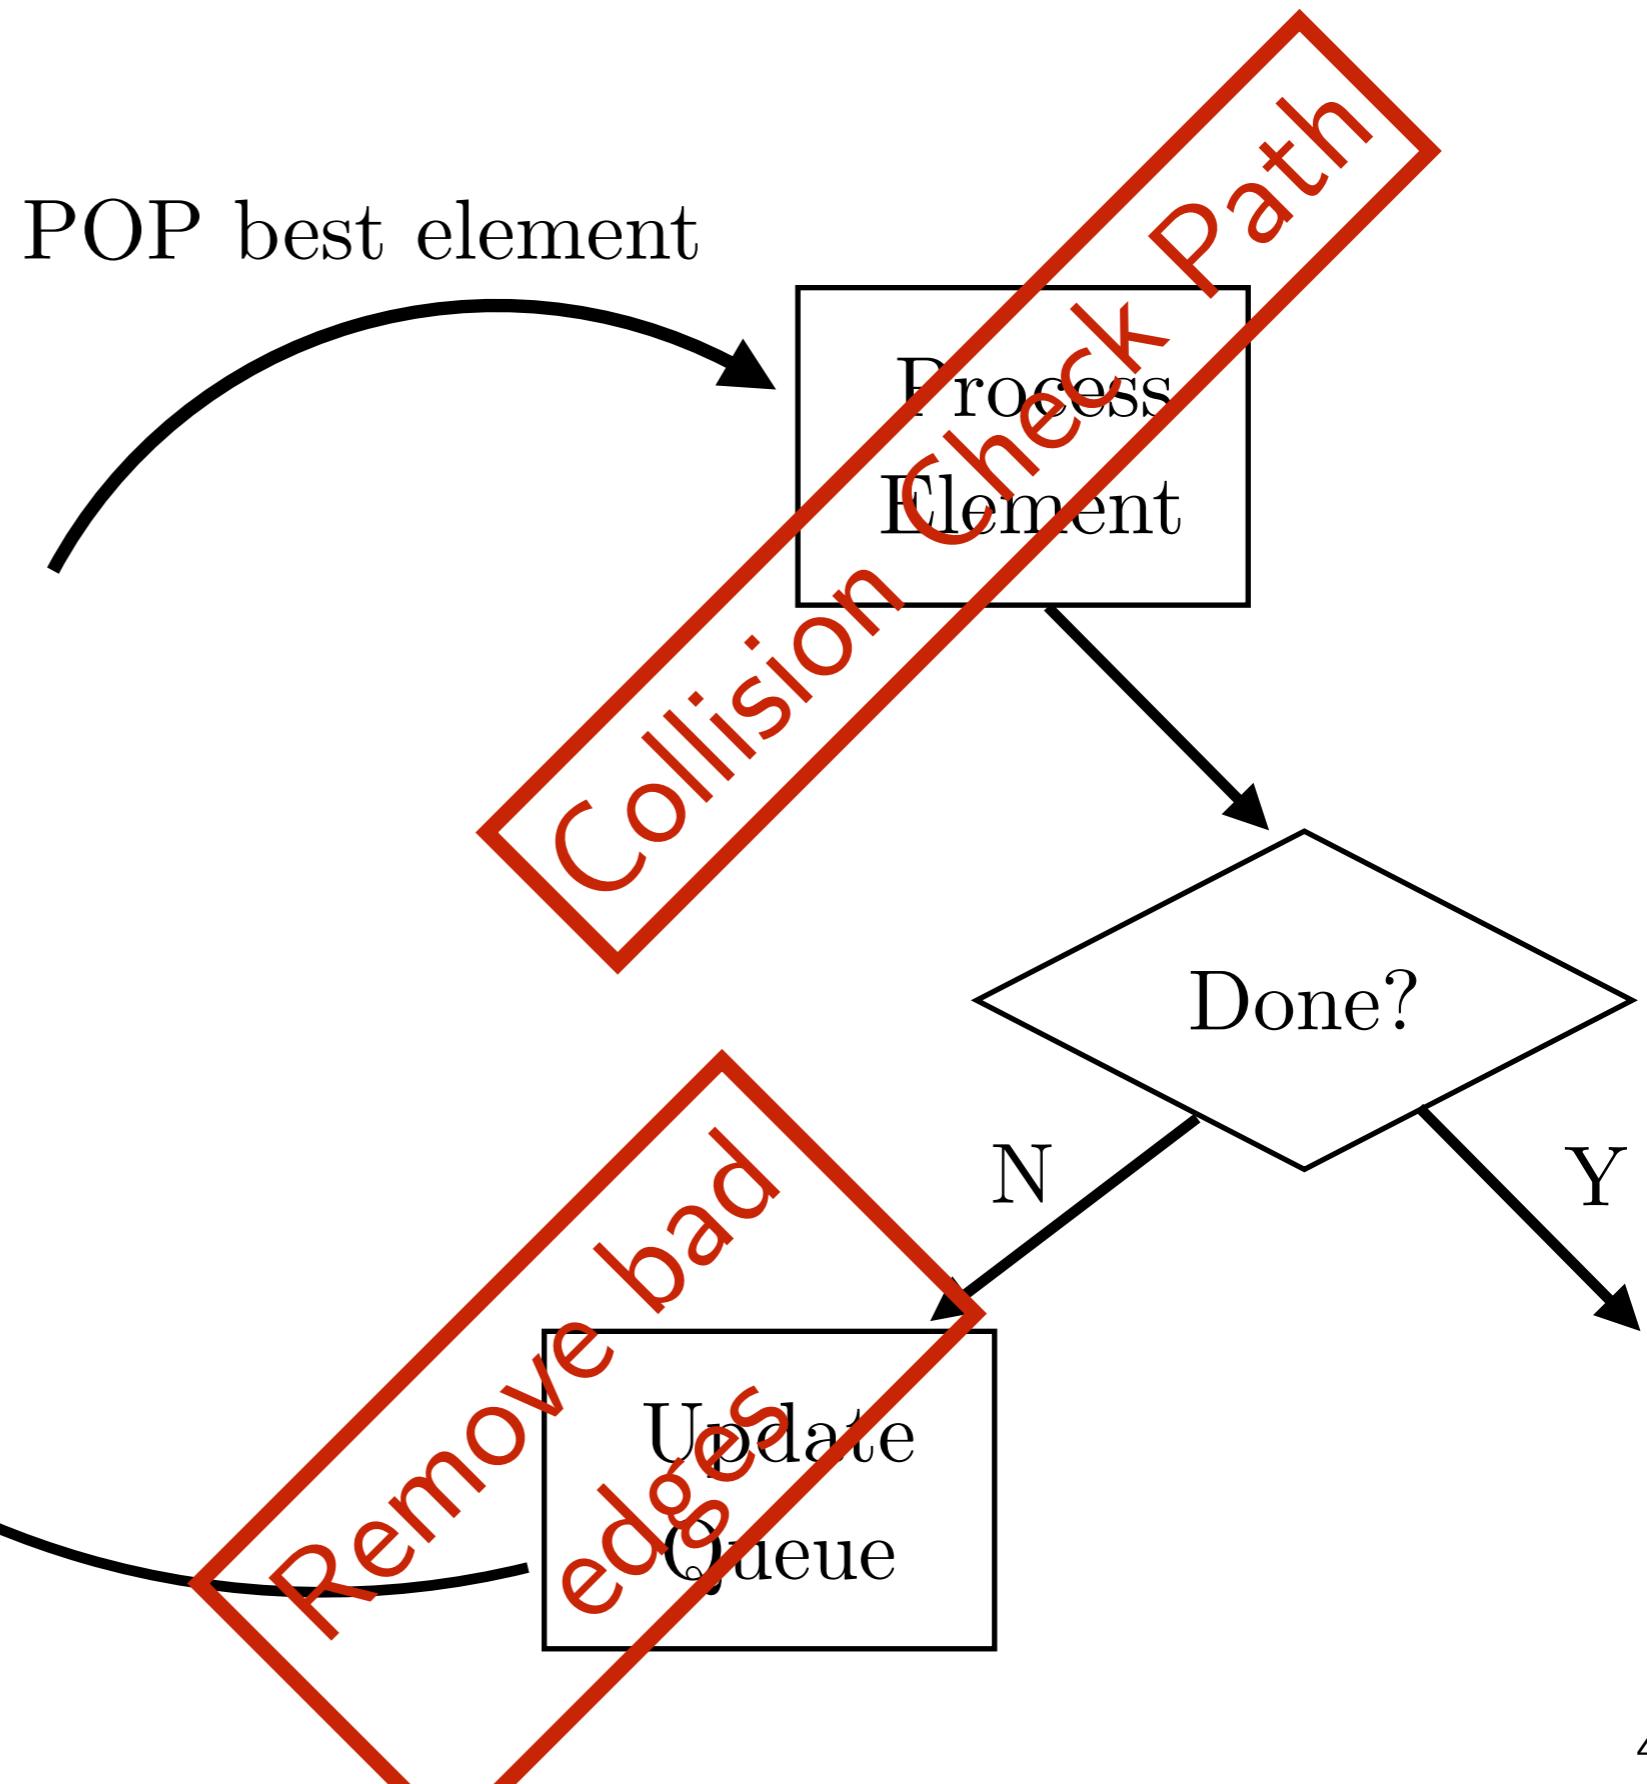
# LazySP

Greedy Best-first Search over Paths

To find the shortest path,  
eliminate all shorter paths!

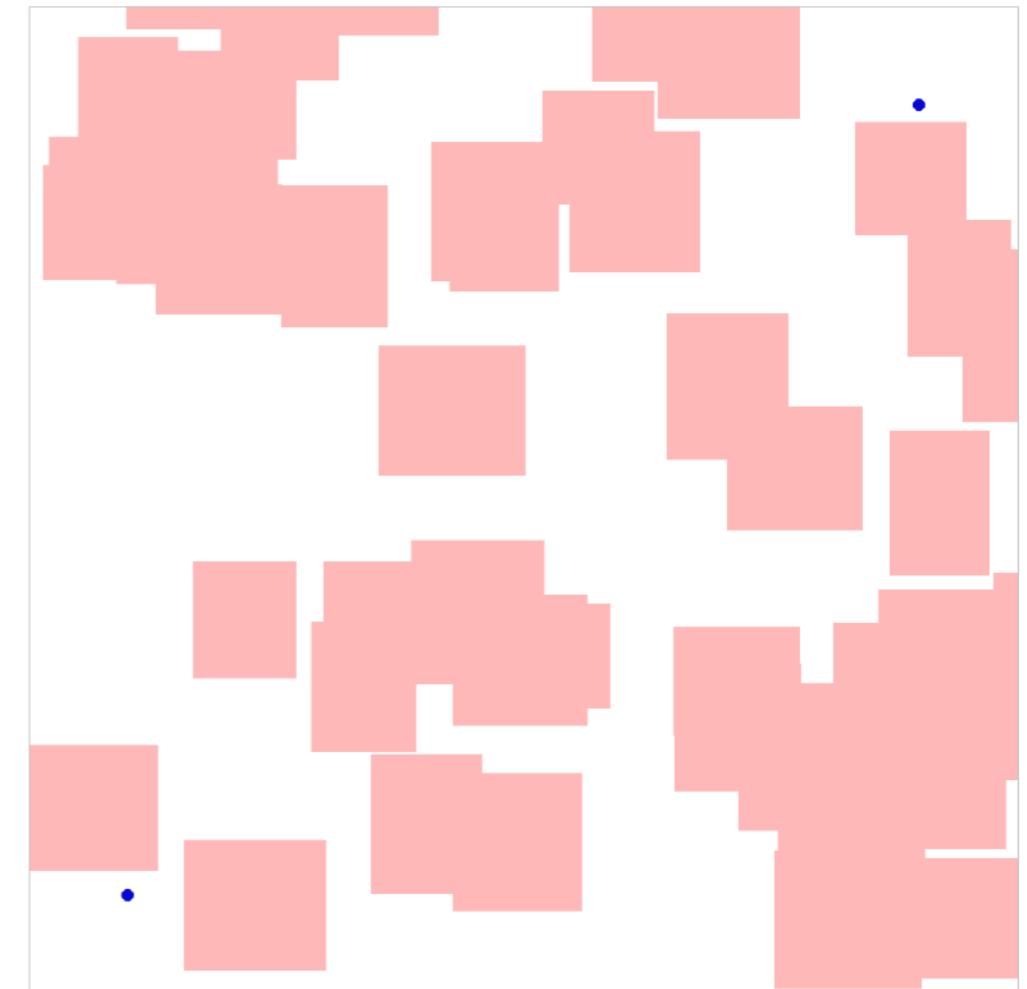
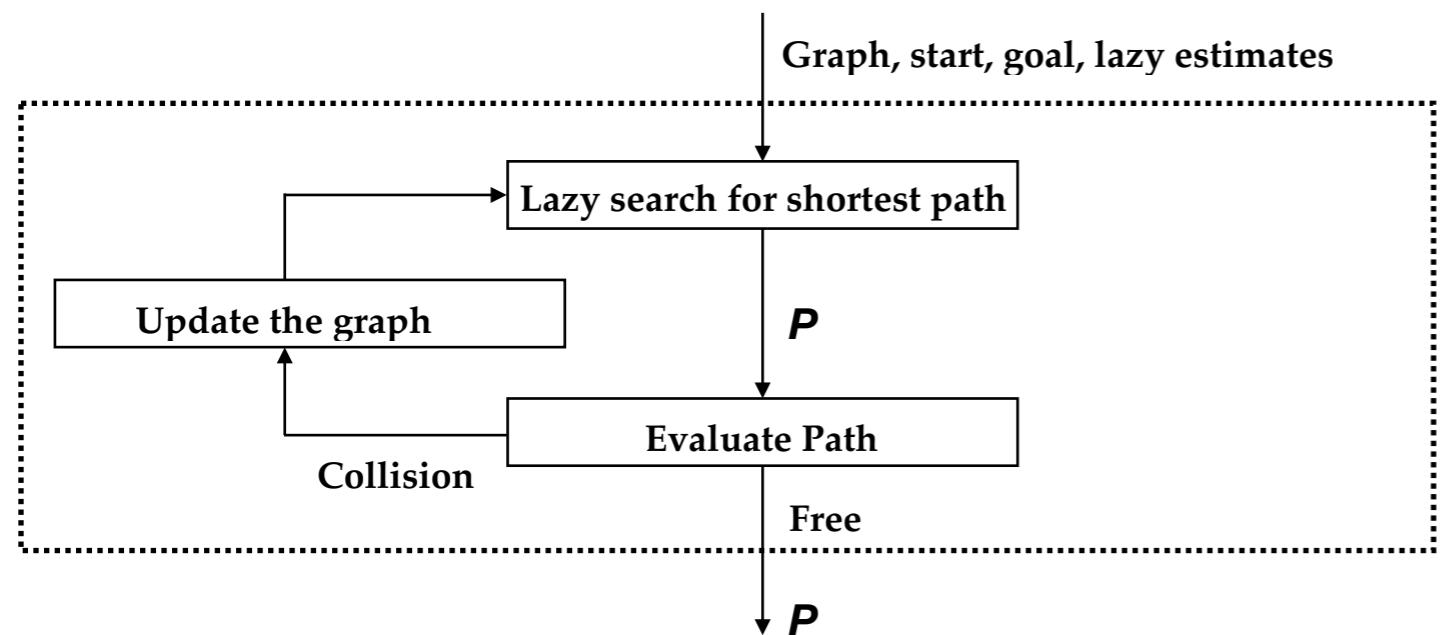
# Lazy A\* as BFS

Element	Ranking value
Element A	$f(A)$
Element B	$f(B)$



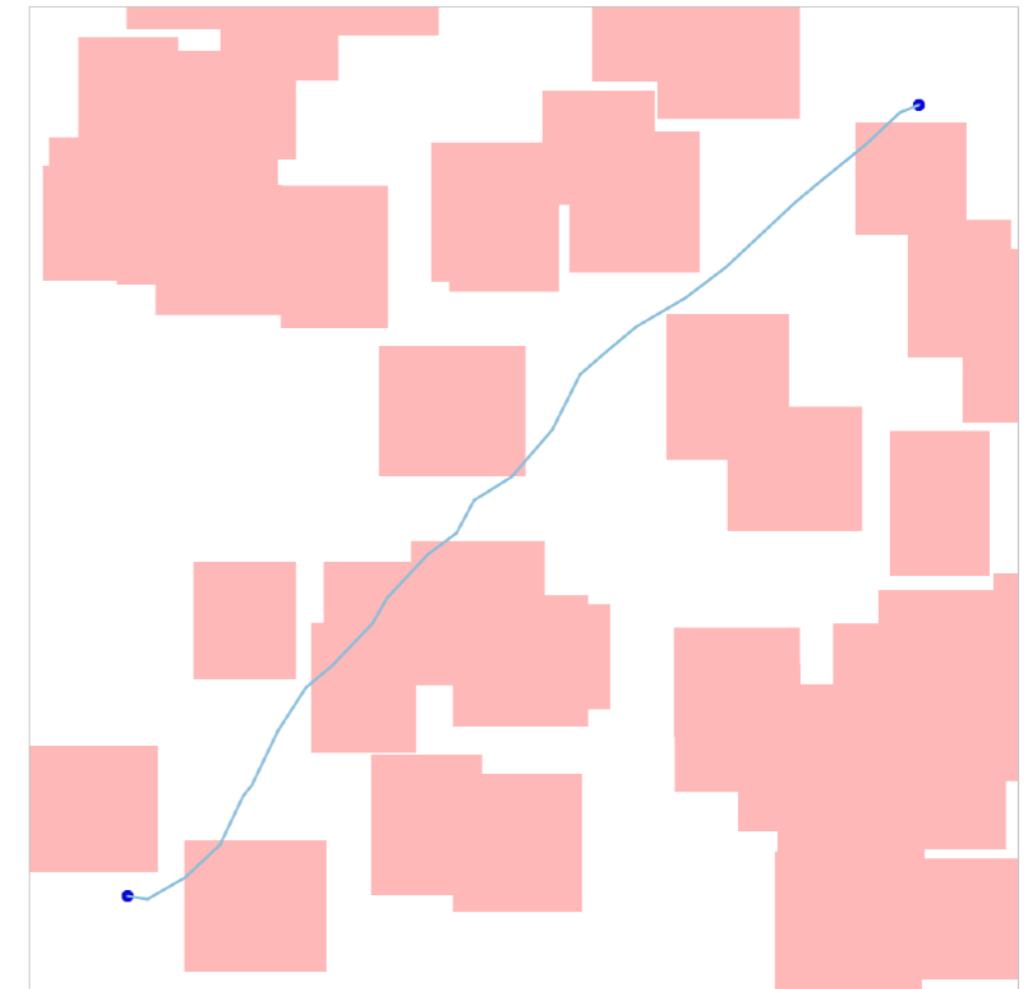
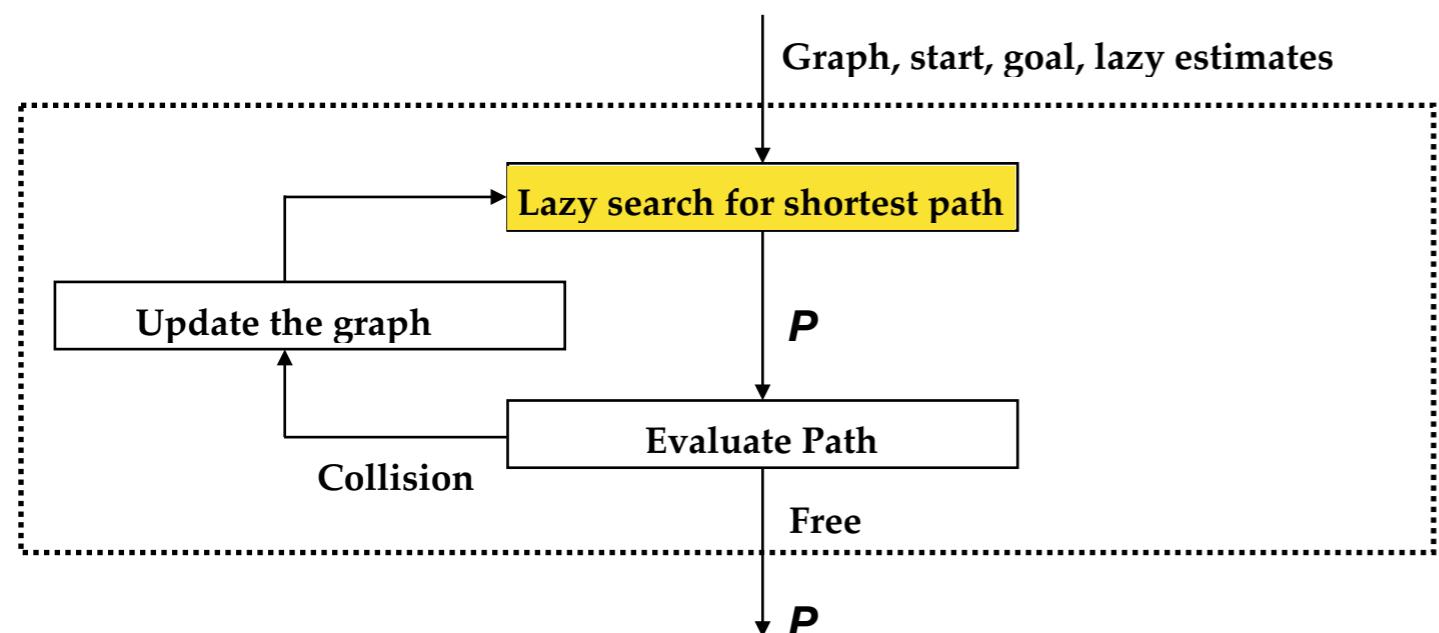
# LazySP

## Optimism Under Uncertainty



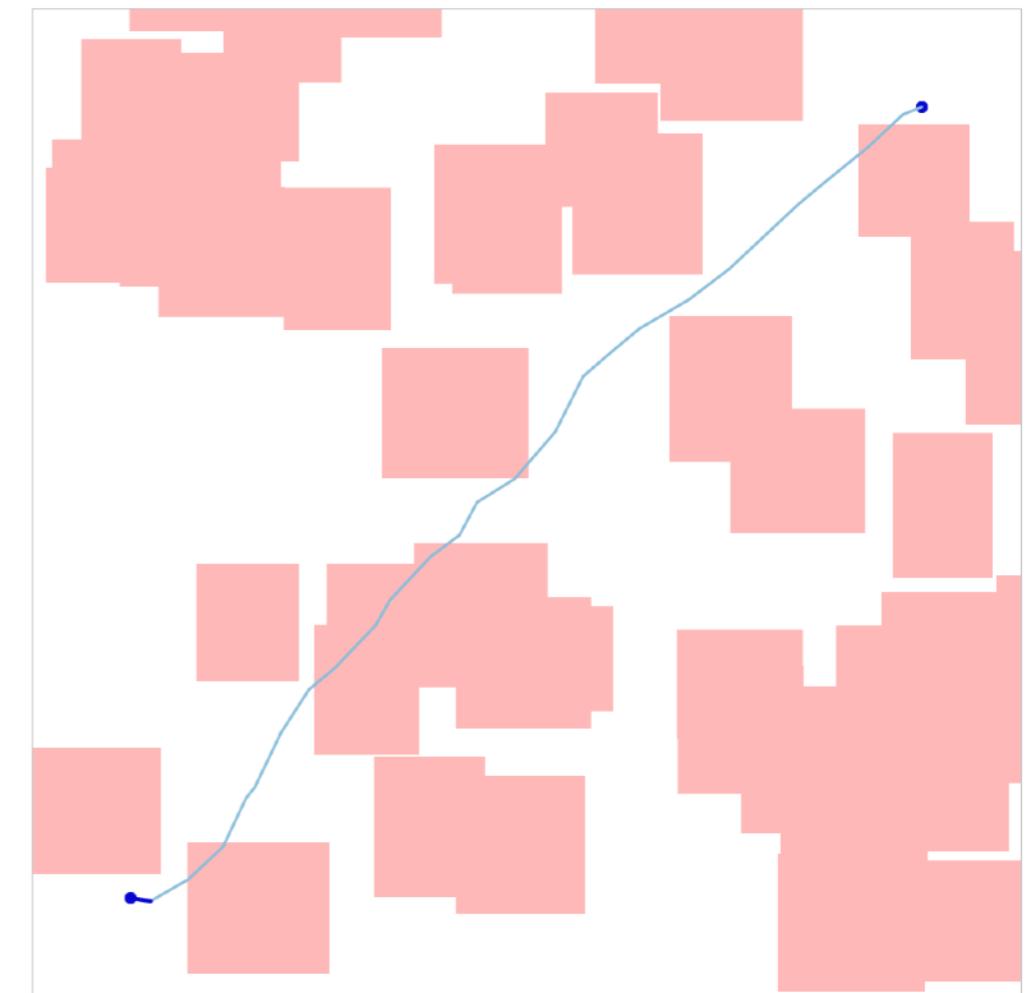
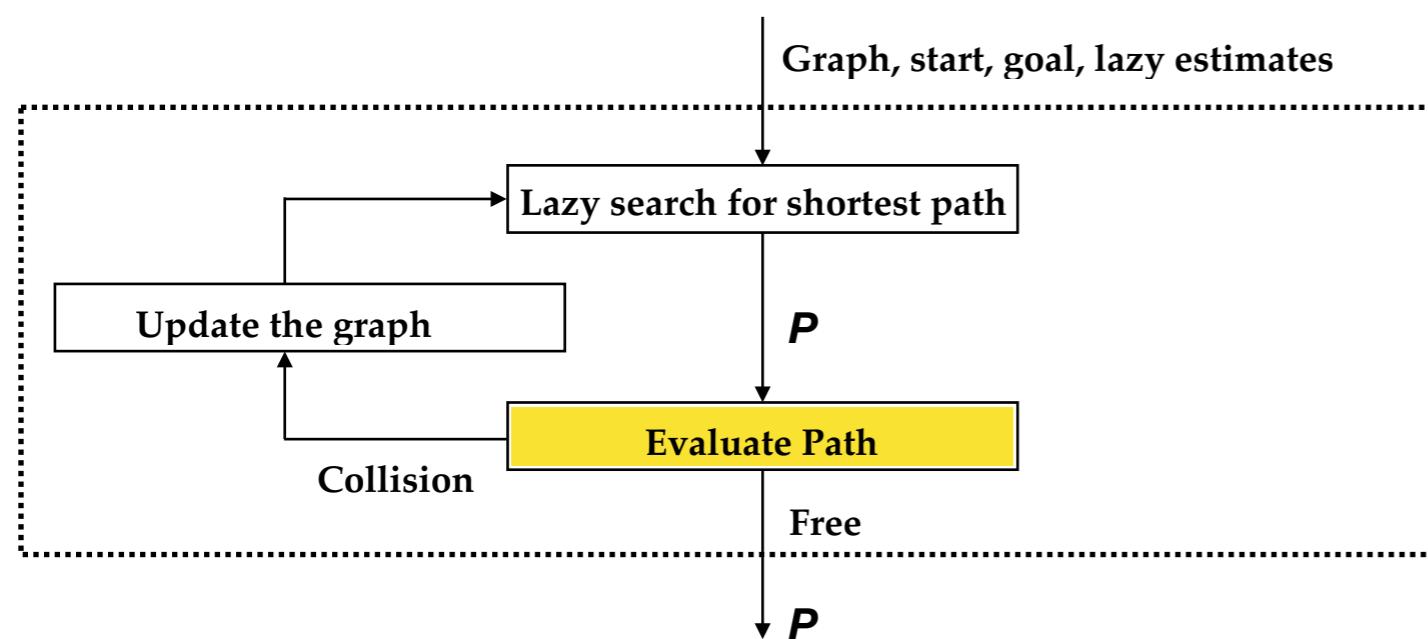
# LazySP

## Optimism Under Uncertainty



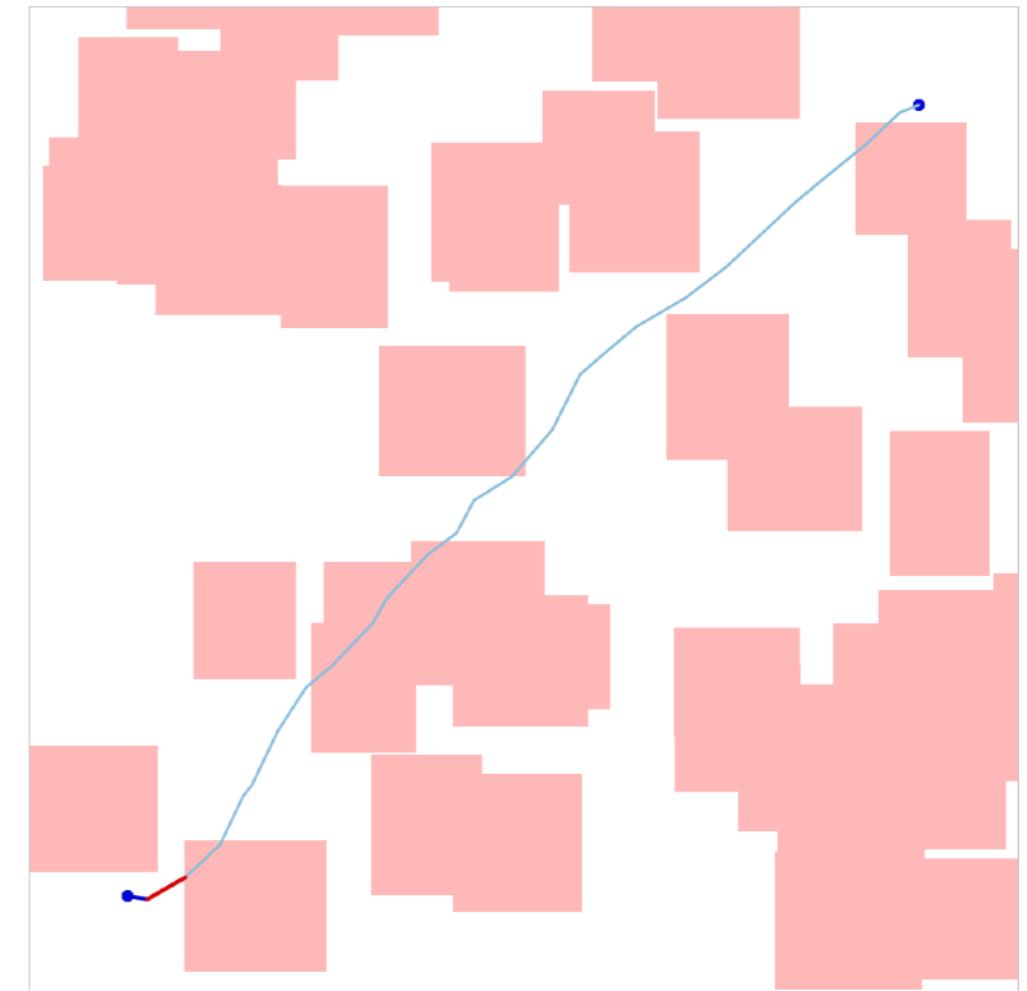
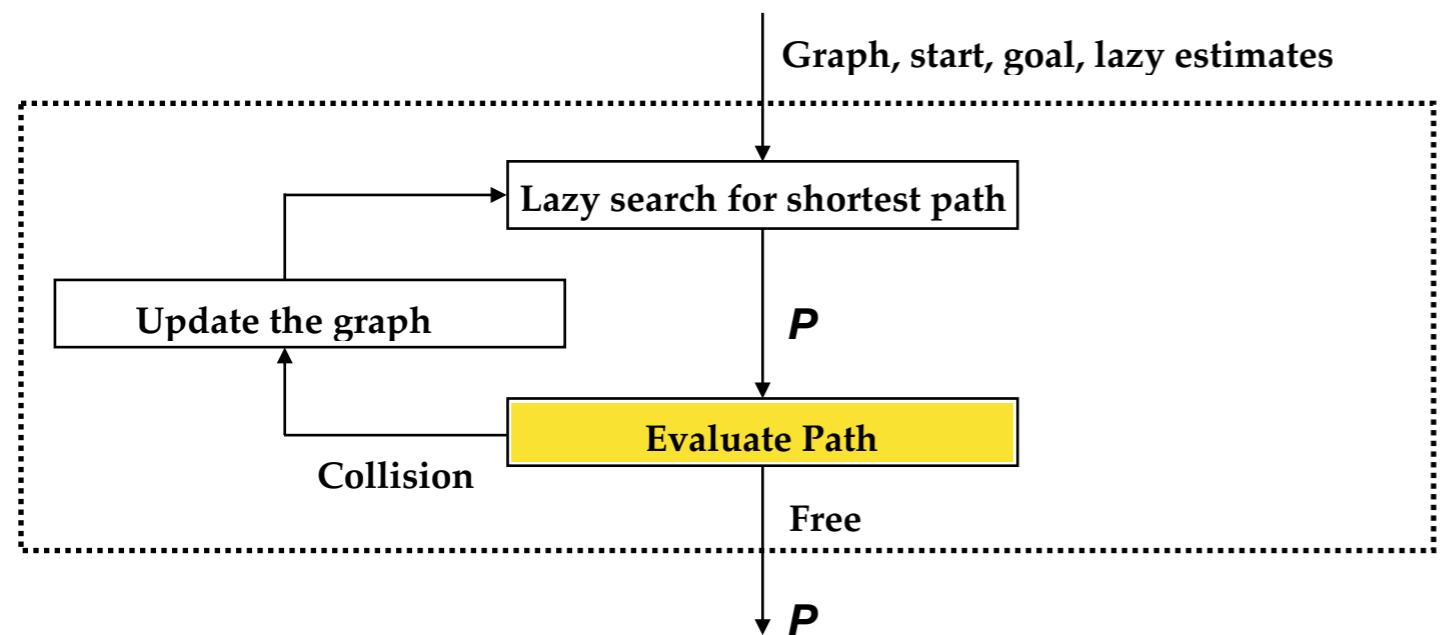
# LazySP

## Optimism Under Uncertainty



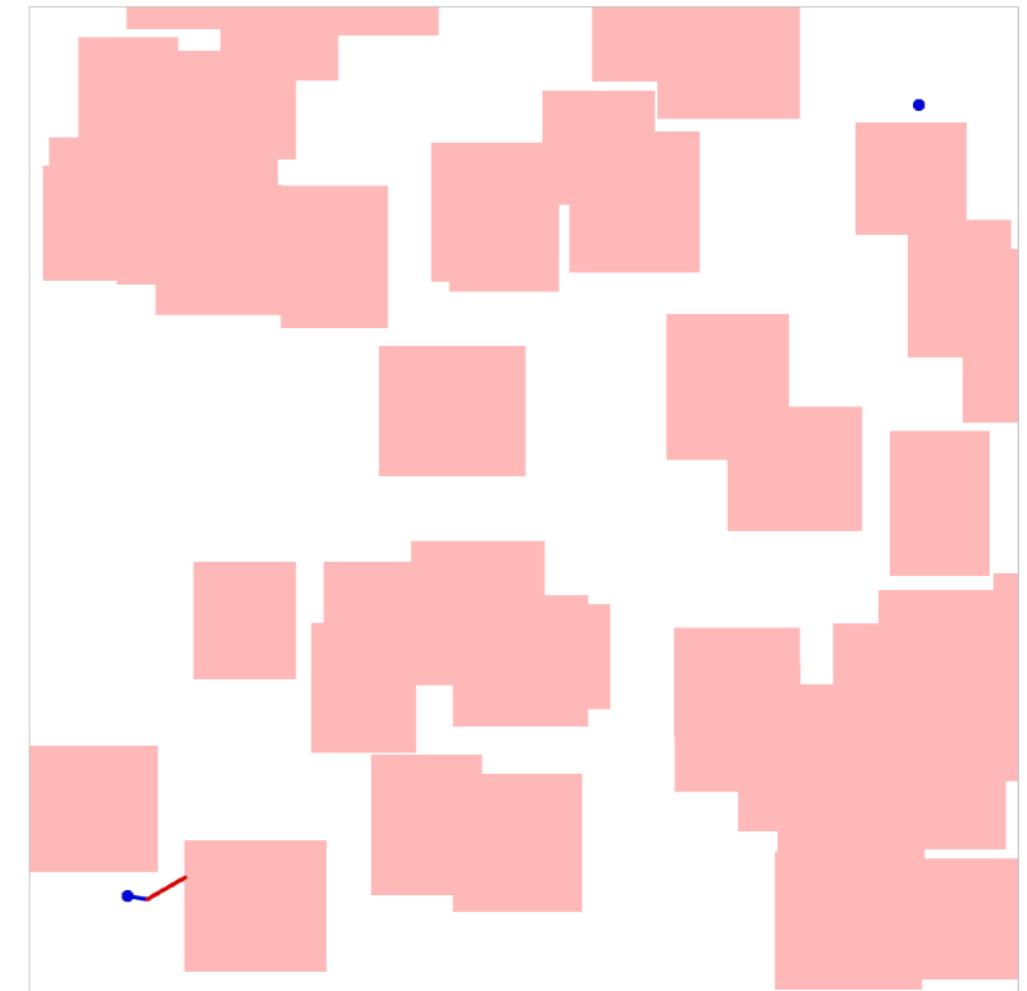
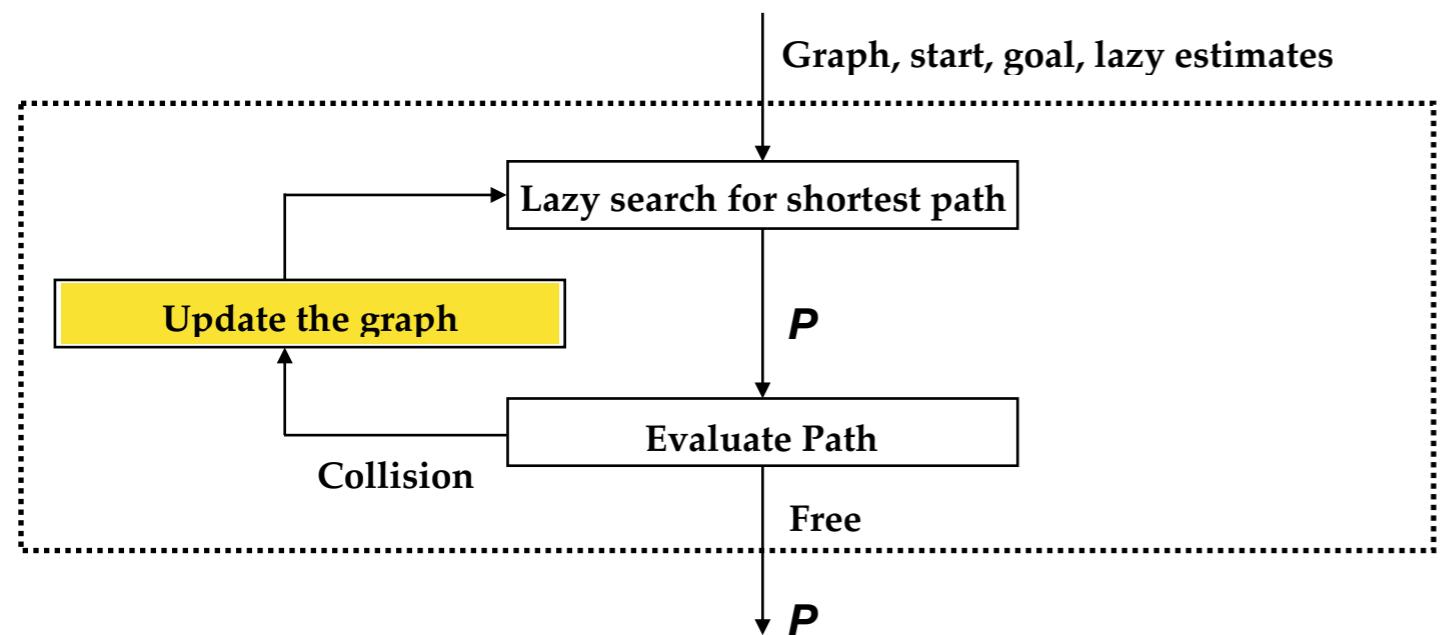
# LazySP

## Optimism Under Uncertainty



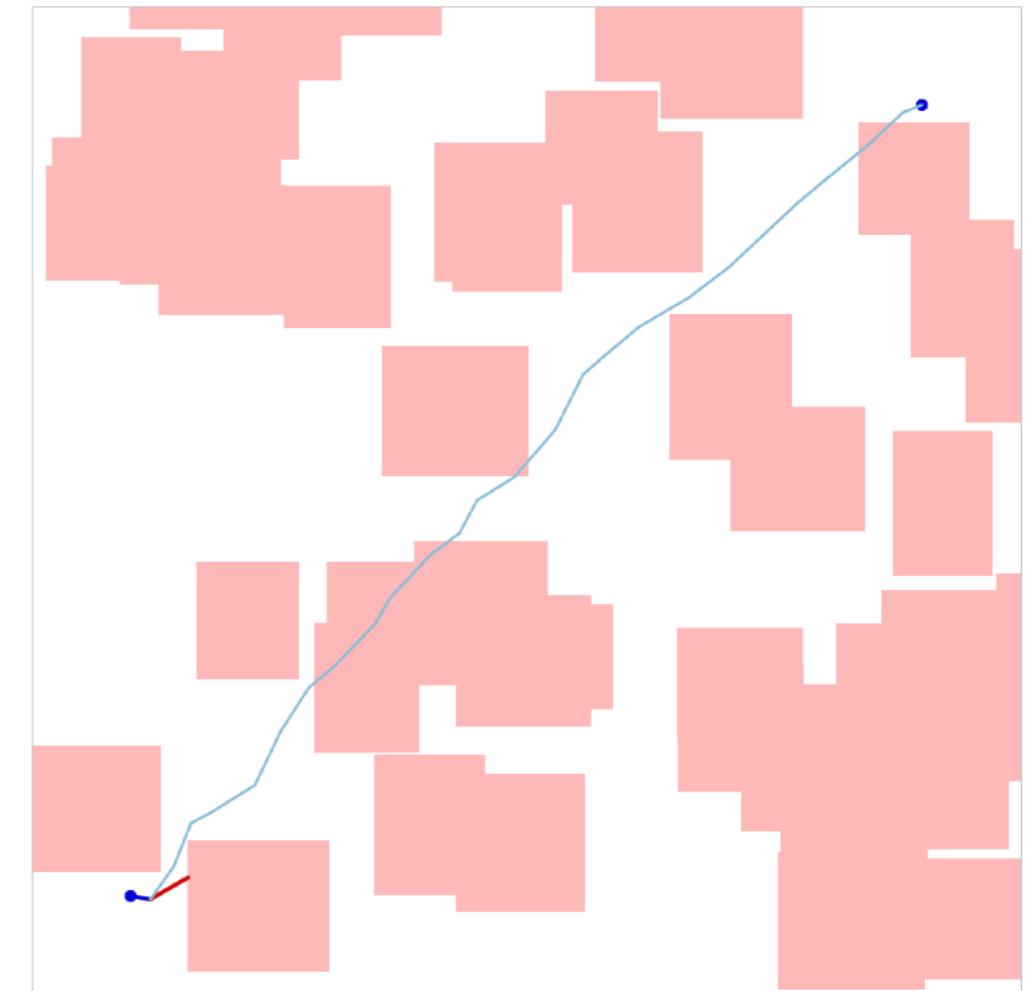
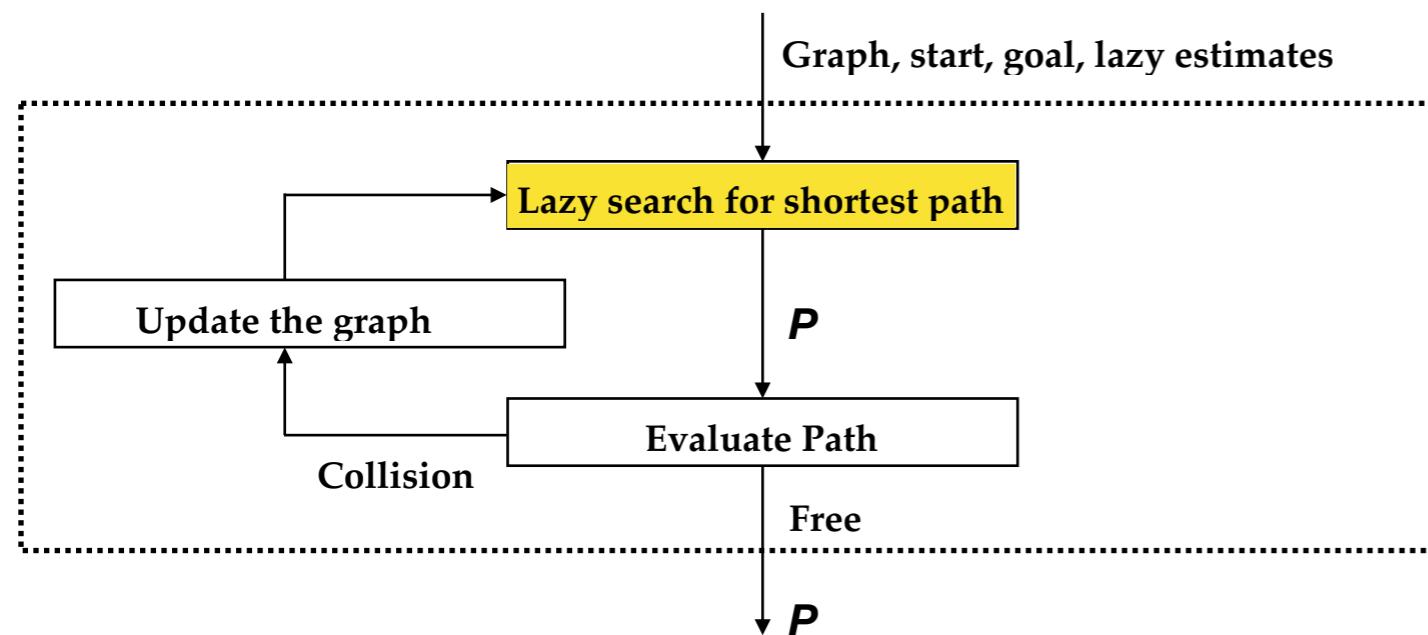
# LazySP

## Optimism Under Uncertainty



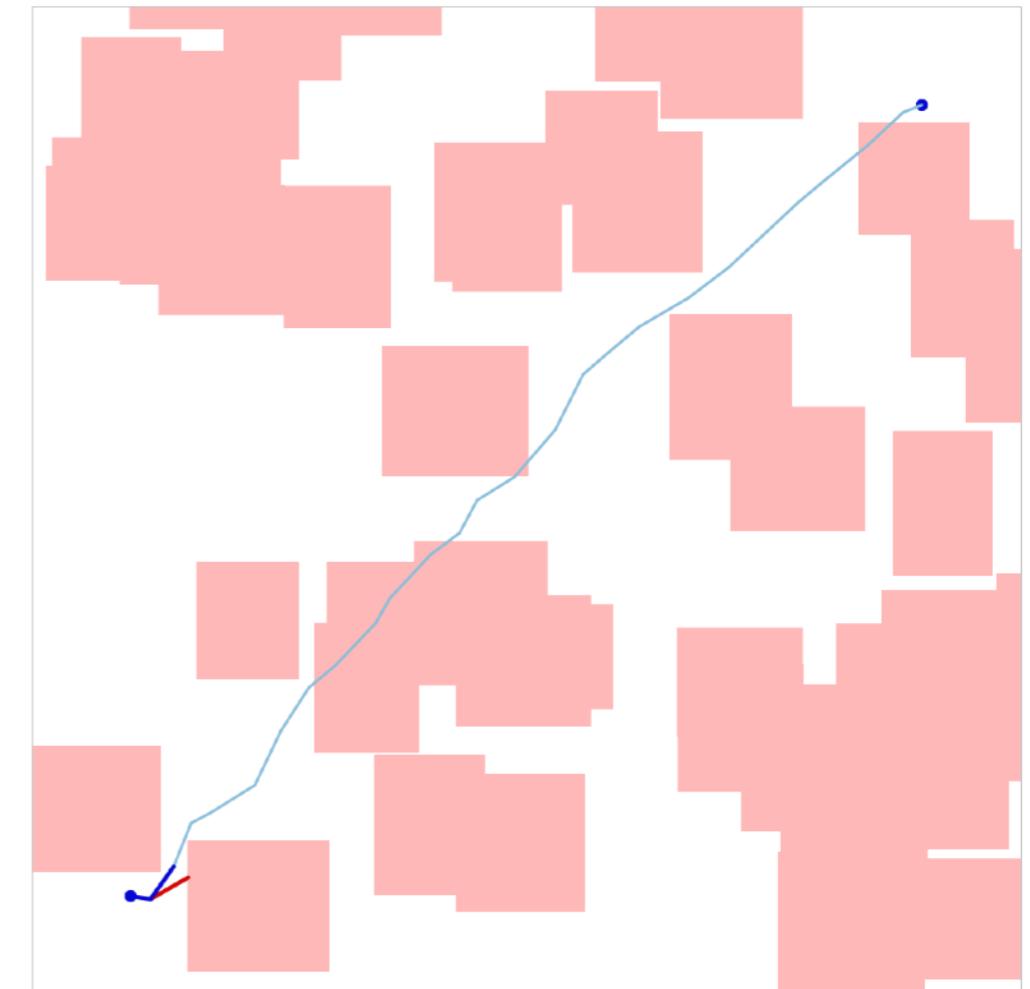
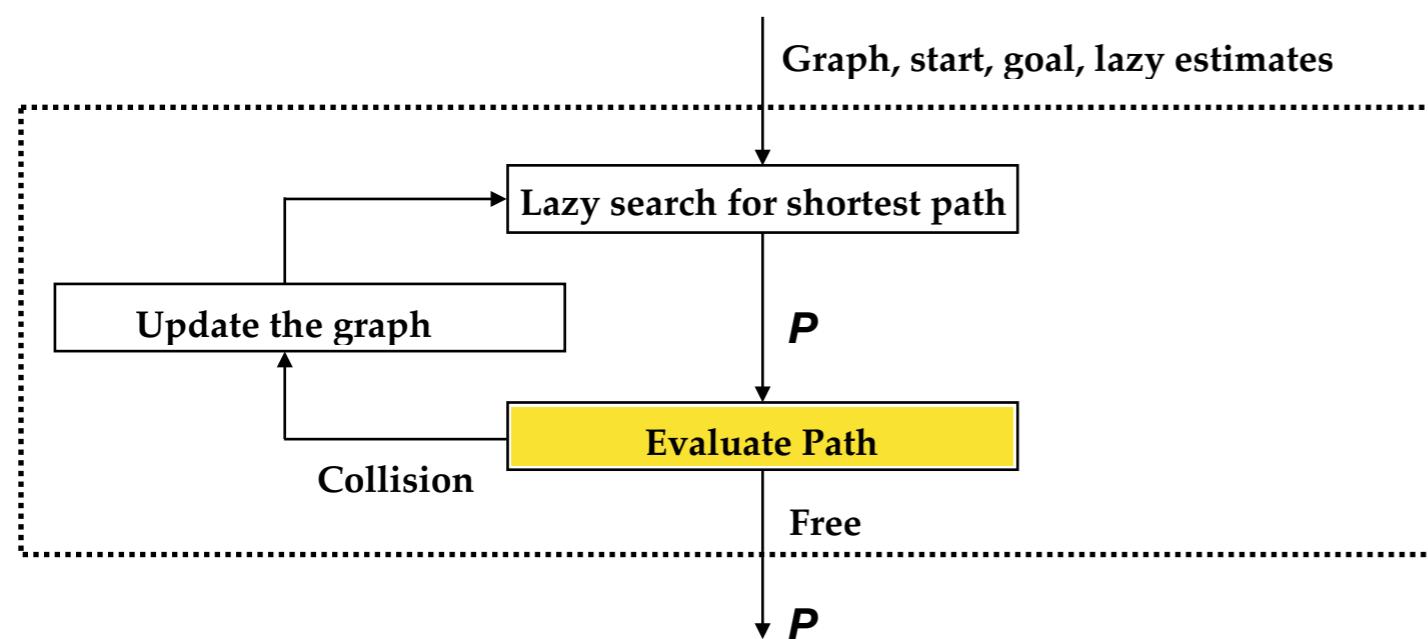
# LazySP

## Optimism Under Uncertainty



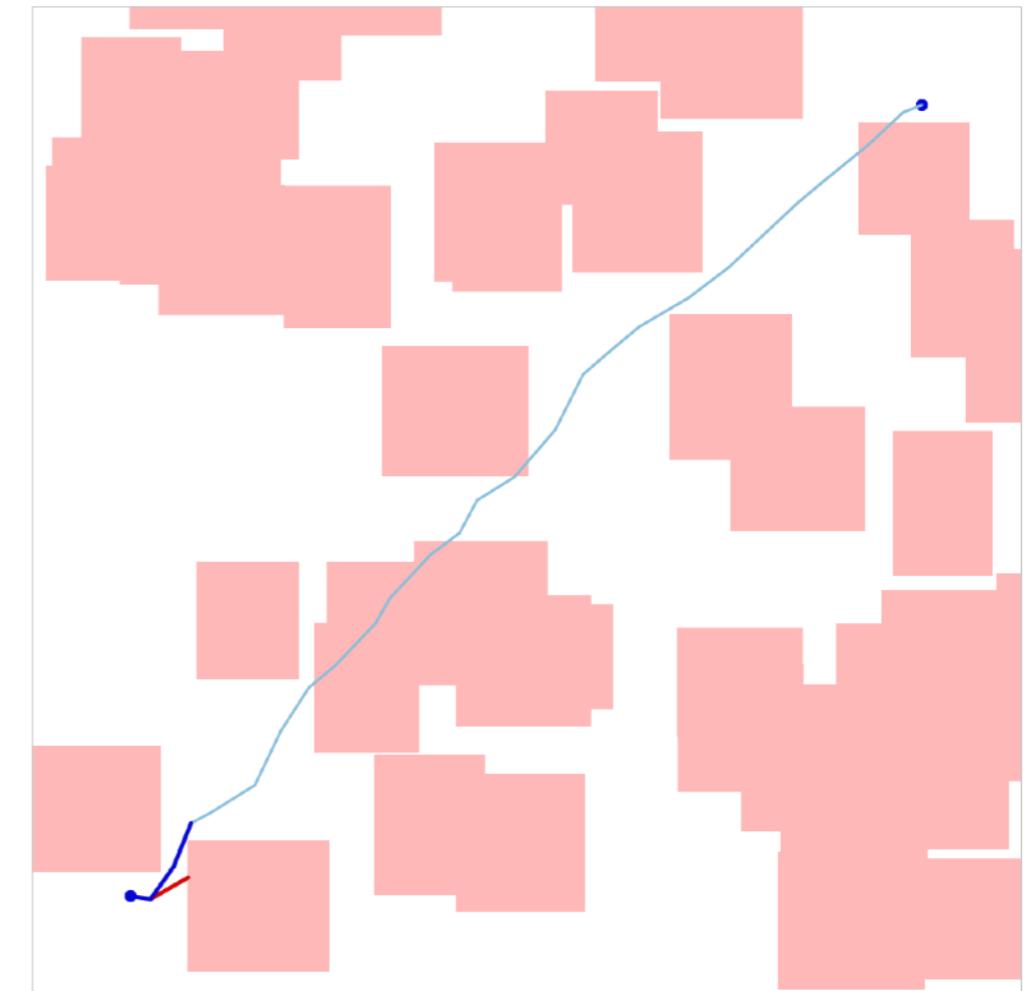
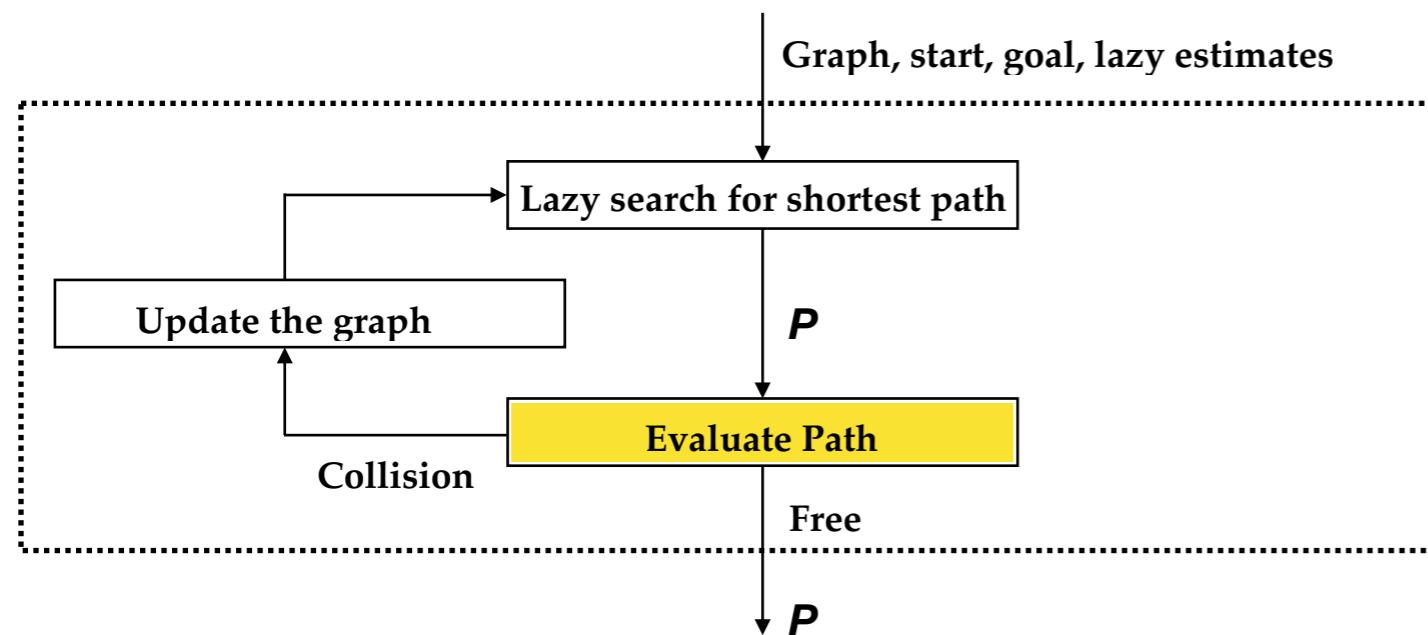
# LazySP

## Optimism Under Uncertainty



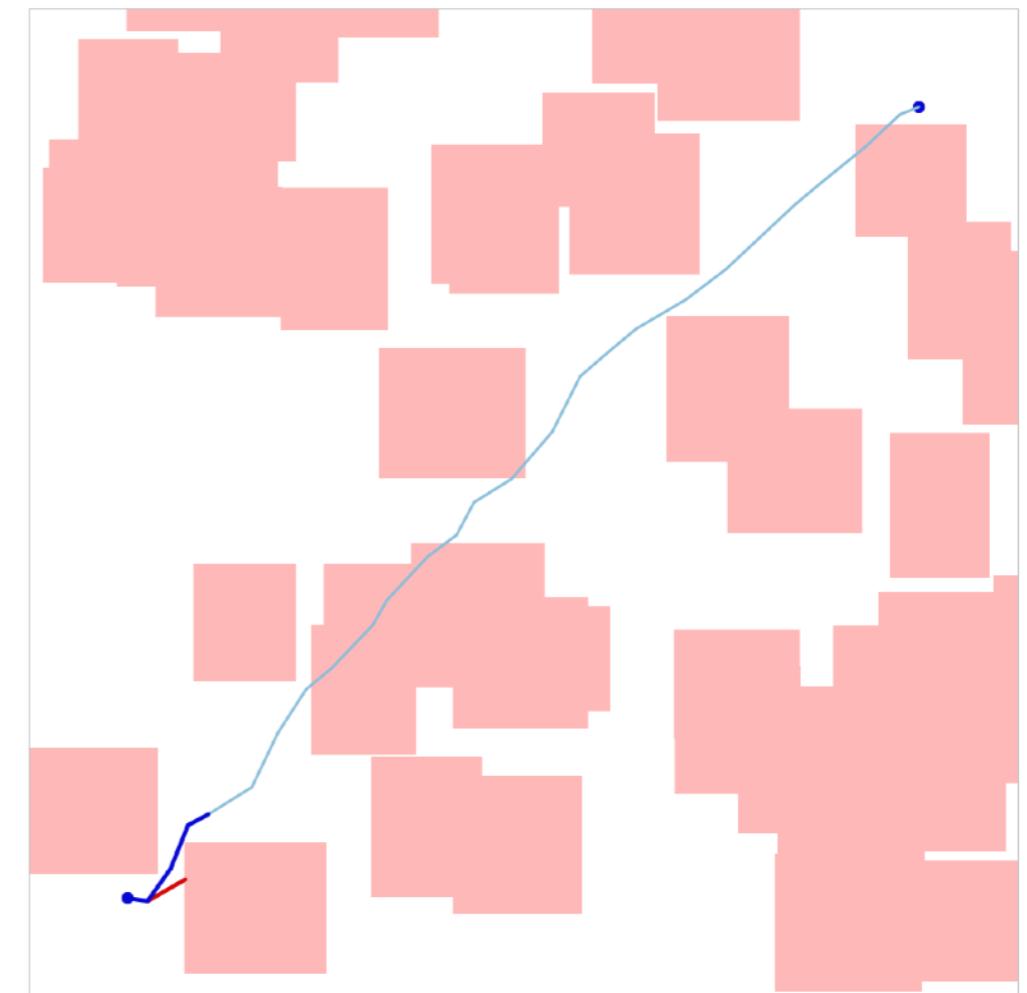
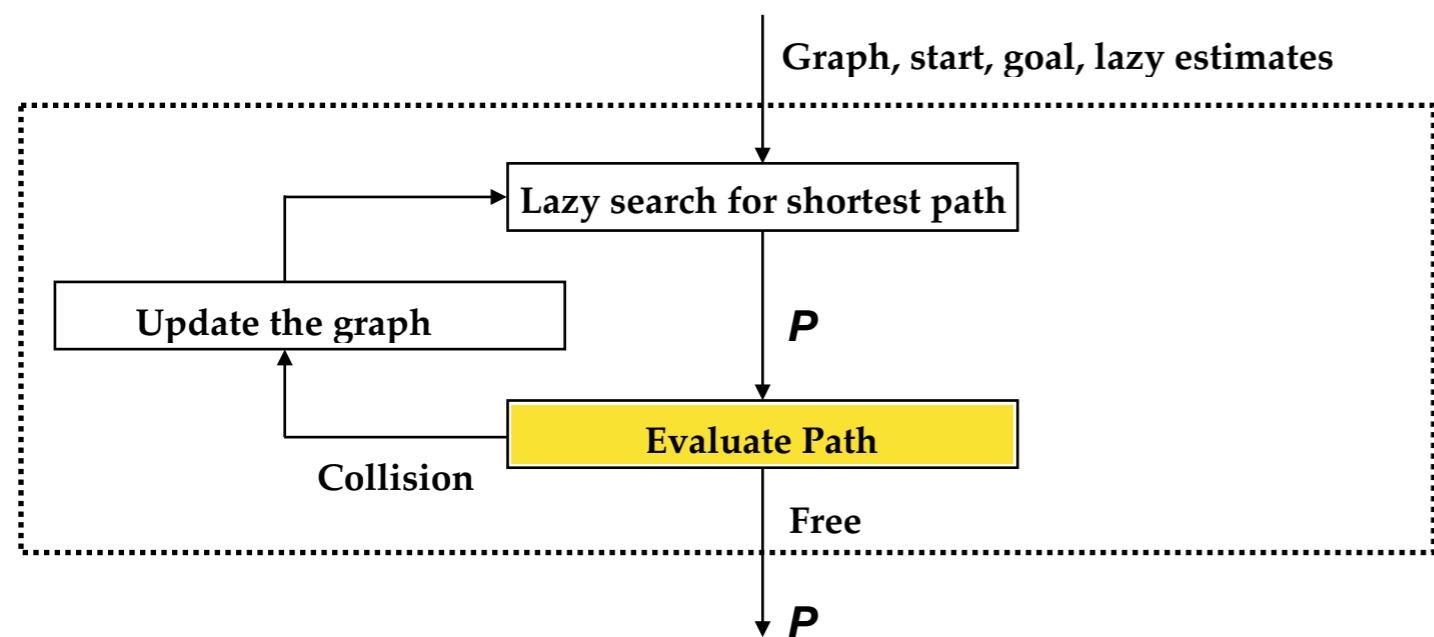
# LazySP

## Optimism Under Uncertainty



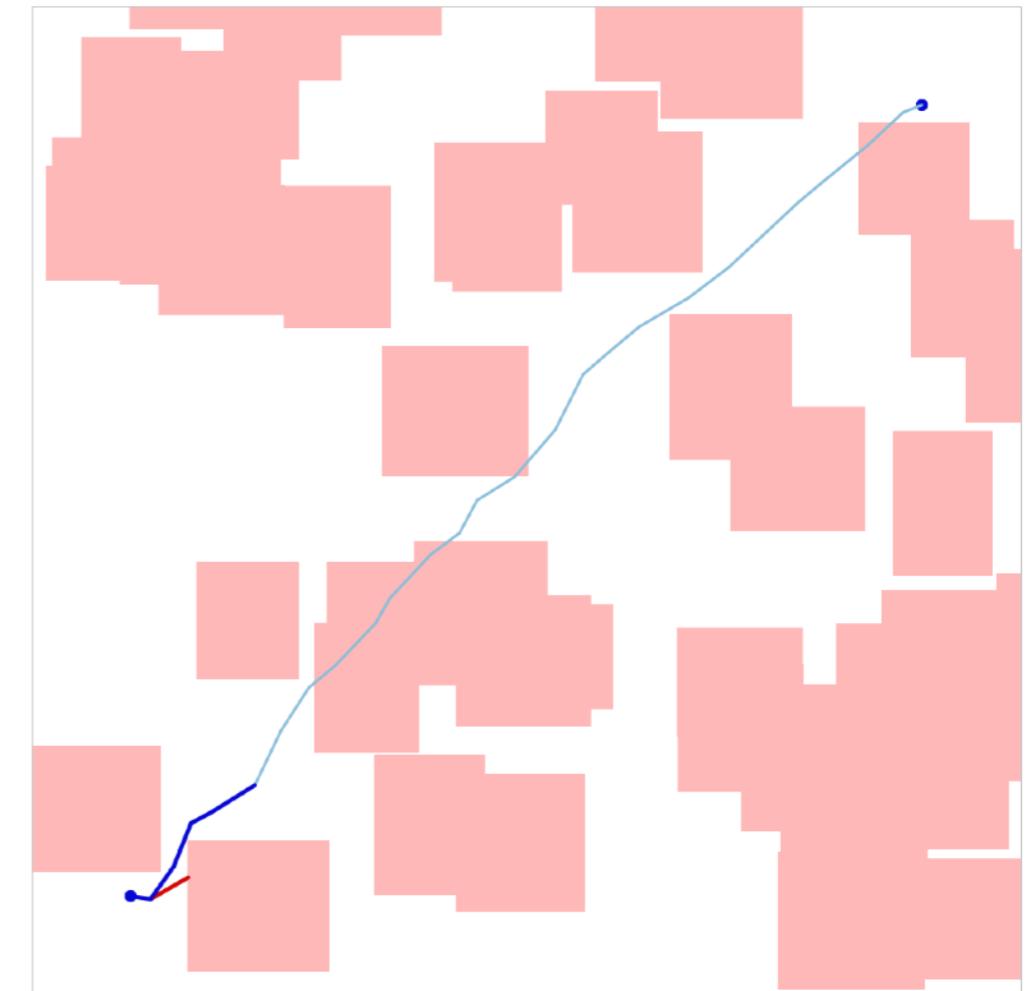
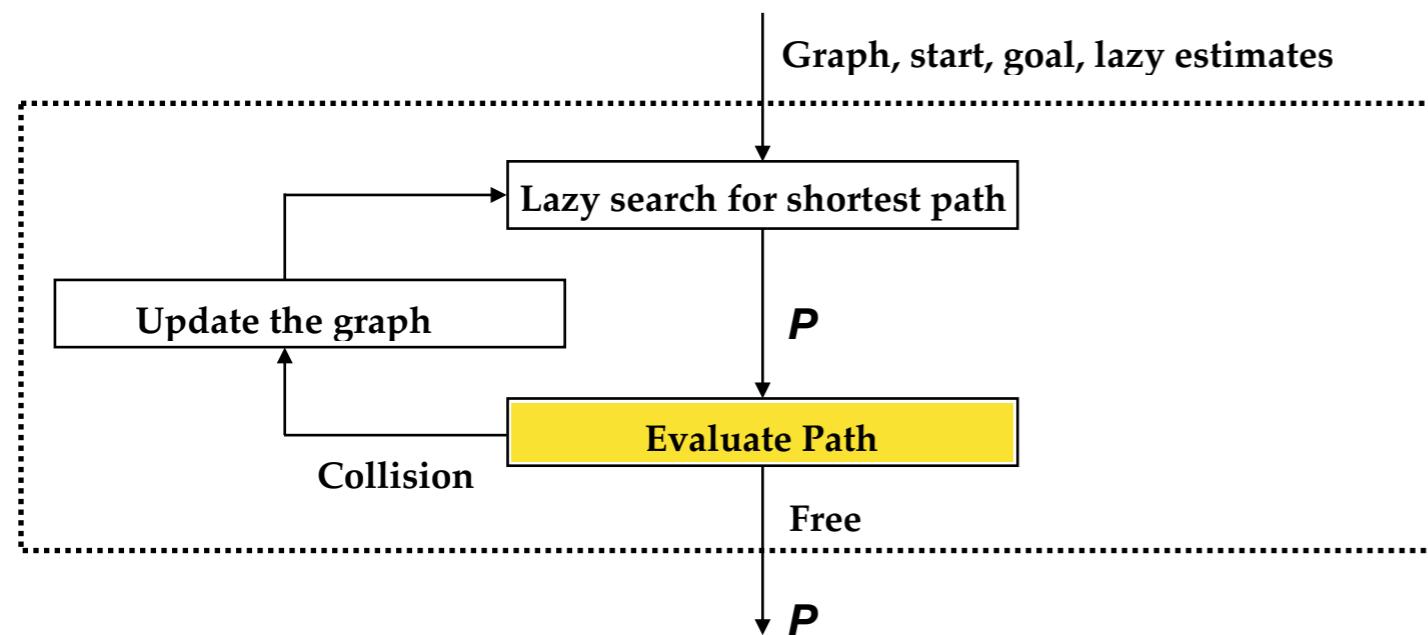
# LazySP

## Optimism Under Uncertainty



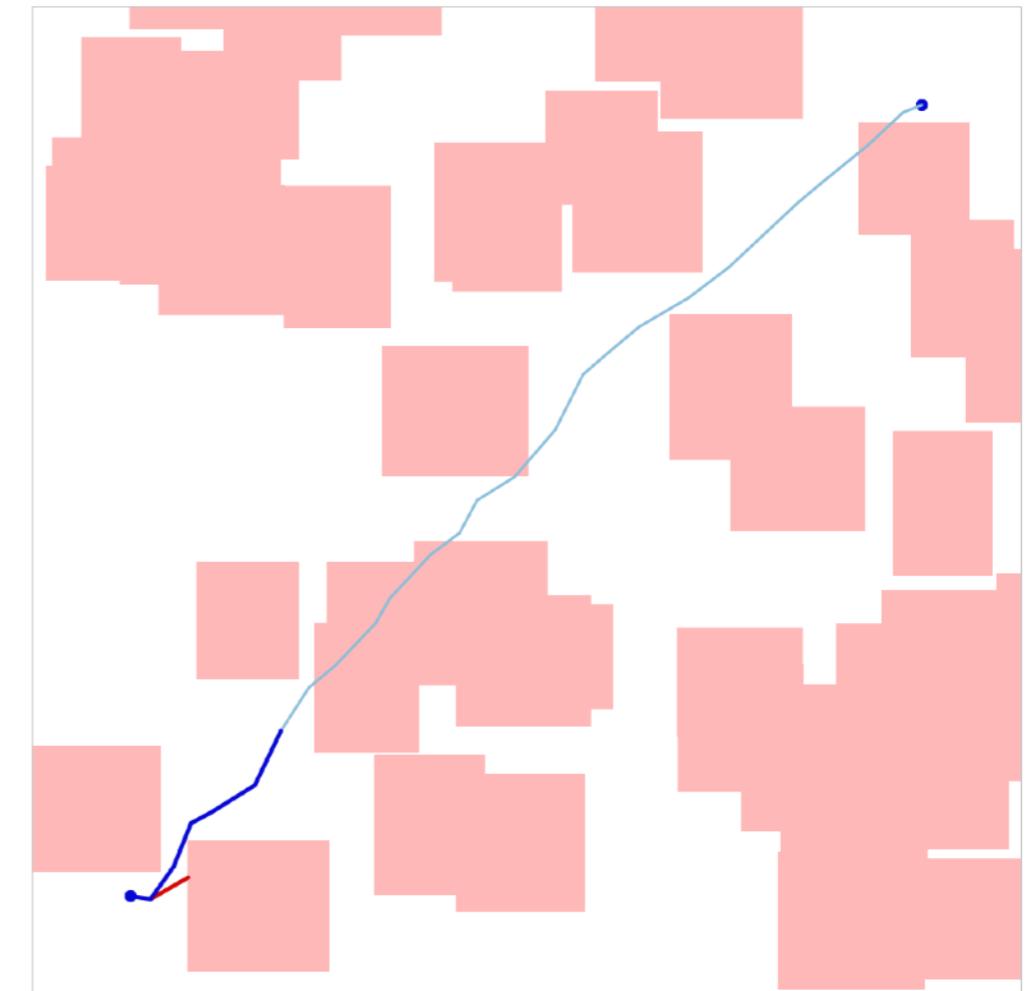
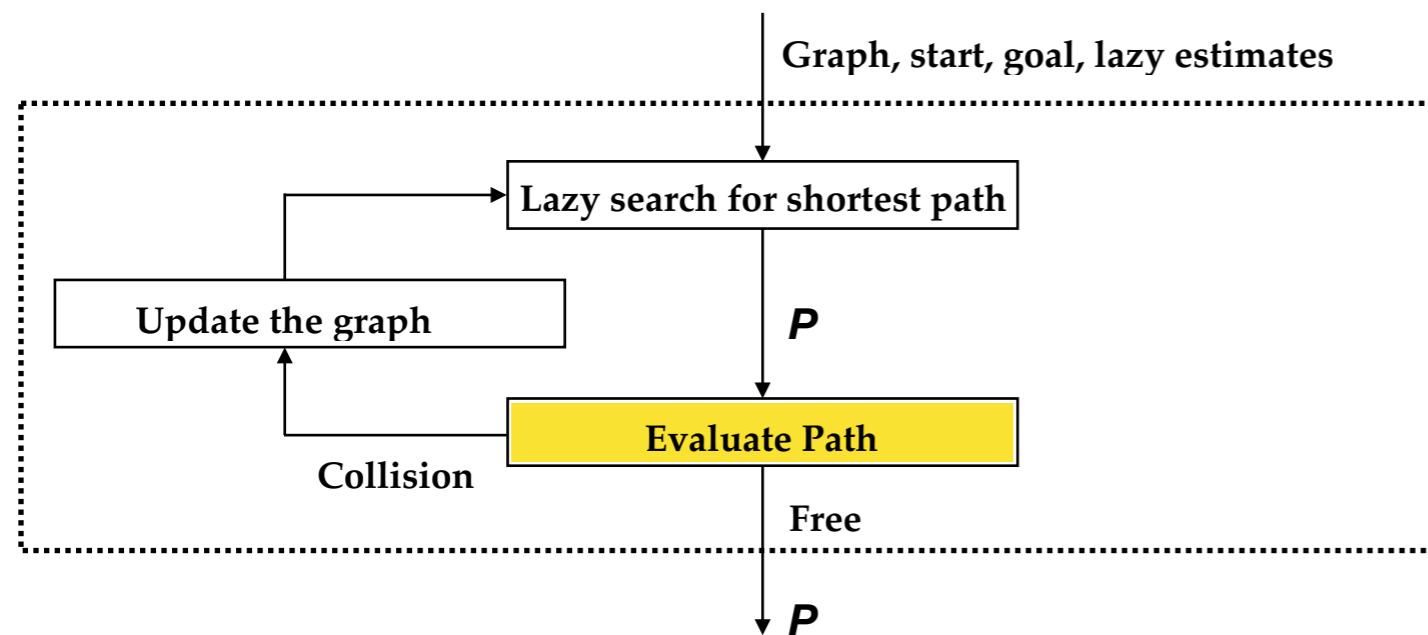
# LazySP

## Optimism Under Uncertainty



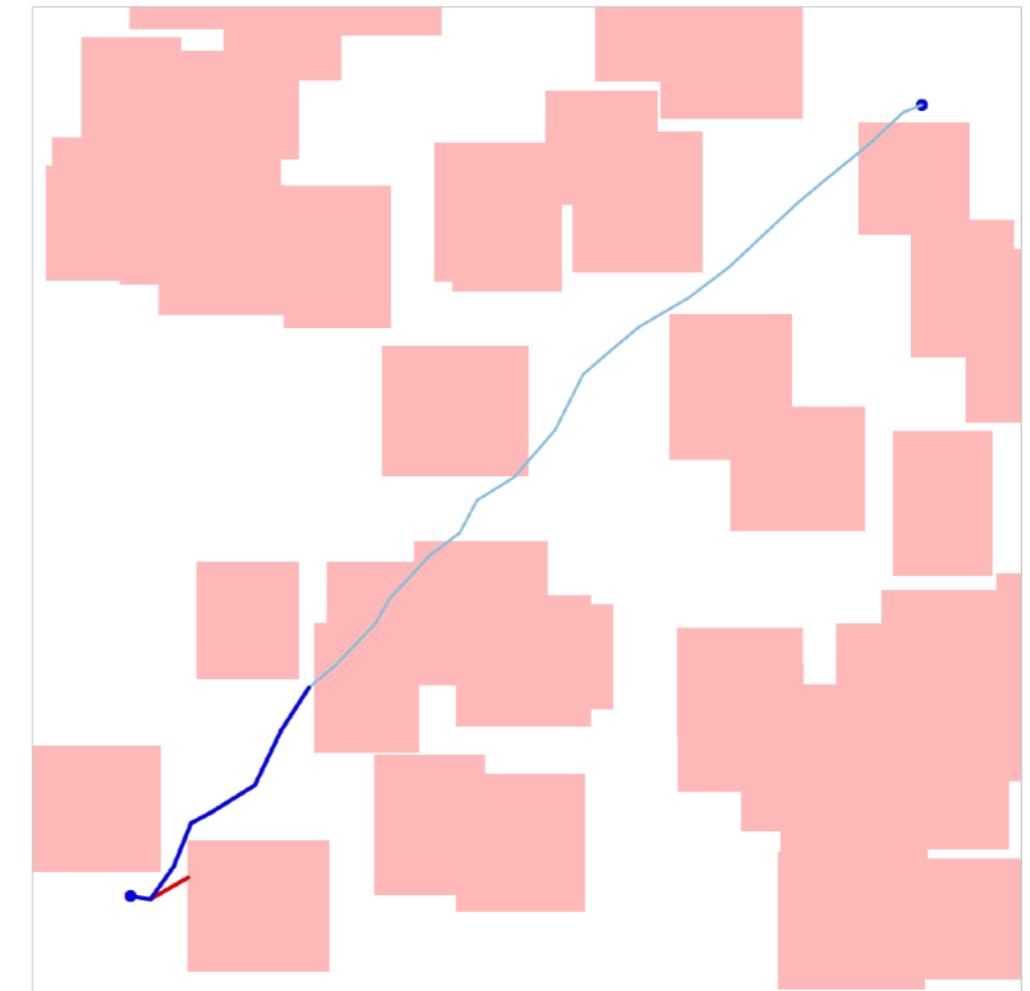
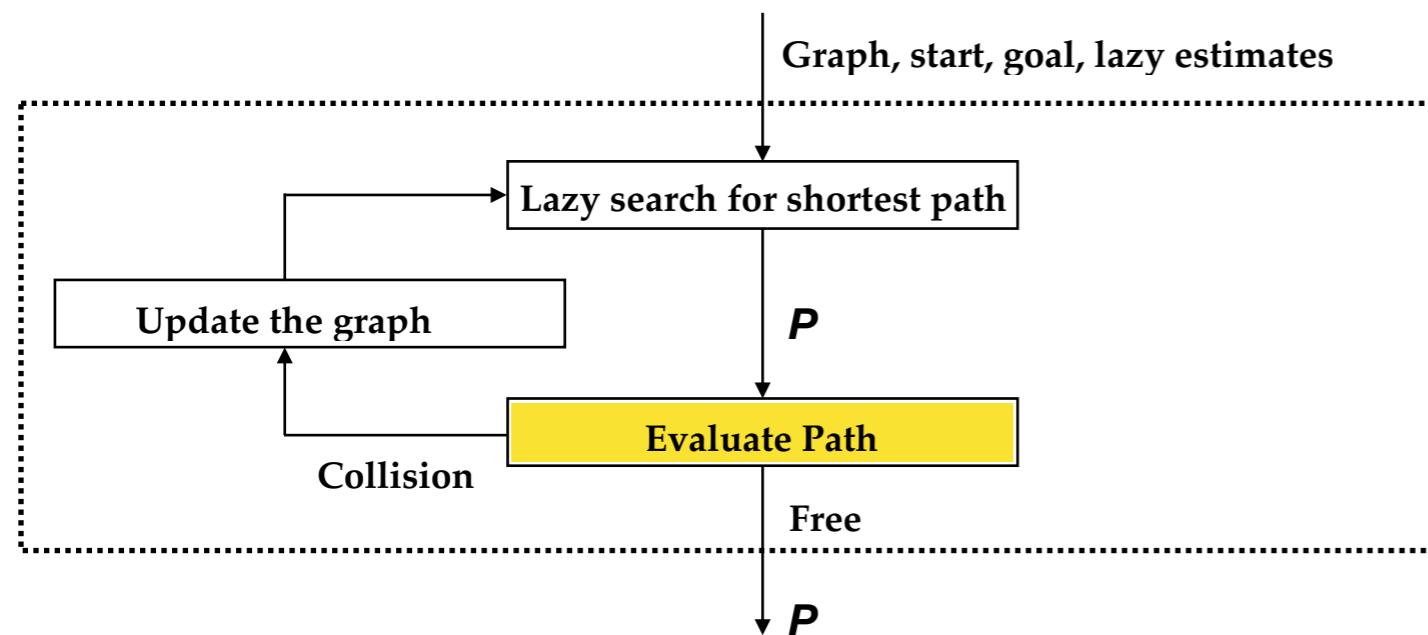
# LazySP

## Optimism Under Uncertainty



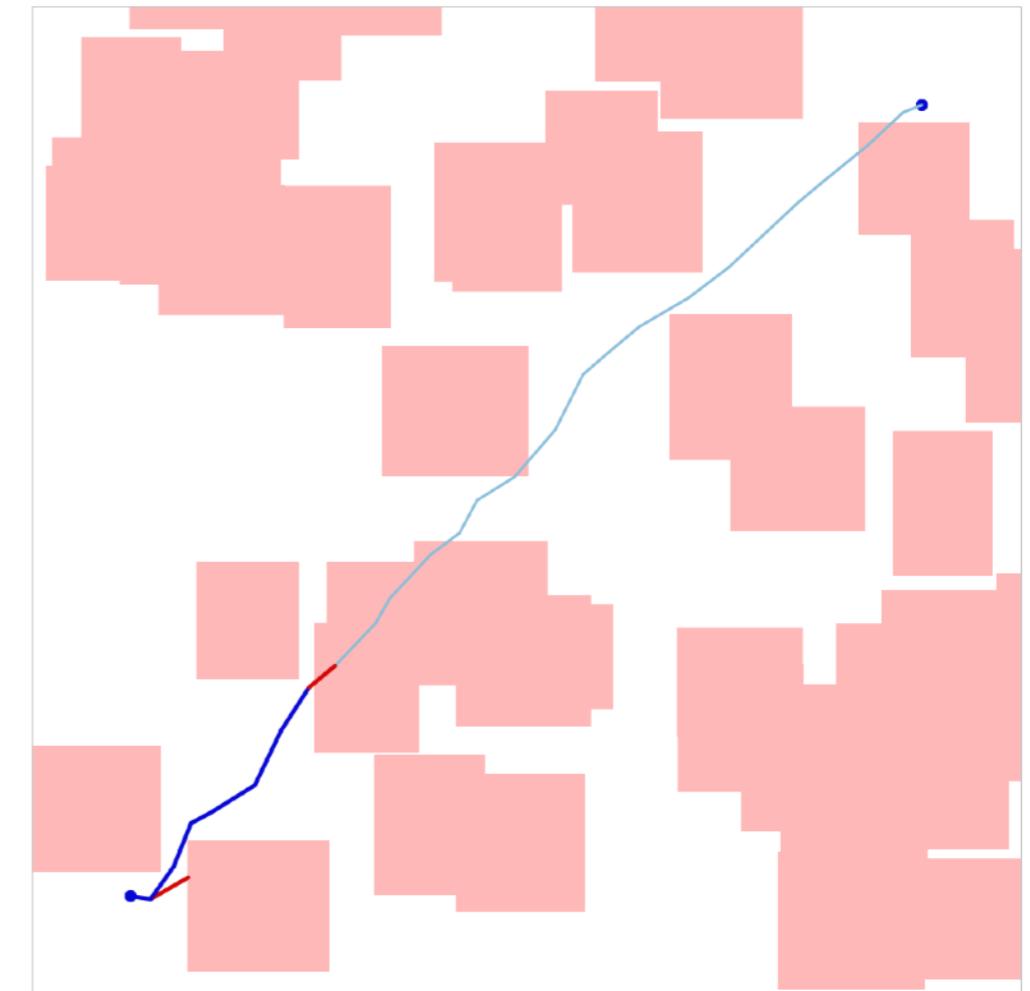
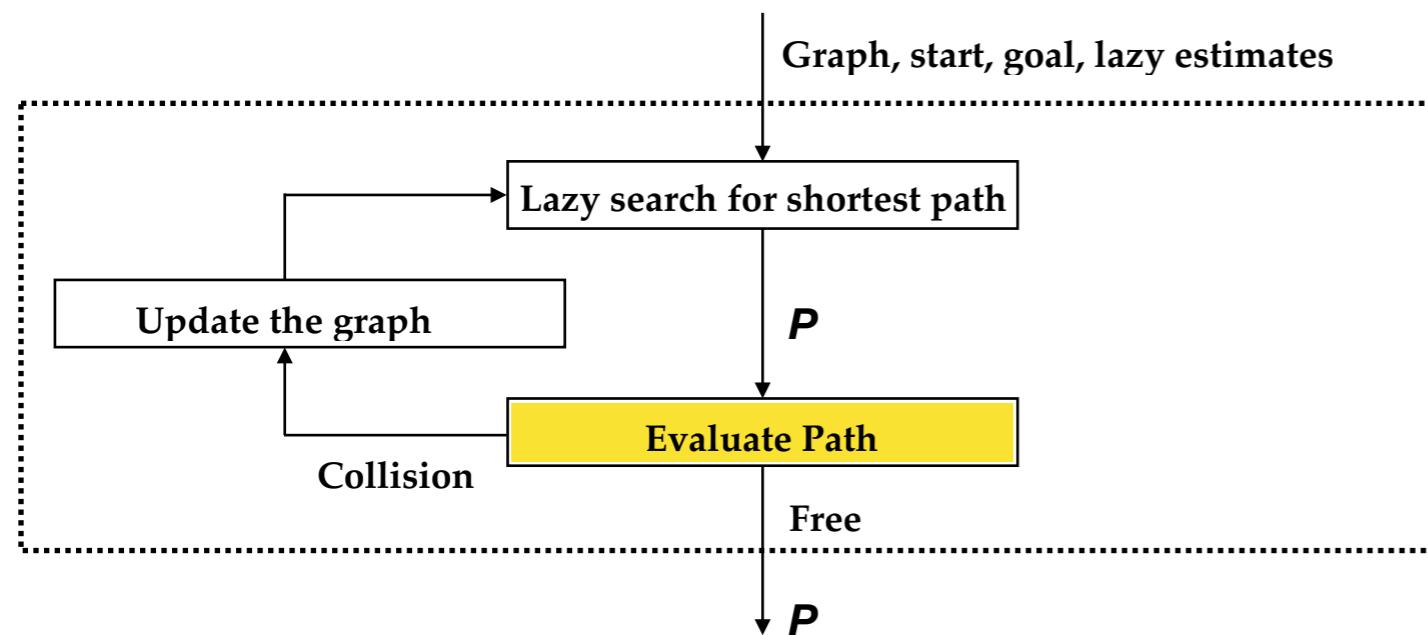
# LazySP

## Optimism Under Uncertainty



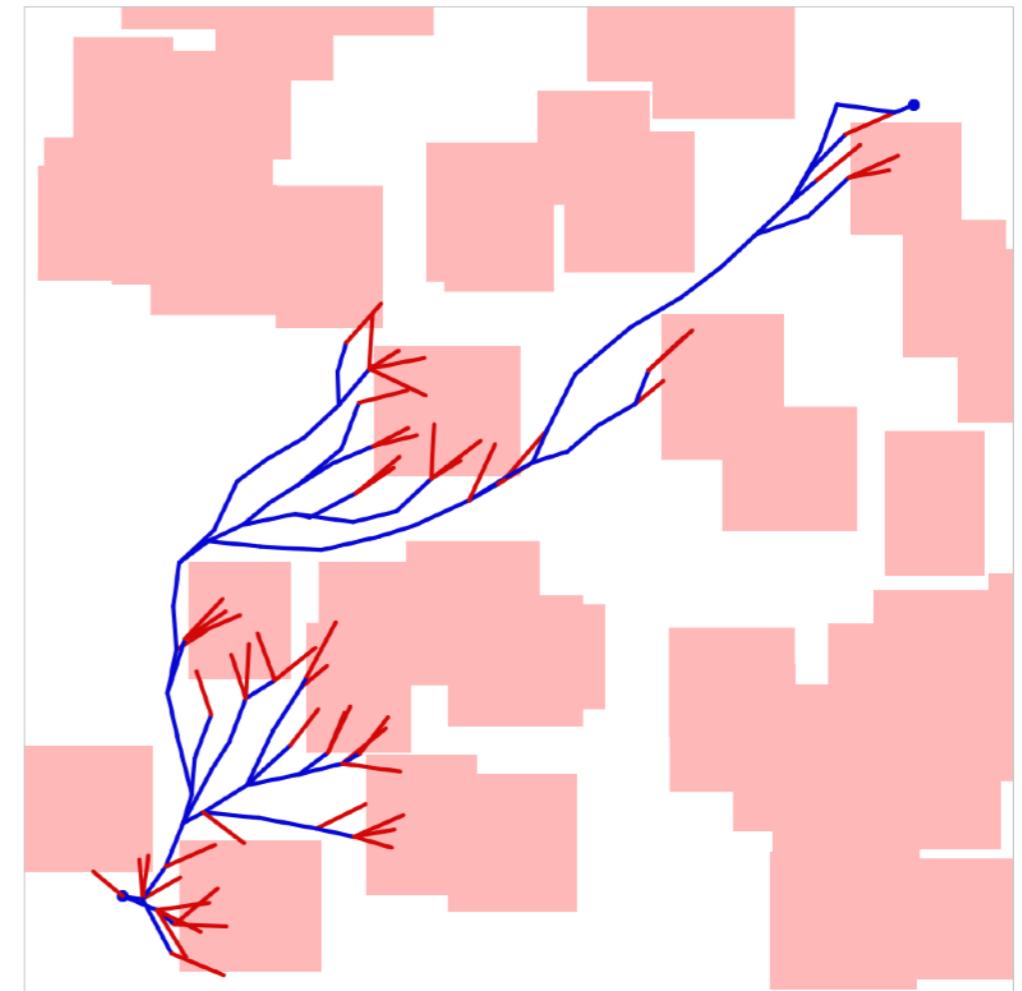
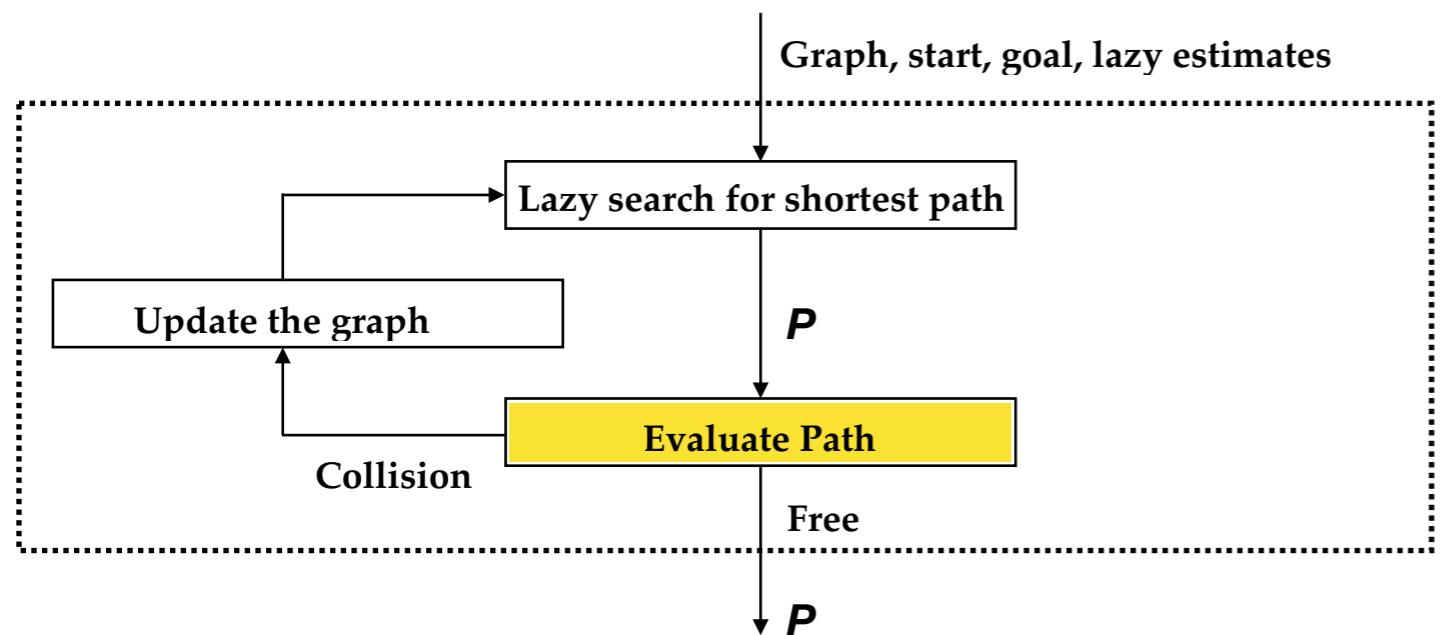
# LazySP

## Optimism Under Uncertainty



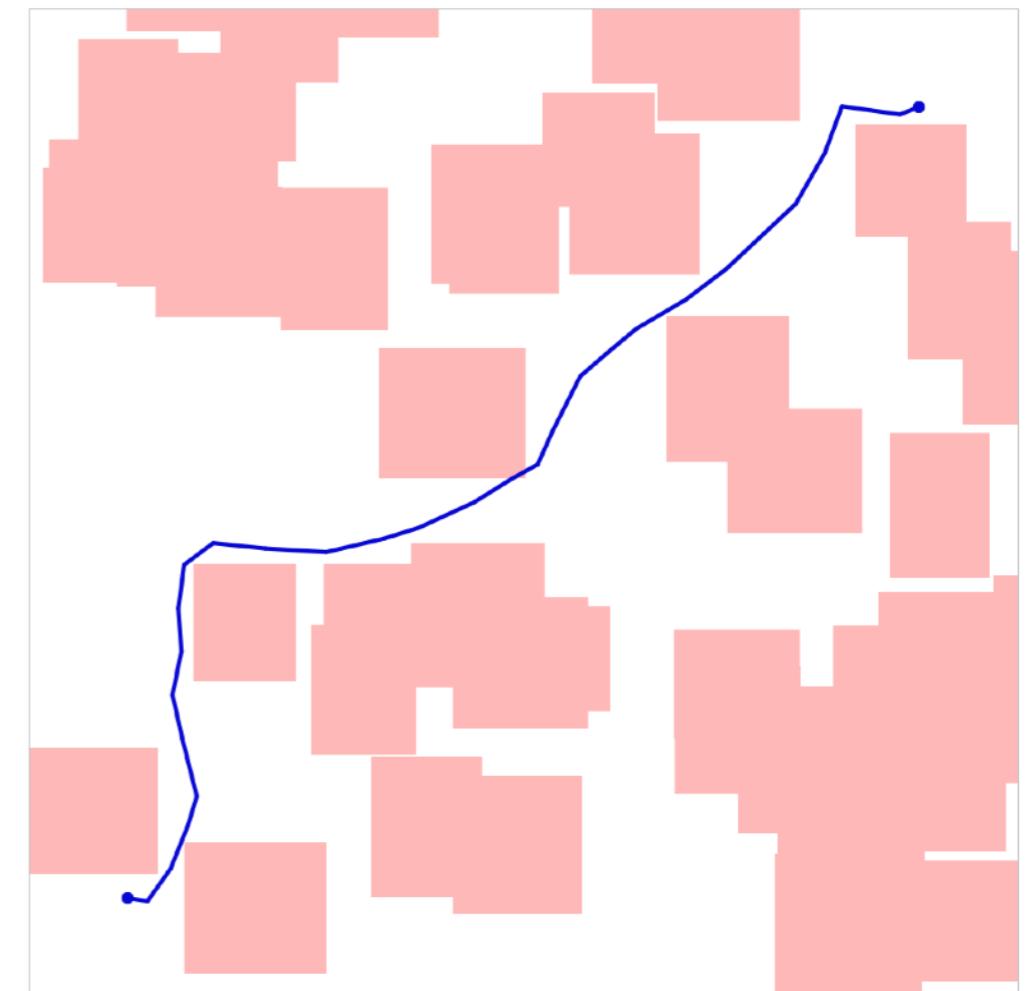
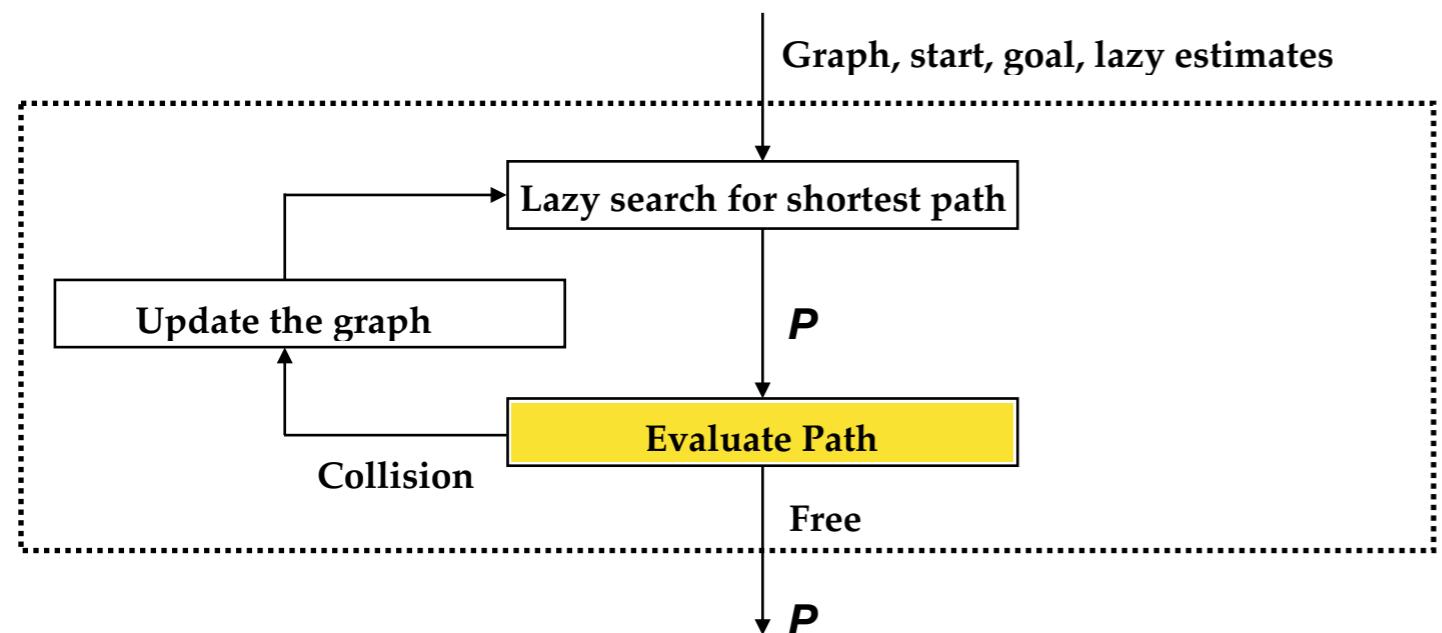
# LazySP

## Optimism Under Uncertainty



# LazySP

## Optimism Under Uncertainty



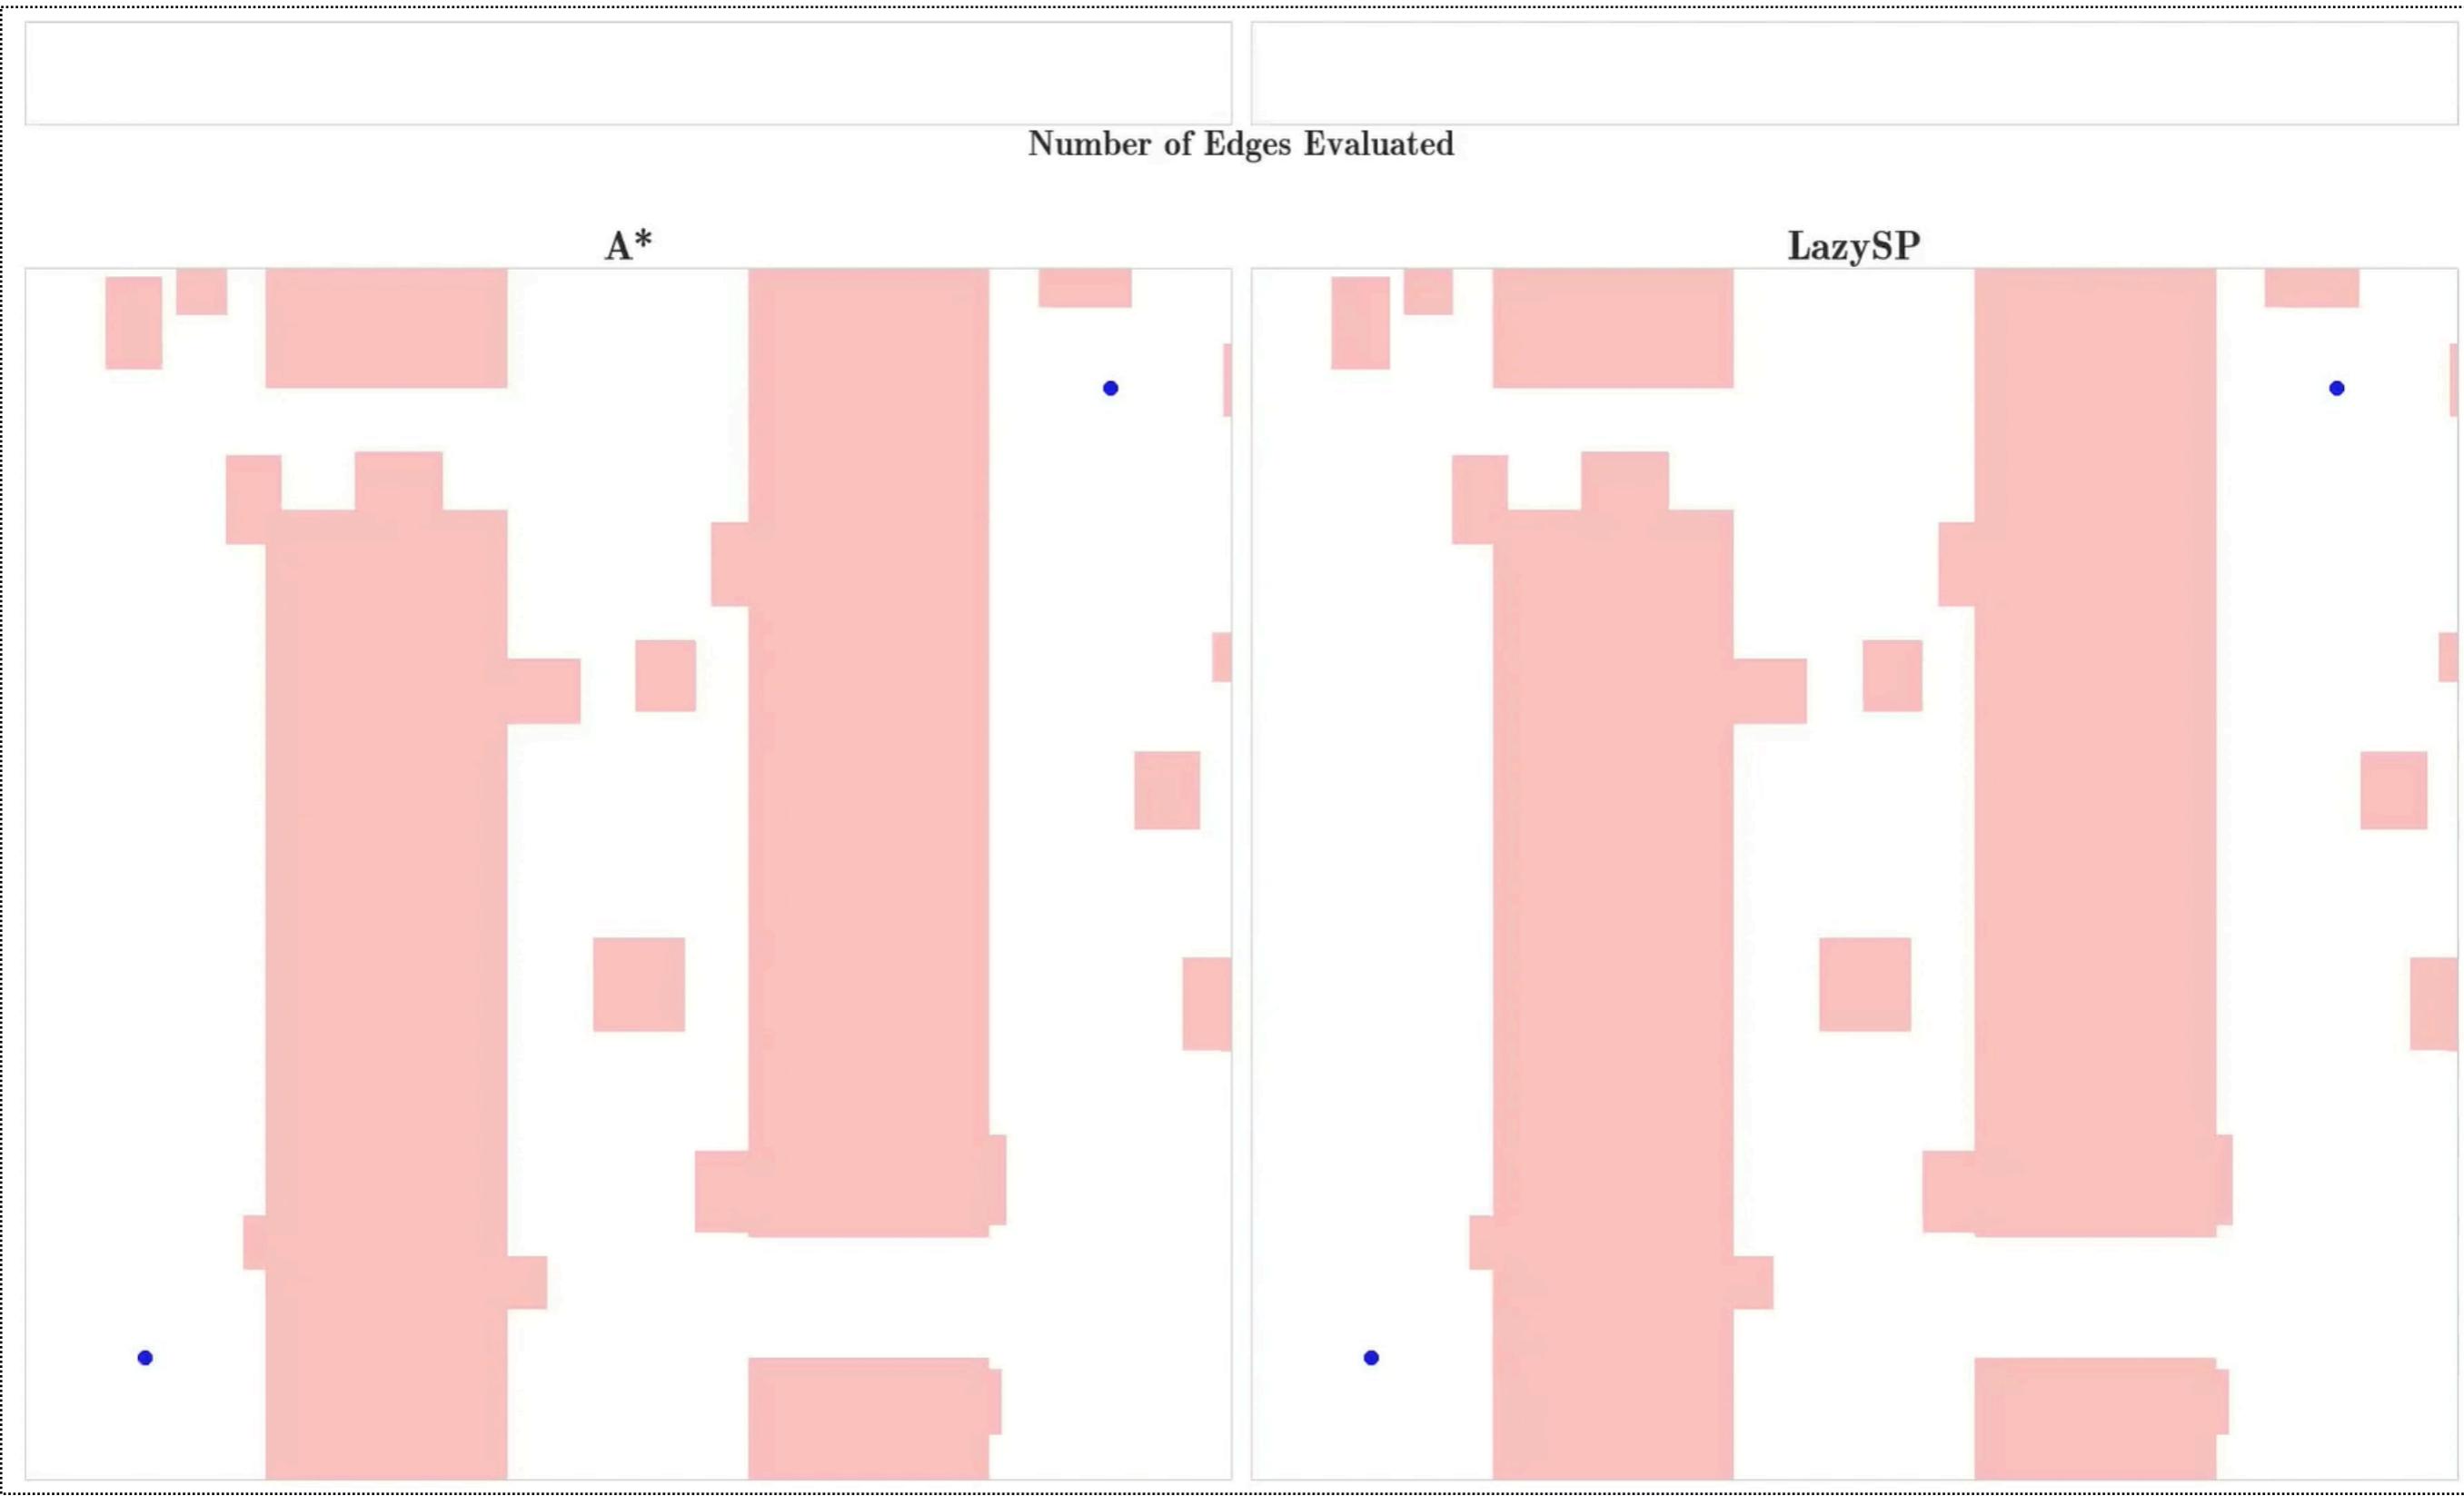
# Comparison across environments



# Comparison across environments



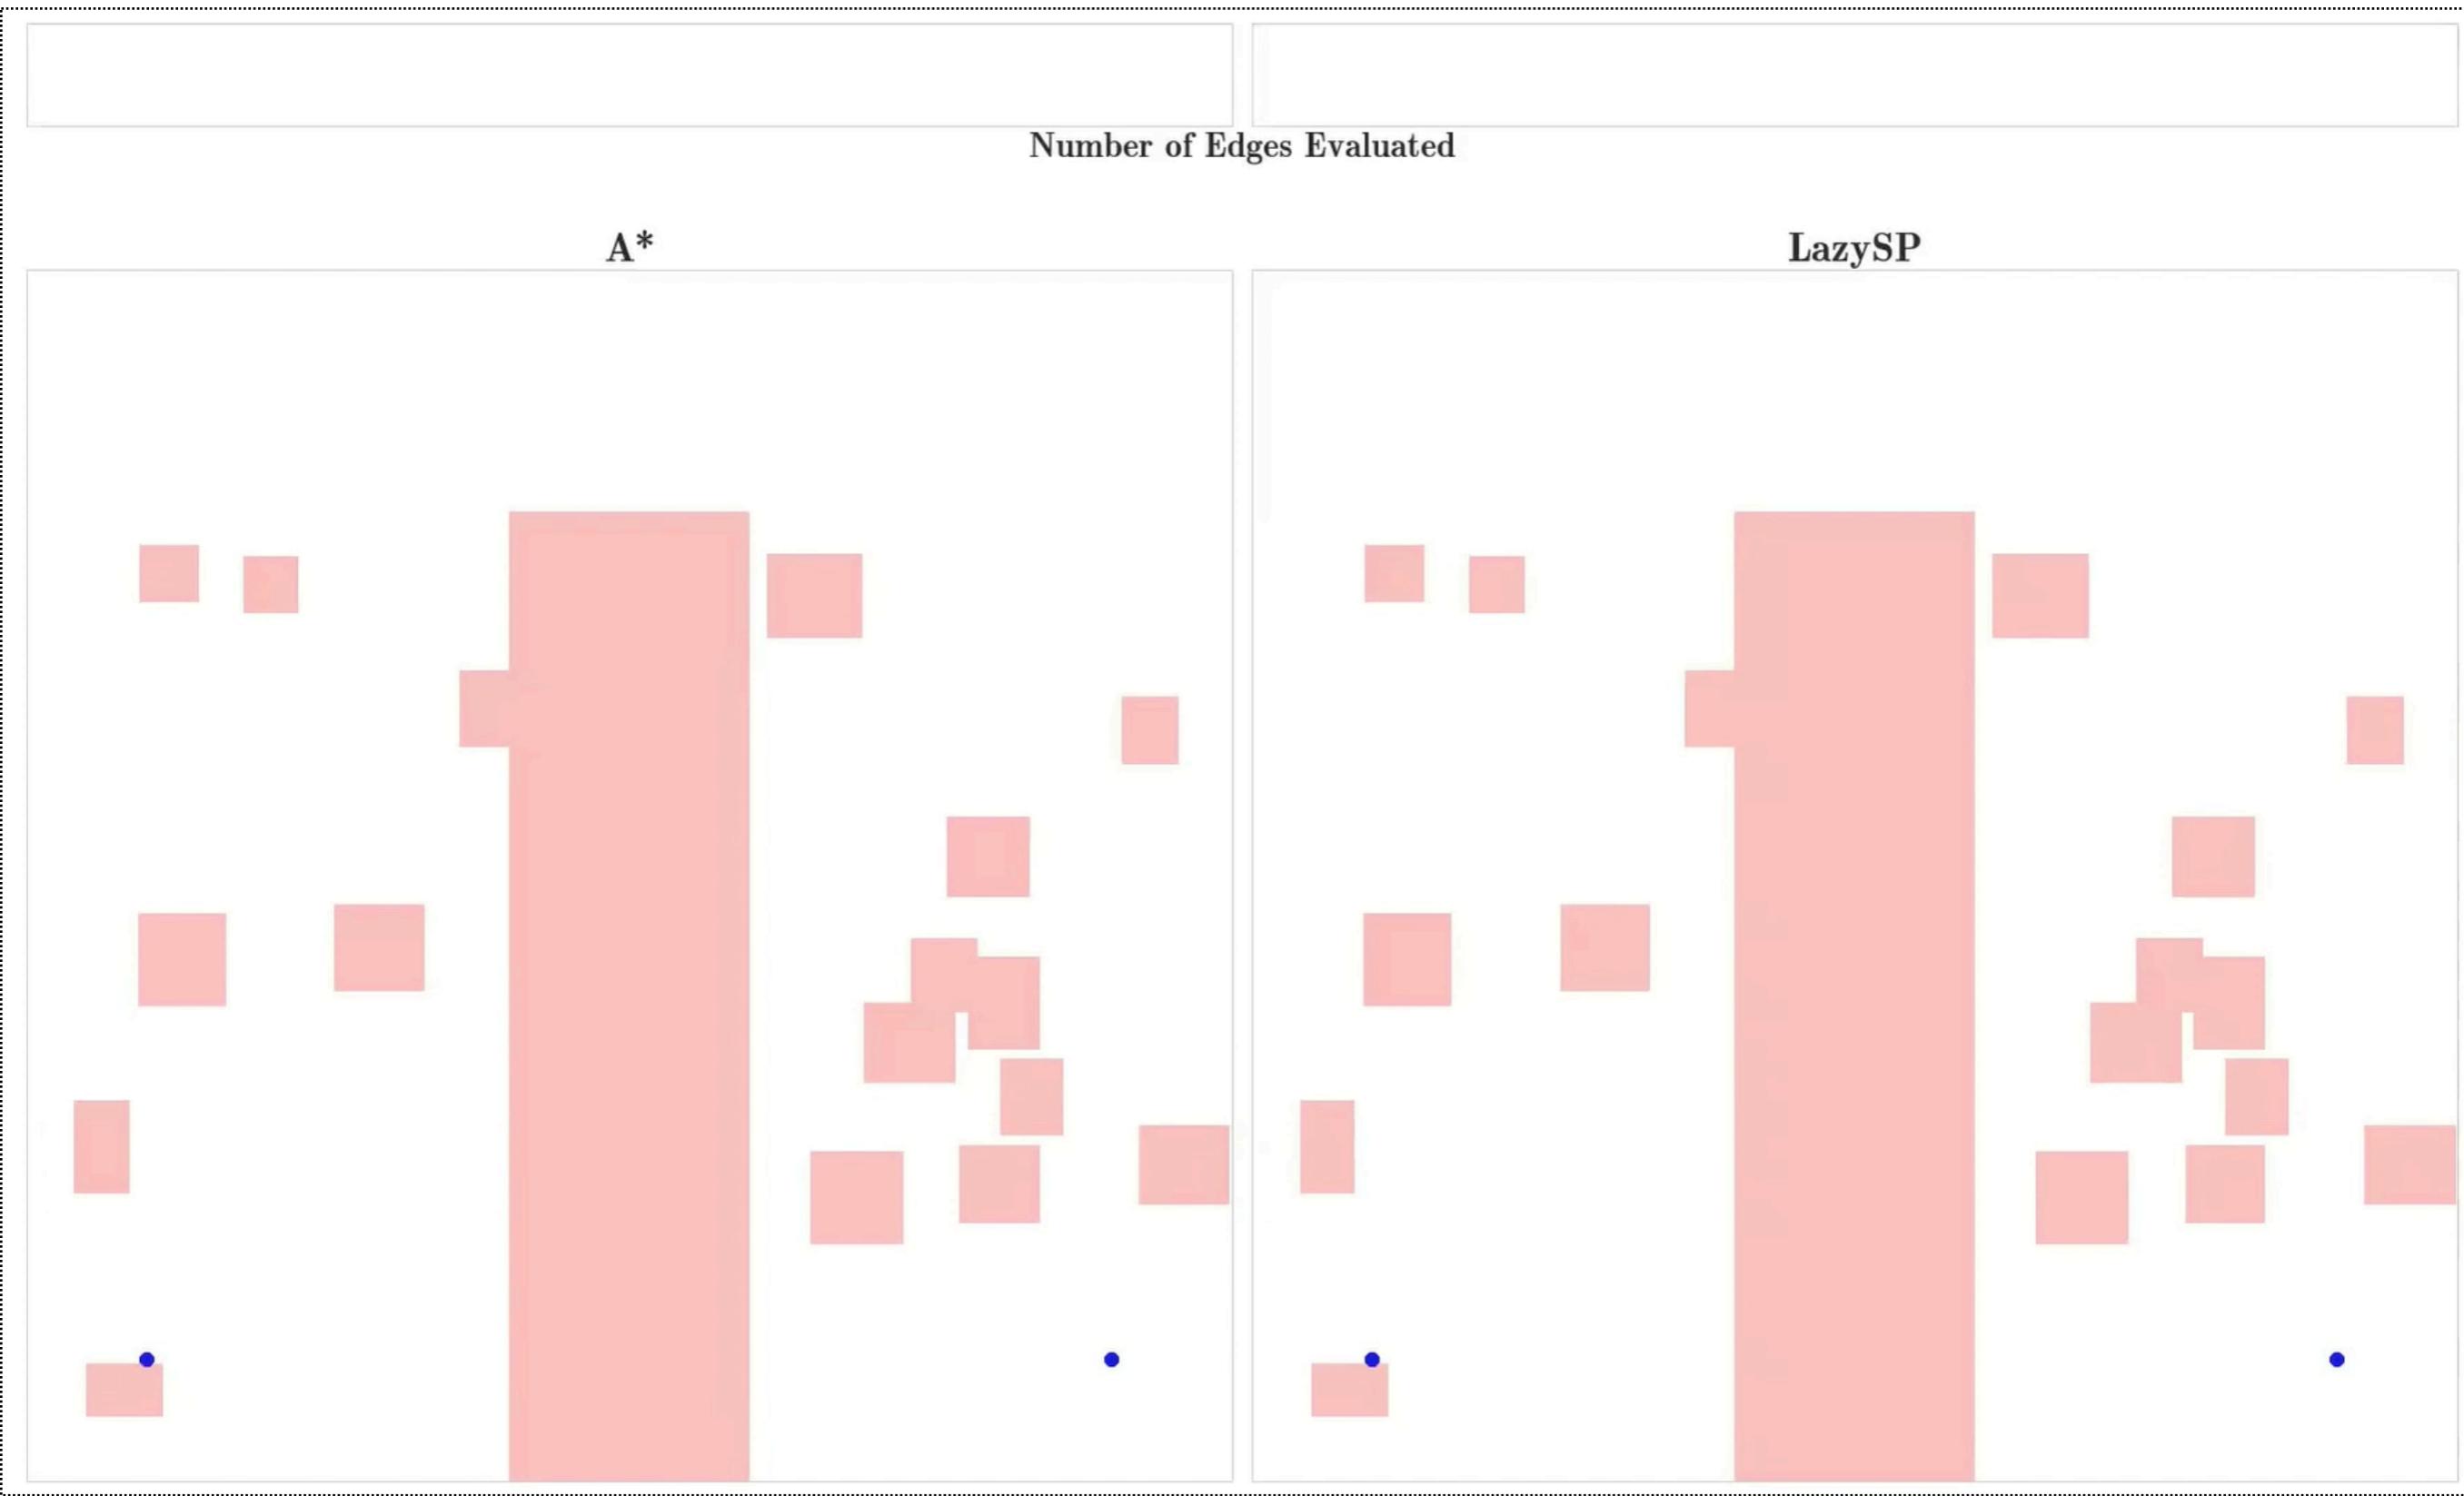
# Comparison across environments



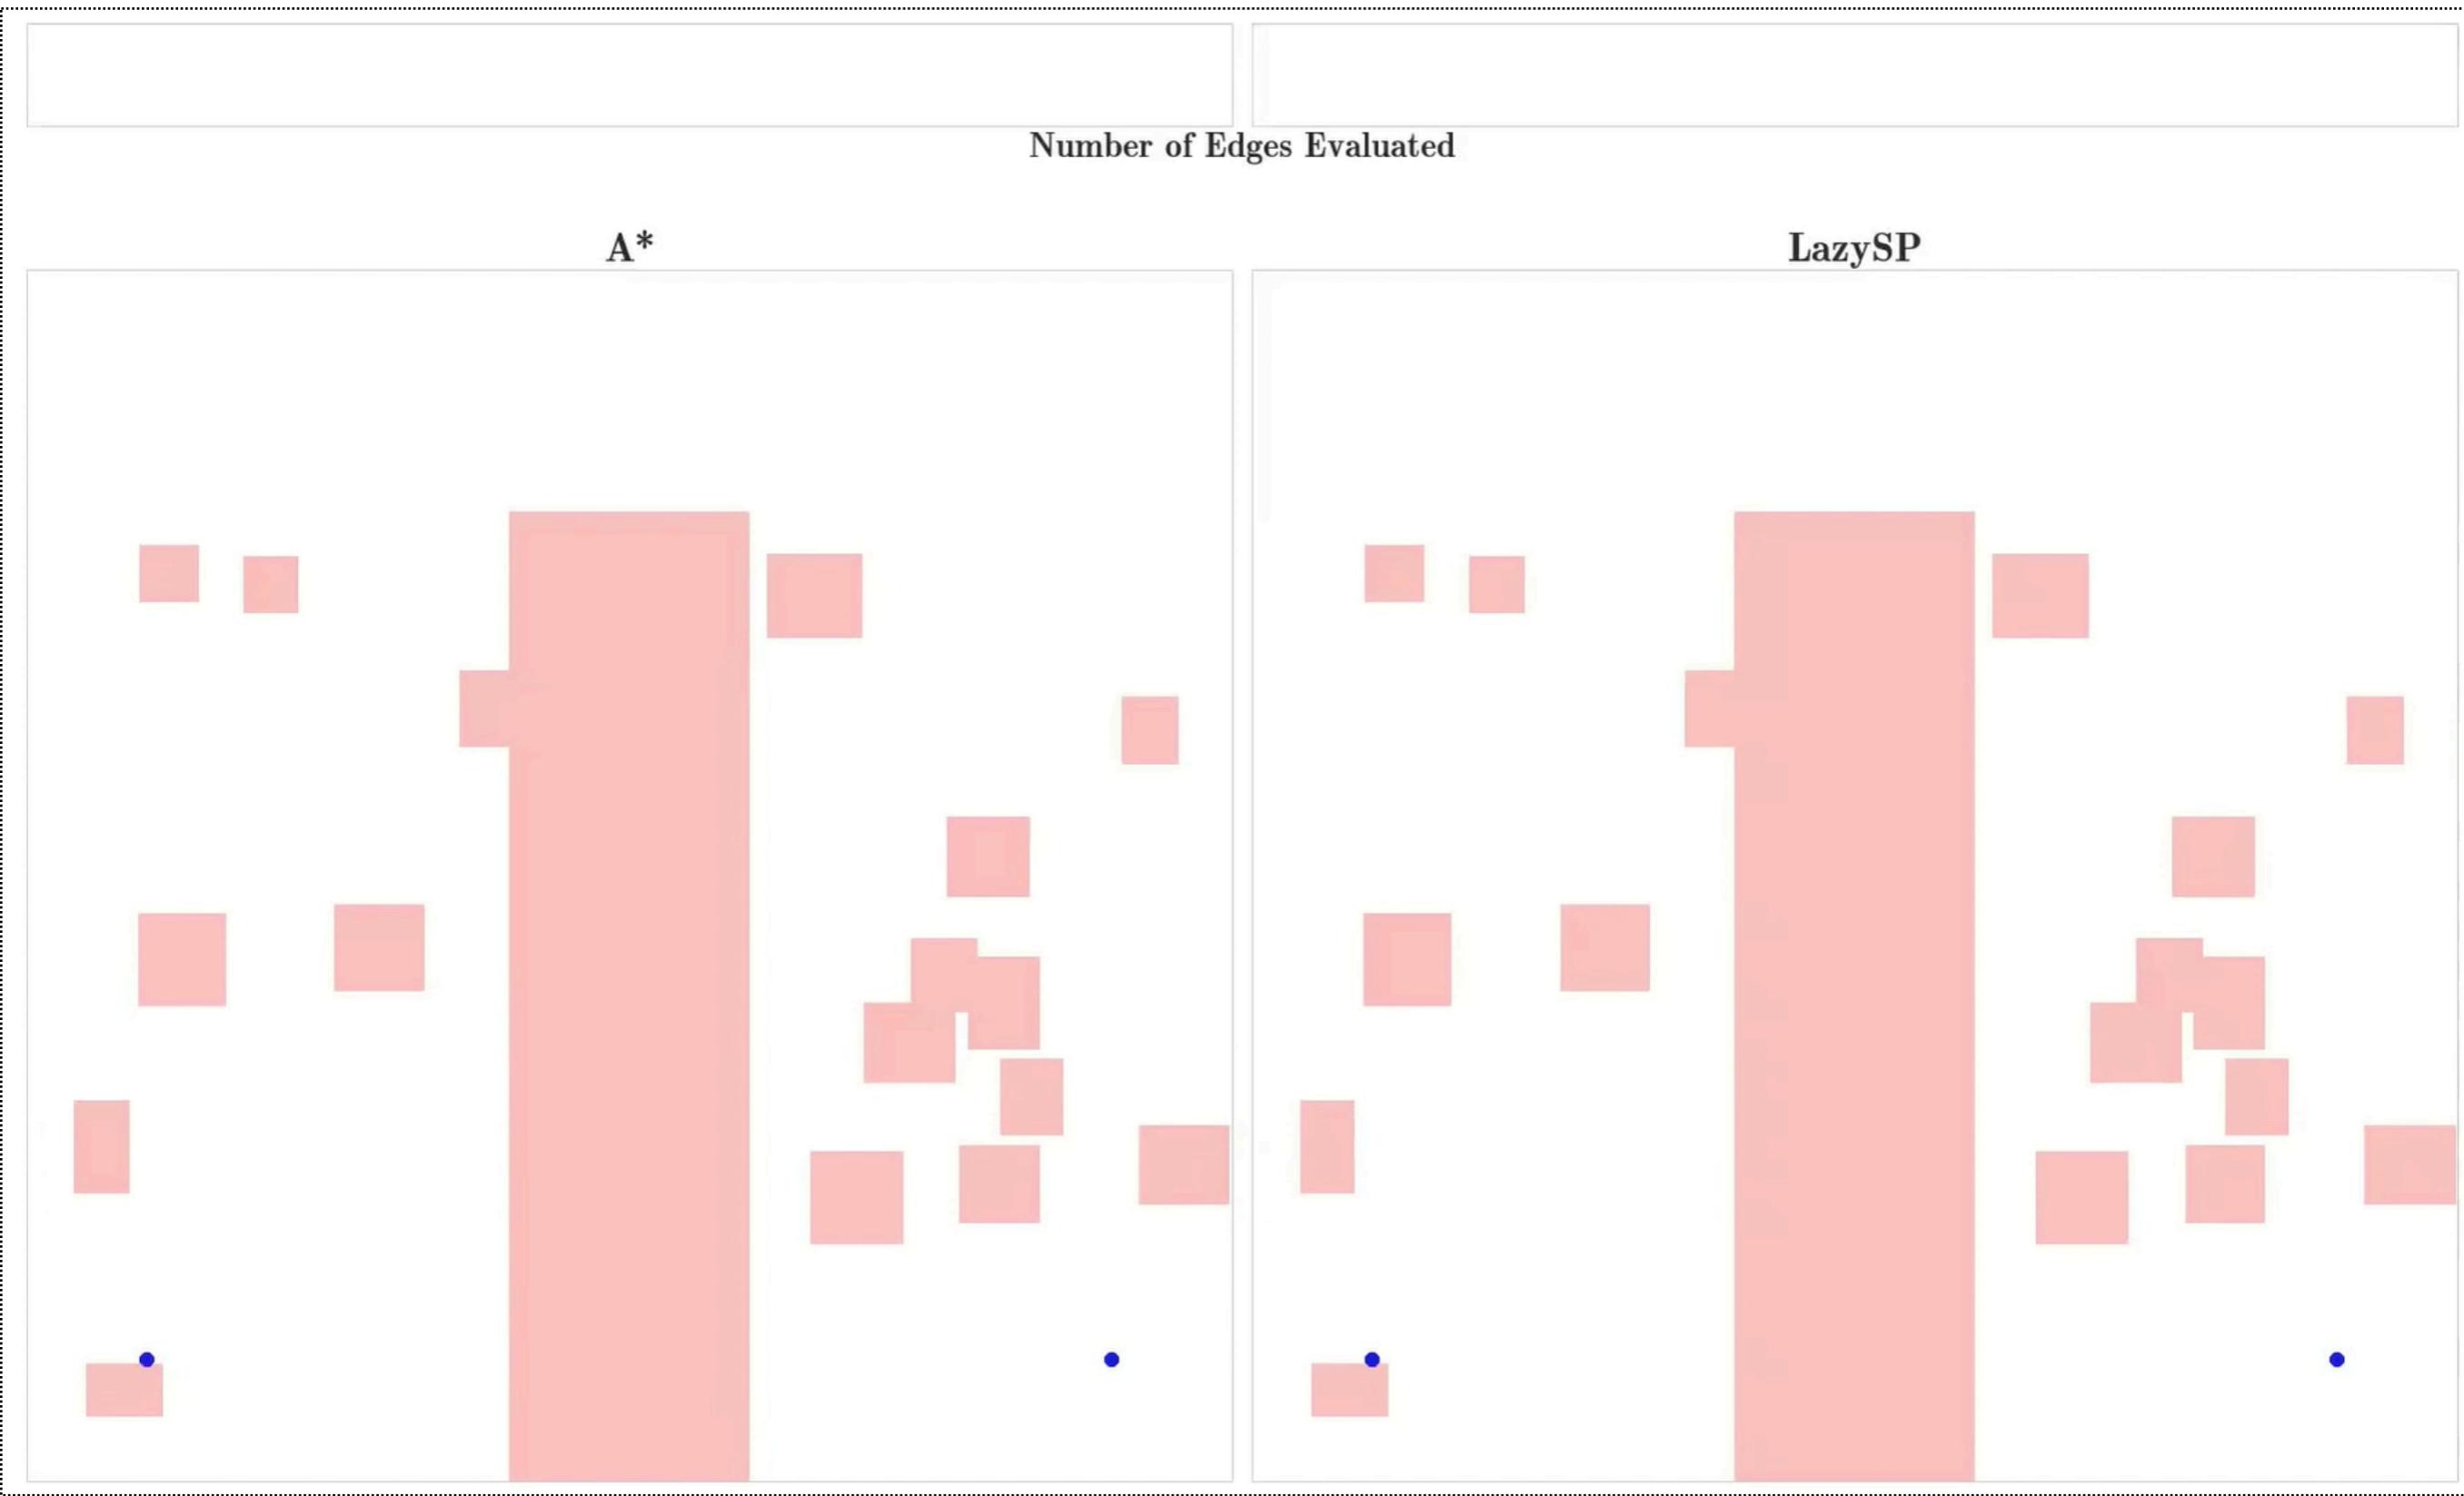
# Comparison across environments



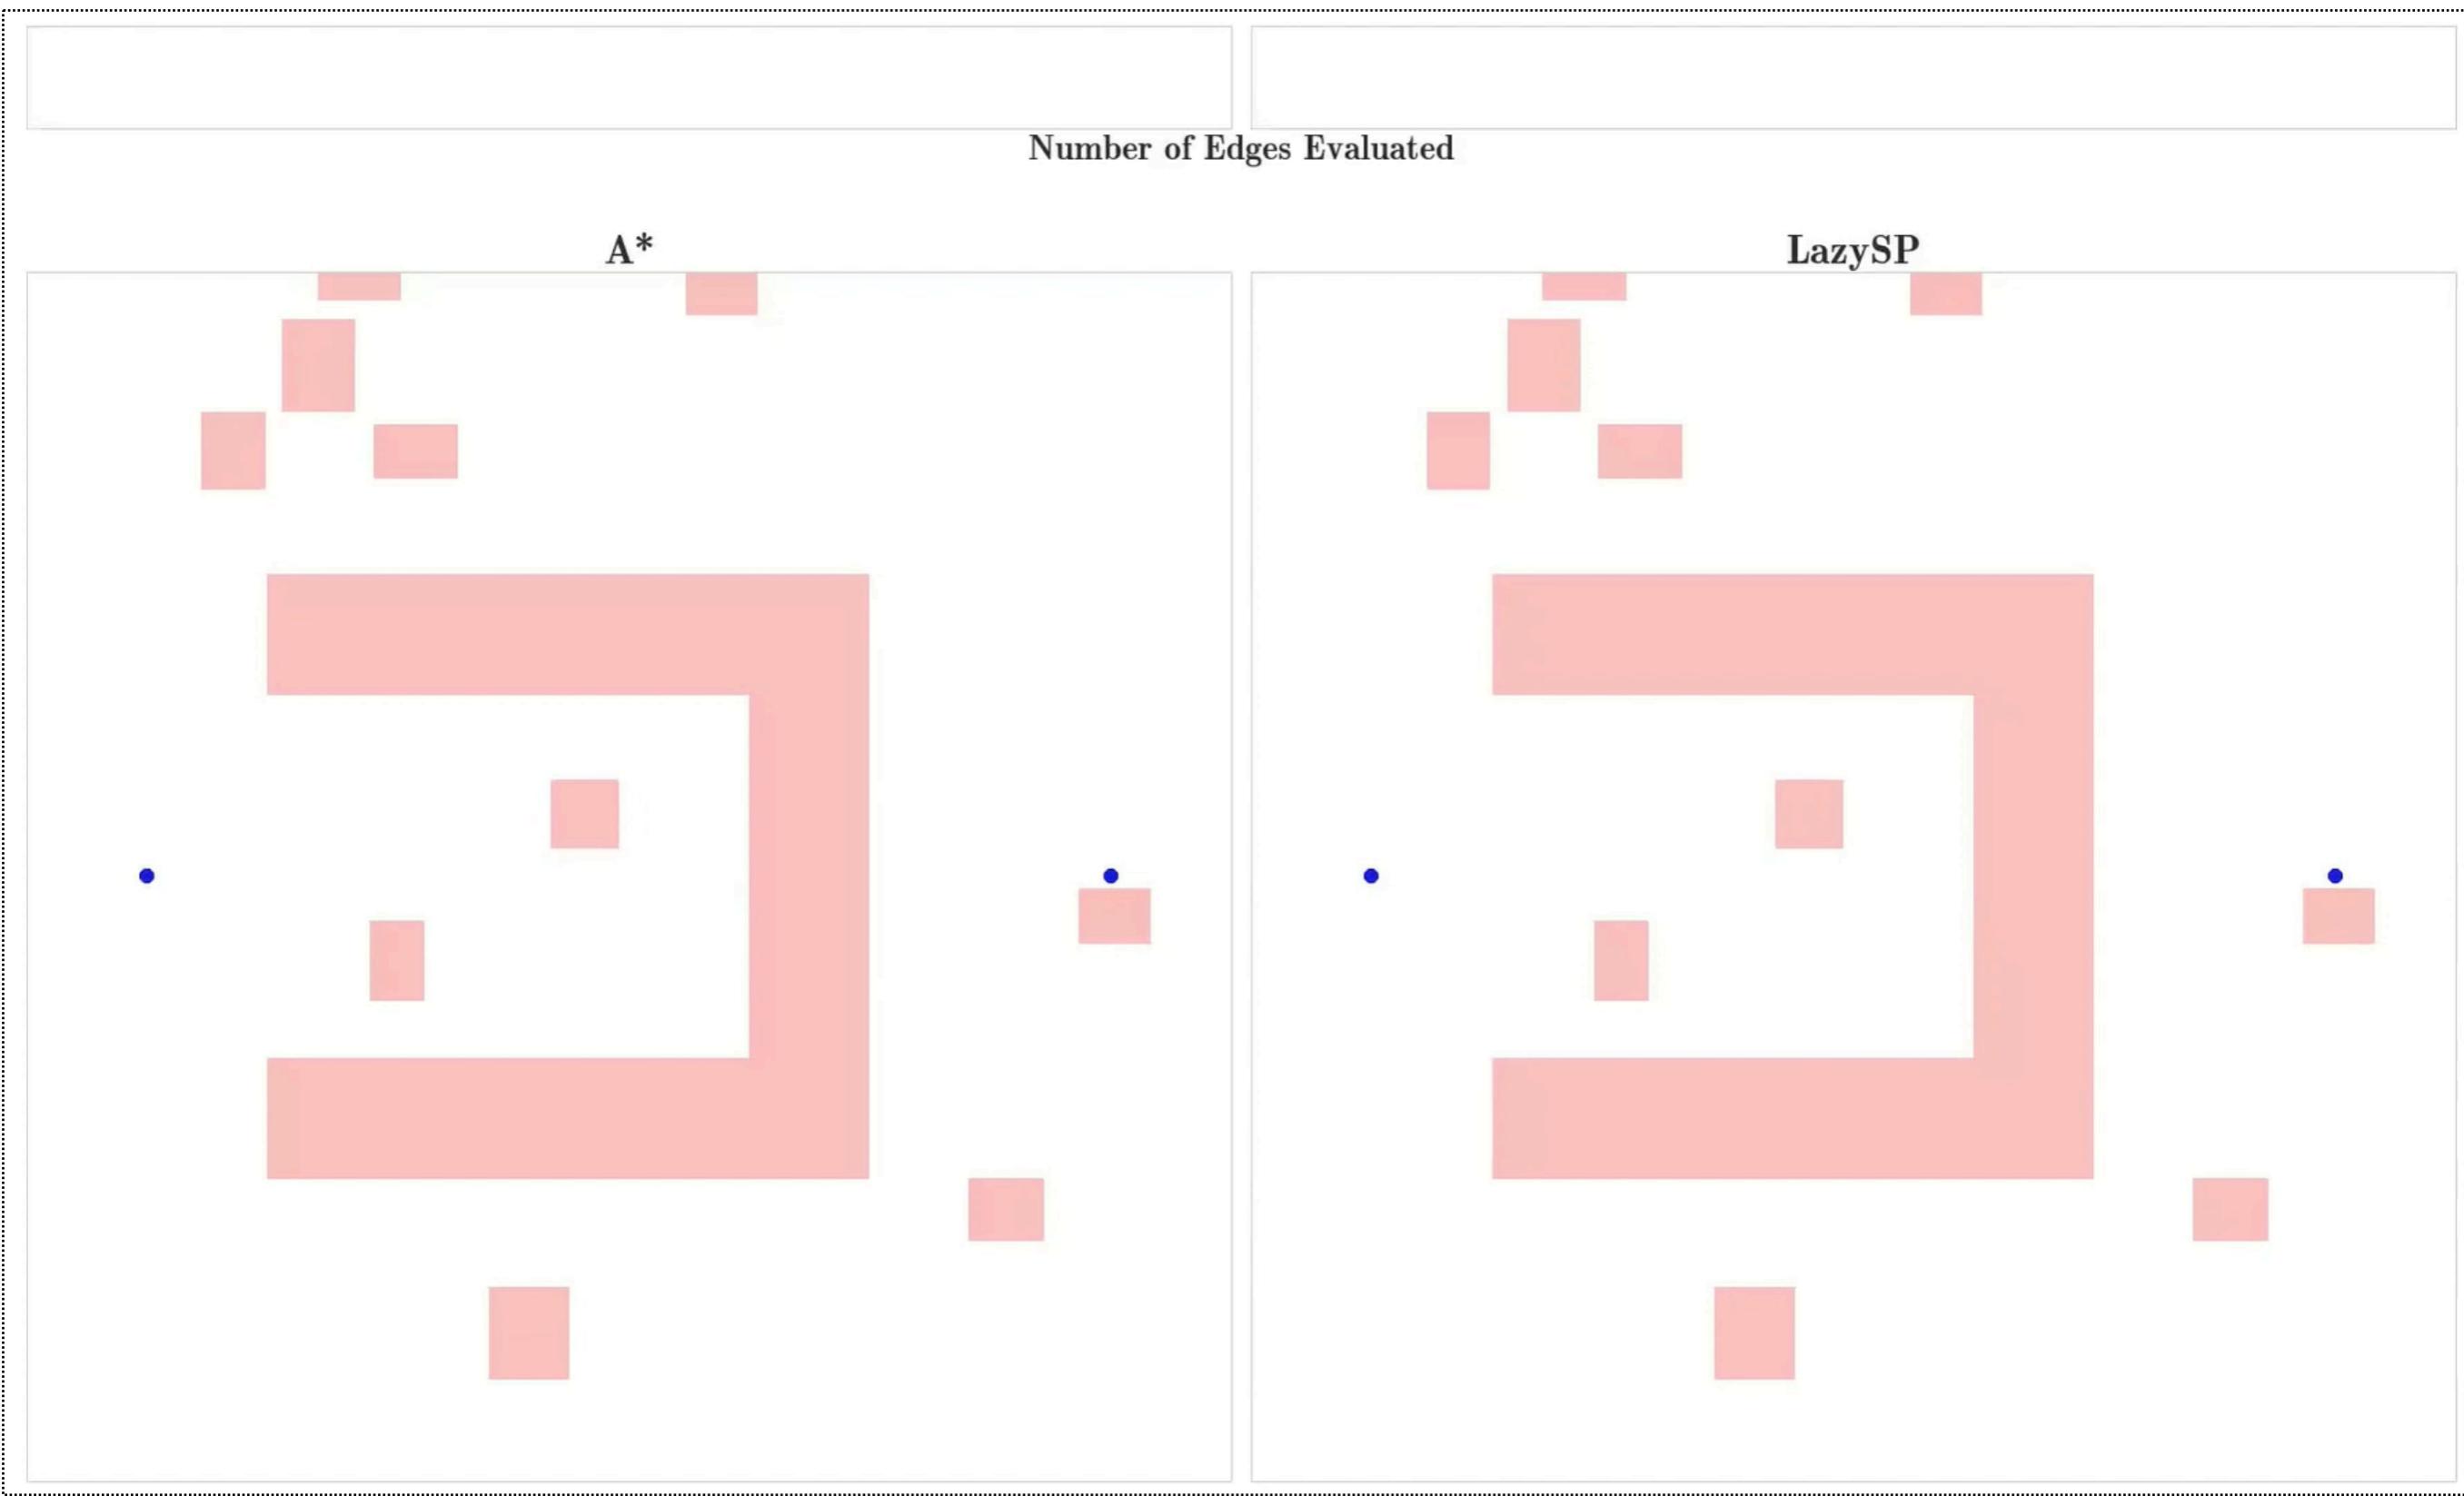
# Comparison across environments



# Comparison across environments



# Comparison across environments



# Comparison across environments

